

Experimental apparatus for universal spin-motion control in multi-ion registers



Donovan Webb
St Peter's College
University of Oxford

A thesis submitted for the degree of
Doctor of Philosophy in Physics

Trinity 2026

Contents

1	Towards Hybrid System Control and Quantum Error Mitigation	1
1.1	Continuous-Variable Control in Trapped Ions	2
1.1.1	Motivation	2
1.1.2	Theory	3
	Computing on a Hybrid System	3
	Gate Implementation	3
	Readout	3
1.1.3	Experimental Demonstration	3
	Two Spin Dependent Forces	4
	Squeezing	4
1.1.4	Outlook	5
1.2	Virtual Distillation	5
1.2.1	Error Mitigation	5
1.2.2	Virtual Distillation Protocol	7
	Purification	7
	Circuit	10
1.2.3	Experimental Demonstration	10
	Implementation	11
	B-matrix Calibration	11
	Injecting Spatial Correlations	11
1.3	Proposed experiments and technical requirements	12
	Extending to > 2 copies	12
	Virtual Distillation on Hybrid Systems	12
1.4	Outlook	13

Appendices

References	15
-------------------	-----------

1

Towards Hybrid System Control and Quantum Error Mitigation

Contents

1.1	Continuous-Variable Control in Trapped Ions	2
1.1.1	Motivation	2
1.1.2	Theory	3
1.1.3	Experimental Demonstration	3
1.1.4	Outlook	5
1.2	Virtual Distillation	5
1.2.1	Error Mitigation	5
1.2.2	Virtual Distillation Protocol	7
1.2.3	Experimental Demonstration	10
1.3	Proposed experiments and technical requirements . .	12
1.4	Outlook	13

Keep in mind: Want to make clear that this is future work, and not specifically QEM on CV! Suggested rough intro: Having established and characterised the experimental platform and demonstrated its application to *k-copy* randomised benchmarking, this chapter considers future research goals utilising the multi-qubit array, and control over both the discrete qubits subsystems, as well as the continuous bosonic motional modes. In broad terms, we are interested in including the motional modes as computational resources, following recent progress in spin-conditioned bosonic operations demonstrated within our

group [bazavan_thesis, saner_thesis, shinbrough_applications_2026, 1, 2]. In section 1.1 we establish a gateset required for continuous-variable control of the ion motion, and demonstrate the generation of a squeezed state on the new apparatus.

The second half of this chapter considers an error mitigation technique applicable to both discrete and continuous systems. Simulation and computation utilising both discrete and continuous resources has been proposed to offer advantages over fully discrete systems due to resource scaling... [<empty citation>]. Demonstrations include [feynman, 3]. However, to extract faithful results, we must reduce the impact of errors. We explore in section 1.2 [virtual_cooling, balint, huggins] the error mitigation technique virtual distillation, which has applications on both discrete [google, 4] and continuous systems [<empty citation>] **virtual cooling**¹. This section begins with the theory behind virtual distillation, and two protocols for its application. We then provide initial demonstrations of this protocol applied to a discrete-variable only system, and explore the performance under both local depolarising and under the correlated noise due to a common fluctuating control field, explored in detail in chapter ???. We end this chapter with an outlook of how VD will be applied to larger system sizes, including the ultimate goal of application to hybrid systems.

1.1 Continuous-Variable Control in Trapped Ions

Short summary of section One paragraph on why this apparatus is well placed for this work. **this section develops the CV gateset that is a prerequisite for hybrid QEM, without yet performing error mitigation on it**

1.1.1 Motivation

One paragraphs on what is CV, and why it is interesting. One paragraph on its conception in ITQC. One paragraph on why error mitigation in this framework is important.

¹In fact, the initial proposal of virtual cooling was on a continuous system

Gate	H_{osc}	Gaussian	modes num.	σ	m
Displacement	$a + a^\dagger$	yes	1	σ_x	0
Squeezing	$a + a^\dagger$	yes	1	σ_z	1
two-mode squeezing	$a + a^\dagger$	yes	2	σ_z	1
dispersive	$a + a^\dagger$	yes	1	σ_z	1
beam splitter	$a + a^\dagger$	yes	2	σ_z	1
3rd order squeezing	$a + a^\dagger$	no	1	σ_x	2

Table 1.1: Gateset, need to check factors and phases for all gates. Check if Hermitian!

1.1.2 Theory

Computing on a Hybrid System

How information is stored in oscillator and spin, and retrieved via spin?

Summary from a couple of reviews (Liu?)

Gate Implementation

Two-SDF toolkit from TS+RS

Could also mention SW and higher order S.B. approaches.

Magnus expansion

point to gateset table of how to generate interaction using TS+RS method

Readout

Define state for CV.



Define characteristic and Wigner function.

$$\mathcal{P}_a = \frac{1 + \text{Re}(\xi(\beta))}{2}, \quad (1.1)$$

$$\mathcal{P}_a = \frac{1 + \text{Im}(\xi(\beta))}{2}, \quad (1.2)$$

1.1.3 Experimental Demonstration

Short description of what we will show

**Figure 1.1: Figure: circuit diagram for tomography.****Figure 1.2: Figure: characteristic function of $|0 + 1\rangle$ Fock state.**


Two Spin Dependent Forces

Experimental sequence, hardware involved, thermal drifts

Four tones on one AOM.

Squeezing

Experimental sequence, dephasing dominated, improved by grounding

**Figure 1.3: Figure: detuning for four tones, Z measurement with ramp**




Figure 1.4: Figure: Sqz characteristic function with varying duration

1.1.4 Outlook

Some next steps for demonstrating full CV control

1.2 Virtual Distillation

Short summary of section Here we develop experimental tools for demonstrating VD on the discrete variable portion of the full hybrid system.

1.2.1 Error Mitigation

Noise coupled to a quantum system results in unintended coherent, and incoherent evolution. Deviations of the realised state from the target state are known as errors, and, if left unchecked, lead to the failure of quantum algorithms. This failure can either be due to a reduction of signal to noise ratio, or a bias in the measured expectation value, resulting in misleading conclusions. Ultimately, the goal of quantum error correction, QEC, is to detect and correct these errors through embedding information in a larger Hilbert space [qec_review]. Full QEC will allow for the execution of arbitrarily long algorithms ensuring errors (of up to a certain distance) do not accumulate over time. However, current QEC methods require large overheads in both qubit number and circuit depth, and are therefore not feasible on current quantum hardware. Furthermore, QEC is

formulated for digital “gate-based” quantum computers, which although promise to be universal, pay for that privilege in often inefficient decompositions of target unitaries into a discrete set of gates. This poses an opportunity for near term devices to show quantum advantage using either short depth, uncorrected circuits, or analog quantum simulators, which although not universal, can provide meaningful results for specific problems [**nisq_review**]. Particularly, phrasing the problem in terms of expectation values of an observable on a physical final state, rather than a sequence of measurements on a set of logical state, allows for a more efficient use of resources. As a tactile example, consider the problem of finding the ground state energy of a Hamiltonian. This can be done by preparing a trial state, and directly measuring the expectation value of the Hamiltonian on that state, building up statistics over many samples [**variational_qc_review**]. Increasing the number of samples will reduce the statistical error, however, if the trial state is not prepared perfectly due to either environmental noise, or to imperfect application of unitary evolutions, the measured expectation value will be biased. To this end, quantum error mitigation (QEM) techniques have been developed to reduce the bias of measured expectation value of an observable at the expense of requiring more samples. Reference [5] defines QEM as “... algorithmic schemes that reduce the noise-induced bias in the expectation value by postprocessing outputs from an ensemble of circuit runs, using circuits at the same noise level as the original unmitigated circuit or above.” This highlights the crux of how information on the system is collected by a QEM implementation, which I shall explore with the following thought experiment on the “perfect” QEM algorithm. Consider the problem of measuring the expectation value of an observable $\hat{O}$ on the target state ρ_{id} . The state is prepared by a quantum circuit, which due to noise, instead prepares the mixed state $\rho = (1 - \lambda)\rho_{id} + \lambda\rho_{err}$, where ρ_{err} is the error state. Each measurement of the observable $\hat{O}$ on the noisy state ρ will return a value o_i , and the average of these measurements will converge to the expectation value $\langle \hat{O} \rangle = \text{Tr}[\hat{O}\rho]$. Given S total samples, the final result will be,

$$\langle \hat{O} \rangle = \frac{1}{S} \left(\sum_{i=1}^{S_{id}} o_i + \sum_{j=1}^{S_{err}} o_j \right),$$




Figure 1.5: Figure: unmitigated vs mitigated expectation value, showing how bias is reduced at the expense of more samples.

where S_{id} and S_{err} are the number of samples that return a measurement outcome from the ideal and error states respectively, with $S = S_{id} + S_{err}$. The bias in the measured expectation value from the ideal is then given by contributions from the samples S_{err} . An ideal QEM algorithm would flag these error samples, and remove them from the final average, resulting in a bias free expectation value using the restricted samples S_{id} .

- Error mitigation motivation. Mainly on contrasting to quantum error correction.
- QEM is framed as finding expectation value of observable, rather than maintaining coherence for a long computation.
- How QEM reduces bias (at the expense of required samples). Look at Cai paper for structured argument.

1.2.2 Virtual Distillation Protocol

Purification

Estimate expectation value of observable on a state without ever physically preparing the state. Purification here refers to approaching the dominant pure state of a

mixed state. The k th degree purified state is given by,

$$\rho_{\text{corrected}}^{(k)} = \frac{\rho^k}{\text{Tr}[\rho^k]} = \sum_{i=1}^d \frac{\lambda_i}{\sum_{i=1}^d \lambda_i^k} |\psi_i\rangle \langle \psi_i|, \quad (1.3)$$

where d is the dimension of the Hilbert space, λ_i are the eigenvalues of ρ , and $|\psi_i\rangle$ are the corresponding eigenvectors. If we assume that one eigenvalue is dominant, $\lambda_1 \gg \lambda_{i>1}$, then the normalised purified state exponentially approaches the dominant eigenvector, $\rho_{\text{corrected}}^{(k)} \rightarrow |\psi_1\rangle \langle \psi_1|$, as k increases. The hope, which we shall verify in section XX on coherent mismatch, is that the dominant eigenvector of the noisy state closely approximates the ideal state, and therefore effective measurements on the purified state yield accurate expectation values of the ideal state. The estimated expectation value of measuring an observable, $\hat{O}$, on this purified state is then,

$$\langle \hat{O} \rangle_{\text{corrected}} = \text{Tr}[\hat{O} \rho_{\text{corrected}}^{(k)}] = \frac{\text{Tr}[\hat{O} \rho^k]}{\text{Tr}[\rho^k]}. \quad (1.4)$$

Next, how to access this corrected state on which to make these measurements? Kozcor and Huggins show that this can be accessed without explicit preparation of the purified state. Instead the numerator and denominator may be found through preparing k copies of the noisy state, applying some derangement operation to entangle all copies together, and measuring the desired observable on one of the copies. This can be understood through the following identity,

$$\text{Tr}[\hat{O} \rho^k] = \text{Tr}[\hat{O}^i S^{(k)} \rho^{\otimes k}], \quad (1.5)$$

where $S^{(k)}$ is the derangement operator, which permutes the k copies of the state, and $\hat{O}^i$ is the observable acting on the i th copy. We shall define the derangement operator explicitly as the cyclic permutation of k objects,

$$S^{(k)} |\psi_1\rangle \otimes |\psi_2\rangle \otimes \dots \otimes |\psi_k\rangle = |\psi_2\rangle \otimes |\psi_3\rangle \otimes \dots \otimes |\psi_k\rangle \otimes |\psi_1\rangle. \quad (1.6)$$

Huggins provides a proof of this identity, which I reproduce here for the simple example of $k = 2$. The derangement operator for two copies is simply the swap

operator, $S^{(2)} = \text{SWAP}$, and we will assume the observable acts only on the first copy, giving $\hat{O}^1 = \hat{O} \otimes \mathbb{I}$. The identity is then,

$$\text{Tr}[\hat{O}^1 S^{(2)} \rho^{\otimes 2}] \quad (1.7)$$

$$= \sum_{\substack{i_1, i_2 \\ j_1, j_2 \\ k_1, k_2}} \langle i_1 | \hat{O} | k_1 \rangle \delta_{i_2, k_2} \langle k_1, k_2 | S^{(2)} | j_1, j_2 \rangle \langle j_1 | \rho | i_1 \rangle \langle j_2 | \rho | i_2 \rangle \quad (1.8)$$

$$= \sum_{\substack{i_1, i_2 \\ j_1, j_2 \\ k_1, k_2}} \hat{O}_{i_1, k_1} \delta_{i_2, k_2} \delta_{k_1, j_2} \delta_{k_2, j_1} \rho_{j_1, i_1} \rho_{j_2, i_2} \quad (1.9)$$

and contracting the Kronecker deltas, we find,

$$\text{Tr}[\hat{O}^1 S^{(2)} \rho^{\otimes 2}] \quad (1.10)$$

$$= \sum_{i_1, j_1, j_1} \hat{O}_{i_1, j_2} \rho_{j_1, i_1} \rho_{j_2, j_1} \quad (1.11)$$

$$= \text{Tr}[\hat{O} \rho^2], \quad (1.12)$$

with the last equality following from the definition of matrix multiplication.

The denominator of the corrected expectation value is found by replacing the observable with the identity and measuring the expectation value of the derangement operator on k copies of the noisy state.

Access to this corrected expectation value is then possible if we can measure the multi-qubit observable $\hat{O}^i S^{(k)}$ on k copies of the noisy state. We will discuss two methods of implementing this measurement. First, a Hadamard test, given in Fig. ??, which uses an ancilla qubit to control the derangement and desired observable operation, and then subsequent measurement of the ancilla only. The second method, given in Fig. ??, uses a B-matrix to implement the derangement operation, and then measures the observable on one of the copies. The Hadamard test is more general, as it can be used to measure any observable, however, it requires an additional ancilla qubit and controlled operations. The B-matrix method has better qubit number and operation depth efficiency, however, requires simultaneous diagonalisation of $S^{(k)}$ and $\hat{O}^i S^{(k)}$, which is not always possible for an arbitrary observable and derangement operator.

Next, need to derive scaling of bias with copy number. And finally, discuss the variance of measurement with shot number (and how this scales with copy number?).




Figure 1.6: **Figure: circuit diagram for ancilla based virtual distillation, Hadamard test, derangement operator by CSwap**




Figure 1.7: **Figure: circuit diagram for B matrix**

Circuit

Ancilla method

B-matrix protocol

If we can diagonalise both S and SO . Can apply transversally

On hybrid systems

It can be used for both spin and bosonic objects

1.2.3 Experimental Demonstration

Short description of following results.




Figure 1.8: Figure: B matrix decomposition.




Figure 1.9: Figure: pi-4 MS gate calibration. Figure: Heuristic signal of B-matrix working. Tomography would be best!

Implementation

Protocol, circuit diagram, hardware involved.

B-matrix Calibration

pi-4 gates, single addressing

Injecting Spatial Correlations

two copy, B matrix demonstration

Analysis of this data.




Figure 1.10: Figure: Cicruit diagram of RB + VD, exp data for noise floor for two copy B-matrix, with and without spatial correlations

1.3 Proposed experiments and technical requirements

How CVDV and VD combine.

Extending to > 2 copies

plan for c-swap implementation.

What is difficult, what numbers we need (gate fidelities, shot number)

Virtual Distillation on Hybrid Systems

Cite Quantum virtual cooling

Why important to develop

What errors does this mitigate?

How it looks on a thermal state and on various quantum states (sqz, coherent, fock)

What is difficult, what numbers we need (gate fidelities, shot number)




Figure 1.11: Figure: Fredkin gate on oscillator, and how to compose derangement transversally.




Figure 1.12: Figure: Simulation of tri-sqz state with VD to show negativity!

1.4 Outlook

Broader perspective of these QEM research directions, and also how we can play with our larger CVDV system

Appendices

References

- [1] O. Băzăvan et al. *Squeezing, trisqueezing, and quadsqueezing in a spin-oscillator system*. en. arXiv:2403.05471 [physics, physics:quant-ph]. 2024.
- [2] S. Saner et al. *Generating arbitrary superpositions of nonclassical quantum harmonic oscillator states*. arXiv:2409.03482 [quant-ph]. 2024.
- [3] S. Saner et al. *Real-Time Observation of Aharonov-Bohm Interference in a $\mathbb{Z}_2$ Lattice Gauge Theory on a Hybrid Qubit-Oscillator Quantum Computer*. en. arXiv:2507.19588 [quant-ph]. 2025.
- [4] William J. Huggins et al. “Virtual Distillation for Quantum Error Mitigation”. In: *Phys. Rev. X* 11.4 (2021).
- [5] Zhenyu Cai et al. “Quantum error mitigation”. In: *Rev. Mod. Phys.* 95.4 (2023).