

An Experimental Platform for Multi-Ion Quantum Control

Scalable Error Characterisation
and Hybrid Resources



Donovan Webb
St Peter's College
University of Oxford

A thesis submitted for the degree of
Doctor of Philosophy in Physics

Trinity 2026

Abstract

Write abstract here.

Contents

List of Abbreviations	vii
1 Introduction	1
1.1 Classical and Quantum Computing	1
1.2 Quantum Computing Platforms	3
1.3 Scaling Trapped Ion Systems	4
Challenge I: Scaling Quantum Resources	4
Challenge II: Scaling Control	5
Challenge III: Scaling with Imperfect Systems	6
1.4 Structure of this Thesis	7
2 Trapped-Ion Quantum Computing	8
2.1 The Ion	9
2.2 Spin	12
2.2.1 The Qubit	12
2.2.2 Coherent Spin Control	12
Rabi Oscillations	14
AC Stark Shifts	15
2.3 Motion	15
2.3.1 Linear Paul Trap	15
2.3.2 The Qumode	18
2.3.3 Driving Motion	19
2.4 Spin-Motion	21
2.4.1 Resolved Sideband Interactions	21
2.4.2 The Spin-Dependent Force	23
2.4.3 Geometric Phase Gates	24
3 Apparatus	27
3.1 Specification	28
3.2 The Ion Trap	29
3.2.1 Trap RF Chain	32
3.2.2 Trap DC Voltages	37

3.3	Magnetic Field	39
3.4	Vacuum System and Beam Geometry	39
3.5	Imaging System	42
3.6	Ca ⁺ Laser Systems	43
3.6.1	Frequency Stabilisation	44
3.6.2	Intensity Stabilisation	46
3.6.3	Diode Lasers	47
3.6.4	Narrow Linewidth 729 Laser	50
3.6.5	Single-ion Addressing System	53
3.6.6	Global Addressing System	57
3.7	Control System	58
3.8	Ion Loading	58
3.9	Outlook	59
3.9.1	Limitations	59
4	Calibration and Characterisation	60
4.1	Spin	61
4.1.1	State Preparation	62
4.1.2	Measurement	63
4.1.3	Single-Qubit Rotations	64
	Robust Phase Estimation	65
	Transition Frequency Calibration	66
	Gate Time Calibration	67
4.1.4	Single-Qubit Randomised Benchmarking	69
4.1.5	Spin Coherence Times	73
4.1.6	Single-ion Addressing	75
4.2	Motion	78
4.2.1	Calibrating Mode Frequencies	79
4.2.2	Cooling	80
	Doppler Cooling	80
	Sideband Cooling	81
4.2.3	Heating Rates	84
4.2.4	Motional Coherence Times	85
4.3	Spin-Motion	88
4.3.1	Calibrating the SDF	88
4.3.2	Two-Qubit Entangling Gates	91
4.4	Summary	95

5	<i>k</i>-copy Randomised Benchmarking	96
5.1	Errors and Characterisation	98
5.2	Existing approaches to correlated-noise characterisation	101
5.3	The <i>k</i> -copy RB scheme	103
5.3.1	Protocol	104
5.3.2	Representations of $Cl_1^{\otimes 2}$	105
5.3.3	Common random unitary noise	106
5.3.4	Local depolarising noise	109
5.3.5	Combined model	110
	Asymptotes	111
	Schematic explanation	113
5.3.6	State preparation errors	114
5.4	Parity contrast: an independent probe of coherence	115
5.5	Sampling strategy	117
5.6	Experimental Results	121
5.6.1	Physical implementation	122
5.6.2	Z-basis measurements	123
5.6.3	Coherence measurements	124
5.6.4	Controlled noise injection	126
	Laser phase jitter	126
	Laser detuning scan	128
	Single addressed Stark shifts	131
5.6.5	Scaling to a multi-ion register	133
5.7	Outlook	134
5.7.1	Limitations	135
5.7.2	Future work	136
6	Towards Hybrid System Control and Quantum Error Mitigation	138
6.1	Hybrid Spin-Oscillator Control in Trapped Ions	140
6.1.1	Theory	141
	Gate set for a Hybrid System	141
	Gate Implementation	143
	Readout	145
6.1.2	Experimental Demonstration	147
	Two Spin Dependent Forces	148
	Squeezing	149
6.1.3	Outlook	151
	Limitations	152
6.2	Virtual Distillation	153

6.2.1	Error Mitigation	154
6.2.2	Virtual Distillation Protocol	156
	Circuits	158
6.2.3	Experimental Demonstration	161
	Virtual Distillation on Two Copies	161
	Virtual Distillation under Correlated Errors	164
6.2.4	Outlook	168
	Extending to > 2 copies	169
	Virtual Distillation on Hybrid Systems	170
	Limitations	171
7	Conclusion	173
7.1	Summary	173
	The Experimental Platform	173
	Scaling Characterisation	175
	Scaling Resources	175
	Scaling Imperfect Systems	176
Appendices		
A	Appendix	178
A.1	Magnus expansion	178
A.1.1	AC Stark Shift	179
A.1.2	Geometric Phase	180
A.1.3	“Two-SDF” Dynamics	180
A.2	Gate Fidelity	183
A.3	Irreducible Representations of Cl_1	185
A.3.1	Twirling	185
A.3.2	Character Table of Cl_1	186
A.3.3	Projectors of $Cl_1^{\otimes 2}$	188
A.3.4	ZZ -type noise	189
A.4	Derangement for $k = 2$	191
	References	193

List of Abbreviations

ADC Analogue-to-Digital Converter	PD Photodiode
AM Amplitude Modulation	PDH Pound–Drever–Hall
AOD Acousto-Optic Deflector	PEEK Polyether Ether Ketone
AOM Acousto-Optic Modulator	PI Photo-Ionisation
AWG Arbitrary Waveform Generator	PID Proportional–Integral–Derivative
BSB Blue Sideband	PM Polarisation Maintaining
cAOD Crossed Acousto-Optic Deflector	PMT Photomultiplier Tube
CPTP Completely Positive Trace Preserving	PSD Power Spectral Density
CV Continuous Variable	PZT Piezo-Controller
DAC Digital-to-Analogue Converter	QC Quantum Computing
DDS Direct Digital Synthesiser	QEC Quantum Error Correction
DP Double-Pass (AOM)	QEM Quantum Error Mitigation
DV Discrete Variable	QIP Quantum Information Processing
EOM Electro-Optic Modulator	RB Randomised Benchmarking
FM Frequency Modulation	RF Radio Frequency
FNC Fibre Noise Cancellation	RPE Robust Phase Estimation
FPGA Field-Programmable Gate Array	RSB Red Sideband
LCC Leadless Chip Carrier	SDF Spin-Dependent Force
MS Mølmer–Sørensen	SP Single-Pass (AOM)
NA Numerical Aperture	SPAM State Preparation and Measurement
NPL National Physical Laboratory	SSB Single Sideband
PBS Polarising Beam Splitter	TTL Transistor–Transistor Logic
	UHV Ultra-High Vacuum

1

Introduction

1.1 Classical and Quantum Computing

Quantum mechanics (QM) successfully describes the behaviour of sufficiently small things. The theory is indispensable in the description of sub-atomic particles, atoms, the molecule they compose, crystals, light, the vacuum, and many other parts of our reality. That is to say, given a choice, bet things obey quantum mechanics. It was recognised in the 1980s that quantum mechanics and information must be unified for fully unlocking the burgeoning field of computation [1, 2]. Motivated with the intuition that if reality is described by QM, then surely quantum objects will offer more efficient strategies to simulate reality, the idea of the *Quantum Computer* was formed. We will elaborate this concept, but first we devote a paragraph to the classical computers that have gone before.

What is the goal of computing? We will define “compute” as modelling some physical system with another, disparate, physical system. This could be an analogue simulator with continuous input and outputs, or it could be some digital model that mimics the system at an arbitrary precision. Apart from the engineering practicalities of realising the device, this analogue and digital distinction may be abstracted away. The *Computer* enters when Turing and Church proposed a class of these simulators that not only mimic one specific system, but rather can model any

physical system, provided enough time and memory [3, 4]. A universal simulator is of course a useful tool. With sufficient understanding of the physical system we wish to model, we may observe *simulated* behaviour without requiring control over the actual system. A profound result in the study of digital universal simulators is the small set of elementary operations required for this universality, for example, the boolean set {AND,OR,NOT}, are sufficient to simulate any function to an arbitrary precision.

However, *universal* does not mean *practical* — although a physical device is able to simulate anything, it may still not be feasible for an operator to “code” some desired model¹, or for a useful result to be returned in reasonable time — i.e. is tractable². Regardless, highly advanced and operable computers have been constructed and have successfully modelled and transformed almost all aspects of modern life.

Unfortunately, many problems intractable using classical computers are of great practical importance, for example accurately modelling chemical systems for industrial, and drug design, and simulating solid-state systems for fundamental material, and physical research.

Both these examples intentionally fall into the simulation of quantum systems. Models of quantum systems are notoriously expensive to encode, fundamentally due to the concepts of superposition and interference — a pure quantum state is generally described by the superposition of all possible d basis states, where d is the system dimension. Given an underlying number, n , of q -level subsystems that compose the system, we find the dimension exponentially scales as $d = q^n$. Each basis state of the n -body system will have some associated amplitude and phase, resulting in $2q^n - 2$ real parameters required to track and accurately model the pure quantum state. This exponential scaling is not kind to a digital computer.

Returning to the 1980s, it was realised that this exponential scaling could provide the very resources required to simulate other quantum systems. Can we construct

¹play with esolangs for a few hours to confirm this for yourself.

²factor a 2048-bit number for a few billion years to confirm this for yourself.

a Turing machine that utilises the resource scaling of a quantum system [5]³ — a quantum computer (QC)?

The ingredients for such a device are well known [8], and have been demonstrated on multiple physical platforms to date. Theoretical work is placing ever-more lenient constraints on resource requirements, as algorithmic strategies are developed, while experimental work is pushing to scale the available resources, and performance of actual hardware. The two frontiers are yet to meet at a realisation of a useful universal QC, however the list of prerequisites is shortening rapidly. We shall give a short overview of the distinct platforms in development for scaling QC, and then concentrate on trapped-ions — the experimental platform we use in this thesis.

1.2 Quantum Computing Platforms

An effective quantum computer requires meticulous control over a collection of simple quantum systems. A good starting point is thus selecting that simple system to build our control around. Nature provides us with a bounty of choices in this regard: atoms. It is definitively difficult to choose a simpler and more reproducible object: atoms of the same species are identical to one another no matter where or when you are. We may use their electronic energy structure, which atomic physicists have scrutinized for over a century, to define our quantum bits, and we can use an assortment of electromagnetic frequencies to manipulate these internal states. This is the strategy of both the “neutral-atom”, and “trapped-ion” platforms, with the latter choosing to remove or add a few electrons to more easily manipulate the position of the atoms in space⁴. The trapped-ion platform also has access to the motion of the ions, which may be used as continuous variable quantum resources — so called qumodes — that can be coherently controlled, and initialised to highly non-classical states.

³As an aside, the implications of quantum mechanics on information extend further than purely mimicking other dynamic systems – using and distributing quantum objects also allows for dominant strategies in cryptography [6], and some competitive (although currently esoteric) games [7].

⁴We note that the neutral-atom platform has made significant progress in the spatial control of atoms in optical tweezer traps, however, the trap depths between the two platforms is incomparable.

Photons are another “natural” choice for the quantum-resource. They may either be treated also as continuous quantum systems, or have some discrete level system encoded, such as a qubit. The photonic qubit is unique in being a “flying-qubit”: photons are able to transmit information over long distances, lending them particular interest in their use for interconnecting distant quantum computers.

The final class of platforms is that of “manufactured” quantum systems, such as superconducting circuits, quantum dots, and NV centers; benefiting from control over the level structure, and owing to their solid state, presumably allowing for easier integration into existing semiconductor fabrication techniques.

At this time, it is unclear whether one of these platforms, or an entirely new architecture, will be the first to reach a scaled device. Each have their own advantages and unique technical challenges, and each are in rapid active development. The next section outlines some technical challenges in scaling current quantum hardware, focussing on the trapped-ion platform studied in this thesis.

1.3 Scaling Trapped Ion Systems

For useful quantum computation and simulation, we require access to a sufficiently large Hilbert space, with sufficiently accurate control with which to traverse it. How do we scale the ion-trap platform to provide those resources? In principle, we know *how*, but the required engineering and control presents unique and technical challenges. I identify three specific challenges associated to scaling, that motivate the work presented in this thesis: increasing the available quantum resources, controlling those resources, and characterising and mitigating the errors that arise as the system grows.

Challenge I: Scaling Quantum Resources

We need larger Hilbert spaces — to solve a “useful” computation, such as Shor’s algorithm [9] applied to factoring RSA-2048 integers, it is currently estimated that $10^4 - 10^5$ physical qubits will be required [10, 11]. Architectures for such

scaling include quantum charge-coupled device (QCCD) [12], to reconfigure an array of qubits using multi-zone surface traps; and optical networked modules [13], using photonic interconnects to combine many small qubit registers into a single traversable Hilbert space. Another disparate strategy — continuous-variable QC (CVQC) — replaces many low-dimensional qubits, with few high-dimensional qumodes. Qumodes store information in the continuous states of a quantum harmonic oscillator, as such, in principle have an infinite-dimensional Hilbert space. This exchanges the technical issues of controlling vast numbers of qubits, to instead requiring truly exquisite control of a few qumodes. The innate utility of this continuous system, is exemplified by a CVQC scheme for factoring RSA-2048 integers using only three qumodes and a single qubit [14]. Although this particular protocol is unfeasible due to technically prohibitive required high-energy oscillator states, scaling accessible Hilbert spaces via qumodes remains a promising strategy. Trapped ions therefore provide two complementary routes to increasing the accessible Hilbert space: increasing the total number of ions in the register, and utilising the motional modes of the ions already present within that register.

Challenge II: Scaling Control

Access to a sufficiently large Hilbert space is useful only if we can reliably and selectively navigate it. What are the elementary operations we require, and how do we physically implement them on our system? The elementary unitary operations often are expressed as minimal gate-sets allowing for universal quantum computing. These are typically comprised of one, and two-body interactions [15–18], which greatly simplifies the elementary operation, and the required number of implementable interactions. This does come at the considerable cost of “addressability” — we must be able to selectively drive interactions locally on either one, or two physical carriers. For trapped-ion qubits, this requires a form of spatial selectivity, while for motional-state qumodes, often requires a combination of frequency and spatial selectivity. This technical constraint is that of “control-signal delivery”, a common problem to both classical and quantum computing. Trapped-ion systems

require hardware for delivering spatially, and spectrally selective electronic control, and optical-beam delivery, both of which are under active development to enable scalability [19, 20].

Challenge III: Scaling with Imperfect Systems

The last specific challenge we discuss is that of unwanted evolution due to *noise*. Qubits stored in the electronic states of ions are sensitive to external electric and magnetic fields, while qumodes are particularly sensitive to electric fields originating from the trap electrodes. The effects of noise on the desired state evolution are varied, but if left unchecked, inevitably lead to information loss. Furthermore, increased system scale and complexity changes the relevant error processes. In particular, noise leading to unintended correlations across the register, and across time can be difficult to quantify and mitigate. To combat this noise, we first require methods to characterise performance, and quantify how this noise leads to logical errors in an executed algorithm. This motivates the wide field of quantum characterisation and verification [21, 22]. A varied collection of characterisation protocols exist due to an unavoidable trade-off between *expressibility* and *scalability* — The systems we wish to characterise are inherently large, and thus full descriptions require vast amounts of information. So we often sit in two regimes: we may either fully characterise few-body interactions [23], or we must use abstracted, and system-wide metrics to gauge performance of the scaled system [24, 25].

When noise is suitably suppressed through hardware engineering, we may then utilise algorithmic techniques to be robust to the remaining imperfections. Foremost, analogous to classical computing, it is possible to detect and correct errors during the execution of an algorithm [26]. This comes at a considerable overhead in both register size and circuit depth, only recently in reach of experimental platforms, but only for small numbers of logical qubits [27]. A complimentary method for suppressing the impact of errors is through myriad error mitigation techniques, using repeat measurements to infer error-suppressed algorithmic outputs [28].

Therefore, even with imperfect physical qubits or qumodes, it appears to still be physically possible to construct useful, universal quantum computers.

1.4 Structure of this Thesis

This thesis describes the development of a trapped-ion platform for multi-ion spin–motion control, and its application to the characterisation of correlated errors, hybrid spin–oscillator control, and quantum error mitigation.

Chapter 2 provides further details on the trapped-ion platform, and the underlying light–matter interactions enabling the use of ions as quantum-information carriers. Chapter 3 details the ion-trap apparatus constructed for this thesis, enabling the study of multi-qubit and qumode registers.

Chapter 4 presents the hardware characterisation, giving details on experimental calibrations, the initialisation of the qubits and qumodes, and elementary single- and two-qubit gate performance.

Chapter 5 presents a novel and scalable randomised benchmarking method for characterising and quantifying spatially correlated errors in multi-qubit devices, including experimental results on up-to an 8-ion chain.

Chapter 6 explores two ongoing research directions enabled by the newly constructed apparatus. First, we outline how to utilise the qumodes in our system, providing a gate-set for full qubit-qumode control. We discuss a practical way to drive this full gate-set, and by this method, demonstrate spin-conditioned motional squeezing. Second, we experimentally investigate virtual distillation, an error mitigation technique with practical application to computation involving both spin and motional degrees of freedom.

Chapter 7 concludes this thesis by summarising its contributions, and providing an outlook for future research enabled by the experimental platform.

2

Trapped-Ion Quantum Computing

Contents

1.1	Classical and Quantum Computing	1
1.2	Quantum Computing Platforms	3
1.3	Scaling Trapped Ion Systems	4
1.4	Structure of this Thesis	7




Figure 2.1: A schematic representation of the hybrid qubit-qumode register in the trapped-ion system. With N physical ions, with qubits encoded in the electronic levels, and qumodes encoded in the secular motion of the ion chain, there are N -qubits and $3N$ -qumodes available.

This thesis considers control over a hybrid spin-motion system, as shown schematically in Figure 2.1. For an N -ion chain, we have access to N -qubits, defined by electronic transitions of the ions, and $3N$ -qumodes, emerging from the

collective motion of the ion chain.

In this chapter we lay out the foundational concepts and interactions that will be used throughout this thesis. We begin by considering the choice of ion, $^{40}\text{Ca}^+$, and its relevant optical transitions. We then abstract the electronic levels of the ion to a single two-level system, the qubit, and discuss coherent control using oscillating electric fields. Next we consider the motion of the ions, how this motion originates using a linear Paul trap, and how to coherently excite it. This chapter ends with a discussion of the interactions utilised for hybrid spin-motion control, with an emphasis on this control for realising two-qubit entangling gates via the shared motional modes.




Figure 2.2: Electronic energy levels of $^{40}\text{Ca}^+$ — the ion species used in this thesis. The levels are split by the Zeeman effect due to a ≈ 5 G external magnetic field (splitting shown explicitly for the $|4S_{1/2}\rangle$ and $|3D_{5/2}\rangle$ manifolds). The transition wavelengths required for cooling and control over the ion are highlighted. The chosen qubit levels are labelled by $| \downarrow \rangle$, $| \uparrow \rangle$.

2.1 The Ion

There are many choices of ion for trapped-ion quantum computing [29]. Selecting a species is a compromise between capability and experimental ease-of-use. For example, ions with hyperfine structure, compared to those without, have many more energy levels with accessible clock-transitions that are insensitive to external

magnetic field fluctuations [30, 31]. This insensitivity is of great importance when envisioning scaled quantum-computing with long coherence times enabling deep circuits. However, practically the extra hyperfine splitting adds experimental complexity, requiring more control for state-preparation, and more channels for leakage errors. Thus, to ease the control requirements, a species without hyperfine structure was selected allowing for focus of the constructed apparatus to be instead spent on the planned demonstrations of hybrid control of a multi-ion chain. The choice of species is further restricted by the desire to use visible wavelengths for the light-matter interactions, and an optical qubit-transition with reasonable excited state lifetimes. We require the wavelength associated with the qubit frequency to have commercially available narrow-linewidth lasers with power outputs of ≈ 5 W. Given those conditions, we selected $^{40}\text{Ca}^+$, with a qubit defined by the optical transition at 729-nm, between the ground manifold $|4S_{1/2}\rangle$ and the excited manifold $|3D_{5/2}\rangle$. This optical transition is dipole forbidden by the symmetries of the $|4S_{1/2}\rangle$ and $|3D_{5/2}\rangle$ electron wavefunctions, as such the excited $|3D_{5/2}\rangle$ manifold is metastable with a natural lifetime of $\tau \approx 1.1$ s [32]. we shall refer to the transition between the two manifolds as either the quadrupole transition, or the 729-nm transition. The energy levels for $^{40}\text{Ca}^+$ are shown in Figure 2.2, with the wavelengths of laser highlighted to drive a selection of transitions.

The 397-nm and 866-nm lasers are used to form a dissipative closed-cycle transition useful for both state-dependent fluorescence (§ 4.1.2), and for laser cooling (§ 4.2.2). The 854-nm laser is used to deshelve the electron from the $|3D_{5/2}\rangle$ manifold for qubit state preparation (§ 4.1.1). Details on these lasers can be found in § 3.6.

An external magnetic field of ≈ 5 G provides the quantisation axis for the ion. This field enables both polarisation and frequency selectivity of the atomic transitions, ensuring unintended off-resonant excitations are suppressed when optically driving interactions. The Zeeman splitting magnitudes are highlighted for the ground and excited manifold used for the qubit. The transition frequencies thus




Figure 2.3: Relative light-matter coupling strengths (given as Rabi frequencies), Ω , and frequency splittings for the $|4S_{1/2}\rangle \leftrightarrow |3D_{5/2}\rangle$ Zeeman sublevels in $^{40}\text{Ca}^+$. The transitions are labelled by their excited state m_j level. Solid bars correspond to transitions with lower state $|4S_{1/2}, m_j = -1/2\rangle$, while open bars correspond to transitions with lower state $|4S_{1/2}, m_j = +1/2\rangle$. (a) Relative Rabi frequency for transition $|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = m_j\rangle$ with angle γ between the polarisation vector and the magnetic field. The k -vector of the light is perpendicular to the magnetic field. (b) Relative frequency splitting and Rabi frequency when $\gamma = \pi/2$ (b) and for $\gamma = \pi/4$ (c).

have a dependence on field strength given by the difference in magnetic moments of the two states. For the quadrupole transition this is given by,

$$\left| \frac{df}{dB} \right| = \Delta\mu_B/h, \quad \Delta = g_D m_D - g_S m_S, \quad (2.1)$$

where μ_B is the Bohr magneton, and g_D and g_S are the g -factors for the manifolds $|3D_{5/2}\rangle$ and $|4S_{1/2}\rangle$ respectively. The magnetic field sensitivities are then, $|\frac{df}{dB}| = \{-2.799, -1.119, 0.561, 2.241, 3.922\}$ MHz G $^{-1}$ between ground state $|\downarrow\rangle$ and excited states with varying m_j , $|3D_{5/2}, m_j = \{-5/2, -3/2, -1/2, 1/2, 3/2\}\rangle$ respectively. The quantising magnetic field is set perpendicular to the 729-nm laser beam access. As such, the coupling strength of the various quadrupole transitions for a linear-polarised optical field is set exclusively by the angle, γ , between the magnetic field axis and the optical polarisation axis. This relationship is plotted in Figure 2.3. Figure 2.3 (a), shows the relative coupling strength (here given as Rabi frequency, which is discussed in § 2.2.2) of each transition as the angle γ is varied. To maximise the coupling of one transition, and suppress other possible transitions,

the polarisation of the 729-nm laser can be set to $\gamma = \frac{\pi}{2}$. This would be a good choice if high-strength coupling interactions are desired to minimise off-resonant driving of spectator interactions. If it is desirable to access multiple quadrupole transitions the setting $\gamma = \frac{\pi}{4}$ is a convenient choice. At the magnetic field strength of ≈ 5 G, the relative frequency splitting and strengths of the interactions are shown for $\gamma = \frac{\pi}{2}$ and $\gamma = \frac{\pi}{4}$ in panels (b) and (c) respectively.

2.2 Spin

2.2.1 The Qubit

The particular choice of qubit, highlighted in Figure 2.2, $|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = -5/2\rangle$ with $|\downarrow\rangle := |4S_{1/2}, m_j = -1/2\rangle$, $|\uparrow\rangle := |3D_{5/2}, m_j = -5/2\rangle$, is convenient as the 729-nm polarisation when set to $\gamma = \frac{\pi}{2}$ can maximise its coupling. Furthermore, due to the dipole selection rule, $|\Delta m_j| \leq 1$, a closed-cycle transition for sideband cooling (details in § 4.2.2) is available using this excited $|3D_{5/2}\rangle$ state.

The computational eigenstates of the qubit, $|\downarrow\rangle$, $|\uparrow\rangle$, have eigenvalues,

$$\hat{\sigma}_z |\downarrow\rangle = -|\downarrow\rangle, \quad \hat{\sigma}_z |\uparrow\rangle = +|\uparrow\rangle, \quad (2.2)$$

and the dynamics of the system are giving by the spin Hamiltonian,

$$\hat{H}_s = \frac{\hbar\omega_0}{2}\hat{\sigma}_z, \quad (2.3)$$

where ω_0 is the qubit transition frequency.

We define transition elements $\hat{\sigma}_+ = |\uparrow\rangle\langle\downarrow|$, and $\hat{\sigma}_- = |\downarrow\rangle\langle\uparrow|$. The next section outlines how to physically realise such transitions.

2.2.2 Coherent Spin Control

To manipulate the spin-state we consider periodic perturbations given by the interaction Hamiltonian $\hat{H}_I(t)$, with the full time dependent Hamiltonian of the

system given by,

$$\hat{H}(t) = \hat{H}_s + \hat{H}_I(t). \quad (2.4)$$

The perturbation is treated semi-classically, consisting of the interaction between the quantised electronic states of the ion and a classical spatially varying electromagnetic field.

For a travelling monochromatic optical field (almost always from a laser) the transition-driving part of the interaction is generally given by,

$$\hat{H}_I = \frac{\hbar\Omega}{2}(\hat{\sigma}_+ + \hat{\sigma}_-) \left(e^{i(\vec{k} \cdot \vec{R} - \omega t + \phi)} + e^{-i(\vec{k} \cdot \vec{R} - \omega t + \phi)} \right), \quad (2.5)$$

where Ω is the Rabi frequency, $\vec{R}$ is the ion position, ω is the optical frequency, and ϕ is the optical phase. The Rabi frequency contains the specifics of the physical light-matter interaction, $\hbar\Omega \propto \langle \uparrow | \hat{H}_I | \downarrow \rangle$. Assuming the phase variation of the optical field across the electrons spatial extent is small, the perturbation can be expanded around the ion position, R , for small relative electron position r , to generate a series of electric and magnetic multipole interactions. For the qubit transition considered here the electric dipole interaction is forbidden, as such in this thesis we exclusively consider the electric quadrupole interaction.

For a plane travelling optical field given by,

$$\vec{E} = E_0 \epsilon \cos(\vec{k} \cdot \vec{R} - \omega t + \phi), \quad (2.6)$$

the quadrupole Rabi frequency is given by,

$$\Omega_{\text{EQ}} = \frac{1}{6\hbar} \sum_{i,j=\{x,y,z\}} \langle \downarrow | Q_{ij} | \uparrow \rangle \partial_i E_j |_R \quad (2.7)$$

where Q_{ij} is the quadrupole tensor [33]¹. Apart from the dependence on polarisation and k -vector discussed in Figure 2.3, the details of the quadrupole derived Rabi frequency are not relevant for the remainder of this section, and a general Rabi frequency, denoted by Ω , is instead considered.

¹For reference, the simpler electric dipole interaction has Rabi frequency is given by $\Omega_{\text{ED}} = \frac{1}{\hbar} \sum_j \langle \downarrow | d_j | \uparrow \rangle E_j$, where d_j is the dipole moment

To evaluate the dynamics of $\hat{H}_I$ we go to a frame rotating at the qubit frequency, ω_0 . This interaction picture with respect to $\hat{H}_s$ is $\tilde{H}_I = \hat{U}_s^\dagger \hat{H}_I \hat{U}_s$ with $\hat{U}_s = e^{-i\hat{H}_s t/\hbar}$. Defining a relative detuning between the qubit and optical frequency as $\delta = \omega - \omega_0$, and neglecting fast oscillating terms (making the rotating wave approximation) with frequency $\omega + \omega_0$, the Hamiltonian is approximated by,

$$\tilde{H}_I \approx \frac{\hbar\Omega}{2} \hat{\sigma}_+ e^{i(\vec{k} \cdot \vec{R} - \delta t + \phi)} + \text{h.c.} \quad (2.8)$$

The dynamics due to the fast oscillating terms are safely neglected in the case of weak coupling $\Omega \ll \omega + \omega_0$.

For the rest of this subsection the ion is assumed to be stationary over time, thus the motional term contributes a simple phase factor which can be absorbed into a global phase, $\phi' = \phi + \vec{k} \cdot \vec{R}$.

Rabi Oscillations

When the optical field frequency is tuned to the qubit resonance, $\delta = 0$, the Hamiltonian becomes time-independent,

$$\tilde{H}_I = \frac{\hbar\Omega}{2} \hat{\sigma}_\phi, \quad (2.9)$$

where $\hat{\sigma}_\phi = \cos(\phi)\hat{\sigma}_x - \sin(\phi)\hat{\sigma}_y$. Exponentiating², this yields the unitary propagator,

$$\hat{R}_\phi(\theta) := \hat{U}(t, 0) = \begin{pmatrix} \cos(\theta/2) & -ie^{i\phi} \sin(\theta/2) \\ -ie^{-i\phi} \sin(\theta/2) & \cos(\theta/2) \end{pmatrix}, \quad (2.10)$$

which is a rotation on the Bloch sphere by angle $\theta = \Omega t$, around the Bloch vector, $\hat{\sigma}_\phi$. Varying phase ϕ then sets the rotation axis in the xy -plane, with $\hat{R}_x(\theta) := e^{-i\theta\hat{\sigma}_x/2}$ defined for $\phi = 0$. Applying single qubit rotations to the initial state $|\downarrow\rangle$ and observing the spin populations results in spin-state transfers, termed Rabi oscillations,

$$P_\uparrow = |\langle \uparrow | \hat{R}_x(\Omega t) | \downarrow \rangle|^2 \quad (2.11)$$

$$= \sin^2\left(\frac{\Omega t}{2}\right). \quad (2.12)$$

where $P_\uparrow$ is the population observed in $|\uparrow\rangle$.

²Using identity $e^{-i\theta\vec{n}\cdot\vec{\sigma}/2} = \hat{I} \cos\left(\frac{\theta}{2}\right) - i(\vec{n} \cdot \vec{\sigma}) \sin\left(\frac{\theta}{2}\right)$

AC Stark Shifts

When considering the optical interaction given in equation 2.8 for a large detuning $|\Delta| \gg \Omega$, the Hamiltonian remains time dependent, and does not self-commute at different times. The Magnus expansion is utilised to evaluate the dynamics. Details are given in Appendix A.1.1, with the resulting dynamics approximating $\hat{\sigma}_z$ rotations of the form,

$$\hat{R}_z := e^{-i\theta\hat{\sigma}_z/2}, \quad \theta = \frac{\Omega^2 t}{2\Delta}. \quad (2.13)$$

In this thesis these $\hat{R}_z$ rotations are often caused unintentionally due to off-resonant coupling of other electronic transitions when addressing the qubit transition. As such, for a given desired interaction and beam power, these Stark shift rotations are calibrated and accounted for by updating the target rotation axis of subsequent pulses.

This concludes the spin-only control of our system. To access two-body interactions, such as two-qubit entangling gates, we must first introduce the shared motion of the qubits.

2.3 Motion

The potential used to trap the ions also dictates their motion. Here we outline the motion of a single trapped ion in a radio-frequency (RF) Paul trap, and go on to define the emergent qumodes, and two representations to describe their state.

2.3.1 Linear Paul Trap

From Earnshaw's theorem, $\nabla^2 V = 0$, a stable stationary point in 3D cannot be realised using only static electric potentials. If the potential is confining in two dimensions it must be anti-confining in the third. Therefore, to achieve stable trapping, either a supplementary oscillating electric field, as used in RF Paul traps [34], or a static magnetic field, as used in Penning traps [35], is required.

In this thesis we use an RF Paul trap (hardware details in § 3.2). Radial confinement is provided by an oscillating radial electric potential, while the axial confinement is provided by a static axial potential. Both potentials are considered to be ideal quadrupolar electric potential. The combined potential is given as,

$$\Phi(x, y, z, t) = \frac{V_{\text{DC}}}{2d^2}(2z^2 - x^2 - y^2) + \frac{V_{\text{RF}}}{2\hat{\rho}^2} \cos(\Omega_{\text{RF}}t)(x^2 - y^2), \quad (2.14)$$

where V_{DC} is the voltage applied to axial “endcap” electrodes which are distance d from the trapping point. V_{RF} is the RF voltage with frequency Ω_{RF} applied to the radial RF electrodes which are distance $\hat{\rho}$ from the trapping point. This potential considers a highly symmetric field configuration with the axial potential confining in the z direction while equally anti-confining in both radial directions. While the radial potential has no axial component, and initially confines in x and anti-confines in y , but switches periodically to confining in y and anti-confining in x .

From this potential the equation of motion in the radial x for a trapped ion with charge $+e$ is,

$$m \frac{d^2x}{dt^2} = \left(\frac{eV_{\text{DC}}}{d^2} - \frac{eV_{\text{RF}}}{\hat{\rho}^2} \cos(\Omega_{\text{RF}}t) \right) x, \quad (2.15)$$

where m is the mass of the ion. This motion is thus described by the Mathieu equation,

$$\frac{d^2x}{d\tau^2} + (a - 2q \cos(2\tau))x = 0, \quad (2.16)$$

where we have defined $\tau = \frac{\Omega_{\text{RF}}t}{2}$, $a = -\frac{4eV_{\text{DC}}}{m\Omega_{\text{RF}}^2 d^2}$, and $q = -\frac{2eV_{\text{RF}}}{m\Omega_{\text{RF}}^2 \hat{\rho}^2}$. Analysis of the two parameters, a , q , of the Mathieu equation yield regions of stability (motional amplitude converges) and instability (motional amplitude diverges). For the “lowest stability region”, the regime relevant to our trap, the motion is stable when $|a|, q^2 \ll 1$ [36]. In this regime the ion dynamics are well approximated by the motion,

$$x(t) \approx x_{\text{eq}} + x_{\text{sec}}(t) \left(1 + \frac{q}{2} \cos(\Omega_{\text{RF}}t) \right), \quad (2.17)$$

where x_{eq} is the equilibrium position, and x_{sec} is a secular motional term,

$$x_{\text{sec}}(t) = A \cos(\omega_x t + \theta_0), \quad (2.18)$$

with phase, θ_0 and amplitude, A , set by initial conditions and a frequency,

$$\omega_x = \frac{\Omega_{\text{RF}}}{2} \sqrt{a + \frac{q^2}{2}}. \quad (2.19)$$

A similar form exists for motion in y . The non-secular component of the motion is termed “micromotion”, and provides a fast oscillation superimposed on the secular motion. Due to the different time-scales of the two motional terms, and the relatively smaller amplitude of the micromotion term, averaging over periods greater than $\frac{1}{\Omega_{\text{RF}}}$ often is sufficient to negate the effects of the fast oscillations. In this case, the motion is well described by simple harmonic motion in 3D with a “pseudopotential” given by,

$$U_{\text{pseudo}}(x, y, z) = \frac{m}{2} (\omega_x^2 x^2 + \omega_y^2 y^2 + \omega_z^2 z^2). \quad (2.20)$$

The axial mode is given exclusively by the DC term,

$$\omega_z = \sqrt{\frac{2eV_{\text{DC}}}{md^2}}. \quad (2.21)$$

The radial terms are given by,

$$\omega_x = \omega_y = \sqrt{\omega_r^2 - \omega_z^2}, \quad (2.22)$$

where ω_r is the radial trapping frequency in the limit of $q^2 \gg a$,

$$\omega_r = \frac{eV_{\text{RF}}}{\sqrt{2}m\Omega_{\text{RF}}\hat{\rho}^2}. \quad (2.23)$$

The axial potential then acts to “defocus” the radial confinement, with an increased axial frequency leading to a reduced radial. This equal defocussing of the two radials is due to the assumed electrode symmetry of the idealised potential in equation 2.14. If the static potential acts asymmetrically on the two radial dimensions, this breaks the degeneracy of the radial modes. In this thesis we enforce a strong DC radial asymmetry to ensure the two radial modes are well frequency separated (electrode configuration details in § 3.2.2).

For N ions there are $3N$ common secular motional modes, with mode frequencies and ion spacing set by the pseudopotential and the Coulomb repulsion between ions [33].

2.3.2 The Qumode

This secular motion is well approximated by simple harmonic oscillations for amplitudes small compared to the length scale of any anharmonicities in the pseudopotential. When the motion is suitably damped to a temperature of order the mode frequency, $T \sim \hbar\omega_x/k_B$, a quantised description of the oscillator is required. The quantum harmonic oscillator (qumode) is given by,

$$\hat{H}_{\text{osc}} = \frac{1}{2m}\hat{p}^2 + \frac{1}{2}m\omega_x^2\hat{x}^2, \quad (2.24)$$

where m is the ion mass, ω_x is the motional mode frequency, and $\hat{x}$, $\hat{p}$ are the position and momentum operators respectively. A description of the oscillator state may be given in phase space using the $\hat{x}$ and $\hat{p}$ operators. In the phonon, or Fock-state, representation, the creation and annihilation operators, $\hat{a}^\dagger$ and $\hat{a}$, are related to this picture by,

$$\hat{x} = x_0(\hat{a}^\dagger + \hat{a}), \quad (2.25)$$

$$\hat{p} = ip_0(\hat{a}^\dagger - \hat{a}), \quad (2.26)$$

where $x_0 = \sqrt{\frac{\hbar}{2m\omega_x}}$ is the ground state wave-packet extent, and $p_0 = \sqrt{\frac{\hbar m\omega_x}{2}}$ is the ground state momenta. These ladder operators provide an alternative basis to describe the oscillator state. They act on the oscillator energy levels as,

$$\hat{a}^\dagger |n\rangle = \sqrt{n+1} |n+1\rangle, \quad \hat{a} |n\rangle = \sqrt{n} |n-1\rangle, \quad (2.27)$$

where n is the Fock state, and the eigenvalue of the number operator, $\hat{n} |n\rangle = n |n\rangle$ with $\hat{n} = \hat{a}^\dagger \hat{a}$. The oscillator Hamiltonian is then given succinctly by,

$$\hat{H}_{\text{osc}} = \hbar\omega_x \left(\hat{n} + \frac{1}{2} \right). \quad (2.28)$$

In this thesis we will denote an oscillators state either in this Fock-state basis, where, $|n\rangle = \frac{\hat{a}^{\dagger n}}{\sqrt{n!}} |0\rangle$, or in a coherent-state basis, $|\beta\rangle = \hat{D}(\beta) |0\rangle$, where $\hat{D}(\beta)$ is the displacement operator which we shall explore in the following section.




Figure 2.4: Phase space trajectories from the displacement operator acting on the initial ground state wavefunction (shown by faded state at origin), $\hat{D}(\beta)|0\rangle$. (a) When displacement is driven resonantly the trajectory is linear. (b) When driven off-resonantly the trajectory traces a circle.

2.3.3 Driving Motion

We now consider driving the qumode by a near-resonant spatially uniform electric field,

$$E(t) = E_0 \cos(\omega t + \phi), \quad (2.29)$$

with the Hamiltonian given by,

$$\hat{H}(t) = \hat{H}_{\text{osc}} + \hat{H}_I, \quad \hat{H}_I = -e\hat{x}E(t), \quad (2.30)$$

where $+e$ is the ion charge. Going into the interaction picture with respect to $\hat{H}_{\text{osc}}$, we find,

$$\tilde{H}_I = -\hbar\Omega_D \cos(\omega t + \phi)(\hat{a}^\dagger e^{i\omega_x t} + \hat{a}e^{-i\omega_x t}), \quad (2.31)$$

where $\Omega_D = \frac{ex_0 E_0}{\hbar}$. Expanding the oscillating field and assuming $\omega + \omega_x \gg \Omega_D$, the fast oscillating terms are neglected giving,

$$\tilde{H}_I \approx \hat{H}_D := -\frac{\hbar\Omega_D}{2}(\hat{a}^\dagger e^{-i(\delta t + \phi)} + \hat{a}e^{i(\delta t + \phi)}), \quad (2.32)$$

where $\delta = \omega - \omega_x$.

In the case of resonant driving, $\delta = 0$, the Hamiltonian becomes time-independent, and we recover the propagator,

$$\hat{D}(\beta) := \hat{U}(t) = e^{i\frac{\Omega_D t}{2}(\hat{a}^\dagger e^{-i\phi} + \hat{a}e^{i\phi})}, \quad (2.33)$$

$$= e^{\beta\hat{a}^\dagger - \beta^*\hat{a}}, \quad (2.34)$$

where the second expression is the conventional displacement operator, $\hat{D}(\beta)$, where here $\beta = \frac{\Omega_D t}{2} i e^{-i\phi}$. The displacement operator is used in the coherent state representation of the oscillator, $|\beta\rangle = \hat{D}(\beta) |0\rangle$, where $|\beta\rangle$ is the eigenstate of the annihilation operator, $\hat{a} |\beta\rangle = \beta |\beta\rangle$. The action of the displacement operator may be visualised in phase space, shown for resonance in Figure 2.4 (a).

Considering the off-resonant case, $\delta \neq 0$, the Hamiltonian $\tilde{H}_I$ does not self-commute over time, generally $[\hat{H}_D(t_1), \hat{H}_D(t_2)] \neq 0$. The resulting dynamics are derived by the Magnus expansion in Appendix A.1.2, yielding the effective unitary at time T ,

$$\hat{U}(T) \approx \hat{D}(\beta(T)) e^{-i\Phi(T)}. \quad (2.35)$$

The first term is the displacement operator, the second is an accumulation of global geometric phase. The displacement parameter β is given by,

$$\beta(T) = \frac{\Omega_D}{\delta} \sin\left(\frac{\delta T}{2}\right) i e^{-i(\delta T/2 + \phi)}. \quad (2.36)$$

with a trajectory tracing a circle of radius $\frac{\Omega_D}{2|\delta|}$, returning to the initial motional state with periodicity $T_n = \frac{2\pi n}{|\delta|}$, for integer n . The phase space dynamics for the off-resonant displacement is shown in Figure 2.4 (b).

The geometric phase term, originating from the second order term in the Magnus expansion, is given by,

$$\Phi(T) = \frac{\pi \Omega_D^2}{4\delta} \left(T - \frac{\sin(\delta T)}{\delta} \right). \quad (2.37)$$

Considering the loop-closure times, T_n , the oscillatory term of the phase disappears, leaving,

$$\Phi(T_n) = \frac{n\pi \Omega_D^2}{2\delta^2}. \quad (2.38)$$

In this example the geometric phase is globally acting, and as such can be safely ignored. This geometric phase is utilised spin-dependently in the final section of this chapter to induce two-qubit entanglement via a shared motional mode coupling the two qubits.

2.4 Spin-Motion

In this section we explore light-matter interactions that couple the spin and motional subsystems. Reconsider the spin transition Hamiltonian, $\tilde{H}_I$, in equation 2.8, but now including quantised secular ion motion in the x -direction given by equation 2.25, around an equilibrium position, $\vec{R} = R_{\text{eq}} + x_0(\hat{a}^\dagger + \hat{a})\hat{e}_x$, where $\hat{e}_x$ is the x unit vector. The interaction Hamiltonian is then given by³,

$$\tilde{H}_I \approx \hat{H}_{\text{osc}} + \frac{\hbar\Omega}{2} \left(\hat{\sigma}_+ e^{i\eta(\hat{a}^\dagger + \hat{a})} e^{-i(\delta t - \phi)} + \text{h.c.} \right), \quad (2.39)$$

where the equilibrium position has been absorbed into the global phase, and we have defined $\eta = |\vec{k}|x_0 \cos(\gamma)$, the Lamb-Dicke parameter, with γ being the angle between the k -vector of optical field and the mode vector $\hat{e}_x$. The Lamb-Dicke parameter relates the excitation wavelength to the extent of the ground state oscillator wavefunction, x_0 . The parameter also quantifies the strength of light coupling to motional change, it may be expressed in terms of photon recoil energy, $E_R = \frac{\hbar^2 k^2}{2m}$, to oscillator energy, $E_{\text{osc}} = \hbar\omega_x$, as $\eta^2 = \frac{E_R}{E_{\text{osc}}}$. For the wavelength (729-nm) and mode frequencies (≈ 3 MHz) considered in this thesis, the Lamb-Dicke parameter is of order $\eta \approx 0.05$.

To further analyse the dynamics, the interaction Hamiltonian is transformed to the frame rotating with the oscillator energy, $\hat{H}_{\text{osc}}$, yielding,

$$\tilde{\tilde{H}}_I \approx \frac{\hbar\Omega}{2} \hat{\sigma}_+ e^{i\eta(\hat{a}^\dagger e^{i\omega_x t} + \hat{a} e^{-i\omega_x t})} e^{-i(\delta t - \phi)} + \text{h.c.} \quad (2.40)$$

2.4.1 Resolved Sideband Interactions

The resolved motional sideband regime is well approximated when the motional wavefunction extent $\sqrt{\langle \hat{x}^2 \rangle}$ is small compared to the optical wavelength, $\eta^2(2\hat{n} + 1) \ll 1$. In this ‘‘Lamb-Dicke’’ regime, the interaction Hamiltonian can be expand in terms of η as (here only to first order),

$$\hat{H}_{\text{LD}} = \frac{\hbar\Omega}{2} \hat{\sigma}_+ \left(\hat{I} + i\eta(\hat{a}^\dagger e^{i\omega_x t} + \hat{a} e^{-i\omega_x t}) + O(\eta^2) \right) e^{-i(\delta t - \phi)} + \text{h.c.}, \quad (2.41)$$

³Keeping in mind this Hamiltonian is already in the interaction picture with respect to the spin, $\hat{H}_s$, and is neglecting fast oscillating terms




Figure 2.5: In the resolved sideband regime the hybrid qubit-qumode system appears as a ladder of possible states. The carrier (green) induces spin flips, $|\downarrow, n\rangle \leftrightarrow |\uparrow, n\rangle$, while the sideband interactions couples the spin and neighbouring oscillator levels. For the red-sideband (red) the coupled states are $|\downarrow, n\rangle \leftrightarrow |\uparrow, n-1\rangle$, and for the blue-sideband (blue) they are $|\downarrow, n\rangle \leftrightarrow |\uparrow, n+1\rangle$.

where $\hat{I}$ is the identity operator of the qumode.

Three resonant regimes are then apparent:

Taking $\delta = 0$, and assuming $\eta\Omega \ll \omega_x$ (adiabatic driving) the non-motional term is brought into resonance, and terms linear in η may be neglected. This “carrier” interaction is identical (up to first order) to the Hamiltonian considered above for Rabi oscillations.

The two remaining resonance conditions are termed sidebands, with $\delta = +\omega_x + \Delta$, and $\delta = -\omega_x - \Delta$, with $\Delta \ll \omega_x$, yielding the blue-sideband (BSB) and red-sideband (RSB) respectively. To ensure off-resonant carrier excitations are negligible, it is required that $\Omega \ll \omega_x$. In this case the sideband Hamiltonians are written as,

$$\hat{H}_{\text{BSB}} = \frac{i\hbar\eta\Omega}{2} \left(\hat{\sigma}_+ \hat{a}^\dagger e^{-i(\Delta t - \phi)} - \hat{\sigma}_- \hat{a} e^{i(\Delta t - \phi)} \right), \quad (2.42)$$

$$\hat{H}_{\text{RSB}} = \frac{i\hbar\eta\Omega}{2} \left(\hat{\sigma}_+ \hat{a} e^{i(\Delta t + \phi)} - \hat{\sigma}_- \hat{a}^\dagger e^{-i(\Delta t + \phi)} \right), \quad (2.43)$$

with the subscript labelling either BSB or RSB interactions. When $\Delta = 0$, the RSB is the Jaynes-Cummings Hamiltonian familiar from cavity QED, driving oscillations of a single quanta of energy between a two-level system and a photon in a cavity mode, coupling the two states $|\downarrow, n\rangle \leftrightarrow |\uparrow, n-1\rangle$, as,

$$R_{\text{RSB}}(\Omega_{\text{RSB}}^{(n)} t) := e^{-i\hat{H}_{\text{RSB}} t/2} |\downarrow, n\rangle = \cos\left(\frac{\Omega_{\text{RSB}}^{(n)} t}{2}\right) |\downarrow, n\rangle + \sin\left(\frac{\Omega_{\text{RSB}}^{(n)} t}{2}\right) |\uparrow, n-1\rangle, \quad (2.44)$$

where we have assumed $\phi = 0$, and $\Omega_{\text{RSB}}^{(n)} = \sqrt{n}\eta\Omega$ is the n -dependent modified RSB Rabi-frequency.

The BSB is described by the anti-Jaynes-Cummings Hamiltonian, where pairs of quanta are simultaneously created or annihilated, akin to pair-creation in two-mode squeezing interactions, coupling the two states $|\downarrow, n\rangle \leftrightarrow |\uparrow, n+1\rangle$, as,

$$R_{\text{BSB}}(\Omega_{\text{BSB}}^{(n)} t) := e^{-i\hat{H}_{\text{BSB}} t/2} |\downarrow, n\rangle = \cos\left(\frac{\Omega_{\text{BSB}}^{(n)} t}{2}\right) |\downarrow, n\rangle + \sin\left(\frac{\Omega_{\text{BSB}}^{(n)} t}{2}\right) |\uparrow, n+1\rangle, \quad (2.45)$$

where $\Omega_{\text{BSB}}^{(n)} = \sqrt{n+1}\eta\Omega$ is the n -dependent modified BSB Rabi-frequency.

The resulting resolved sideband structure is shown schematically in Figure 2.5 with the carrier, RSB, and BSB transitions highlighted for one state, $|\downarrow, 1\rangle$. In contrast to the coherent displacements in § 2.3.3, these sideband interactions offer a method for coherent control in the discrete Fock-state basis, selectively coupling quantised states of the oscillator. They also provide a method to directly entangle the spin and motional subsystems of the hybrid system.

2.4.2 The Spin-Dependent Force

Here we introduce a fundamental interaction used throughout this thesis — the spin dependent force, or SDF. The SDF is used in this thesis for three applications: For two-qubit entanglement via a shared motional mode (§ 2.4.3 and § 4.3.2), for motional state tomography (§ 6.1.1), and for driving effective nonlinear bosonic interactions (§ 6.1.1).

As the name suggests, we consider displacements, of the form given in equation 2.32, with conditioning dependent on an ancillary qubit state, described by the Hamiltonian,

$$\hat{H}_{\text{SDF}} = \hat{\sigma}_{\phi_s} \hat{H}_{\text{D}}, \quad (2.46)$$

with the spin-conditioning given by spin phase, ϕ_s , as $\hat{\sigma}_{\phi_s} = \cos(\phi_s)\hat{\sigma}_x - \sin(\phi_s)\hat{\sigma}_y$. In this section we analyse how to generate an effective SDF using a combination of RSB and BSB fields. In the next section we discuss using the SDF and the resulting spin-dependent geometric phase to realise two-qubit entangling gates.




Figure 2.6: Motional phase space trajectories for off-resonant spin-conditioned displacements acting on the two-qubit, single-qumode state initialised as $|\downarrow\downarrow, 0\rangle$. Both $|++\rangle$ (red) and $|--\rangle$ (blue) trace circular trajectories, accumulating a geometric phase, whilst $|+-\rangle$ and $|-+\rangle$ (green) remain static at the origin.

In this thesis the SDF is realised by the optical Mølmer-Sørensen (MS) scheme [37] via a bichromatic laser field. Bichromatic refers to the simultaneous application of two frequency tones. The tones are set symmetrically detuned around the carrier transition, $\omega_1 = \omega_0 + \omega_x + \Delta$, $\omega_2 = \omega_0 - \omega_x - \Delta$. The bichromatic Hamiltonian is then given by,

$$\hat{H}_{\text{bi}}(t) = \hat{H}_{\text{RSB}}(t) + \hat{H}_{\text{BSB}}(t). \quad (2.47)$$

The parameters for the sidebands are labelled with subscript r , b , for RSB and BSB respectively. To realise the SDF, the Rabi frequencies are balanced, $\Omega_r = \Omega_b = \Omega$. Expanding the Hamiltonian, and grouping spin and motional terms, we find,

$$\hat{H}_{\text{bi}} = \hat{\sigma}_{\phi_s} \frac{\hbar \Omega_{\text{SDF}}}{2} \left(\hat{a}^\dagger e^{-i(\Delta t + \phi_m)} + \hat{a} e^{i(\Delta t + \phi_m)} \right), \quad (2.48)$$

which is equivalent to equation 2.46 with $\Omega_D \rightarrow \Omega_{\text{SDF}} = \eta\Omega$, spin phase given by $\phi_s = \frac{\phi_b + \phi_r + \pi}{2}$, and a motional phase defined as $\phi_m = \frac{\phi_r - \phi_b}{2}$.

2.4.3 Geometric Phase Gates

To generate two-qubit entanglement we utilise the “MS gate” method [38]: the SDF is applied simultaneously to two qubits, and a common motional mode. The MS gate relies on the spin dependent geometric phase accumulated during the motional displacement [39]. To ensure there is no remaining spin-motion entanglement the SDF is driven off-resonantly by Δ for an integer loop duration T_n , as described

in § 2.3.3, such that the motional state returns to the initial state no matter the two-qubit spin state.

Here we assume the spin phase, $\phi_s = 0$, so the SDF is conditioned on $\hat{\sigma}_x$, and to simplify the analysis we assume the motional phase, $\phi_m = 0$. If applied to two qubits, the spin conditioning is the collective spin operator $\hat{S}_x = \hat{\sigma}_x \otimes \hat{I} + \hat{I} \otimes \hat{\sigma}_x$, yielding the two-qubit and single-qumode Hamiltonian,

$$\hat{H}_{\text{MS}}(t) = \frac{\hbar\Omega_{\text{SDF}}}{2} \hat{S}_x (\hat{a}^\dagger e^{-i\Delta t} + \hat{a} e^{i\Delta t}). \quad (2.49)$$

As with the analysis of the geometric phase of the non-spin-conditioned displacement, we analyse this time dependent Hamiltonian by the Magnus expansion, to find,

$$\hat{U}_{\text{MS}}(t) = \hat{D}(\beta(t)\hat{S}_x) e^{-i\Phi(t)\hat{S}_x^2}, \quad (2.50)$$

where both β and ϕ are the same as in equations 2.36 and 2.37, respectively, but with $\delta \rightarrow \Delta$ and $\Omega_D \rightarrow \Omega_{\text{SDF}}$. Considering the loop closure time $T_n = \frac{2\pi n}{|\delta|}$, the dynamics reduce to the time-independent effective Hamiltonian, with no motional component,

$$\hat{H}_{\text{eff}} = \frac{\pi\hbar\Omega_{\text{SDF}}^2}{4\Delta} \hat{S}_x^2. \quad (2.51)$$

Taking $\hat{S}_x^2 = 2(\hat{I} + \sigma_x \otimes \sigma_x)$, the spin-conditioning has a global and two-body phase component. The global phase may be ignored, leaving the final propagator as,

$$R_{xx}(\theta) := e^{-i\theta\hat{\sigma}_x \otimes \hat{\sigma}_x / 2}, \quad (2.52)$$

with angle,

$$\theta = \frac{\Omega_{\text{SDF}}^2 T_n}{\Delta}. \quad (2.53)$$

As such, the phase accumulated is independent of the initial motional state (assuming we are in the resolved sideband regime), meaning the gate is robust to imperfect cooling of the participating motional mode.

Considering the initial spin-state $|\downarrow\downarrow\rangle$, the maximally entangled two-qubit Bell state is generated after applying $R_{xx}\left(\frac{\pi}{2}\right)$,

$$\hat{R}_{xx}\left(\frac{\pi}{2}\right) |\downarrow\downarrow\rangle = \frac{1}{\sqrt{2}}(|\downarrow\downarrow\rangle - i|\uparrow\uparrow\rangle). \quad (2.54)$$

Setting $\theta = \pi/2$ and ensuring the duration is an integer number of loop periods restricts the gate detuning to $\Delta = \delta_g$ and $T = T_n$ with,

$$\delta_g = 2\sqrt{n}\Omega_{\text{SDF}}, \quad (2.55)$$

$$T_n = \frac{\pi\sqrt{n}}{\Omega_{\text{SDF}}}. \quad (2.56)$$

The phase space dynamics for this entangling gate are shown in Figure 2.6. The initial state is expressed in the x -basis as $|\downarrow\downarrow\rangle = \frac{1}{2}(|++\rangle + |+-\rangle + |-+\rangle + |--\rangle)$, with the trajectories associated with $|++\rangle$ and $|--\rangle$ tracing closed loops in phase space, while spin-component $|+-\rangle$ and $|-+\rangle$ remain at the origin.

3

Apparatus

Contents

2.1	The Ion	9
2.2	Spin	12
2.2.1	The Qubit	12
2.2.2	Coherent Spin Control	12
2.3	Motion	15
2.3.1	Linear Paul Trap	15
2.3.2	The Qumode	18
2.3.3	Driving Motion	19
2.4	Spin-Motion	21
2.4.1	Resolved Sideband Interactions	21
2.4.2	The Spin-Dependent Force	23
2.4.3	Geometric Phase Gates	24

After 70 years of ion confinement development (since the quadrupole mass spectrometer in the 1950s [34]), and the myriad use cases from time standards, to analytic chemistry, to quantum computing, there are naturally many systems and theses’ apparatus chapters in existence today. As such, we structure this chapter in the following way: **Outdated** Following on from the standard requirements of an ion trap system, provided at the beginning of Chapter 2, we discuss the unique motivation and design specifications of our system; we then present an abstract schematic of the system in its current form, highlighting the constituent subsystems,

and the optical and electrical signal paths; the remaining sections of this chapter contain technical details, perhaps more relevant to the practicing experimental student, and the hardware enthusiastic.

3.1 Specification

Ion-trap experiments are exceptionally versatile in their application, however in the vast search space of quantum-information applications, we want a system particularly tailored for three avenues of study:

1. Error channels on multi-qubit chains.
2. Hybrid system quantum computing and simulation utilising the spin and motional degrees of freedom of the confined ions.
3. Fast two-qubit entangling gates via phase stabilised optical standing waves.

These latter two research directions were explored within our group on a previous apparatus, leading to state-of-the-art two-qubit entanglement rates on the ion trap platform [40, 41], creation of highly non-classical motional states [42, 43], and a simulation of $\mathbb{Z}_2$ gauge theories encoded via the spin and motional degrees of freedom of a two-ion chain [44, 45]. However, hard limitations of both scaling these demonstrations to longer ion chains, and of interaction strength due to available optical intensity, were often restrictive. To solve both these issues, a new apparatus has been designed, constructed, and characterised, motivated by these applications. Colloquially the system is named “*FastGates*”, owing to its experimental objectives, and we shall refer to it as such throughout this chapter.

These objectives give the following specific control requirements:

- Storage of medium length (5-10) qubit registers, § 3.2.
- Stable motional modes for storing quantum information, § 3.2.1.

- Dual high-numerical-aperture (NA) optical access for high-intensity standing waves, § 3.4.
- Individual internal state control of each qubit, § 3.6.5.

As the lifetime of such an apparatus extends past a single PhD generation, this thesis serves as a point-in-time description of the system, and where applicable, the planned future extensions will be highlighted as such. An abstract schematic of the current configuration is shown in Figure 3.1 with blocks resembling either electronic, mechanical, or optical subsystems.

3.2 The Ion Trap

Recently, the micro-fabricated surface-style linear Paul trap has gained popularity due to the maturity of chip fabrication technologies [47] and the potential route to scalability this offers. In the surface trap, the 3D radial and axial electrodes of a “macro” trap are effectively projected onto a 2D surface. The stable confining point of such a trap is typically designed to be on the order of $50 - 100 \mu\text{m}$ above the chip surface. The maturity of chip fabrication has enabled the manufacturing of complex multi-zone devices with many DC electrodes, amenable to ion shuttling [48, 49], and hence furthering progress towards Quantum CCD type architectures [12]. However, this surface-style geometry typically comes with two costs: the depth of the trapping potential is often greatly reduced due to anharmonicity [50], and the close proximity of the surface to the ion can be a large contributor to motional heating rates [51–54].

Our trap, provided by the *National Physical Laboratory* in the UK, brings together the advantages of chip fabrication as well as the low heating rates and high trapping depths of a macro 3D style trap through utilising a micro-fabricated 3D trap geometry with greater ion-surface distances than typical planar traps. Details on its design and characterisations can be found in Ref. [46, 55, 56]. Figure 3.2 shows the electrode geometry of the trap and the relevant length scales. The trap package consists of an outer leadless chip carrier (LCC) with outer pads routing DC, RF,




Figure 3.1: **point to glossary?** Block diagram schematic of the *FastGates* apparatus. Blocks resemble discrete subsystems, which shall be discussed within this chapter. Lines resemble relevant optical (free-space or in-fibre), RF, DC, and communication paths between subsystems. Block location signifies whether the device is in-vacuum, within the magnetic shielded (MuMetal) box, within the “trap lab”, or within the “laser lab”.




Figure 3.2: Schematics of the *NPL* trap chip used in this experiment. (a) A photograph of the trap chip used in the *FastGates* experiment before being placed under vacuum. The top “atomic shield” can be seen here. (b) A front view of the trap chip. Axial confinement is provided by a subset of the segmented DC electrodes. Radial confinement is provided by the radio-frequency voltage (RF) rails. The trap chip and leadless chip carrier (LCC) are highlighted. (c) A schematic of the electrode geometries and labelling for the front-face of the trap chip. A similar arrangement exists for the back-face however with the RF rail on the diagonally opposite electrode. (d) A cross-sectional view of the trap. The red point here represents the ions, the gold are the RF and DC electrodes, and the grey is the silica substrate. It can be seen that the ion is $350 \times \sqrt{2}/2 \simeq 250 \mu\text{m}$ away from the nearest electrode. The numerical aperture of the ion plane due to the electrodes is ~ 0.71 . Figures (b) and (d) adapted from Reference [46], other panels adapted from private communication with the *NPL* team.

and ground connections to an intermediate Al_3N_4 substrate. The connection to the front and back surfaces of the trap chip is made through vias and wire bonds to the intermediate substrate [57]. The trap chip is a monolithic double-sided silicon device with gold plated electrodes on the front (*Fr*) and back (*Ba*) surfaces. These two layers are shown schematically in the bottom panel of (c). The electrode schematic for the front surface is shown, and the electrode indexing convention is highlighted. Electrodes recessed from the ion due to the RF rail are named


“compensation”-electrodes (*co*), while electrodes on the opposite side provide the strong DC potential required for axial confinement¹. We typically trap ions centered on electrode *dc5*. Panel (d) shows a cross-sectional view of the trapping region, highlighting the ion position between the front and back surface electrodes. The dual RF rails and DC electrodes are shown in gold, and the undercut silicon bulk is highlighted in dark grey.

The next two subsections describe the source of RF and DC voltages to the trap electrodes. The physical mounting of the trap in the vacuum system is discussed in § 3.4, and the optical access to the trap is discussed in § 3.4.

3.2.1 Trap RF Chain

To utilise radial motional modes for low error quantum gates, we require the RF field that produces the desired trapping pseudopotential to be both frequency and amplitude stable. Furthermore, from equation 2.22, it can be seen that radial motional mode frequency is proportional to RF voltage, and inversely proportional to the RF frequency. If using the motion to enable two-qubit gates (§ 2.4.3), adiabatic excitation with respect to the motional mode period yields a characteristic time scale of $\tau = 2\pi/\omega_m$ for the gate duration. Therefore higher RF voltages are desirable for faster entangling operations.

The *NPL* trap used in the *FastGates* apparatus is rated for a max peak-to-peak voltage of 400 Vpp on the RF electrodes. To target high motional mode frequencies, saturating this maximum voltage is desirable, however, for safe operation, typically the trap is provided a maximum RF voltage $V_{RF} = 360$ Vpp. The RF chain which provides this signal is schematically shown as:



¹Note: Compensation electrodes *co1* and *co11* on both front and back surfaces are grounded on the trap package, and so are not available for use. There is a continuity fault on front *dc11* and so this electrode is also not available for use.




Figure 3.3: **d should be FM not PM** Characterisation of the RF chain. (a) SSB spectra for AM and FM, given as fractional noise, measured after the *Urukul*. (b) Measured after the *Squareatron*. (c) Measured after the amplifier.




Figure 3.4: Expected SSB after the resonator, and estimated mode frequency noise. (a) Resonator circuit diagram. (b) Resonator lineshape, given by magnitude of the transfer function $|H(f)|$. (c) Expected AM and FM fractional noise spectra after the resonator. (d) Integrated RMS noise of radial motional mode frequency due to the RF chain.

and embedded within the schematic 3.1, beginning with the *Trap FPGA*. Each element of the chain is designed to provide a low-noise, high-voltage RF signal to the trap electrodes, and is discussed in turn below.

The frequency source is a direct-digital-synthesiser (DDS), named *Urukul* [58], part of the *Artiq Sinara* [59] hardware ecosystem. *Urukul* is used extensively in our lab for driving acousto-optic modulators (AOMs) and here used as the source drive of our trap RF. The *Urukul* is operated at maximum output power, +10 dBm, and with a frequency ~ 35 MHz (set by resonator, discussed below). As can be seen in

equation 2.22, amplitude and frequency instability will directly lead to radial mode instability. Expanding for small amplitude, dV_{RF} , and drive frequency, $d\Omega_{RF}$, errors we find the motional mode frequency, ω'_r , is related to the ideal frequency, ω_r , by

$$\omega'_r = \omega_r(1 + dV_{RF}/V_{RF} - d\Omega_{RF}/\Omega_{RF}), \quad (3.1)$$

where we can see both errors affect the motional mode frequency linearly. The instability of the oscillator can limit motional mode coherence time (see § 4.2.4 for characterisation with the ion). We measure noise originating from both amplitude modulation (AM) and frequency modulation (FM) across the RF chain using a phase noise analyser². This was done to verify low-noise operation of the individual devices, and to estimate expected motional mode errors due to the inherent noise on the RF chain.

We quote FM as the fractional frequency stability power spectral density (PSD),

$$S_y(f) = \frac{f^2}{f_0^2} S_\phi(f), \quad (3.2)$$

where $S_\phi(f)$ is the phase noise PSD (which is the Fourier transform of the Autocorrelation function) f is the offset frequency, and f_0 is the carrier frequency. $S_y(f)$ has units of $\left[\frac{1}{\text{Hz}}\right]$. AM noise is quoted as the single sideband (SSB) power, S_α , in a 1 Hz bandwidth at an offset frequency, f , from the carrier, normalised to the total power of the carrier, and is here given in units of $\left[\frac{\text{dBc}}{\text{Hz}}\right]$.

Figure 3.3, panel (a) displays the *Urukul* AM and FM spectra. The *Urukul* is phase referenced to the master FPGA board (*Kasli*), which is in turn phase locked to an external 10 MHz atomic clock³. It can be seen that AM noise dominates up ~ 1 MHz, suggesting that this will be the limiting error source for motional mode stability. The AM noise may be reduced by using a nonlinear voltage clamp. We use an “ultra-low noise limiting amplifier” named *Squareatron* developed previously in the group, and further refined by *Oxford Ionics*. *Squareatron* takes an input RF signal and converts to a stable amplitude square-wave output which is subsequently bandpassed around the desired RF frequency to recover the sinusoidal signal, but

²Rhode and Schwartz FSWP PNA

³**ToDo: Abaqus atomic clock PN**

with significantly reduced AM noise. The AM and FM after the *Squareatron* is shown in panel (b). It can be seen that between 1 kHz and 1 MHz, AM is considerably suppressed. There appears to be an increase in AM at higher frequencies **ToDo: explore this further**. The *Squareatron* has a 17 dBm output power, which is attenuated by -20 dB before an amplification stage to reach the desired power for the trap. Highlighted in the above chain schematic is the option to input a parametric excitation signal (P.E.), which we shall discuss further in § ???. For the following measurements, the P.E. branch was removed. AM and FM introduced by the low-noise amplifier⁴, can be seen in panel (c). Most notably an increase in AM is seen between 1 kHz and 10 kHz.

To impedance match the $50\ \Omega$ line from the amplifier to the small capacitance of the trap electrodes and vacuum feedthrough, a resonant matching circuit is used. This “resonator” has three main effects on our system:

1. The source $50\ \Omega$ is matched to the (almost) purely capacitive load.
2. The voltage step up provides the necessary amplitude for high motional mode frequencies.
3. The input RF is bandpassed around the desired trapping frequency.

The latter two points are practically useful, especially the bandpass effect of the resonator attenuating noise on the RF line, however, changes to the resonators Q factor and central frequency due to thermal and mechanical noise can be a significant source of low frequency motional mode instability. Opting for good passive stability, a PCB based design⁵ with good thermal management was seen as preferable to large helical resonators, which are often used in ion trap experiments for their high Q factors. Aiming for a moderate Q factor of ~ 70 allows for a good voltage step up, and some bandpass filtering, without being overly sensitive to thermal and mechanical noise. The circuit diagram of the constructed resonator⁶

⁴Mini-Circuits ZHL-1-2W-S+

⁵This style PCB resonator has been used around the Oxford Ion Trapping Group for some time, and is based originally on a design by Tom Harty [60].

⁶Credit to Emily Hirsch for construction of the recent resonator.

is shown in figure 3.4 (a). The inductor, $L \sim 0.8 \mu\text{H}$, is hand-wound copper wire with a toroidal iron powder core⁷, and the capacitors are low tempco C0G rated ceramic, $C = 5.4 \text{ pF}$. The vacuum can feedthrough and trap were measured to present an $\sim 18 \text{ pF}$ load. The trap itself presents a $C_T = 4 \text{ pF}$ [57], the cable between the resonator and the vacuum feedthrough contributes $\sim 4 \text{ pF}$, and the remaining presented capacitance is expected to originate from the feedthrough and in-vacuum RF line. In vacuum, the RF line is a *Kapton* coated 0.4 mm diameter single strand copper wire, which connects to an in-vacuum distribution PCB. The RF connection to the trap LCC is made by *Fuzz button*⁸ vias.

Outdated - or need plot Measuring the motional mode frequency of a trapped ion while varying the trap RF frequency (see § ??), yields a resonator quality factor $Q = 57.88(6)$, and central frequency $34.659(1) \text{ MHz}$. The effect of the resonator on the AM and FM noise can be calculated through the transfer function, $H(f)$, of the resonator circuit. Capacitors ($C2$ and $C3$) are used to provide a pick-off (p.o.) voltage which can be monitored, and if required, used in a closed feedback loop. We however have not attempted this as of yet, and shall ignore the effect of the pick-off capacitors for this analysis. The transfer function of the simplified circuit is then given by,

$$H(f) = \frac{V_{out}}{V_{in}} = \frac{Z_p}{Z_s + Z_p}, \quad (3.3)$$

where

$$\begin{aligned} Z_p &= \left(\frac{1}{j2\pi f L} + \frac{1}{Z_L} \right)^{-1}, \\ Z_s &= R_s + \frac{1}{j2\pi f C_1}. \end{aligned} \quad (3.4)$$

Given the approximate capacitance and inductance values of the resonator and trap, the analytic $|H(f)|$ is plotted in Figure 3.4 panel (b), and shows acceptable agreement with the measured quality factor and central frequency (here $Q = 69$ and $f_0 = 36.8 \text{ MHz}$) with the discrepancy of quality factor expected to originate

⁷Micrometals T80-6


⁸Custom Interconnects

from excess resistance in the inductor not accounted for in the model. Final AM and FM spectra are calculated through considering both upper ($+\Delta f$) and lower sidebands ($-\Delta f$), as shown in panel (b), following reference [61]. We find that for both AM $\rightarrow$ AM and FM $\rightarrow$ FM (where “ $\rightarrow$ ” represents the transform from input noise type to output noise type) the input PSD are low pass filtered with a cut-off frequency of ~ 0.25 MHz. We also consider the cross term, FM $\rightarrow$ AM, which here we expect to be non-negligible at higher offset frequencies due to the asymmetry of the resonator lineshape. Panel (c) shows the expected fractional noise after the resonator, with total FM noise (orange line), total AM noise (blue line), and contributions from AM $\rightarrow$ AM (dashed green), FM $\rightarrow$ AM (dashed red). Integrating the AM and FM fractional noise spectra between some chosen range of offset frequencies, f_{min} and f_{max} , gives the expected fractional RMS noise on the radial mode frequency of the ions. Assuming a mode frequency of $w_r = 3$ MHz the RMS, σ_r , is shown in Figure 3.4 panel (d) as a function of f_{max} , with $f_{min} = 10$ Hz. It can be seen that the expected noise is dominated by the AM contribution, and in the frequency range considered, corresponds to a mode frequency instability of $\sigma_r = 60$ Hz. This analysis excludes any noise less than 10 Hz, and environmental noise coupled in by the resonator, which may be a significant contribution at low frequencies due to thermal and mechanical noise. The motional modes will also be subject to noise from other coupling mechanisms, such as voltage noise on the DC electrodes, or fluctuating stray electric fields, which we do not consider here.

Here we have characterised both AM and FM noise on the RF chain and estimated the expected contribution to motional mode instability.

3.2.2 Trap DC Voltages

Static fields are required for reshaping the trapping pseudopotential. This is useful for setting desired mode frequency splittings, for shuttling ions, and for compensating stray electric fields. As shown in Figure 3.2, the *NPL* trap has 40 independently controllable DC electrodes. A further DC bias voltage can be applied directly to the silicon bulk. DC voltage is supplied to the LCC through the following chain:



and embedded within the schematic 3.1. Two multi-channel DACs, known as *Fastino*, part of the *Artiq Sinara* hardware ecosystem, provide the initial 41 DC channels. The *Fastino* outputs are filtered by a passive cascaded RC low pass filters with resistances, $R = 2k$, and capacitances $C_1 = 2.2$ nF and $C_2 = 4.7$ nF⁹ yielding a -3 dB “cut-off” frequency of 7.4 kHz. The low pass filter circuit is shown in the above schematic, with the first stage on an external-to-vacuum PCB (Ext. Filt), and the final shunt capacitor on an in-vacuum PCB (Vac. Filt)¹⁰ within close proximity of the trap chip to mitigate RF pickup on the DC lines. The feedthrough to the vacuum system has 51 pins¹¹ with non-DC pins used to provide ground connections. On the external filter board, one channel, corresponding to electrode *Ba co4*, has a capacitively coupled in AC signal. This originates from a DDS (*Urukul*) and is used to resonantly excite the motional modes of the ion chain for characterisation and calibration purposes (see § 4.2.1). The in-vacuum PCB is connected to the trap chip via a custom interposer made of PEEK polymer, with embedded *Fuzz button* vias¹².

The ion is confined in the “operation” zone, seen in Figure 3.2 (c). The axial trapping is provided by electrodes *Ba dc5* and *Fr dc5*. The DAC can provide up to 10 V to these electrodes to produce an axial frequency of 1.6 MHz.

⁹For in-vacuum use. 1206 NPO 4.7nF 100V, SRT 1206A472JBCB

¹⁰Vacuum compatible PCB using Rogers RT6002. Filter boards were designed by Mariella Minder.

¹¹*Allectra* 51-way Micro-D

¹²*Custom Interconnects* Au/BeCu 0.025" diameter x 0.104" length Fuzz Button

3.3 Magnetic Field

Due to the fragility of the internal states of the ion (these are *literally* state-of-the-art magnetic field sensors), we must take great care in isolating the ion from a noisy environment. We do not have access to magnetic field insensitive transitions in $^{40}\text{Ca}^+$, therefore uncontrolled Zeeman shifts of the internal atomic energy levels must be suppressed through enforcing a stable external magnetic field. This magnetic field stability is required for long spin coherence times (see § 4.1.5), and for low error single- and two-qubit gates. A permanent magnet array of Samarium Cobalt¹³, $\text{Sm}_2\text{Co}_{17}$, is used in a Helmholtz configuration to create a stable magnetic field of $\simeq 5 \text{ G}$ ¹⁴. Samarium Cobalt was chosen for its low temperature coefficient of remanence, $-0.03\%/K$.

The ion is shielded from unwanted external magnetic fields by a custom designed MuMetal enclosure¹⁵, with an internal dimension of $650 \times 650 \times 650 \text{ mm}$. The shielding consists of two layers of MuMetal shielding with internal shield thickness of 1.5mm and external shield thickness of 2.0mm. The factory quoted shielding factor at the center of the box is $\times 546$ for DC fields.

3.4 Vacuum System and Beam Geometry

Ultra High Vacuum (UHV) is required to prevent ion loss. However, UHV equipment is often bulky, and puts constraints on the access and visibility of the ion chain. Here, we describe the custom vacuum system, allowing for multiple beam access geometries including dual high-numerical-aperture (NA) objectives required for single addressed optical standing waves. The vacuum system was designed by Sebastian Saner, Christopher Ballance, and Mariella Minder, and constructed by Sebastian Saner, Fabian Pokorny, and myself.

A residual pressure of $< 10^{-11} \text{ mbar}$ is desired. For this strict requirement, care

¹³*Bunting Magnets* E643SM and E644SM.

¹⁴Designed and constructed by Sophie Howarth.

¹⁵*Magnetic Shields* custom design.




Figure 3.5: (a) Laser and atomic beam geometries with respect to the ion chain. (b) Compact trap housing with dual NA=0.6 access. The size of the high-NA objectives restricts access for other beams, and so in-vacuum mirrors are used to redirect low-NA beams (373, 397, 423, 854, 866-nm) to the trap center. (c) Exploded view of the trap housing with components labelled.




Figure 3.6: Cross section of the vacuum chamber and dual-high-NA objective sandwich. Lower right inset shows the vacuum chamber housed in the MuMetal enclosure, with dashed lines indicating the cross section region. Custom recessed CF-100 viewports are used to provide optical access for the high-NA objectives (working distance = 24.4 mm).

must be taken in selecting in-vacuum materials, and thorough cleaning and baking procedures must be followed. A summary of tactics that were useful in the construction of the vacuum system can be found in [62, 63].

The system consists of a 6-inch spherical-octagonal experimental chamber¹⁶ connected to a spherical chamber with ion pump¹⁷, ion gauge¹⁸ and Titanium sublimator pump (TSP)¹⁹ attached. After the final system bake, a He leak test was performed and no external leaks were found. We do observe a gradual pressure increase that we suspect is either due to the in-vacuum optics, the optics epoxy adhesive²⁰, or the PCB solder. Regardless, the ion pump and TSP maintain UHV at the desired $< 10^{-11}$ mbar (the TSP is fired every 4 weeks at 41 A for 60 seconds to provide a fresh Ti layer).

Figure 3.5 (a) shows the beam geometries with respect to the ion chain axis. The quadrupole 729-nm beam has high-NA (NA = 0.6) access perpendicular to the ion chain axis, while 373-nm, 397-nm, 423-nm, 854-nm and 866-nm beams have low-NA (NA < 0.05) access at an angle of 45° to the ion chain axis. The atomic beam from the oven is at an angle of $\sim 57^\circ$ to the ion chain axis. As such, there is a doppler shift of the atomic beam with respect to the 423-nm photoionisation laser, which is accounted for in the set laser frequency (given 12° angle between atomic and 423-nm beam, and 500 K mean temperature, expect +126 MHz frequency shift). Figure 3.5 (b) shows the custom trap housing including NA=0.6 access for the 729-nm beams, and in-vacuum mirrors (prisms) to redirect the low-NA beams to the trap center. The low-NA beams enter the housing via one of the four beam access channels parallel to the trap chip plane, and are reflected through the trap slit. The exiting beam is reflected by a second mirror out of the housing to prevent scattered light incident on the ion chain. Figure 3.5 (c) shows an exploded view of the trap housing with components labelled. The trap chip is interfaced to an in-vacuum PCB (described in § 3.2.2) via a custom PEEK interposer with the embedded *Fuzz*

¹⁶Kimball MCF600-SphOct-F2C8

¹⁷Agilent VacIon Plus 20 Pump

¹⁸Agilent UHV-24P Ion gauge

¹⁹Scanwel custom housing

²⁰EPO-TEK 353ND




Figure 3.7: Optical path for fluorescent imaging of the ions using both a CMOS camera and a Photon-multiplier tube. Insets on right hand side show typical signal for: spatial resolution of eight ions in a linear chain (top), time resolution of detected photons from the ion chain (bottom).

button vias (one for each DC and RF channel). The trap chip is held in place by 4 steel clamp arms, which attach to a steel base plate which provides mechanical stability and houses the lower in-vacuum prisms. A top steel plate houses the upper prisms. The tube oven is mounted on two steel clamps which are electrically isolated from the base plate by a PEEK spacer. The trap housing is fastened to the vacuum can by 4 steel arms attached to alternating side ports of the octagonal chamber²¹. Figure 3.6 shows a cross section of the vacuum chamber and recessed dual high-NA objectives. Two recessed CF100 viewports²² are required to bring the custom objectives²³ to the designed working distance of 24.4 mm from the ion chain.

3.5 Imaging System

To image the ion chain, and readout the qubit states, fluorescent scatter of incident 397-nm light is collected. Here we briefly describe the imaging system, while § 4.1.2 describes the qubit readout procedure and fidelity. The imaging system is designed to provide high spatial resolution of the ion chain, while also providing a large field of view to accommodate long ion chains. The imaging system must also be compatible with the high-NA objectives used for single addressed optical standing waves.

²¹Using *Kimball Groove grabbers* P/N 53-482610

²²*UK Atomic Energy Authority* P/N VPR100015

²³*Photon Gear, inc.* Custom 18WD Atom Imager

Figure 3.7 shows the optical path schematic of the imaging system, with both a high quantum efficiency ($\simeq 70\%$ at 397 nm) CMOS camera²⁴ for spatially selective qubit readout on a multi-ion chain, and a photon multiplier tube (PMT)²⁵, for experiments requiring time resolved photon counting (e.g. micromotion compensation, § ??). Typical signals of both devices can be seen in the right panels of Figure 3.7. The optical path consists of a high numerical aperture objective²⁶ (effective focal length, $f = 33$ mm), NA = 0.6, a dichroic mirror²⁷ (transmission: 397 nm; reflection: 729 nm), a relay lens²⁸ ($f = 200$ mm), a 90 : 10 plate beamsplitter²⁹, and the camera and PMT.

The objective lens is custom designed to compensate for aberrations due to the recessed vacuum viewport, however, due to the large NA, accurate alignment of all degrees of freedom of the objective is critical. As such, the NA is attached to a 5-axis alignment stage³⁰, with z-axis focus controlled by a piezo micrometer³¹ for fine adjustment of the ion focus with the MuMetal box closed (requires $< \mu\text{m}$ precision).

Given the chosen lens focal lengths, we expect a magnification of $\times 6$ from the ion plane to the imaging plane. With an expected ion separation of $5 \mu\text{m}$, and the camera pixel pitch of $6.5 \mu\text{m}$, we expect around 4.5 pixels between ions, which, if the system is well aligned, is sufficient for low cross-talk spatial readout of qubits.

3.6 Ca^+ Laser Systems

Singly charged group 2 elements are popular in ion trap experiments due to their single outer electron resulting in Hydrogen-like energy levels. In this thesis we use $^{40}\text{Ca}^+$.

$^{40}\text{Ca}^+$ has no nuclear spin giving the (relatively) simple level structure shown in

²⁴Hamamatsu Orca-Fusion BT C15440-20UP

²⁵Hamamatsu Photon counting head H10682-210

²⁶Photon Gear, inc. Custom 18WD Atom Imager

²⁷**Thorlabs TODO**

²⁸Thorlabs Achromat Doublet, ACT508-200-A-ML

²⁹**Thorlabs TODO**

³⁰Newport ULTRAlign M-562-XYZ, M-36, SM-13

³¹Newport CONEX-TRA12CC

Figure 2.2 with the relevant laser transitions used in our apparatus indicated. Details on the ion choice, and utility of the lasers is given in § 2.1, here we only describe the hardware and feedback control to provide this stabilised light.

The relevant wavelengths are the following: The qubit is defined by the quadrupole transition at 729-nm. Access to other excited levels, outside of our defined two level system, are crucial for ion trap quantum logic to enable qubit readout via state selective fluorescence (397-nm, and 866-nm transitions), and state preparation via optical pumping (866-nm, and 854-nm transitions). Access to two additional transitions in neutral calcium for isotope selective ionisation, 423-nm and 373-nm, are also required [64].

This section begins with a description of laser frequency and power stabilisation techniques common to multiple laser systems used in this thesis. Following this, the specific systems are described, starting with the diode lasers, which drive all but the 729-nm transition. Subsequently, the Ti:Sapph 729-nm source is discussed. Finally, the global and single addressed beam arrangements for the 729-nm system are described.

3.6.1 Frequency Stabilisation

All lasers exhibit frequency noise due to unavoidable environmental coupling (e.g. temperature fluctuations, and acoustics), and to instabilities in control signals (e.g. voltage and current controls). For the lasers corresponding to dipole transitions used in this thesis, we are relatively lenient and require slow frequency excursions of < 1 MHz with respect to the qubit transition, and laser linewidths less than the natural linewidth of the dipole transitions, which are often multiple MHz. However, for the quadrupole transition at 729-nm, we require significantly more stringent frequency stability, with laser linewidths of < 100 Hz, and negligible phase modulation of the laser light over a few MHz of the qubit transition to ensure spectator transitions are not excited during qubit operations. Although the order of magnitude of the required frequency stability is different for the dipole and




Figure 3.8: PDH optical and control schematic. Green optical path represents the carrier frequency, blue and orange lines represent the upper and lower sidebands respectively, generated by an EOM. The reflected light is directed onto a fast photodiode, and the signal is demodulated with the same oscillator signal used to drive the EOM, but delayed by some chosen phase. The signal is finally low pass filtered to recover the error signal, an example of which is seen in the upper right panel for a finesse $F \sim 4000$, a free spectral range of $FSR = 1.5$ GHz, and a modulation frequency of $f_{mod} = 5$ MHz.

quadrupole transitions, the concepts and techniques used to achieve this stability are similar. Generally, a closed loop consisting of measurement of the optical frequency of the laser and feedback to the laser control is required. In this work, stable external optical cavities are used for referencing the laser frequency. Specifically, we use the Pound-Drever-Hall (PDH) technique to generate the error signal for the feedback loop. The PDH technique is well described in the literature [65–67], and so here we only give a brief overview of the technique to contextualise the required hardware introduced in the following sections.

The conceptual goal of the PDH technique is to measure the frequency difference between the laser and a resonant frequency of a stable optical cavity to generate an error signal, preferably linear over typical frequency excursions (while *locked*). This error signal can then be modified by some controller (typically PID), to directly infer how to tune the laser controls to minimise this error signal, often termed *locking* the laser to the cavity. Further, the error signal should be insensitive to amplitude noise on the laser light, to avoid amplitude to frequency noise conversion due to the feedback loop (in contrast, side-of-fringe stabilisation is sensitive to amplitude drifts).

Practically this scheme consists of phase modulation of the incident laser light, here by electro-optical modulators (EOMs), to generate sidebands, which are used as phase references to compare the phase of the incoming laser field to that of the phase of the field inside the cavity. Figure 3.8 shows a schematic of the PDH scheme, highlighting the beam paths and closed loop. The phase modulated light is directed onto the cavity, with the total return signal being some combination of the reflected sidebands, carrier, and the cavity leakage field, with the latter providing the phase of the electric field inside the cavity. The reflected light is directed onto a fast photodetector, and the signal is demodulated with the same oscillator signal used to drive the EOM, but delayed by some chosen phase. The resulting signal is low pass filtered to remove high frequency components, yielding the final error signal, an example of which is seen in the upper right panel of Figure 3.8. The central region is anti-symmetric and linear, centered on the desired lock point. Using a tuned controller for feedback to the laser controller, this provides a stable lock of the laser to the cavity.

As the cavity is used as a frequency reference, both a high finesse and stable optical lengths are desirable. The gradient of the PDH error signal is proportional to finesse, leading to increased signal to noise with higher finesse (see Ref. [66] for more details). Optical length deviations are minimised through mounting the cavity in a vacuum chamber, and taking care in isolating it from temperature fluctuations and vibrations.

Particular details of the PDH hardware and control loops are given in the following sections for each laser system.

3.6.2 Intensity Stabilisation

The light-matter interactions considered in this thesis are all proportional either to intensity or amplitude of the optical electric field at the ion's position. Intensity fluctuations over time therefore lead to fluctuations in the interaction strength, which are detrimental to the fidelity of quantum operations. Intensity noise originates

either from total beam power fluctuations, or due to beam pointing drift with respect to the ion position. The former is mitigated through a closed feedback loop between a pick-off photodiode signal before the ion monitoring the optical power, and the drive RF amplitude of an upstream AOM. For all beams, this optical pick off is one of the last elements before the beam enters the vacuum system to reduce the impact of lossy elements between the stabilise point and the ion position. The low-noise photodiodes used were designed by David Nadlinger (*QLNPD*), and the feedback system controller, known as *SUServo*, was written by Peter Drmota.

outdated Beam pointing drifts are mitigated through careful mechanical design of the optical paths for passive stability, and through automated intermittent beam pointing calibrations using the ion as a sensor (see § ??).

3.6.3 Diode Lasers

For transitions 397-nm, 854-nm, and 866-nm, we require light that is frequency stable to < 1 MHz level, well below the natural linewidths of the dipole transitions, and power stable to < 1 μ W. All beam paths and relevant control loops for the diode laser systems are shown in Figure 3.9. The laser diode heads³², controllers³³, and splitterboards³⁴ (fig 3.9 (a)) are housed in a single 19" rack. The splitterboard distributes the light between two ion-trap experiments, the PDH lock cavity (fig 3.9 (c)), and a diagnostic wavemeter for checking wavelength and lock status³⁵. The splitterboard also houses an Electro-optic modulator (EOM) module for applying both offset and PDH sidebands for an offset PDH lock. The offset sidebands may be arbitrarily set (within the GHz range of the EOM) to provide a continuous choice of laser lock frequencies between cavity modes, rather than maintaining a fixed offset between the laser output frequency and the discrete cavity mode frequencies. For 397-nm and 423-nm, resonant free space EOMs³⁶ are used, for 854-nm and 866-nm

³²*Toptica* diodes. 854/866: MDL DL pro; 397: MDL DL pro HP; 423: DL pro

³³*Toptica* DLC Pro with PDH boards

³⁴Designed by *OITG Cavities* group

³⁵*HighFinesse* WS7

³⁶*Qubig* PM7-VIS-25 and PM9-VIS-2.25




Figure 3.9: All beam paths for diode laser systems. Dotted boxes resemble modular optical paths with fibre and control lines in and out only. Frequency and optical power control loops are included.

broadband fibre EOMs³⁷ are used for convenience. The path on the reference cavity board (fig 3.9 (c)) consists of a four-bore reference cavity³⁸ under vacuum and temperature control, two lenses for matching the collimated input beam to the cavity mode, and a beamsplitter and $\lambda/4$ waveplate for directing the reflected cavity light to a fast photodiode³⁹. The beatnote due to the carrier and PDH sidebands are mixed down and low-passed by the *Toptica* PDH module to produce the signature antisymmetric PDH trace. Gains for PID controllers of both the diode current and voltage are set to stabilise the optical frequency.

The light from the “laser” laboratory is carried to the “trap” laboratory using 16 m polarisation maintaining fibres⁴⁰. In the “trap” laboratory we use AOMs in a double pass configuration for both frequency shifting and power stabilising the light (fig 3.9 (d) and (e)).

As space inside the MuMetal enclosure is constrained, it is advantageous to couple multiple beams in a single fibre before the enclosure. The second stage photo-ionisation light, at 373-nm⁴¹, is combined with the 423-nm beam by a dichroic (fig 3.9 (f)) before being coupled to a 16 m fibre that takes the light directly into the MuMetal enclosure. The respective refractive indices of fused silica for 373-nm and 423-nm differ to the extent that the two beams were separated by more than one waist at the ion plane after initial rough alignment. To ensure the two beams overlap at the ion plane, care was taken for the beam propagation to be perpendicular to both the lens and vacuum viewport window, and a parabolic mirror outcoupler⁴² (rather than a lens) was used inside the MuMetal box.

The 854-nm and 866-nm light is combined after their respective double and single pass AOMs, and coupled into a 3 m PM⁴³ fibre carrying the light into the enclosure. Light from both the PI and the 397/854/866-nm paths are directed onto the ions via the in-vacuum prisms, discussed in § 3.4. The final mirrors prior to the vacuum

³⁷*EOSPACE* PM-0s5-05-PFA-PFA-850

³⁸*Stable Laser Systems* SLS-6010-1-4bore

³⁹*Thorlabs* PDA10A2

⁴⁰*Schafter-Kirchoff ...*

⁴¹*Toptica* iBEAM-SMART-375-S.

⁴²*Thorlabs* RC02APC-F01

⁴³*SK ...*

system are motorised⁴⁴, so that beam alignment to the ions can be optimised when the MuMetal enclosure is closed. Practically we find that these beams rarely need to be realigned (few months), unless there is a significant ambient temperature change of a few degrees Celsius.

3.6.4 Narrow Linewidth 729 Laser

Lasers are a key tool for creating the highly localised, strong electric field amplitudes and gradients needed to drive both carrier and sideband transitions of the trapped ion.

As shown in Figure 2.2, two sublevels within the $4S_{1/2}$ to $3D_{5/2}$ manifolds define the qubit. This is an electric quadrupole transition as $\Delta l = 2$. For the Calcium ion this transition is at 729-nm, and so a near resonance 729-nm laser is used to implement single- and multi-qubit gates (§ 4.1.3 and 4.3.2). This transition is also used after Doppler cooling for resolved sideband cooling to prepare the motional mode close to its ground state (as discussed in § 4.2.2). This transition is narrow linewidth due to the long lived $3D_{5/2}$ level, and so, for power efficiency, a narrow linewidth laser must be used. For the “fast entangling gates” use case, high intensity electric-fields are required at the ion to drive sideband transitions, this will be discussed further in the single-addressing and two-qubit gate sections, but here it is sufficient to say we require >100 mW of light at the ion plane. Here we describe the 729-nm system consisting of a Ti:Sapph laser system pumped by an Nd:YAG 532-nm laser.

An *M Squared Solstis* Ti:Sapph is pumped via 18 W of 532-nm light from a *Coherent Verdi-V* system to produce around 5 W of 729-nm light. The Ti:Sapph is engineered to operate with a stable < 50 kHz linewidth. Ti:Sapph crystals have broadband gain profiles which make them amenable in research environments to create frequency tunable laser systems. They also often have lower intrinsic phase noise than diode lasers at the same wavelength, which is favourable for our application. For single mode operation, the *Solstis* employs multiple intracavity frequency selective elements

⁴⁴*Newport Picomotor* piezo mirror mount P/N 8821

which consist of (in order of coarse frequency selectivity), a birefringent filter, a tunable Fabry-Pérot etalon, and the surrounding bow-tie cavity. To lock the tunable etalon to one longitudinal cavity mode, an actively stabilised “dither” phase-lock loop is used. This dither consists of periodically varying the etalon spacing with a frequency of around 20 kHz. This dither frequency directly phase modulates the light, which leads to the creation of sidebands which can interact with the ion causing errors in gates. This dither frequency (and harmonics of it), can be observed via composite pulse experiments on the ion, and are discussed in § 4.1.5.

Referencing the *Solstis* output with an external cavity and using PDH feedback can greatly suppress phase noise. Here an ultra high finesse cavity⁴⁵ is used, with a schematic of the 729-nm system shown in Figure 3.10. As opposed to the simpler controller implemented for the diode lasers, the 729-nm uses a cascaded PID lock with a fast feedback loop to the AOM situated after the *Solstis*, and an intermediate and slow feedback loop to piezos controlling the *Solstis* cavity length. With this system, a linewidth of $\ll 1$ kHz for the 729-nm light is expected. The electronics for this control loop are provided by *Stable Laser Systems* in the form of their *FPGA Servo* lock box. For an effective PDH lock, both short and long term stability of the reference cavity is required. To ensure the cavity is insensitive to the environment, it is made of an ultra-low expansion material, and mechanically isolated from the environment by an acoustic dampening enclosure and a vibration isolation platform. The cavity is housed in a vacuum system to prevent refractive index fluctuations due to air, and to further reduce thermal coupling to the environment. The cavity is temperature stabilised at the zero crossing temperature of 30.6°C. Long term cavity drifts of 309(1) Hz/hr are measured over 125 days, seen in Figure 3.11. These measurements use the ion at a known magnetic field as a frequency reference.

Figure 3.10 displays the other beam paths for our 729-nm system. Some light is picked off and sent to a wavemeter to continuously monitor the frequency. However, the majority is coupled to two output fibres for our experiment and another within the group. The 729-nm light is transported from a dedicated laser lab to a the

⁴⁵*Stable Laser Systems* P/N 6020-4




Figure 3.10: The 729-nm system. A Ti:Sapph laser tuned to 729-nm is pumped by a 532-nm source. Light is picked off at the first beam sampler to stabilise by PDH locking to a cavity. Light is distributed to two experiments, with the *FastGates* line including a fibre noise cancellation system.




Figure 3.11: Cavity drift over 125 days measured by reference to the ion transition. The laser offset here is with respect to running all 729-nm AOMs at their central frequencies. Outliers are due to bad fits of the ion transition frequency, not due to changes in the cavity resonance.

trap apparatus lab by a 10 m single-mode-polarisation-maintaining fibre⁴⁶. The fibre introduces phase noise due to mechanical and thermal perturbations along the 10 m length. To remove this introduced noise, passive stabilisation in the form of thick foam tubing along the fibre length, as well as active stabilisation by a fibre-noise-cancellation (FNC) technique [68], is used⁴⁷. The efficacy of the FNC system is characterised in section 4.1.5 by measuring the spin coherence time of the ion.

3.6.5 Single-ion Addressing System

Selectively driving qubits and qumodes in a multi-ion register is a requirement for scalable QC. Our apparatus uses an optical single-ion addressing system to spatially select one or two ions in the chain. Spatial selectivity using narrow diameter beams is further advantageous as the focussed beam has greatly increased optical intensity compared to a globally addressing field, and so drives much stronger interactions,

⁴⁶*SK ...*

⁴⁷FNC control consists of *Sinara Stabilizer* and *Pounder* mezzanine board with PID software developed by Ayush Agrawal.




Figure 3.12: Schematic of the constructed compact single-ion addressing apparatus using crossed acousto-optic deflectors. The red beam represents the 729-nm field, with Gaussian waist changes (not to scale) indicating where in the optical train the beam is expanded or focussed. An approximately $1.5 \mu\text{m}$ waist is formed at the ion-position. This schematic is reproduced with permission from I. R. Øvergaard transfer of status [69].

and makes available much faster gates. A unique design feature of our system is the ability to produce an array of single ion addressing standing waves with individually controllable spatial phase, thus enabling any-to-any connectivity for fast entangling gates based on the scheme demonstrated in reference [41].

The design and construction of the single addressing system was led by Iver Romvig Øvergaard and, as of writing, has been in constant use in the experiment for around 1 year. Here we shall quote the design specification and relevant details of the single addressing system. In § 4.1.6 we present the characterisation data upon installation. Further technical information can be found in I. R. Øvergaard’s Transfer report [69], available upon request.

Multiple methods for single ion addressing have been developed in the ion trap community, including waveguide arrays [20, 70], micro-electronic-mechanical mirror arrays [71, 72], spatial-light-modulators [73], and acousto-optic deflectors [74, 75]. Each method has its own advantages and disadvantages in terms of switching speed, resolution, ease of hardware integration, and scalability, and so the choice of hardware is a matter of experimental preference, and constraints from the remaining

ion-trap apparatus. We require:

- Addressing of up to two ions in a ≈ 10 ion chain with low crosstalk to neighbouring ions.
- Fast reconfigureability, $\approx 1 \mu\text{s}$, of the beam position within a single experimental shot (i.e. between different gates within a single circuit execution).
- Good optical power efficiency to ensure high intensity electric fields at the ion.
- Flexibility in beam positioning to allow for fine tuning of the overlap of two counterpropagating beams used to create the standing wave.
- Individual phase control of each addressed beam to stabilise the standing wave phase at the ion plane.
- Beam position stability to avoid power fluctuations.
- A compact design to fit within the magnetic shielded region of the experiment.
- Good telecentricity (fixed beam k -vector across the scanned area) to avoid unwanted travelling wave component in a standing wave gate scheme.
- Zero frequency shift of the spatially tunable beam at the ion plane.

Crossed Acousto-Optic Deflectors (cAODs) are a popular choice for single ion addressing, offering fast rise times ($< 1 \mu\text{s}$), and good resolution over a large scanning area (≈ 500 resolvable spots with 2° angular range) with zero frequency shift along the addressing axis. AODs operate through the same mechanism as AOMs, but are optimised for beam deflection rather than frequency modulation. As an addressing system cAODs have poor scaling of power efficiency with the number of addressed spots, and so are not ideal for large scale systems. As we are interested in addressing ≤ 2 ions at a time, this is not a significant issue for our system.

Figure 3.12 shows a schematic of the constructed addressing system. Due to space constraints within the magnetic shielded region, the optical elements are




Figure 3.13: Spots generated by cAOD with wedge angle α , when supplying two frequencies to each AOD. The black points have no over-all frequency shift, and are aligned to the ion-chain. This schematic is reproduced with permission from I. R. Øvergaard transfer of status [69].

located on two height levels with a periscope steering the beam between them. The beam is delivered via single-mode polarisation maintaining fibre to the upper level where it is steered through the first AOD, with beam angle optimised for maximum diffraction efficiency into the -1^{st} order. Two lenses⁴⁸ in a $4f$ configuration re-image the deflection origin of the first AOD⁴⁹ onto the second AOD with the second AOD angled for maximum diffraction efficiency into the $+1^{st}$ order. The cAOD arrangement is termed “crossed” as the second AOD is rotated by some angle α with respect to the first, so that the deflection from the second AOD has some orthogonal component to that of the first. When the two AODs are driven with RF tones of the same frequency, f , the high efficiency diffracted order combinations are: $(0,0)$, $(-f,0)$, $(0,+f)$, $(-f,+f)$. The last combination of $(-f,+f)$ has zero overall frequency shift yet retains some frequency dependent deflection angle, which is suitable for addressing ions in a 1D line. When two frequencies, f_1 and f_2 , are supplied to the cAODs, two spots are formed at the ion plane. This arrangement is required for addressed two-qubit entangling gates. The spot configuration when two frequencies are present is shown schematically in Figure 3.13. In this case, there are two spots close to the ion chain with frequencies of $f_x = \pm(f_1 - f_2)$ compared to the zero-frequency shifted spots aligned to the ion-plane. These nearby spots can lead to unintended crosstalk errors, particularly when addressing neighbouring ions. The choice of wedge angle α then acts as a trade-off between total scan range of the

⁴⁸**P/N**

⁴⁹**P/N**

steerable beam, and the crosstalk due to the off-ion-axis beams, with a decreasing α effectively further displacing the off-axis beams from the ion chain, as shown schematically in Figure 3.13.

To translate this deflection angle into some physical displacement at the ion plane, a telecentric scanning lens⁵⁰ is used. The telecentricity ensures that the beam k -vector is fixed across the scanned area, which is important for standing wave gate schemes to avoid unwanted travelling wave components. At the intermediate focus after the scanning lens, the unwanted orders are blocked by an adjustable slit aperture and the beam polarisation is cleaned and selected using a polarisation filter and half-wave plate. A plate beam splitter⁵¹ is used within the persicope to pick off $\approx 5\%$ of the light for monitoring on a photodiode⁵². This photodiode is used for an intensity-stabilising feedback loop as discussed in § 3.6.2. The scan lens, in combination with a relay lens, forms a beam expander to increase the beam NA for the final objective lens to focus to a beam waist radius of $\approx 1.5 \mu\text{m}$ at the ion plane.

Characterisation of the system (details given in § 4.1.6) determine that experimentally the addressing beam is focused to a Gaussian waist of $1.50(1) \mu\text{m}$, measured through Rabi-oscillations while scanning the beam position across a single trapped ion. The same beam was then tested in a five-ion chain, where the average nearest-neighbour intensity crosstalk was measured to be $4.16(2) \times 10^{-4}$, with crosstalk defined as the relative intensity at the probe position compared to the addressed position.

3.6.6 Global Addressing System

The single addressing provides the utility of spatially selecting a few ions to drive with the 729-nm light, however, the power efficiency scales as $O(n^2)$, with n being the number of addressed spots. As such, we limit the single addressing to simultaneously illuminate only one or two ions at a time. It is therefore convenient to also have a

⁵⁰**P/N**

⁵¹**P/N**

⁵²**P/N**

global beam that has dimensions larger than the ion chain to address all the ions in parallel. This global branch is currently counterpropagating to the single addressing path, and shares the same optical path as the imaging system (see Figure 3.7). To increase the beam intensity at the ions, the global branch uses an anamorphic prism pair⁵³ to elliptically shape the beam, with a 4:1 aspect ratio, with major axis along the ion chain direction. The beam is focussed onto the ion plane using an $f = xx$ lens⁵⁴, resulting in a major axis beam waist radius of $52(3) \mu\text{m}$ ⁵⁵. As with the other beams directed to the ion, we use a final piezo controlled mirror⁵⁶ for fine alignment and control when the MuMetal enclosure is closed.

3.7 Control System

Could tie up loose end here of control FPGA blocks.

3.8 Ion Loading

Atoms are provided to the trapping region by heating Calcium flakes housed in a resistive tube oven with the oven mount shown in Figure 3.5. A thermocouple is tag welded onto the tube oven to provide a temperature reading for current feedback for fast stabilisation to the target temperature of $\simeq 200^\circ\text{C}$. A steady state of $\simeq 2 \text{ A}$ across the 0.5Ω load of the oven tube is required to maintain 200°C . The tube oven has a pinhole that is aligned to the ion flux aperture, visible in Figure 3.5 with atom flux shown in green, and the atomic shield pinhole shown in Figure 3.2. These apertures were included in the design to improve the operational lifetime of the trap by ensuring neutral Calcium metal cannot accumulate on the trap surface and cause shorts between RF, ground, and DC electrodes. The aperture design *luckily* failed. It failed as we observe atomic flux at all electrode locations suggesting the aperture is too large. It was *lucky* as the planned trapping region of *dc10* became

⁵³ThorLabs PS883-B

⁵⁴**P/N**

⁵⁵The beam profile was measured by extracting per-ion Rabi frequencies while driving on resonance quadrupole transitions (see § 4.1.3) and assuming a Gaussian cross section.

⁵⁶ThorLabs POLARIS-K2S2P

non-viable after electrode *Fr dc11* developed a continuity fault. As such, ions are loaded at *dc5*.

The neutral atom flux is ionised using a two-step photoionisation scheme [64] that can selectively load the desired isotope of Calcium. Natural abundance Calcium, with typical isotope proportions of 96.9% : ^{40}Ca , 2.09% : ^{44}Ca , 0.647% : ^{43}Ca [64], was appropriate as we aim to exclusively load $^{40}\text{Ca}^+$. The first stage involves a frequency stabilised 423-nm beam that couples the $4s^2 \rightarrow 4s4p$ transition in neutral calcium. The natural linewidth of this transition ($\simeq 35$ MHz) is much smaller than the isotope shifts between abundant isotopes of Calcium (typically > 100 MHz shifts), as such, through frequency selectivity, only the desired isotope can be excited. The second ionisation stage couples $4s4p \rightarrow \text{continuum}$ through a 373-nm UV laser source (although the exact wavelength does not matter as long as < 389.8 -nm). We practically see no loading of other Calcium isotopes, and only see rare non-fluorescent ions (once every few weeks) after collisions with residual gas (we suspect this is CaH^+).

The 373-nm laser causes charge build-up on the trap electrodes, which can affect the trapping potential, and so a physical shutter, controlled by a TTL signal, is used to block the beam when not actively loading ions.

3.9 Outlook

Does this section need a bit of a closing word?

3.9.1 Limitations

4

Calibration and Characterisation

Contents

3.1	Specification	28
3.2	The Ion Trap	29
3.2.1	Trap RF Chain	32
3.2.2	Trap DC Voltages	37
3.3	Magnetic Field	39
3.4	Vacuum System and Beam Geometry	39
3.5	Imaging System	42
3.6	Ca⁺ Laser Systems	43
3.6.1	Frequency Stabilisation	44
3.6.2	Intensity Stabilisation	46
3.6.3	Diode Lasers	47
3.6.4	Narrow Linewidth 729 Laser	50
3.6.5	Single-ion Addressing System	53
3.6.6	Global Addressing System	57
3.7	Control System	58
3.8	Ion Loading	58
3.9	Outlook	59
3.9.1	Limitations	59

Before we can proceed with our planned experiments involving the motion and spin of the ions, the apparatus must be characterised. This characterisation serves two purposes: it allows us to benchmark the performance of our system, and it reveals any current limitations that must be addressed before we can successfully carry out the experiments.




Figure 4.1: State preparation error as a function of the number of state preparation pulses, N . The error is defined as population not in the $|\downarrow\rangle$ state after the preparation sequence has been applied. The error is reduced to below 10^{-3} after $N = 7$ pulses, which is used in our experiments.

We split this chapter into interactions involving only the spin, then those focussed only on the motion, and finally we describe spin-motion interactions. This chapter demonstrates the full control required for discrete-variable quantum computation. Section 6.1 in Chapter 6 applies the spin-dependent motional interactions demonstrated here for continuous-variable readout, and to realise the remaining nonlinear and two-mode operations required for full hybrid CV-DV control.

4.1 Spin

Discrete variable quantum computing consists of control over many two-level systems, which we refer to either as spins, or qubits. In this section the experimental methods used to prepare, manipulate, and measure these spins are described, and the performance of these operations are benchmarked.

Parameter	Value
729-nm laser power	1000 μW
729-nm pulse duration	5 μs
854-nm laser power	60 μW , $250 \times I_{\text{SAT}}$
854-nm pulse duration	40 μs
866-nm laser power	15 μW , $\sim 70 \times I_{\text{SAT}}$
866-nm pulse duration	30 μs

Table 4.1: Experimental parameters used for state preparation of the ion into the $|\downarrow\rangle = |4S_{1/2}, m_j = -1/2\rangle$ state. All lasers frequencies are set to be resonant with the corresponding transition frequencies.

4.1.1 State Preparation

To utilise two levels of the ion as a qubit, the electronic state must first be prepared into the qubit manifold. Here, the ion is prepared into the $|\downarrow\rangle = |4S_{1/2}, m_j = -1/2\rangle$ state. This is done via optical pumping by pulsing the quadrupole transition $|4S_{1/2}, m_j = +1/2\rangle \leftrightarrow |3D_{5/2}, m_j = -3/2\rangle$, followed by deshelling pulses using the 854-nm, 866-nm lasers on resonance. The 729-nm laser pulse duration and intensity is set to invert the populations, i.e. perform a π -pulse. Details for laser parameters are summarised in Table 4.1. These three pulses are repeated N times sequentially. Figure 4.1 shows the state preparation error with increasing N . The error is found by applying a π -pulse on the state preparation transition directly after the state preparation sequence and measuring the population in the $|\downarrow\rangle$ state, $P_\downarrow$. Any population not in $P_\downarrow$ indicates either residual population in the $|4S_{1/2}, m_j = +1/2\rangle$ state, or trapped in the metastable D levels. After $N = 7$ pulses there is a state preparation error of $< 10^{-3}$. The state preparation error is likely limited by accidental excitation of the $|4S_{1/2}, m_j = -1/2\rangle$ state to the $D_{5/2}$ level. This may be due to off-resonant excitation of another transition, or due to spectral purity of the 729-nm laser.




Figure 4.2: Readout histogram with 100,000 individual measurements for bright (red) and dark (blue) populations. The camera counts are presented as a histogram and fitted to Poissonian distributions to calibrate a threshold for discerning if the ion is bright or dark with the lowest probability for false measurement. Dark counts are mainly due to 397-nm light scattered from the nearest trap electrodes. With 100 μ s readout duration, and laser parameters set to Doppler cooling parameters, a threshold of 137 photons is found to be sufficient to discern the two populations.

Parameter	Value
397-nm laser power	3 μ W, $1 \times I_{\text{SAT}}$
397-nm laser detuning	-17.5 MHz
866-nm laser power	15 μ W, $\sim 70 \times I_{\text{SAT}}$
866-nm laser detuning	0 MHz
Doppler cooling duration	1 ms
Readout duration	100 μ s
Camera threshold	137 counts

Table 4.2: Experimental parameters used for Doppler cooling and state-selective fluorescence of the ion.

4.1.2 Measurement

To measure the qubit state, the 397-nm and 866-nm lasers are applied, and the scattered 397-nm photons are imaged by the objective onto a camera (see § 3.5). 397-nm light will only be scattered if the qubit is in $|\downarrow\rangle$. A typical image generated by the camera is seen in Figure 3.7. A small window of pixels is set around each ion (typically a 3×3 px region), and the photon counts in this region are summed. Use

of a high numerical aperture ($\text{NA} = 0.6$) objective greatly increases the collection efficiency. This efficiency allows the 397-nm beam to be red-detuned to Doppler cooling parameters (see § 4.2.2), while still maintaining low readout duration (100 μs) and low readout error. Cooling during readout increases the experiment’s repetition rate by reducing the amount of dedicated cooling time required.

The parameters used for readout are summarised in Table 4.2, and a typical histogram of readout counts for one ion can be seen in Figure 4.2. The readout threshold is calibrated by taking bright counts with the 397-nm and 866-nm lasers on, and dark counts with the 397-nm laser on and the 866-nm laser off. This assumes dark counts are due to 397-nm scattering from the trap electrodes. Fitting Poissonian distributions to the bright and dark measurements, an optimal count threshold is determined, shown by black dashed line. Assuming Poissonian statistics this threshold would give an expected statistical readout error of $\epsilon_{\downarrow} = \epsilon_{\uparrow} = 6 \times 10^{-9}$, where $\epsilon_{\downarrow, \uparrow}$ are the expected error associated with measuring bright or dark. This measurement ignores errors due to the finite lifetime of the metastable $3D_{5/2}$ level, which dominate $\epsilon_{\uparrow}$. The expected error due to the finite lifetime, τ , is $\epsilon_{\text{decay}} = 1 - e^{-t/\tau}$, where t is the readout duration. For $^{40}\text{Ca}^+$, $\tau \sim 1.1$ s, $\epsilon_{\uparrow} \approx \epsilon_{\text{decay}} = 10^{-4}$. Assuming the metastable lifetime is limited only by spontaneous decay overlooks leakage from the 854-nm shelving beam, which could significantly shorten τ [70]. Further characterisation of the combined state-preparation and measurement fidelity is provided in § 4.1.4.

4.1.3 Single-Qubit Rotations

Single-qubit rotations driven by optical fields, as described in § ??, along with two-qubit gates are required for universal discrete variable quantum computation.

We divide errors affecting single-qubit gates into two categories: unitary, and non-unitary. The unitary error component is labelled “systematic” if there exists simple (and accessible) control parameters to suppress such errors. We outline in the next section robust phase estimation (RPE), the calibration routine used to

reduce systematic errors. In contrast, reducing non-unitary errors require improved isolation between the qubit and environment, or suppressed noise on the single-qubit-driving field. At the end of this section we characterise the performance of the single-qubit gates using randomised benchmarking.

Robust Phase Estimation

Many gate calibrations can be formulated as phase estimation tasks. A single-qubit rotation may be generally represented by $R_{\vec{n}}(\theta)$, where $\vec{n}$ is some given axis around which the state evolves with rotation angle $\theta = \omega t$, where the angle is accumulated over time by rate ω . For a gate we specify this axis and total angle, with deviations to the ideal setting being a systematic error.

Suppose we can set axis $\vec{n}$, and the interaction duration t arbitrarily well, the task of calibration is then limited to how precisely we can measure ω . We further assume some statistical uncertainty for a practical phase measurement, σ_θ , that is independent of the magnitude of θ . The uncertainty of angular frequency $\sigma_\omega = \frac{\sigma_\theta}{t}$ is thus inversely proportional to duration t — longer probe durations provide more information. This improved signal to noise is understood by the coherent accumulation of phase over time due to an underlying ω .

However, there are two issues with going to large t : first, incoherent errors also accumulate, attenuating the measured signal, and violating our assumption of a fixed statistical uncertainty σ_ω . Second, quantifying θ for a fixed duration t is ambiguous — the measurement only provides the wrapped angle, θ_w , for example in the interval $(-\pi, \pi]$. The first limitation restricts what practical maximum t we can measure, heuristically often set to the $\frac{1}{e}$ -coherence time, τ . The second limitation may be mitigated through supplemental measurements. The RPE protocol provides an algorithm for determining which branch the measured θ is on through an optimal number of supplemental measurements [76].

The protocol is as follows:

1. Select some duration t_0 such that $\theta_0 = \omega t_0 \in (-\pi, \pi]$, and determine parameter $k_{\max} = \left\lfloor \log_2 \left(\frac{\tau}{t_0} \right) \right\rfloor$, to set the maximum probe duration, $t_{\max} = 2^k_{\max} t_0$, below the coherence time, τ .
2. For $k \in \{0, 1, \dots, k_{\max}\}$, measure total phase $\theta_{k,\text{tot}}$ wrapped to $(-\pi, \pi]$.
3. Infer candidate base-duration phases for each k as $\theta_{k,0}^j = \frac{\theta_{k,\text{tot}} + 2\pi j}{2^k}$ for $j \in \{0, 1, \dots, 2^k - 1\}$.
4. The branch for $\theta_{k,0}$ is then recursively selected to be the closest $\theta_{k,0}^{(j)}$ to the determined $\theta_{k-1,0}$.
5. Angular frequency, ω , is inferred from $\omega = \frac{\theta_{\max,0}}{t_0}$ with uncertainty $\sigma_\omega = \frac{\sigma_0}{2^k t_0}$ where σ_0 is the phase uncertainty from a measurement at t_0 .

We use this protocol for both frequency (ω) and gate-duration, t , calibrations, which we outline below.

Transition Frequency Calibration

Rotations around the $\hat{\sigma}_z$ axis, correspond to a relative detuning, δ , between the qubit precession and a the reference frequency of a near resonant drive field, given by the idle Hamiltonian, $\hat{H}_z = -\hbar \frac{\delta}{2} \hat{\sigma}_z$. The phase accumulated due to the application of $\hat{H}_z$ for a fixed time, T , is $\theta = \delta T$. This phase may be measured through RPE using Ramsey spectroscopy. The sequence starts with an initial $R_x(\frac{\pi}{2})$ pulse, then a delay with duration $2^k t_0$, followed by a final $R_\phi(\frac{\pi}{2})$ pulse with variable rotation axis $\vec{n} = \cos(\phi) \hat{\sigma}_x + \sin(\phi) \hat{\sigma}_y$. Varying the final axis angle, ϕ , for a given probe duration reveals the familiar Ramsey fringes, with phase offset given by θ . For sample-efficient data collection, the phases $\phi \in [0, \frac{\pi}{2}, \pi, \frac{3\pi}{2}]$ are chosen, which provide enough information to determine the phase angle in the xy -plane, amplitude, and any asymmetric qubit readout offset of the measured fringes.

By this method, we calibrate the reference laser-frequency with respect to the bare-qubit frequency. The Stark shifts due to the two $\frac{\pi}{2}$ -pulses will give some systematic offset between the bare-qubit frequency and the calibrated laser frequency, however




Figure 4.3: Rabi oscillations due to on-resonant 729-nm beam. Estimated probability $P_{\downarrow}$ at each time step is an average of 200 individual experimental shots. The solid line is a fit using decaying oscillation model given by equation 4.1.

are negligible for long Ramsey probe durations.

Resonant $R_x(\theta)$ gates are then driven by setting the relative detuning to $\delta = 0$. This works well for gates using low optical intensity, however, if high intensity is required (e.g. for motional-sideband pulses), a further calibration of the beam-induced Stark-shift is required. In the case that the shift is dominated by far-off-resonant dipole transitions, the described Ramsey experiment may be simply modified to include the optical dressing field to be present during the delay period. The dressing field is detuned sufficiently to not drive resonant interactions, but still provide near-identical Stark-shifts on the qubit transition. Using a Hahn-echo sequence with the dressing field present only in one of the arms allows for extracting the dressing field induced Stark shift while being insensitive to any systematic frequency offsets between the bare qubit and the calibrated resonant $\frac{\pi}{2}$ -pulses.

Gate Time Calibration

With the resonant frequency well calibrated, the pulse durations for various gates must be set. Here we outline the calibration routine for a $R_x(\frac{\pi}{2})$ -gate, although this may be adapted for gate angles correspond to integer fractions of 2π .

We use a modified version of the RPE protocol to determine optimal gate durations. This protocol converges with a sufficiently close initial value for pulse duration $T_{\pi/2}$, and so a simple calibration using resonant Rabi oscillations is first described.

The qubit is prepared in the computational eigenstate $|\downarrow\rangle$, followed by an on-resonant 729-nm carrier pulse inducing $R_x(\Omega t)$ rotations. The final state population is measured as a function of pulse duration to observe familiar Rabi oscillations. Figure 4.3 displays these oscillations when using a global-addressing beam with a total power of 1 mW. Here each point corresponds to the average of 200 samples. The Rabi frequency, Ω , and decay rate, λ , are determined through fitting the recovered populations to the model,

$$P_{\downarrow} = \frac{1 + e^{-\lambda t} \cos(\Omega t)}{2}. \quad (4.1)$$

In this case, the Rabi frequency is found to be $\Omega = 2\pi \times 0.1433(1)$ MHz, and the decay rate $\lambda = 0.0104(6) \mu\text{s}^{-1}$. From this a first-guess pulse duration of $T_{\pi/2} \approx 1.75 \mu\text{s}$ can be estimated. This is a sufficient initial value to calibrate further, however it does not account for any amplitude ramps or dead-time of the optical-pulse compared to the digital control signal.

RPE provides a higher-precision calibration, that may account for non-square-pulse profiles. The protocol is simply adapted from the RPE method described above by setting the initial base probe as a single $R_x(\frac{\pi}{2})$ -pulse, with duration $t_0 = T_{\pi/2}$. The pulse is repeated 2^k times up to k_{max} (inferred from characteristic decay time $\tau = \frac{1}{\lambda}$) applications of the $R_x(\frac{\pi}{2})$ -pulse. The phase angle in the yz -plane is determined through applying an additional $R_x(\theta_f)$ rotation with angle $\theta_f \in [0, \frac{\pi}{2}, \pi, \frac{3\pi}{2}]$, before measurement. The final base phase accumulated, θ_0 , is determined after applying the phase unwrapping protocol given above. Any deviation between this measured phase and the target, $\frac{\pi}{2}$, may be subsequently calibrated by rescaling the original pulse in time by factor $s = \frac{\pi}{2\theta_0}$.

This method, along with the transition frequency calibrations given above, provide an efficient method to calibrate systematic offsets of single qubit rotation gates driven through on-resonant carrier excitation. Next we discuss randomised benchmarking, a characterisation technique that provides a measure of remaining incoherent and systematic errors for the calibrated gates.




Figure 4.4: Standard randomised benchmarking applied to a single qubit. Panel a) shows the required circuit for a random sequence of Clifford gates, C_i , of total length $m-1$, with a final inverting Clifford, C_{inv} , and a random Pauli gate, P , before measurement in the computational basis. Panel b) shows the survival probability, $\mathcal{P}_S$, versus Clifford sequence length, m . Each point is found by averaging over $r = 50$ and $s = 50$ repeats for each m . The solid line is a fit using the decay model given by equation 4.4, the dashed line is the expected asymptote for a fully depolarised single qubit.

4.1.4 Single-Qubit Randomised Benchmarking

A further characterisation protocol is required to quantify the performance of single-qubit gates under both systematic and incoherent noise. A scalable class of methods, prevalent in quantum computation, is that of randomised benchmarking (RB). We will give a short description of RB here, but further details of the goals of characterisation, and complementary techniques to RB are discussed in Chapter 5 § 5.1.

The purpose of RB is to simplify a general completely positive, trace preserving (CPTP) error channel by enforcing the symmetry of a chosen group on the channel [77]. The general error channel is described by the map acting on the space of density matrices,

$$\mathcal{E}(\hat{\rho}) : \mathbb{C}^{d \times d} \rightarrow \mathbb{C}^{d \times d} \quad (4.2)$$

where d is the Hilbert space dimension. The symmetry is enforced on the channel through twirling — averaging a channel $\mathcal{E}$ over a group $\mathcal{G}$. A treatment of the action of twirling on a CPTP channel is given in § A.3.1 in the context of *k-copy*

RB, an extension to standard RB methods for a single qubit [78].

The resulting simplified channel is then quantified to obtain metrics of the performance of the executed circuit. This protocol thus can reduce the problem to measuring $O(1)$ parameters, rather than the $O(d^4)$ parameters required for a general quantum channel acting on a Hilbert space with dimension d . Furthermore, RB is particularly robust to state-preparation and measurement errors (SPAM), allowing for the isolation of errors originating from imperfect SPAM, to errors originating from the qubits idling or from the execution of gates.

Practically RB protocols are simple to implement, approximating the twirling procedure to the generation and execution of an ensemble of random gate sequences of varying lengths, m . We focus on the “standard” RB sequence, enforcing the single-qubit Clifford group symmetry onto a single-qubit error channel [78]. The aim of standard RB is to twirl over the single-qubit unitary group, $SU(2)$, which takes the Haar measure, and transforms single-qubit error channels to highly-symmetric depolarising type channels, fully expressed by a single parameter, λ ,

$$\mathcal{E}_{dp}(\hat{\rho}) = (1 - 2\lambda)\hat{\rho} + \lambda I_2 \text{Tr}(\hat{\rho}), \quad (4.3)$$

where $\mathcal{E}_{dp}$ is the resulting twirled channel, and I_2 is the 2×2 identity matrix. Unfortunately, $SU(2)$ is continuous, as such it is impractical to average over the entire group. However, there are certain subgroups of $SU(2)$ that can be averaged over to recover an approximation of the Haar measure over $SU(2)$, up to the moment t . These subgroups are known as unitary t -designs of $SU(2)$. The single-qubit Clifford group, Cl_1 , is a known 2-design, and consists of only 24 elements [79].

The standard RB implementation for one sequence is presented in Figure 4.4 (a), and the full sequence is as follows:

1. The qubit is initialised to $|\downarrow\rangle$.
2. A sequence of $m - 1$ single qubit Clifford gates, C_i , are applied, with each gate randomly sampled from the Clifford gate-set.

3. A final “inverting” Clifford gate, C_{inv} , is applied which is chosen such that the full Clifford gate sequence is equivalent to the identity operation.
4. A random Pauli gate, $P \in \{I, X, Y, Z\}$, is applied with equal probabilities¹.
5. The population is measured in the computational basis, with shots resulting in the expected outcome (accounting for the random Pauli) labelled as survivals, S , and opposite outcomes labelled as failures, F .
6. Steps 1-5 are repeated s times to find the average survival population, $\mathcal{P}_S$.
7. Steps 1-6 are repeated r times with different random sequences for sequence with length m .
8. Steps 1-7 are repeated for a chosen set of sequence lengths, m , here log-sampling m up to some m_{max} .
9. The average survival populations versus sequence length, m , are fit to a decay model to extract a depolarising error per Clifford gate, ε_c , and a SPAM error ε_{sp} .

The decay model used to find the fidelity versus number of Clifford gates is given by,

$$\mathcal{P}_S(m) = \frac{(1 - 2\varepsilon_{\text{sp}})(1 - 2\varepsilon_c)^m + 1}{2}, \quad (4.4)$$

where the first term in brackets is the probability of no error during SPAM, and the second, exponential term provides the probability of not depolarising during a sequence of length m .

This scheme is applied experimentally to characterise the performance of the calibrated single-qubit gates. The 24 Clifford gates are decomposed into sequences of $R_x(\pm\frac{\pi}{2})$ and $R_x(\pm\frac{\pi}{2})$, with an average of 2.17 $\frac{\pi}{2}$ -pulses per Clifford gate. Figure 4.4 (b) shows the average survival populations for a log-sampled range

¹This acts to randomise the final outcome between computational states $|\downarrow\rangle$ and $|\uparrow\rangle$, enforcing a further symmetry on any state-preparation and measurement error component.




Figure 4.5: Calibrating the relative detuning between the 729-nm laser and the qubit during the application of $\frac{\pi}{2}$ -pulses. Each point corresponds to the extracted error-per-Clifford from standard RB. The solid line is the fit of a quadratic model, the dashed line is the extracted center frequency.

of sequence lengths up to $m_{\max} = 2048$ Clifford gates. Each point corresponds to an average over $r = 50$ randomisations, and $s = 50$ sample repeats per randomisation. The solid line follows equation 4.4, with fit parameters, $\varepsilon_c = 1.9(1) \times 10^{-3}$, and $\varepsilon_{\text{sp}} = 3(1) \times 10^{-3}$. The dashed line is the expected asymptote for a fully depolarised qubit. Error bars show the standard error of the average survivals for each random sequence for a given length m . From ε_c , we find a single $R_x(\frac{\pi}{2})$ average fidelity of $\mathcal{F}_{\pi/2} = 1 - \frac{\varepsilon_c}{2.17} = 99.913(2)\%$. The measured SPAM is consistent with the independently characterised state preparation (§ 4.1.1) and measurement (§ 4.1.2) errors of $\sim 10^{-3}$, and $\sim 10^{-4}$ respectively.

We attribute the single-qubit rotation error to the finite-temperature of the radial modes. Expanding the light-matter interaction Hamiltonian given in equation 2.41 to second order in the Lamb-Dicke parameter, the carrier Rabi frequency is modulated depending on motional Fock state, n , as [36],

$$\Omega_n \approx \Omega \left(1 - \eta^2 \left(n + \frac{1}{2} \right) \right), \quad (4.5)$$

where η is the Lamb-Dicke parameter. For a thermal distribution each shot will result in a different rotation angle leading to a thermal dephasing error. For small rotation errors, the fidelity averaged over the Haar measure (see Appendix A.2) is

given to first order by $\mathcal{F} = 1 - \frac{1}{3}\langle\theta^2\rangle$. If targetting an $R_x(\frac{\pi}{2})$ -gate, the expected fidelity due to thermal errors is then,

$$\mathcal{F}_{\pi/2}(\bar{n}) = 1 - \frac{1}{12}\pi^2\eta^4(2\bar{n}^2 + \bar{n}). \quad (4.6)$$

The radial modes for our target trapping frequencies (≈ 2.5 MHz) have Lamb-Dicke parameters of $\eta \approx 0.05$. For these experiments the radial modes are only Doppler cooled. In the following section on cooling, § 4.2.2, we will determine the Doppler limit for these modes to be $\bar{n}_d \approx 4$. Using these two estimated parameters, we expect a fidelity due to thermal errors of $\mathcal{F}_{\pi/2,\text{th}} = 99.96\%$. This is similar in magnitude to the measured value and makes the assumption cooling is optimised to reach the Doppler limit. The single-qubit gate fidelity measured here is sufficient for the projects in this thesis, however if a lower error is required, we will be required to sub-Doppler cool all radial modes.

RB can be used to calibrate systematic errors, however requiring many more samples than the RPE protocol described above. Figure 4.5 presents the detuning of the 729-nm laser with respect to the resonant frequency extracted by RPE. The points correspond to the extracted ε_c from RB sequences with the same parameters as used for Figure 4.4 (b). The solid line is a quadratic fit, with center frequency $\delta = 2\pi \times 0.76(8)$ kHz. The RPE protocol does not account for the Stark-shift due to the 729-nm beam, as such we see a frequency shift between the qubit and laser when probing by RB.

We use standard RB as a daily characterisation experiment to verify consistent performance of the single-qubit gates. In Chapter 5 we show further utility of single-qubit Clifford RB using a novel extension to parallel single-qubit Clifford application in a multi-qubit register.

4.1.5 Spin Coherence Times

Qubit decoherence may originate from uncontrolled fluctuations of the qubit transition frequency or from phase noise on the laser driving the qubit transition.




Figure 4.6: Coherence times of three quadrupole transitions with varying magnetic field sensitivities. Points correspond to the extracted contrasts of Ramsey fringes. Solid lines are exponential fits to determine the $\frac{1}{e}$ coherence time, τ . The inset plot shows the measured coherence time as a function of the magnetic field sensitivity, $\Delta\mu/h$, of the transition being probed.

Decoherence directly contributes to loss of fidelity of qubit operations, and will limit the circuit-depths possible on our apparatus.

Both electric and magnetic field noise can lead to qubit transition frequency fluctuations. The quadrupole transitions between the ground state $|4S_{1/2}, m_j = -1/2\rangle$ and the metastable manifold $|3D_{5/2}\rangle$ in $^{40}\text{Ca}^+$ are first order sensitive to any external magnetic fields. As such, even with the MuMetal magnetic shield in place, external magnetic fields are expected to be the dominant noise source. To verify this, and benchmark the current decoherence time, the Ramsey contrast is measured as a function of delay duration. The Ramsey sequence is as follows:

1. Prepare the qubit in $|\downarrow\rangle$.
2. A $R_y(\pi/2)$ -pulse is applied. preparing the state $|+\rangle = \frac{1}{\sqrt{2}}(|\downarrow\rangle + |\uparrow\rangle)$.
3. The system is left to evolve for Ramsey delay, t_{RAM} .
4. A variable $R_\phi(\pi/2)$ -pulse is applied.
5. The population of $|\downarrow\rangle$ is measured.
6. Repeat steps 1-5 for different values of ϕ to find the contrast of the Ramsey fringe.

7. Repeat steps 1-6 for different values of t_{RAM} to find the decay in contrast with time.

This Ramsey fringe contrast is directly proportional to the coherence of the resulting noisy state, $\hat{\rho}$, given by matrix element, $\rho_{\downarrow,\uparrow}$.

Figure 4.6 shows the coherence time of three quadrupole transitions:

$|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = \{-5/2, -3/2, +3/2\}\rangle$. These transitions have magnetic field sensitivities of $\Delta\mu/h = \{2.799, 1.119, 3.922\}$ MHz G⁻¹ respectively (calculated using equation 2.1). The contrast decays for all three qubits are shown in the figure. The $\frac{1}{e}$ -coherence time is extracted by fitting exponential decays to the contrasts. Generally, exponential decays indicate the timescale of the noise is fast compared to the Ramsey probe durations, whilst Gaussian decays indicates comparatively slow drifts. Neither exponential nor Gaussian well describe the measured decays here due to the strong periodic structure above the exponential decay. This suggests a strong narrow-band coherent component to the noise. At the time scale of ≈ 20 ms we expect to see the effects of UK mains hum, which likely cause this structure. Regardless, fitting the exponential envelope suggests a coherence time of 17(4) ms for the qubit transition, $|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = -5/2\rangle$. The inset plot shows the extracted coherence time as a function of transition magnetic field sensitivity. There is a clear dependency between the coherence time and the magnetic field sensitivity, suggesting that magnetic noise is dominant.

4.1.6 Single-ion Addressing

The single-ion addressing system, introduced in § 3.6.5, is calibrated and characterised using the ion as a sensor. We verify that the beam waist at the ion plane approximately matches the designed specification, and we measure the crosstalk due to an addressed neighbouring ion. Crosstalk can cause structured coherent errors acting on the qubit register, which can have particularly adverse effects to error correction and mitigation protocols that assume independent, uncorrelated noise on each qubit. As such, designing an addressing system for low-crosstalk, and




Figure 4.7: Characterisation of the single-ion addressing system using the ion as a probe. a) The normalised intensity of the focussed 729-nm beam at the ion-plane. The $1/e^2$ -waist is highlighted by the contour. b) Demonstration of spatial selective addressing by scanning the beam across a 5-ion chain while applying $R_x(\pi)$ -gates calibrated for ion 3. The probability of observing $|\downarrow\rangle$, $P_\downarrow$ is shown for each ion.

verifying its performance is critical.

The intensity profile of the beam in the ion plane is measured by Rabi oscillations on a single-trapped ion while scanning the beam position in a 2D grid. To efficiently measure this 2D array, a fixed Rabi pulse duration is used, although we note that similar measurements can be made with adaptive pulse durations, or through using RPE type protocols to improve the sensitivity of these measurement. To avoid ambiguity of the accumulated phase, we ensure that at the point of maximum field amplitude the accumulated phase does not wrap, $\theta_{\max} < \pi$. The relative optical intensity may then be extracted by considering equation 4.1, assuming a negligible decay rate, λ , and recalling that $\Omega \propto \sqrt{I}$, yielding,

$$\frac{I(x, y)}{I_{\max}} = \left(\frac{\arccos(\sqrt{P_\downarrow(x, y)})}{\theta_{\max}} \right)^2. \quad (4.7)$$

The normalised intensity profile extracted by this method is shown in Figure 4.7 (a). The qubit is initialised in $|\downarrow\rangle$ using the global addressing beam, and the total

single-addressing beam power is stabilised to 20 μW . The black contour signifies the $\frac{1}{e^2}$ waist. The spot closely resembles the expected Gaussian shape, with some aberrations, especially leading to a shoulder in the $-x$ direction. Fitting a 2D Gaussian distribution, we find $\frac{1}{e^2}$ -waists of $\omega_x = 1.584(3) \mu\text{m}$, and $\omega_y = 1.386(2) \mu\text{m}$. These closely match the target waist of $1.5 \mu\text{m}$.

The crosstalk is measured by scanning a single beam across a 5-ion chain. The axial frequency of the trap is set to 0.7 MHz, to provide an ion spacing of $\approx 5 \mu\text{m}$. When addressing the central ion of the chain, ion 3, we may observe Rabi oscillations on all qubits simultaneously to extract the relative Rabi frequency crosstalk. The Rabi frequency is determined through fitting equation 4.1, finding frequencies of $2\pi \times \{1.0(1), 1.67(5), 281.7(5), 7.96(1), 5.29(1)\}$ kHz for ions 1 to 5 respectively. This corresponds to intensity crosstalks of $\{0.12(4), 0.35(2), -, 7.99(3), 3.53(1)\} \times 10^{-4}$ respectively.

To highlight this spatial selectivity, we show the beam scanning across the 5-ion chain with a fixed pulse duration. For reference, above this plot we show fluorescence of the 5-ion chain as viewed by the imaging camera. The fixed Rabi-pulse duration is calibrated to be a $R_x(\pi)$ -pulse when acting on the central ion. Figure 4.7 (b) shows the measured populations, $P_{\downarrow}$, as a function of beam position along the axial ion-chain direction. The points correspond to populations averaged over 200 samples, with the population for each ion gathered simultaneously using the spatially resolved readout, as discussed in § 4.1.2. The solid lines follow the simple model inferred from the expected Gaussian waist of the beams,

$$P_{\downarrow,i} = 1 - \sin \left(A_i \exp \left(\frac{-(x - x_i)^2}{\omega^2} \right) \right)^2, \quad (4.8)$$

where x_i is the central position for ion, i , and A_i is included to account for a drop in field amplitude due to any misalignment between the addressing axis and the ion-chain, or due to changes in diffraction efficiency of the AODs when driving them at different frequencies. Practically we may calibrate the beam power set-point to balance any Rabi-frequencies for every addressing configuration to mitigate these

small misalignments and limited bandwidths of the AODs.

The achieved waist of $\approx 1.5 \mu\text{m}$ allows for high intensity fields at the qubits, exemplified here by the low $20 \mu\text{W}$ 729-nm beam power required for driving $\approx 1 \mu\text{s}$ $R_x(\frac{\pi}{2})$ gates. The low intensity crosstalk of $< 10^{-3}$ for neighbouring ions is sufficient for the projects demonstrated in this thesis. If lower crosstalk errors are required, it is likely this may be reduced to $< 10^{-4}$ solely through further optimisation of beam alignment through the optical train. This is suggested by the current asymmetry of crosstalk errors either side of the addressed qubit. Furthermore, as the crosstalk error is coherent, it may be further suppressed using composite pulse sequences [80].

4.2 Motion

We outline methods to track the mode frequencies, cool the motion of the qubit, quantify thermal heating of the motion, and characterise the qumode coherence time. These methods are necessary for control over qumodes for information, or for using the common motion as a means to entangle two qubits.

Motion of ions trapped in the same potential can be described by a set of normal modes [33]. These modes are labelled with respect to the ion chain geometry: axial modes oscillate parallel with the chain, and radial modes oscillate perpendicular to the chain. Due to the geometry of the 729-nm beam access with respect to the ion chain (see § 3.4), only the radial modes can be addressed in our system. For the *NPL* trap (see § 3.2), the radial modes frequencies are non-degenerate. In our system, for a single ion, the mode frequencies are approximately $\omega_{\text{upper}} = 2\pi \times 2.9 \text{ MHz}$, $\omega_{\text{lower}} = 2\pi \times 2.0 \text{ MHz}$, and $\omega_{\text{axial}} = 2\pi \times 1.0 \text{ MHz}$, where “upper” and “lower” are the two radial modes. All results in the following section are using the upper radial mode.




Figure 4.8: Motional mode resonance measured by RF “tickle” experiment. Points correspond to fluorescence rate measured by the camera while the solid line is a Lorentzian model fit.

4.2.1 Calibrating Mode Frequencies

The radial mode frequencies are proportional to the RF voltage (equation 2.22), as such any amplitude drift on the RF chain translates to mode frequency drift. The efficiency of the resonator, as discussed in § 3.2.1, is dependent on ambient temperature, and is currently not actively stabilised. As such the mode frequencies are routinely (order minutes) measured to track and update the frequency set point for interactions involving the mode.

The motional mode frequency is measured via an RF “tickle” experiment. This consists of short (≈ 1 ms) RF field pulses applied to one of the trap DC electrodes which has a non-zero electric field projection to the mode we wish to probe. The frequency of the RF field is scanned around the expected motional mode frequency, and when resonant drives the motion of the ion. Excess motional amplitude is measured by observing an increase in fluorescence when driving a closed dipole transition detuned to a nonlinear region of the resonance. The motional induced sidebands at this nonlinear point are unbalanced leading to a modulation of fluorescence. Here the 397-nm laser is set to provide approximately 1 saturation intensity of light to avoid power broadening the dipole transition. The 397-nm is red-detuned 100 MHz from resonance. The 866-nm is set to typical readout parameters to repump any population trapped in the D -manifold. The modulated fluorescence is integrated over ≈ 1 ms leading to an increased average count rate

when the RF tone is resonant with the mode frequency.

Figure 4.8 shows the typical fluorescence count rate for this mode frequency calibration scan on the upper radial mode. A Lorentzian profile is fit to extract central mode frequency $\omega_{\text{upper}} = 2\pi \times 2.87034(2)$ MHz, and FWHM of 280(50) Hz. The uncertainties of these parameters are determined by the bootstrapping method using 1000 re-samplings.

4.2.2 Cooling

For any interaction that involves the motion of the ion, it is essential to both prepare and measure the motional state with high fidelity. For all motional interactions considered in this thesis, the ion is initially prepared as close as possible to the motional ground state, with a mean phonon number $\bar{n} \approx 0$. This is done by an initial stage of Doppler cooling, followed by pulsed sideband cooling on the 729-nm transition. The final motional state is characterised by measuring the population in Fock states using RSB and effective BSB pulses.

Doppler Cooling

Doppler cooling relies on light incident on a moving ion appearing frequency shifted in the ion's rest frame. For Doppler cooling of $^{40}\text{Ca}^+$, both the 397-nm and 866-nm lasers are applied. The 397-nm laser is red-detuned from the dipole $|4S_{1/2}\rangle \leftrightarrow |4P_{1/2}\rangle$ transition, leading to preferential scattering of 397-nm photons that have a net cooling effect [81]. The excited $4P_{1/2}$ level decays via two channels: emission of a 397-nm photon back to $|4S_{1/2}\rangle$, or an 866-nm photon to $|3D_{3/2}\rangle$, with a branching ratio of 14.5 [82]. Since many photon scattering events are required, the cooling cycle must be closed. To prevent population trapping in the metastable $|3D_{3/2}\rangle$ level, the 866-nm laser is applied on resonance during cooling.

Equilibrium is reached when the cooling rate balances photon recoil heating. Assuming the 397-nm transition is broadened only by its natural linewidth γ ,

Parameter	Value
729-nm laser power	3.0 mW
729-nm laser detuning	-4.0 MHz*
729-nm pulse duration $N = 1$	2.5 μ s
729-nm pulse duration $N = N_{\max}$	33 μ s
Repumping 866-nm duration	14 μ s
Deshelving 854-nm duration	5 μ s
Number of cooling cycles $N_{\max}$	50

Table 4.3: *729-nm detuning is relative to the $|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = -5/2\rangle$ transition frequency.

Experimental parameters used for sideband cooling. Other parameters for the repump and deshelve pulses are the same as for state preparation, see Table 4.1. The 729-nm laser is detuned to the RSB resonance for the upper radial mode, $\omega_{\text{upper}}/2\pi = 4.0$ MHz, and the pulse duration for $N = N_{\max}$ is chosen to be the expected RSB $R(\pi)$ -pulse duration for $n = 1$.

the minimum Doppler temperature is,

$$T_{\text{Doppler}} \approx \frac{\hbar\gamma}{2k_B}, \quad (4.9)$$

where $\hbar$ is the reduced Planck constant, and k_B is Boltzmann's constant [81].

For $^{40}\text{Ca}^+$, the natural linewidth of the 397-nm transition is $\frac{\gamma}{2\pi} = 21$ MHz[83], giving $T_{\text{Doppler}} = 0.5$ mK. Given a mode frequency, the expected mean phonon number is,

$$\bar{n} = \frac{1}{e^{\hbar\omega/k_B T} - 1}. \quad (4.10)$$

For a radial mode with frequency $\omega \approx 2\pi \times 2.5$ MHz, the Doppler limit is approximately $\bar{n} \approx 4$. The parameters used here for Doppler cooling are summarised in Table 4.2.

Sideband Cooling

To further cool the ions toward their motional ground state, resolved sideband cooling is used [84]. As discussed in § 2.4, red and blue sidebands (RSB and BSB) of the quadrupole transition are accessible, which may both flip the spin and add or remove a single phonon from the motional mode. For the $|4S_{1/2}\rangle \leftrightarrow |3D_{5/2}\rangle$ transition, at appropriate motional mode frequencies, interaction strengths, and using a




Figure 4.9: (a) Example pulse sequence for one experimental shot including pulsed sideband cooling. The *Exp.* block represents an arbitrary sequence of operations, for the experiment under study, such as a thermometry sequence. The blue blocks represent Doppler cooling and readout (§ 4.1.2), the green blocks represent qubit state preparation (§ 4.1.1), the red blocks are RSB pulses with varying durations $t(N)$, and the purple blocks are deshelling pulses. After every 5 sideband cooling pulses, there is an interleaved state preparation sequence to ensure there is no population trapped outside of the closed cooling cycle. The RSB length is linearly increased with pulse number N , up to the expected RSB $R(\pi)$ -pulse duration for $n = 1$. (b, c) Thermometry scans after Doppler and sideband cooling. The red points are the populations measured via on-resonance RSB pulses, and the blue points are the populations after a sequence of $R(\pi)$ -RSB- $R(\pi)$ with varying RSB probe durations. The solid lines are fits to a thermal Fock state distribution, with truncation at $n = 100$. The extracted mean occupation number is $\bar{n} = 0.03(1)$, and $\eta\Omega/(2\pi) = 11.6(1)$ kHz.

narrow linewidth laser (§ 3.6.4), these sidebands can be resolved spectroscopically. The pulsed sideband technique employed consists of RSB pulses on the 729-nm transition, followed by deshelling, and repumping pulses on the 854-nm and 866-nm transitions respectively. For efficient cooling, initial short 729-nm pulses are applied to preferentially excite the thermal population with high phonon number, as the pulse duration for RSB $R(\pi)$ is proportional to $\sqrt{n}$, where n is the phonon number. The RSB pulse length is linearly increased to the expected RSB $R(\pi)$ -pulse duration for $n = 1$.

An example pulse sequence for a single experimental shot using sideband cooling, can be seen in Figure 4.9 (a). Experimental parameters used are summarised in Table 4.3.

We use the RSB of the quadrupole transition $|4S_{1/2}, m_j = -1/2\rangle \leftrightarrow |3D_{5/2}, m_j = -5/2\rangle$ as the subsequent deshelve and spontaneous decay from $|4P_{3/2}, m_j = -3/2\rangle$ to $|4S_{1/2}, m_j = -1/2\rangle$ forms a closed cycle, lessening the need for state-preparation pulses throughout the sideband-cooling sequence.

To verify the efficacy of our sideband cooling, thermometry experiments are performed by probing Fock state populations. This is done using on resonance RSB pulses and effective BSB pulses composed of a $R_x(\pi)$ -RSB- $R_x(\pi)$ pulse sequences. The spin population dynamics are measured as RSB pulse length, t , is varied. In the case of Fock state $|0\rangle$, there should be full contrast “BSB” oscillations, and no visible oscillations when probed with RSB pulses. A thermal Fock state distribution [85] (with truncation at Fock state = 100) is fitted to these signals to extract the mean occupation number, and $\eta\Omega$, the carrier Rabi frequency multiplied by the Lamb-Dicke parameter. A typical thermometry scan after Doppler and sideband cooling can be seen in Figure 4.9 (b, c), where (b) shows the population after effective BSB pulses, and (c) shows the population after RSB(t) pulses with duration t . The mean occupation number after sideband cooling is found to be $\bar{n} = 0.03(1)$, and $\eta\Omega = 11.6(1)$ kHz.




Figure 4.10: Heating rates of upper radial mode, measured by varying the delay between cooling and thermometry pulses. The solid line is a linear fit to the data, with slope $33(3)$ quanta/s. The error bars are given by the standard deviation of the fitted $\bar{n}$.

4.2.3 Heating Rates

After cooling the motion, the ion will return to a temperature in equilibrium with the environment. However, this process can be slowed by adequate isolation of the ion. In UHV ion trap systems (see § 3.4), motional heating is predominantly caused by electric field noise from the ion trap [51]. Noise due to thermal charge fluctuations at the trap surface can be mitigated by increasing ion-electrode distances. The *NPL* trap (§ 3.2), used here, has an ion-electrode distance of $250\text{ }\mu\text{m}$, larger than most surface traps, but less than that of a macroscopic blade or rod style trap [86]. Heating rates were probed by a series of thermometry experiments, explained above, with varying delay times, t_{delay} , between cooling and thermometry pulses. A typical plot can be seen in Figure 4.10. The heating rate of the system was found by a linear fit, shown by the solid line, to be $33(3)$ quanta per second on the upper radial with mode frequency $\omega_{\text{upper}} \approx 3\text{ MHz}$ on one ion. As will be discussed in § 4.2.4, the heating rate sets a limit for the motional coherence time.




Figure 4.11: a) Motional coherence time measured by Ramsey interferometry using Fock states $|0\rangle$ and $|1\rangle$. Points are populations measured in $|\downarrow\rangle$ after the Ramsey sequence with varying delay time. The black dashed line is an exponential decay fit, with $\frac{1}{e}$ coherence time $\tau = 18.0(8)$ ms, and initial contrast $C_0 = 0.90(2)$. The red line corresponds to the expected decay due to the independently measured heating rate of $\dot{n} = 33(3)$ q/s. b) Displays the residual oscillations after subtracting the expected heating-rate profile. c) Displays the PSD of the residual oscillations, with the dashed red line signifying 50 Hz.

4.2.4 Motional Coherence Times

For both hybrid computation, and for two-qubit entangling gates via geometric phase gates, we require motional modes with long coherence times compared to the desired circuit length. Electric field noise on the trap RF, or fluctuating external fields cause decoherence of the qumode. The field noise can cause coherent errors, motional frequency dephasing, or thermal heating of the mode, depending on its coherence and underlying frequency spectrum. Here the coherence time is quantified using Fock state interferometry — a Ramsey technique with initial superposition states $|\psi_{0m}\rangle = (|\downarrow, 0\rangle + |\downarrow, m\rangle)/\sqrt{2}$, where m is a pure Fock state [87]. The acquired Ramsey phase is mapped to the optical qubit through sideband pulses. Fock state interferometry is easy to implement through driving resonant carrier and RSB gates, and is sensitive to both dephasing and mode-heating noise. We note that similar interferometry techniques based on cat states [88] also quantify decoherence, however with different sensitivities to both dephasing and heating.

The Ramsey contrast decays is measured by the following sequence:

1. Prepare the ground state of both qubit and qumode, $|\downarrow, 0\rangle$. The qubit is first prepared via optical pumping, and the qumode is subsequently prepared by Doppler and resolved sideband cooling.
2. A $R_y(\pi/2)$ -pulse is applied on the carrier transition, preparing the state $(|\downarrow, 0\rangle + |\uparrow, 0\rangle)/\sqrt{2}$.
3. A RSB $R_{\text{RSB}}(\pi)$ -pulse is applied, preparing the state $|\psi_{0m}\rangle = (|\downarrow, 0\rangle + |\downarrow, 1\rangle)/\sqrt{2}$.
4. The system is left to evolve for Ramsey delay, t_{RAM} .
5. A RSB $R_{\text{RSB}}(\pi)$ -pulse is applied to map accumulated phase onto the qubit state.
6. A variable carrier $R_\phi(\pi/2)$ -pulse is applied.
7. The population of $|\downarrow\rangle$ is measured.
8. Repeat steps 1-7 for different values of ϕ to find the contrast of the Ramsey fringe.
9. Repeat steps 1-8 for different values of t_{RAM} to find the decay in contrast with time.

Figure 4.11 (a) shows the recovered coherence contrast for variable Ramsey delays for the upper radial mode of a single ion, with frequency $\omega_{\text{upper}} \approx 2\pi \times 2.9$ MHz. The black solid line indicates an exponential decay with a fit $\frac{1}{e}$ coherence time of $\tau = 18.0(8)$ ms, and initial contrast $C_0 = 0.90(2)$. The low initial contrast is attributed to imperfect RSB pulses due to a combination of finite residual mode temperature (expected around $\bar{n} = 0.03$), and imperfect calibration of the sideband frequency and π -gate duration. The exponential, rather than Gaussian, profile of the decay suggests that the dominating noise source is correlated for durations much shorter than the Ramsey delay times [89].

If motional-mode heating is dominant, the expected decay in the coherence between Fock states $|0\rangle$ and $|m\rangle$ for an equal superposition is given by [87, 90],

$$\rho_{0m}(t) = \frac{1}{2(1 + \dot{n}t)^{m+1}}, \quad (4.11)$$

where ρ_{0m} is the off-diagonal density matrix element, and $\dot{n}$ is the heating rate in phonons per second. Taking $m = 1$, and using the above protocol mapping the motional coherence to the spin-qubit, we find spin populations,

$$P_{\downarrow}(t) = \rho_{01}(t) + \rho_{10}(t) = \frac{C_0}{(1 + \dot{n}t)^2}, \quad (4.12)$$

where C_0 is included to account for imperfect mapping pulses. The dashed red line overlaid on Figure 4.11 (a) is the heating induced decoherence using the independently measured heating rate of $\dot{n} = 33(3)$ q/s, found in § 4.2.3. The measured contrast closely follows this decay, giving strong evidence that the coherence of the qumode is dominated by the underlying heating rate.

Further periodic structure on the decaying contrast suggests the presence of coherent noise sources. The frequencies of these oscillations are recovered by finding the power spectral density (PSD) of the residual contrast above the expected decay profile. Panel (b) displays the residual contrast amplitudes when subtracting the decay due to the heating rate from the observed populations $P_{\downarrow}$. The PSD, shown in panel (c), is found using Welch's method with a Hann window and zero-padded to visually smooth the spectrum. The vertical dashed red line indicates 50 Hz, the electric hum from the UK AC mains supply. The PSD suggests that this hum is the dominant coherent noise source affecting the contrast decay profiles.

The effect of dephasing and heating noise depends strongly on the state of the motional mode, as such, this metric of Fock state coherence time provides evidence of adequate performance of the qumode, but does not rigourously determine the error channel acting on the oscillator. However, this measurement suggests that the coherence of the qumode is sufficient for the short hybrid sequences we plan to explore with the apparatus.

If this coherence time is not adequate for future projects, there are three possible mitigations. First, we may explore the dependence of motional mode frequency on coherence time, effectively probing the electric-field noise spectrum to find quieter regions. Second, we can decrease the impact of any incoherent noise by reducing the duration of motional interactions, however emergent off-resonant coherent errors will then require mitigation. Third, we may explore algorithmic error mitigation techniques that reduce the impact of noise on estimated expectation values, for example, using virtual distillation. This technique was first developed explicitly to suppress the bias of an estimated observable due to finite-temperature effects [91], but has been generalised to mitigate errors occurring in quantum circuits [92, 93]. We discuss this method in the context of hybrid systems in Chapter 6, and provide an initial demonstration on a spin-only system.

4.3 Spin-Motion

For full control of the spin-motion hybrid system, we require interactions that couple the two. Perhaps the simplest of this class of interactions are the red- and blue-sidebands (RSB and BSB respectively, see § 2.4). The spin-dependent force (SDF), as described in § 2.4.2, is another highly useful spin-motion interaction. The calibration and characterisation of the SDF is presented here. Using the SDF applied simultaneously to two qubits, we demonstrate the Mølmer Sørensen two-qubit entangling gate [38].

4.3.1 Calibrating the SDF

Spin-dependent forces (SDF), as described in § 2.4.2, are used in this thesis for both two-qubit entangling gates (following section), and for generating and characterising nonlinear bosonic states (§ 6.1.1 and § 6.1.1). Here we describe the calibration strategy and characterisation of a single SDF.

The SDF is generated using the bichromatic field method introduced by Mølmer and Sørensen [38], using simultaneous detuned red- and blue-sidebands. In practice,




Figure 4.12: Detuning and duration scans of the spin-dependent force applied to an initial thermal state with $\bar{n} = 0.03(1)$. (a) Detuning scan, where the detuning, δ_g , is varied, and the duration is fixed at $t = 100$ μ s. (b) Detuning scan with a mis-set qubit frequency of 0.5 kHz. (c) Duration scan, where the duration, t , is varied, and the detuning is fixed at $\delta_g = 0$. The solid line is a fit to the expected behaviour given by equation 4.13 with only Ω_{SDF} as a free parameter. The extracted Ω_{SDF} is used with equation 4.13 to plot the dashed red lines in (a) and (b) to guide the eye.

this SDF interaction involves relatively few degrees of freedom to calibrate: the central qubit “carrier” frequency, the motional mode frequency, the power in each optical tone, and ensuring negligible off-resonant excitation of other spin, and spin-motion interactions.

The bichromatic field is generated by providing two simultaneous RF tones to the single-pass (SP) AOM, as seen in the *TrapLab* panel of Figure 3.10. The SDF may be driven by either the single-ion addressing system (path 729#1), or by a global field (path 729#2). The SP AOMs are driven by an arbitrary waveform generator² to provide phase, frequency, and amplitude control over the two tones. The bichromatic field is measured using a high-bandwidth photodiode after an optical pick-off before addressing the ions. The intensity measured by the photodiode

²Phaser Sinara board

has a beat-note with frequency $\omega_b \approx 2\omega_{\text{upper}}$. The contrast of this beat-note is optimised to calibrate the power balance between the two optical frequencies. The power balance is varied by updating the RF amplitude supplied to the SP AOMs. AOMs generate frequency shifts in the laser beam at the expense of small frequency-dependent angular shifts. The beam pointing of the diffracted order also has some temperature dependence. After the AOM, the beam is coupled into a single mode fibre, as such the AOMs introduce a frequency- and temperature-dependent loss. These losses can lead to significant errors in power-balance between the two tones. We found that optimising for short optical path-lengths between the SP AOM and the fibre, careful alignment to the fibre, and ensuring the SP AOM is supplied with a constant duty cycle were all required for stable power-balance.

Off-resonant excitations of other interactions are negated through selecting trap parameters to spectrally isolate a single-motional mode, and ensuring that gates are driven adiabatically — the SDF Rabi frequency, Ω , is significantly less than the frequency spacing to nearby interactions. Ensuring the interaction is smoothly ramped on and off also reduces excess off-resonant excitations.

The bichromatic field leads to a Stark shift, δ_s , of the central qubit frequency due to off-resonant coupling of both dipole and nearby quadrupole transitions. As such, the bare-qubit frequency measured by RPE in § 4.1.3 is insufficient for low-error SDFs. The Stark-shift calibration is thus done through observing the application of the SDF to the ion when all other parameters are well calibrated. When the qubit frequency is set incorrectly, by Δ , the two tones no longer have the same absolute detuning from their respective RSB or BSB. The impact of this error depends on the detuning of the SDF compared to the motional mode, δ_g . When $\delta_g \gg \Delta$ the induced phase space dynamics are similar to the desired interaction, however when $\delta_g \approx \Delta$, we see a strong deviation from ideal behaviour due to the resonant driving of only one of the sidebands. We utilise this frequency dependent non-ideal behaviour of the uncalibrated SDF to manually calibrate the induced qubit Stark shift. The phase space dynamics can be observed through projective spin measurements. An SDF square pulse of duration t_{pulse} with spin conditioning,

$\hat{\sigma}_x$, and variable detuning, δ_g , is applied to the initial prepared state $|\downarrow, 0\rangle$. The spin population returns to $P_\downarrow = 1$ only when the pulse duration and SDF detuning are related by $\delta_g = \frac{2\pi}{nt_{\text{pulse}}}$, for integer n . These spin dynamics are described by [94],

$$P_{\downarrow, \text{th}} = \frac{1}{2} \left(1 + e^{-4(\bar{n} + \frac{1}{2})|\beta(t, \delta_g)|^2} \right), \quad (4.13)$$

where $\bar{n}$ is the initial thermal phonon occupation, and $|\beta(t, \delta_g)| = \Omega_{\text{SDF}} \sin(\delta_g t/2)/\delta_g$, is the absolute displacement of the initial state. Deviations from this behaviour indicate either a miscalibration of Stark shift, or unintended off-resonant excitation of nearby transitions.

To give an example of this calibration procedure we show experimental data of both a well-calibrated, and a miscalibrated SDF, in Figure 4.12 (a), and (b) respectively. For these detuning scans we set $t_{\text{pulse}} = 100 \mu\text{s}$, and use the global addressing beam with 2 mW of total power. The points indicate the spin population averaged from 200 samples. The calibrated scan assumes a Stark-shift $\delta_s = 2\pi \times 4\text{kHz}$, while the miscalibrated scan assumes $\delta_s = 2\pi \times 3.5\text{kHz}$. The dashed red lines in both (a) and (b) are the expected population dynamics inferred using an independently measured SDF strength of Ω_{SDF} . The well-calibrated SDF shows good agreement to the expected dynamics, while the qubit-frequency mis-calibrated SDF displays a clear asymmetry near $\delta_g = 0$.

Panel (c) displays a “splitting” experiment, applying the SDF on resonance $\delta_g = 0$, and varying the duration. The solid black line is a fit of the spin populations using equation 4.13, and assuming $\bar{n} = 0$. We find here $\Omega_{\text{SDF}} = 2\pi \times 6.6(2) \text{ kHz}$. Splitting of the motional wave-packet by a resonant SDF is exactly the procedure used for motional state tomography, which we give further details for in § 6.1.1.

4.3.2 Two-Qubit Entangling Gates

To complete the demonstration of elementary operations required for universal qubit control, we characterise the performance of two-qubit entangling gates, $R_{xx}(\theta)$, via the MS interaction on our apparatus.

This interaction, as discuss in § 2.4.3, is implemented through a detuned SDF. To




Figure 4.13: Experimental calibration and characterisation of a Walsh-modulated Mølmer-Sørensen (MS) two-qubit entangling gate with fidelity $\mathcal{F} = 0.980(6)$. (a) Populations versus duration scan of the MS interaction. Each point corresponds to the average of 200 shots, and the error bars are given by the standard deviation of these averages. The desired entangling gate corresponds here to $t_g = 95\mu\text{s}$. (b) Observed parity oscillations after the MS interaction. The fitted contrast here is $C = 0.97(1)$. (c) Populations at the desired entangling gate duration, $t_g = 95$ us.

reduce the sensitivity to incoherent motional errors such as heating and dephasing, a multi-loop sequence is used. Specifically, we use a “Walsh-modulated two-loop” sequence consisting of two complete loops in phase space with the motional phase of the SDF flipped by π for the second loop [95]. The two-loop gate requires gate duration and detuning related by equation 2.55, yielding, $T_2 = \frac{\pi\sqrt{2}}{\Omega_{\text{SDF}}}$ and $\delta_g = 2\sqrt{2}\Omega_{\text{SDF}}$. The change of SDF sign after the first loop reduces sensitivity to systematic detuning offsets, for example due to mode-frequency drift.

The two-qubit gate performance can be characterised many different ways. Ideally a process tomography method would be used to fully characterise the CPTP map implemented by the noisy two-qubit gate [23]. However such a method requires near-perfect state-initialisation to correctly discern dynamics originating from the gate under test. Random circuit methods such as two-qubit Clifford RB [96], or

cyclic benchmarking [24] could also be used to gain single-value performance metrics insensitive to state-preparation errors. In this thesis we use a state tomographic technique [97], which although does not rigorously characterise the action of the gate on arbitrary initial states, is quick to implement, and only requires a single additional $R_x(\frac{\pi}{2})$ -pulse.

A single $R_{\phi_s\phi_s}(\frac{\pi}{2})$ two-qubit entangling gate, where ϕ_s is some fixed angle in the xy -plane, is applied to the initial state $|\downarrow\downarrow\rangle$. This gate targets the creation of the maximally-entangled state $|\Phi_{\phi_0}^+\rangle = 1/\sqrt{2}(|\downarrow\downarrow\rangle + e^{i\phi_0}|\uparrow\uparrow\rangle)$, where ϕ_0 , is some fixed phase³. Practically, due to systematic errors and noise on both the qubit and motional state, the final mixed state $\hat{\rho}$ is produced. The state-fidelity between $\hat{\rho}$, and the pure Bell state is given by,

$$\mathcal{F} = \langle \Phi_{\phi_0}^+ | \hat{\rho} | \Phi_{\phi_0}^+ \rangle = \frac{1}{2} (\rho_{\downarrow\downarrow,\downarrow\downarrow} + \rho_{\uparrow\uparrow,\uparrow\uparrow}) + \frac{1}{2} (e^{i\phi_0} \rho_{\uparrow\uparrow,\downarrow\downarrow} + e^{-i\phi_0} \rho_{\downarrow\downarrow,\uparrow\uparrow}), \quad (4.14)$$

where the matrix elements are given by $\rho_{ij,kl} = \langle ij | \hat{\rho} | kl \rangle$. Thus to extract this fidelity, only partial state tomography is required [97]. The first bracketed term of equation 4.14, denoted A , is measured by performing projective measurements of the two ions after the applied $R_{\phi_s\phi_s}(\frac{\pi}{2})$. The extracted two-qubit joint populations directly recover this population term, $A = P_{\downarrow\downarrow} + P_{\uparrow\uparrow}$, where P_{ij} is the population of the two ions in state $|i\rangle |j\rangle$. The coherence terms in the second bracket are measured by applying a global analysis $R_\phi(\pi/2)$ pulse to the two ions with variable phase ϕ . The final spin populations are combined to a parity measurement, defined as $\Pi = P_{\downarrow\downarrow} + P_{\uparrow\uparrow} - P_{\downarrow\uparrow} - P_{\uparrow\downarrow}$. Varying ϕ , the observed parity relates to the matrix elements of $\hat{\rho}$ as,

$$\Pi(\phi) = \rho_{\downarrow\uparrow,\uparrow\downarrow} + \rho_{\uparrow\downarrow,\downarrow\uparrow} + e^{-2i\phi} \rho_{\uparrow\uparrow,\downarrow\downarrow} + e^{+2i\phi} \rho_{\downarrow\downarrow,\uparrow\uparrow}. \quad (4.15)$$

In the case that there are no coherences between single-excitation states $|01\rangle \leftrightarrow |10\rangle$, the parity reduces to simple oscillations of the form,

$$\Pi(\phi) = 2\text{Re} [e^{-2i\phi} e^{i\phi_0} \rho_{\uparrow\uparrow,\downarrow\downarrow}] = C \cos(2\phi - \phi_0), \quad (4.16)$$

³ ϕ_0 is left unspecified as the target of a specific phase is not required, we only require that it is stable over time, and measured precisely.

where we have taken $\rho_{\uparrow\uparrow,\downarrow\downarrow} := e^{i\phi_0}|\rho_{\uparrow\uparrow,\downarrow\downarrow}|$, and $C = 2|\rho_{\uparrow\uparrow,\downarrow\downarrow}|$. The contrast of the parity oscillations thus provide the magnitude of the coherence term. The Bell-state fidelity is thus given by $\mathcal{F} = \frac{A+C}{2}$.

Using this method, the fidelity of a globally addressed two-qubit entangling gate by the Walsh method is experimentally measured. Figure 4.13 (a) shows the spin-population dynamics as a function of total Walsh-1 duration for an MS gate acting on the upper radial common mode, with gate parameters, $\delta_g = -19.5$ kHz, and 35 mW of total beam power, causing a Stark shift of $\delta_s = 14$ kHz. To suppress off-resonant excitation of nearby spin-motion interactions, the SDF field is ramped on and off with a rise envelope of $\sin^2(\pi t/(2t_{\text{ramp}}))$ applied via the AWG, where $t_{\text{ramp}} = 9.6$ μs is the ramp duration. The desired geometric phase is accumulated by $t_g = 95$ μs . The parity oscillations and two-qubit populations are measured at this point, and shown in panels (b) and (c) respectively. A contrast of $C = 0.97(1)$, and population of $A = 0.987(2)$ is found, resulting in a state fidelity, $\mathcal{F} = 0.980(6)$. Each population point in these figures are found by taking the average of 200 shots of the gate sequence.

Similar fidelities were measured when applying the interaction by the single-ion addressing system on the upper radial common mode. Using gate parameters $\delta_g = -15$ kHz, $t_g = 110$ μs , and 0.148 mW of total beam power, a fidelity of $\mathcal{F} = 0.983(4)$ was measured, with contrast $C = 0.976(7)$ and population contribution $A = 0.991(1)$.

We believe we are currently limited either by systematic calibration errors, or by laser phase noise of the 729-nm system. We have anecdotally observed a strong dependence on the gate-fidelity with PID parameters for the 729-nm PDH lock. To increase the two-qubit gate fidelity, we believe that further time must be devoted to characterising the noise-spectrum of the 729-nm, and optimising lock parameters to suppress phase noise at the frequencies relevant to the two-qubit entangling gates. This, along with developing an automated gate-tuning algorithm

Operation	Fidelity	Section
State Prep	99.96(2)	§ 4.1.1
Measurement	< 99.7(3)	§ 4.1.4
$R_x(\pi/2)$	99.913(2)	§ 4.1.4
$R_{xx}(\pi/2)$	98.3(4)	§ 4.3.2

Table 4.4: Summary of qubit-operation fidelities presented in this chapter.

for our apparatus, is left to future work. The current fidelities observed here are sufficient for the projects outlined in this thesis. In § 6.2.3 we calibrate $R_{xx}(\frac{\pi}{4})$ gates using this state-fidelity approach.

4.4 Summary

These results serve as both a proof of principle for full discrete-variable control on our system, but also as a benchmark to which we can compare future system improvements. The fidelities of the operations demonstrated in this chapter are summarised in Table 4.4, with relevant sections highlighted.

5

k-copy Randomised Benchmarking

Contents

4.1	Spin	61
4.1.1	State Preparation	62
4.1.2	Measurement	63
4.1.3	Single-Qubit Rotations	64
4.1.4	Single-Qubit Randomised Benchmarking	69
4.1.5	Spin Coherence Times	73
4.1.6	Single-ion Addressing	75
4.2	Motion	78
4.2.1	Calibrating Mode Frequencies	79
4.2.2	Cooling	80
4.2.3	Heating Rates	84
4.2.4	Motional Coherence Times	85
4.3	Spin-Motion	88
4.3.1	Calibrating the SDF	88
4.3.2	Two-Qubit Entangling Gates	91
4.4	Summary	95

To construct useful scale quantum processors we must house many qubits physically close together, and must ensure the full register operates with “low error rates”. But how should we define our errors in a large system? Naively, one might hope that the error rate of a multi-qubit system could be inferred from the error rates of its constituent qubits, but unfortunately, not all errors are commensurable — there are a zoo of possible error channels. They may be

coherent, incoherent, continuous, stochastic, local, correlated, time dependent or independent (and I am sure a host of further adjectives exist). In response, a zoo of error correction and mitigation techniques have been developed each with their own scaling overheads and assumptions on these noise models they treat [28, 98, 99]. It is a large and intimidating search space, and so, to better equip ourselves for navigation, we require tools to characterise what errors are relevant in our physical systems. Crucially, errors which induce correlations across multiple qubits, termed spatial correlations or crosstalk errors [100, 101], require mitigation and correction strategies built on noise models that faithfully represent their correlated structure [102]; assuming globally depolarising noise models (white noise), can lead to systematic underestimation of logical error rates [26, 103], or systematic bias of the corrected observable outcome [104]. In reference [105] the threshold for several error correcting codes was shown to be correlation-dependent, with some larger codes having expected failure probabilities increased by several orders of magnitude when exposed to highly correlated noise. Furthermore, in § 6.2.3, we experimentally demonstrate that virtual distillation, an error mitigation technique, is entirely ineffective against perfectly correlated errors.

In this chapter we introduce a scalable method for characterising correlated errors in multi-qubit systems: *k-copy randomised benchmarking*. This work extends the widely adopted single-qubit Clifford randomised benchmarking (RB) protocol [78, 79, 106] (discussed in § 4.1.4) to capture and quantify correlated errors in scaled multi-qubit registers. We introduce a novel method to parallelise RB, and utilise the extracted symmetry resolved dynamics to provide evidence of, and quantify various multi-qubit error channels.

This chapter is structured as follows: we start by giving an overview on the problem of *characterisation*, and what techniques already exist to address it. We then outline the *k-copy* protocol and discuss the relevant symmetric subspace structure. We derive the population dynamics when subject to simple noise models and discuss variations of the protocol that isolate specific subspaces. We provide experimental demonstrations of the protocol on a trapped-ion processor through injecting known

correlated and uncorrelated errors. We end this chapter by demonstrating the scalability of the protocol on an 8-qubit register, and discuss the limitations and potential of this method for further scaling to larger systems.

5.1 Errors and Characterisation

Errors, which we shall give a broad definition as any deviations between the realised and ideal state of a quantum information processor (QIP), are due to unintentional environmental coupling and control imperfections. They are unavoidable, and if not carefully considered or mitigated, lead to unreliable or unverified computation. The process of modelling and quantifying these errors is known as *characterisation*. Characterising errors has two practical advantages: 1) for the experimentalist, quantifying an error allows for a clear signal to improve hardware and calibration routines, and 2) for the theorist, understanding physically motivated error channels allows for the development of error correction and mitigation techniques that are robust to real world noise. Both users are concerned ultimately with the same question: How faithfully can we execute arbitrary circuits on our physical devices? Extracting metrics to this end is a wide research area in quantum information, with techniques of varying scalability and implicit assumptions on underlying noise models [21]. On one end of the spectrum sits full tomography [23, 107–109] which reconstructs quantum channels, here represented as completely positive, trace-preserving (CPTP) maps on the space of density matrices,

$$\mathcal{E}(\hat{\rho}) : \mathbb{C}^{d \times d} \rightarrow \mathbb{C}^{d \times d}, \quad (5.1)$$

where the dimension of the Hilbert space is $d = 2^n$ for n qubits. The CPTP map is rich in information, capturing all Markovian (those which do not depend on history) error channels, but has poor scaling with system size, generally requiring $O(d^4)$ parameters to describe, and $O(d^4)$ measurement settings to reconstruct. This scaling fundamentally limits tomography to small system sizes. The goal then shifts to how to extract useful information in predicting the behaviour of a system executing a circuit, without full channel reconstruction. Now, two strategies emerge:

1) we make assumptions on the underlying noise model, or 2) we actively simplify the structure of the error channel through averaging. The first strategy leads to techniques such as noise learning [110, 111] and the latter leads to random circuit methods, such as randomised benchmarking (RB) [24, 106, 112]. Both approaches dramatically reduce the number of parameters required to describe the system, and the number of measurements required to extract them.

The assumptions on the noise model are cleanly motivated by considering the modularity of quantum circuits, and how errors may violate this modularity. Following the description of imperfect computation and crosstalk errors given in references [77], and [101], respectively, we define the following models of QIPs: *Ideal QIPs*, *Markovian QIPs* (containing the subset of *Crosstalk-Free Markovian QIPs*), and *Non-Markovian QIPs*.

The most desirable, yet technically unfeasible, is the *Ideal QIP* — described by a closed quantum system where the qubits comprise a quantum state represented by a single vector in Hilbert space. The state is isolated from the environment and may be decomposed into a collection of subsystems that remain uncoupled unless explicit multi-qubit operations are applied. Operations are all unitary evolutions, conveniently described by single and two-qubit unitaries, which we term *gates*. Errors may exist, but are limited to reversible unitary interactions.

Markovian QIPs include external and control field noise due to an environment with no history, i.e. a source of Markovian noise. We must depart from the quantum state vector and unitary evolution description, and instead use density matrices and CPTP maps to describe the state and dynamics of the QIP. We highlight a subset of these models, termed *Crosstalk-Free Markovian QIP*, where the dynamics of distinct subsystems are governed by locally acting CPTP maps. To satisfy this condition, operations are said to be *local* and *independent* [101].

- *Local* is the assumption that only explicitly applied multi-body operations lead to correlations between qubits that participate in the ideal multi-body operation. For a bipartite system, a local Markovian QIP executing a circuit

with no explicit coupling between subsystems, A , and B , can be described by the product channel CPTP map,

$$\mathcal{E}_{AB}(\hat{\rho}_A \otimes \hat{\rho}_B) = \mathcal{E}_A(\hat{\rho}_A) \otimes \mathcal{E}_B(\hat{\rho}_B), \quad (5.2)$$

where $\mathcal{E}_A$ and $\mathcal{E}_B$ are CPTP maps acting locally on the A and B subsystems respectively.

- *Independent* is the assumption that the action of an applied gate on a subset of qubits is independent of any other gates applied in parallel on the other qubits. This implies that the CPTP map of gate G , acting on subsystem A in layer i , given by $\mathcal{E}_G^{(i)}(\hat{\rho}_A)$, is identical to the action of the same gate acting on subsystem A in layer j ,

$$\mathcal{E}_G^{(i)}(\hat{\rho}_A) = \mathcal{E}_G^{(j)}(\hat{\rho}_A), \quad (5.3)$$

where parallel applied gates on the other subsystems may be different between layers.

With these two criteria satisfied, we may retain the composability of gates, conveniently allowing us a good prediction of the behaviour of a many-qubit deep circuit, using only characterisations of each gate in a universal gate set applied to small subsystems. If these criteria are not met, to accurately predict the behaviour of a deep circuit, each unique layer must be characterised, which is intractable.

The last class of models, included for completeness, are *Non-Markovian QIP*, where we relax the assumption that the environment exhibits no hysteresis. In this case, any instantaneous state may still be able to be described by a density matrix, however the dynamics can no longer be described by the application of a sequence of CPTP maps. To accurately characterise the system, full process tomography would need to be implemented for each possible circuit.

Practically, to ensure we can reasonably characterise our physical systems, and predict their performance on novel circuits, we must design and appropriately isolate

hardware to be well approximated by the *Crosstalk-Free Markovian* models. The next section provides some background on existing techniques to verify crosstalk-free performance, and provide some quantification when the assumption is violated.

5.2 Existing approaches to correlated-noise characterisation

Existing approaches to correlated-noise characterisation span a trade-off between expressiveness and scalability. The most expressive, and arguably least scalable, are tomographic techniques, such as gate set tomography [107] or the crosstalk-sensitive variation [113]. Here, the implicit assumption that the noise environment is Markovian is sufficient, and no other assumptions of underlying noise models are required. These techniques are valuable in producing detailed CPTP maps that inform not only the presence of crosstalk, but also offer evidence on the underlying noise process. Demonstrations are however limited to small system sizes of two-qubits in [113].

Noise learning protocols are substantially more scalable by imposing a restricted set of noise models, and fitting experimental data to these models. Correlated-noise specific variations of this approach [111, 114] apply repeated random circuits that act to simplify the action of noise channels to Pauli type errors, and then through measurements and statistical inference, likely Pauli error models are found, and crosstalk-free violating errors may be quantified. Noise learning can also be adaptable by first fitting simple models, for example when twirling over random circuits yields structured error channels, and progressively increasing the complexity of the inferred model when arbitrary circuits are run on the same device [115]. This nested-model approach provides a framework for robust characterisation of the system, however, demonstrations are still currently limited to small systems.

Similar to enforcing Pauli-type errors, randomised benchmarking techniques average the action of noise channels to impose structure on general error channels. This is achieved by twirling the action of the channel over deep circuits comprised of a chosen gate group. Clifford twirling techniques, notably the standard RB protocol [78, 79,

106], have been adapted for detecting crosstalk induced correlations by simultaneous RB [116], and correlated RB [117] protocols. Correlated RB uses independent local twirling of multiple system copies to extract multi-qubit survival decay rates with Clifford sequence length. The weight and locality of correlated errors may then be inferred through a distance of the measured channel from a product (uncorrelated) channel. Other, non-Clifford twirled RB protocols, such as cycle-benchmarking [24] and the Error Per Layered Gate protocol (EPLG) [25] apply parallel gate sequences on multiple subsystems to extract some error applied per layer, that can then be used to infer error contributions due to correlated-noise channels. These methods have been shown to scale to 10s of qubits, and provide system level diagnostics of device performance.

We note two further streams of current research, first the application of RB in determining the unitarity of noise [118], and the extension of this work to quantifying bipartite unitary noise [119], which can induce correlations across the qubit register. Second, the use of custom circuits designed to probe crosstalk-free violations [101, 120, 121].

The present protocol we have developed, and demonstrated experimentally, sits firmly in the RB methodology. It however differs from simultaneous and correlated RB by providing not only a quantification of crosstalk-free violating noise, but also providing evidence for the action of the error channel on a multi-qubit system. This extra information is inferred through direct observation of population decay and coupling between the subspaces enforced by twirling over the diagonally applied single-qubit Clifford group. The inspiration of this work comes from correlation spectroscopy [122, 123] and subspace-leakage techniques [124, 125]. Correlation spectroscopy demonstrates that common-mode noise can be revealed through its action on collective and differential degrees of freedom, while subspace-leakage approaches provide a framework for defining error models not by full CPTP maps, but rather by their action on a discrete set of physically relevant subspaces. In such settings, symmetry is not merely a simplifying assumption providing structure to arbitrary channels, but rather it is the very mechanism that separates crosstalk-free




Figure 5.1: Circuit diagram for k -copy randomised benchmarking. A sequence of $m - 1$ random Clifford gates are applied in parallel to k qubits, with a final inverting Clifford applied to ideally recover the initial $|\downarrow\rangle$ state. All qubits are measured, and correlations between qubits that successfully return to the initial state are recorded.

violating errors from local noise. The next section details the subspaces enforced by diagonally applied single-qubit Clifford twirling, and provides the protocol for k -copy RB.

5.3 The k -copy RB scheme

This section provides the theoretical background to the k -copy RB protocol. First, we provide the experimental sequence and analysis required to execute the protocol, with a focus on pairwise analysis of correlations. We then study the symmetries enforced by the random circuits. This provides insight into how k -copy RB reveals correlated errors through the asymptotic structure of the resulting population decay profiles. In this chapter we specifically focus on correlated errors due to globally acting external noisy fields, and local depolarising errors. As such we analyse how these explicit error models can be quantified by k -copy RB, and highlight three emergent parameters of interest in the two-copy case: $\varepsilon_L, \varepsilon_R, \varepsilon_{\text{corr}}$, which respectively represent local depolarising error components on the “left-qubit”, on the “right-qubit”, and a correlated error component acting on the pair. This explicit model is required to provide the analytic decay curves used in the experimental section 5.6. Details of this analytic decay, and an intuition for the observed asymptotic structure of the RB decay curves are given in § 5.3.5.

5.3.1 Protocol

The *k*-copy RB protocol is implemented through standard single-qubit Clifford RB sequence (see § 4.1.4) applied simultaneously to multiple qubits, with the same randomly sampled Clifford applied to each qubit at every step, as shown in Figure 5.1. A single sequence is composed of $m - 1$ parallel gates sampled randomly from the Clifford group. As in standard RB, the final applied layer is an inverting Clifford, which is selected to rotate the final state back to the initial computational eigenstate. The full sequence of random Clifford gates is then equivalent to the identity operation with some unknown noise overlaid¹. Practically, each gate layer can be executed using a global field, or through parallel local control lines addressing each qubit. A final $\langle Z \rangle$ measurement is made in parallel on all qubits to probe the successful execution of the Clifford sequence. Qubits measured in the target state are labelled *S*, and qubits measured in the opposite state are labelled *F*, for success and failure respectively. To approximate ideal twirling, r random circuits are trialled for each length m , and each circuit is repeated s times to reduce shot noise. To retain information on multi-qubit correlations, for every experimental shot, data from each qubit pair (i, j) is assigned to one of four outcomes, $\mathcal{P}_{SS}^{(i,j)}$, $\mathcal{P}_{SF}^{(i,j)}$, $\mathcal{P}_{FS}^{(i,j)}$, or $\mathcal{P}_{FF}^{(i,j)}$, where *S* and *F* denote success and failure. As with other RB protocols, the success probabilities versus sequence length are then fit to explicitly defined decay curves inferred by the action of the twirled error models.

By limiting the analysis to pairwise comparisons, the *k*-copy protocol is both efficient in the number of trials required on the device, and in the postprocessing of the recovered data. For a register of k qubits, the same shot record contains outcomes for all $k(k - 1)/2$ pairs. The two-body error models are derived in the following sections for prevalent noise mechanisms in both ion trapping and superconducting information processors.

¹Note: To further symmetrise state preparation and measurement errors, we may allow the final target state to be randomised between $|\downarrow\rangle$ and $|\uparrow\rangle$ for each sequence [78].

Sub-representation	irrep	Basis operators
Identity	1	II
Trace	1	$(XX + YY + ZZ)/\sqrt{3}$
Left qubit	3	XI, YI, ZI
Right qubit	3	IX, IY, IZ
Antisymmetric	3	$(XY - YX)/\sqrt{2},$ $(XZ - ZX)/\sqrt{2},$ $(YZ - ZY)/\sqrt{2}$
Diagonal traceless	2	$(XX + YY - 2ZZ)/\sqrt{6},$ $(XX - YY)/\sqrt{2}$
Symmetric off-diagonal	3'	$(XY + YX)/\sqrt{2},$ $(XZ + ZX)/\sqrt{2},$ $(YZ + ZY)/\sqrt{2}$

Table 5.1: Sub-representations of the $Cl_1^{\otimes 2}$ group. The irrep, dimension of each sub-representation, and a set of operators in the Pauli basis which span the subspace dimension is given. The irrep basis vectors are from equation 56 in reference [126].

5.3.2 Representations of $Cl_1^{\otimes 2}$

To understand the symmetries enforced by twirling over the $Cl_1^{\otimes 2}$ gate-set, it is first useful to understand the structure of the Clifford group. The “structure” of a group can be efficiently described by the concepts of representations, and irreducible representations (irreps). In Appendix A.3 we outline why twirling over the group leads to this natural language of irreps, however, here we limit its application to that of a convenient basis in which to describe CPTP error maps. We shall refer to this as working in the Schur basis.

The 2-copy single-qubit Clifford group, analysed in detail in references [126, 127], decomposes the general 16-dimensional two qubit space into 4 irreps with dimensions $\{1, 2, 3, 3'\}$, (the notation $3'$ is used to contrast the other 3-dimensional irrep), and multiplicities $\{2, 1, 3, 1\}$. The sub-representations are summarised in Table 5.1. In Appendix § A.3.2 we confirm the multiplicity of these sub-representations by their character on the conjugacy classes of the Clifford group.

Given that the structure of a general twirled map is enforced to be block diagonal in the Schur basis (see equation A.40), we can do a simple count to find generally how many parameters are required to fully describe the CPTP map. Accounting

for the multiplicities, and the constraint that the map is trace preserving, we find that at most $2^2 + 3^2 + 1 + 1 - 2 = 13$ independent scalars are needed to describe $\tilde{\Lambda}$. The diagonal terms, here labelled $\vec{\alpha} = (\alpha_L, \alpha_R, \alpha_{\text{tr}}, \alpha_{\text{dt}}, \alpha_{\text{so}}, \alpha_{\text{as}})$, describe a decay rate contribution from each sub-representation, whilst off-diagonal terms, $\vec{m} = (m_{\text{id} \rightarrow \text{tr}}, m_{L \rightarrow R}, m_{R \rightarrow L}, m_{L \rightarrow \text{as}}, m_{R \rightarrow \text{as}}, m_{\text{as} \rightarrow L}, m_{\text{as} \rightarrow R})$ describe coupling between the multiplicity spaces.

In practice, different physical error mechanisms populate only subsets of the coefficients in $\vec{\alpha}$ and $\vec{m}$, leading to simplified dynamics with characteristic signatures of the underlying process. Table 5.2 provides possible physical mechanisms capable of coupling isotypic representation subspaces. We devote some time in the next two sections to describing noise due to fluctuations of a common field, and to local depolarising type errors. In Appendix A.3.4 we discuss a ZZ -type noise model, relevant to the superconducting qubit platform.

5.3.3 Common random unitary noise

It can be seen that the error channels that lead to couplings neatly discern the physical origin of the noise channel. For example, in transmon superconducting qubit platforms, always on ZZ -couplings are a dominant crosstalk error source [128–130], and so the couplings between single qubit subspaces and the antisymmetric subspace can be quantified; in the neutral atom platform, correlated non-unital errors may be observed if the physical qubit separation is of similar order to the qubit transition wavelength [131]; however, in ion trap platforms, environmental noise is often dominated by magnetic and electric field fluctuations, which may act globally on all qubits if they are stored in the same trapping potential. These field fluctuations induce both common² and independent single-qubit rotations, but do not couple irrep subspaces together (i.e $m = 0$ for $m \in \vec{m}$). Thus, the general error map of this type is now parametrised by at most the 6 scalars, $\vec{\alpha}$, and results in Liouville form diagonal in the Schur basis, $\mathcal{S}$, (ordered by Table 5.1),

$$\tilde{\Lambda}^{\mathcal{S}}(\mathcal{E}) = \text{diag} (I_1, \alpha_{\text{tr}} I_1, \alpha_L I_3, \alpha_R I_3, \alpha_{\text{as}} I_3, \alpha_{\text{dt}} I_2, \alpha_{\text{so}} I_3), \quad (5.4)$$

²Here “common” explicitly means applying identical rotations to each qubit.

Coupling	irrep	Mechanism
$m_{\text{id} \rightarrow \text{tr}}$	1	Decays (non-unital)
$m_{L \rightarrow R},$ $m_{R \rightarrow L}$	3	Qubit Swap, Correlated decays (non-unital)
$m_{L \rightarrow \text{as}},$ $m_{R \rightarrow \text{as}},$ $m_{\text{as} \rightarrow L},$ $m_{\text{as} \rightarrow R}$	3	Entanglement

Table 5.2: Possible mechanisms that result in the coupling of sub-representation sectors of the $Cl_1^{\otimes 2}$ twirled error channel. This list is non-exhaustive, but acts to link the irrep structure to physical channels.

where I_d is the d -dimensional identity.

We term noise due to common small-angle single-qubit rotations acting on k -qubits *common noise*. For two qubits the general form is given by the separable channel,

$$\mathcal{E}_{\mathcal{A}}(\hat{\rho}) = \sum_{a \in \mathcal{A}} p_a (U_a \otimes U_a) \hat{\rho} (U_a \otimes U_a)^\dagger, \quad (5.5)$$

where $\mathcal{A}$ specifies the ensemble of unitaries applied with probabilities p_a , with $\sum_{a \in \mathcal{A}} p_a = 1$, and $U_a(\theta_a, \hat{n}_a) = e^{-i(\theta_a/2)\hat{n}_a \cdot \hat{\sigma}}$, with rotation angle, θ_a , axis, $\hat{n}_a$, and Pauli vector, $\hat{\sigma}$. A subset of error channels following this description cannot be described by a product channel map on a bipartite system, and so violate the *local* condition given in equation 5.2. These errors therefore cannot be modelled by a crosstalk-free Markovian QIP although originating from separable channels. This is an important distinction: non-product channels are not limited to direct entangling noise.

To find the twirled error channel of the common unitary channel, $\tilde{\Lambda}^S(\mathcal{E}_{\mathcal{A}})$, we first consider the effect of twirling on a single common unitary channel, $\mathcal{E}_a(\hat{\rho}) = (U_a \otimes U_a) \hat{\rho} (U_a \otimes U_a)$. We outline two approaches to find $\tilde{\Lambda}^S(\mathcal{E}_a)$:

First, as the basis sets for the relevant sub-representations are known (see Table 5.1), and as the single qubit Clifford group only contains 24 elements, we may construct the common unitary channel in the Pauli transfer matrix (PTM) representation[77, 132], and, using equation A.37, we may explicitly calculate the twirled channel $\tilde{\Lambda}^{\mathcal{P}}(\mathcal{E}_a)$ utilising a symbolic computation program³. Here the superscript $\mathcal{P}$ denotes

³Such as *Mathematica* by *Wolfram Research, Inc.*, or the *Python* package *SymPy*.

we are working in the Pauli basis. The α parameters are then simply read from the diagonal elements after transforming the PTM representation to the known basis operators for the sub-representations, i.e. transforming to $\tilde{\Lambda}^{\mathcal{S}}(\mathcal{E}_a)$.

The second approach utilises the symmetries of the irrep spaces, and the separability of the error channel to construct the PTM projectors onto the irrep spaces. This method is outlined in Appendix A.3.3, and provides further insight into the symmetries each irrep possesses, and the naming convention for the sub-representations.

Both approaches ultimately yield the $\vec{\alpha}$ parameters,

$$\alpha_L = \alpha_R = \alpha_{\text{as}} = \frac{1}{3}(1 + 2c), \quad (5.6)$$

$$\alpha_{\text{tr}} = 1, \quad (5.7)$$

$$\alpha_{\text{dt}} = \frac{1}{2}(c^2 + 2c - 1 + (1 - c)^2\kappa), \quad (5.8)$$

$$\alpha_{\text{so}} = \frac{1}{3}(3c^2 - (1 - c)^2\kappa), \quad (5.9)$$

where $c = \cos(\theta_a)$, and $\kappa = n_x^4 + n_y^4 + n_z^4$ is an anisotropy parameter reflecting the fact that Clifford twirling enforces symmetry under permutations of the Pauli axes, rather than under arbitrary rotations. Taking these scalar products to second order in error angle θ_a , we find that the anisotropic term may be ignored for small errors,

$$\alpha_L = \alpha_R = \alpha_{\text{as}} = 1 - \theta_a^2/3, \quad (5.10)$$

$$\alpha_{\text{tr}} = 1, \quad (5.11)$$

$$\alpha_{\text{dt}} = \alpha_{\text{so}} = 1 - \theta_a^2. \quad (5.12)$$

Reintroducing the stochastic application of common unitaries, given by equation 5.5, we find,

$$\tilde{\Lambda}^{\mathcal{S}}(\mathcal{E}_{\mathcal{A}}) = \sum_{a \in \mathcal{A}} p_a \tilde{\Lambda}^{\mathcal{S}}(\mathcal{E}_a). \quad (5.13)$$

As such, the dynamics for the twirled channel of an ensemble of common unitaries is entirely described by the error angle magnitudes, with alpha parameters,

$$\alpha_i = \sum_{a \in \mathcal{A}} p_a \alpha_a. \quad (5.14)$$

This definition leads to a convenient (and conforming to previous RB work) definition for the correlated error per Clifford given by,

$$\varepsilon_{\text{corr}} = \frac{1}{6} \sum_{a \in \mathcal{A}} p_a \theta_a^2, \quad (5.15)$$

where the correlated error per Clifford, $\varepsilon_{\text{corr}}$, quantifies the single-qubit error contribution due to *common noise*. Comparing to equation 5.4, the general ensemble of common single-qubit rotations channel, after twirling over $Cl_1^{\otimes 2}$, is given by the diagonal matrix in the Schur basis,

$$\tilde{\Lambda}^{\mathcal{S}}(\mathcal{E}_{\mathcal{A}}) = \text{diag}\left(I_2, (1 - 2\varepsilon_{\text{corr}})I_9, (1 - 6\varepsilon_{\text{corr}})I_5\right). \quad (5.16)$$

5.3.4 Local depolarising noise

Ultimately, environmental noise will not perfectly couple symmetrically to both qubits, and so independent local single-qubit error channels must also be considered. Assuming an independent noise channel acting on each qubit, Clifford twirling leads to a fully depolarising noise channel of the form,

$$\mathcal{E}_{dp,i}(\hat{\rho}) = (1 - 2\varepsilon_i)\hat{\rho} + \varepsilon_i I_2 \text{Tr}(\hat{\rho}), \quad (5.17)$$

where $\hat{\rho}$ is the density matrix for a single qubit, and the subscript i signifies which qubit we are considering. For the “left” and “right” qubits we describe the decay rates by ε_L and ε_R , respectively. Parallel local depolarisation is given by the channel, $\mathcal{E}_{dp} = \mathcal{E}_{dp,L} \otimes \mathcal{I} + \mathcal{I} \otimes \mathcal{E}_{dp,R}$, with $\mathcal{I}$ being the identity channel on a single qubit. Comparing to equation 5.4, this channel may be expressed in PTM form as the diagonal matrix,

$$\Lambda^{\mathcal{P}}(\mathcal{E}_{dp}) = \text{diag}\left(1, (1 - 2\varepsilon_L)I_3, (1 - 2\varepsilon_R)I_3, (1 - 2\varepsilon_L)(1 - 2\varepsilon_R)I_9\right). \quad (5.18)$$

We find that $\Lambda(\mathcal{E}_{dp})$ already commutes with the representations $r(g)$ for all elements g in the single qubit Clifford group, as such the channel already obeys the symmetries

enforced by twirling over $Cl_1^{\otimes 2}$ (see equation A.38). Transforming to the Schur basis, we extract α parameters,

$$\alpha_L = 1 - 2\varepsilon_L, \quad (5.19)$$

$$\alpha_R = 1 - 2\varepsilon_R, \quad (5.20)$$

$$\alpha_{\text{tr}} = \alpha_{\text{dt}} = \alpha_{\text{so}} = \alpha_{\text{as}} = 1 - 2(\varepsilon_L + \varepsilon_R). \quad (5.21)$$

5.3.5 Combined model

As both the depolarisation channels and the twirled common single-qubit rotation channel are diagonal, they may be combined by simple composition,

$$\Lambda_T(\varepsilon_L, \varepsilon_R, \varepsilon_{\text{corr}}) = \Lambda_{dp}(\varepsilon_L, \varepsilon_R) \tilde{\Lambda}_{\text{corr}}(\varepsilon_{\text{corr}}). \quad (5.22)$$

This results in scalar parameters to first order in the error terms,

$$\alpha_{\text{tr}} = 1 - 2(\varepsilon_L + \varepsilon_R), \quad (5.23)$$

$$\alpha_L = 1 - 2(\varepsilon_L + \varepsilon_{\text{corr}}), \quad (5.24)$$

$$\alpha_R = 1 - 2(\varepsilon_R + \varepsilon_{\text{corr}}), \quad (5.25)$$

$$\alpha_{\text{as}} = 1 - 2(\varepsilon_L + \varepsilon_R + \varepsilon_{\text{corr}}), \quad (5.26)$$

$$\alpha_{\text{dt}} = \alpha_{\text{so}} = 1 - 2(\varepsilon_L + \varepsilon_R + 3\varepsilon_{\text{corr}}). \quad (5.27)$$

To reproduce the decay dynamics observed in the RB sequence, we must consider an initial prepared state, and the final measurement basis. For the RB sequence given in Figure 5.1, the initial state is prepared as $|\downarrow\downarrow\rangle$, and the final correlated populations are measured in the computational eigenstates, $\{|\downarrow\downarrow\rangle \leftrightarrow SS, |\downarrow\uparrow\rangle \leftrightarrow SF, |\uparrow\downarrow\rangle \leftrightarrow FS, |\uparrow\uparrow\rangle \leftrightarrow FF\}$, where S and F denote “success” and “failure” of recovering the target state outcome for each qubit. The correlated populations may be given analytically using the derived Liouville-space superoperators, and by the vectorized initial and final states,

$$\mathcal{P}_{LR} = \langle\langle L, R | \Lambda_T^m | \downarrow\downarrow \rangle\rangle, \quad (5.28)$$

where L, R denote the successful recovery of the target final state for the “left” and “right” qubit respectively, and m is the number of Clifford gates in the

applied sequence. Using this equation, and the explicit transfer matrices, the final populations decays are found,

$$\mathcal{P}_{SS} = \frac{1}{4} + \frac{1}{4}(\alpha_L^m + \alpha_R^m) + \frac{1}{6}\alpha_{dt}^m + \frac{1}{12}\alpha_{tr}^m, \quad (5.29)$$

$$\mathcal{P}_{SF} = \frac{1}{4} + \frac{1}{4}(\alpha_L^m - \alpha_R^m) - \frac{1}{6}\alpha_{dt}^m - \frac{1}{12}\alpha_{tr}^m, \quad (5.30)$$

$$\mathcal{P}_{SF} = \frac{1}{4} - \frac{1}{4}(\alpha_L^m - \alpha_R^m) - \frac{1}{6}\alpha_{dt}^m - \frac{1}{12}\alpha_{tr}^m, \quad (5.31)$$

$$\mathcal{P}_{FF} = \frac{1}{4} - \frac{1}{4}(\alpha_L^m + \alpha_R^m) + \frac{1}{6}\alpha_{dt}^m + \frac{1}{12}\alpha_{tr}^m. \quad (5.32)$$

Considering the final state of the single qubit subspace, as probed by standard RB, we find the expected single exponential decay,

$$\mathcal{P}_S^{(L)} = \mathcal{P}_{SS} + \mathcal{P}_{SF} = \frac{1}{2} (1 + (1 - 2(\varepsilon_L + \varepsilon_{\text{corr}}))^m), \quad (5.33)$$

where the superscript (L) refers to considering the left qubit only. Comparing to standard single qubit RB parameterisation,

$$\mathcal{P}_S = \frac{1}{2} (1 + (1 - 2\varepsilon_c)^m), \quad (5.34)$$

where ε_c is the “error-per-Clifford”, we find the observed decay of each qubit is now separable into a local, and correlated error component,

$$\varepsilon_c^{(L)} = \varepsilon_L + \varepsilon_{\text{corr}}. \quad (5.35)$$

Asymptotes

Here we consider the asymptotic structure of the correlated decay curves given in 5.29. Under dominant local depolarisation, the two-qubit state will approach the maximally mixed state, $\hat{\rho} \rightarrow I_4/4$, indicated by the asymptote $\mathcal{P}_{SS}|_{m \rightarrow \infty} = 1/4$, as seen in Figure 5.2 (a). Under dominant *common noise*, with decay profile seen in (b), an asymptote of $\mathcal{P}_{SS}|_{m \rightarrow \infty} = 1/3$ is approached due to the invariance of the trace sector, α_{tr} . Panel (c) shows the case where both local depolarising and *common noise* is present. When $\varepsilon_{\text{corr}} > \varepsilon_L = \varepsilon_R$ a clear intermediate asymptote is seen before a final decay to the fully depolarised state. This intermediate asymptote is the central experimental signature of correlated noise exploited throughout this




Figure 5.2: The expected multi-exponential decay profiles for two-qubit correlated success and failures recovered during k -copy RB. Panels (a), (b), and (c) display the structured decays for local depolarising only, *common noise* only, and the combined case, respectively. Dashed lines indicate the expected asymptotes for depolarising (blue) and for *common noise* (red).

chapter, and is explored experimentally in § 5.6.2 below.

The mixed state reached after the application of *common noise* to the initial state $|\downarrow\downarrow\rangle$ can be understood by directly inspecting the space of possible pure two qubit states reached through common only rotations. The general pure state after common rotations is the separable identical two-qubit states,

$$|\Psi\rangle = |\psi\rangle \otimes |\psi\rangle, \quad (5.36)$$

with $|\psi\rangle = \cos\left(\frac{\theta}{2}\right)|\downarrow\rangle + e^{i\phi}\sin\left(\frac{\theta}{2}\right)|\uparrow\rangle$ for a single qubit general state. Depolarising into the symmetric separable two-qubit states is equivalent to averaging over the full Bloch sphere,

$$\begin{aligned} \hat{\rho}_{\text{sym}} &= \frac{1}{4\pi} \int_0^\pi \int_0^{2\pi} (\sin(\theta) |\Psi\rangle \langle\Psi|) d\phi d\theta \\ &= \frac{1}{4} \left(II + \frac{1}{3}(XX + YY + ZZ) \right), \\ &= \frac{1}{3} |\downarrow\downarrow\rangle \langle\downarrow\downarrow| + \frac{1}{3} |\uparrow\uparrow\rangle \langle\uparrow\uparrow| + \frac{1}{3} |\Psi^+\rangle \langle\Psi^+|, \end{aligned} \quad (5.37)$$

where the third term is the Bell state $|\Psi^+\rangle = \frac{1}{\sqrt{2}}(|\downarrow\uparrow\rangle + |\uparrow\downarrow\rangle)$. The presence of this Bell-state component reflects residual coherence induced between the two qubits by the shared rotation error, and motivates the alternative strategy of a coherence-sensitive probe developed in § 5.4.




Figure 5.3: *k*-copy RB applies identical single-qubit Clifford sequences in parallel to all qubits. (a) shows the circuit diagram for a single RB random sequence with random unitary errors included explicitly. (b-d) illustrate the distinction between uncorrelated and *common noise* (a type of correlated noise) in the final-state trajectories.

Schematic explanation

The distinction between correlated and uncorrelated noise is illustrated schematically in Figure 5.3. We consider a particularly simple unitary noise mechanism to highlight the qualitative difference of information gained between standard and *k*-copy RB. After each Clifford gate we consider an independent unitary rotation errors on each qubit, and a common unitary rotation error on qubits $i = 3$ and $i = 4$. Each unitary rotation axis is randomly drawn (isotropically from $\text{SO}(3)$), and the angle of the rotation is small. The error channels are explicitly included in the circuit diagram in (a). The state of the qubits after the twirling procedure, and before measurement, are shown in (b). We see that all qubits are no longer in the eigenstate initial eigenstate (depicted as qubit pointing up) due to the unitary noise. Panels (c) and (d) display the final vectors resulting from 5 random sequences. All qubits here exhibit identical single-qubit RB error rates, which are probed by standard RB, and shown here through the identical depolarised Bloch vectors (black arrows). However, looking at pairwise correlations, we see: independent noise produces uncorrelated sequence-to-sequence deviations, whereas common noise produces shared deviations

between qubits acted on by the same random Clifford sequence. The parallel twirling during *k*-copy RB preserves exactly this symmetry information that would be averaged away under independently sampled local Clifford sequences.

5.3.6 State preparation errors

Similar to previous work on subspace leakage benchmarking [124], state preparation errors can lead to non-trivial modifications of the correlated population decay curves. First, we consider independent state preparation on two qubits. We model independent errors as probabilistic bit flips. To first order in error, on a single qubit we have,

$$\hat{\rho}_{sp}^{(i)} = (1 - \varepsilon_{sp}^{(i)}) |\downarrow\rangle \langle\downarrow| + \varepsilon_{sp}^{(i)} |\uparrow\rangle \langle\uparrow|, \quad (5.38)$$

where the superscript refers to which qubit, L , or R , the preparation error is relevant to. This error may be modelled by $\mathcal{E}_{sp}^{(i)}$, a probabilistic X channel applied to an arbitrary state,

$$\mathcal{E}_{sp}^{(i)}(\hat{\rho}) = (1 - \varepsilon_{sp}^{(i)}) \hat{\rho} + \varepsilon_{sp}^{(i)} X \hat{\rho} X. \quad (5.39)$$

Transforming to the Liouville form, denoted by $\Lambda_{sp}^{(i)}$, this error is simply incorporated into the population dynamics given by equation 5.28, by acting on the initial state,

$$\mathcal{P}_{LR} = \langle\langle L, R | \Lambda_T^m \Lambda_{sp}^{(L)} \Lambda_{sp}^{(R)} | \downarrow\downarrow \rangle\rangle. \quad (5.40)$$

In this picture, the state preparation errors scale only the overlap of the initial state with the sub-representations. The overlaps are multiplied by,

$$A_{sp}^{(i)} = (1 - 2\varepsilon_{sp}^{(i)}), \quad (5.41)$$

if the qubit is involved non-trivially in the subspace projector, i.e. for the left and right qubit sectors $A_{sp} = A_{sp}^{(i)}$. Relevant to this work, the correlated survival populations with state preparation errors included are,

$$\begin{aligned} \mathcal{P}_{SS} = & \frac{1}{4} + \frac{A_{sp}^{(L)}}{4} \alpha_L^m + \frac{A_{sp}^{(R)}}{4} \alpha_R^m \\ & + \frac{A_{sp}^{(L)} A_{sp}^{(R)}}{6} \alpha_{dt}^m + \frac{A_{sp}^{(L)} A_{sp}^{(R)}}{12} \alpha_{tr}^m. \end{aligned} \quad (5.42)$$

Similar forms exist for populations, $\mathcal{P}_{SF}$, $\mathcal{P}_{FS}$, and $\mathcal{P}_{FF}$. Significantly, this state preparation error does not add projection on the antisymmetric or symmetric off-diagonal sectors, which would introduce another exponential component into the observed decay curves.

Correlated state preparation errors may be considered by adapting equation 5.39 to a two-qubit stochastic error channel, such as correlated spin-flips given by

$$\mathcal{E}_{sp}^{\text{corr}}(\hat{\rho}) = (1 - \varepsilon_{sp}^{\text{corr}})\hat{\rho} + \varepsilon_{sp}^{\text{corr}}(X \otimes X)\hat{\rho}(X \otimes X). \quad (5.43)$$

The subspaces involving both qubits non-trivially are now unaffected by the state preparation error, shown explicitly by,

$$\mathcal{P}_{SS} = \frac{1}{4} + \frac{A_{sp}^{\text{corr}}}{4}\alpha_L^m + \frac{A_{sp}^{\text{corr}}}{4}\alpha_R^m + \frac{1}{6}\alpha_{\text{dt}}^m + \frac{1}{12}\alpha_{\text{tr}}^m, \quad (5.44)$$

where $A_{sp}^{\text{corr}} = 1 - 2\varepsilon_{sp}^{\text{corr}}$. These correlations thus appear similar to *common noise* errors in the invariance of α_{tr} to this channel. Therefore correlated state prep errors can lead to a systematic bias in quantification of correlated error per Clifford if not accounted for in the fit by an initial state preparation error offset.

5.4 Parity contrast: an independent probe of coherence

Correlated errors lead to multi-qubit correlations. In the two-qubit case these correlations are clearly visible in the mixed state $\hat{\rho}_{\text{sym}}$, given by equation 5.37, which is approached when *common noise* is dominant. Here, we propose directly measuring induced two-qubit correlations to independently verify the coherent structure seen in $\hat{\rho}_{\text{sym}}$. In § 5.5 we shall see that such a measurement gives an improvement in the uncertainty of the estimated error components.

To quantify the two-qubit correlations, a partial state-tomography method based on [97] is utilised. This method is common in fidelity estimation of two-qubit entangling gates (further details given in § 4.3.2). The protocol quantifies the overlap of the “coherence-terms” between the final mixed state and a chosen Bell state. To quantify the overlap to the mixed state $\hat{\rho}_{\text{sym}}$, the Bell state $|\Psi^+\rangle$ is chosen.




Figure 5.4: Probing coherence of the mixed state after *k*-copy RB. Panel (a) shows a circuit that first twirls the state over $Cl_1^{\otimes 2}$, and then applies differential $\pi/2$ -pulses with variable phase, ϕ , to recover parity oscillations. Panel (b) shows the expected contrast of the oscillations as a function of sequence length when both local depolarising and *common noise* is present. The dashed lines indicate expected asymptotes for depolarising (blue) and for *common noise* (red).

The circuit for probing this coherence is shown in Figure 5.4 (a): differential rotations of the form $\hat{R}_{+\phi}(\pi/2) \otimes \hat{R}_{-\phi}(\pi/2)$, with rotation axis $\hat{\sigma}_\phi = \cos(\phi)\hat{\sigma}_x + \sin(\phi)\hat{\sigma}_y$, are applied prior to measurement, and the contrast of the resulting fringe (as ϕ is varied) is extracted.

Figure 5.4(b) show the expected contrast of parity measurements as a function of Clifford sequence length. The analytic model of the parity contrast can be simply calculated by extending equation 5.28 to include the final differential rotation, in transfer matrix form, Λ_Π ,

$$\begin{aligned} \Pi(\phi) &= \langle \langle \pi | \Lambda_\Pi(\phi) \Lambda_T^m | \downarrow\downarrow \rangle \rangle, \\ &= C \cos(2\phi), \end{aligned} \tag{5.45}$$

where Π is the parity signal, $|\pi\rangle = |\downarrow\downarrow\rangle + |\uparrow\uparrow\rangle - |\downarrow\uparrow\rangle - |\uparrow\downarrow\rangle$, C is the parity contrast, and perfect execution of the probe $\hat{R}_\phi(\pi/2)$ is assumed. Evaluating $\Pi(\phi)$ for *common* and depolarising noise, we find the parity contrast follows,

$$C = \frac{1}{3}(\alpha_{\text{tr}}^m - \alpha_{\text{dt}}^m), \tag{5.46}$$

with α_{tr} and α_{dt} given by equation 5.23. This method is demonstrated experimentally in § 5.6.3 below.

Furthermore, we may now isolate some α decay parameters by combining the




Figure 5.5: The performance of the *k*-copy fitting as a function of maximum Clifford sequence length (probe depth). Probe depth, γ , is relative to the characteristic decay length for the slowest mode, α_{tr}^m . Panel (a) uses the Cramér-Rao derived relative uncertainty for each extracted ε , while (b) shows the minimum eigenvalue (higher is better) of the Fisher information matrix, with associated eigenvectors highlighted.

contrast measurements with the computational state measurements made in *k*-copy RB. Namely, if we define the computational parity, Π_Z , as

$$\begin{aligned} \Pi_Z &= \mathcal{P}_{SS} + \mathcal{P}_{FF} - \mathcal{P}_{SF} - \mathcal{P}_{FS} \\ &= \frac{1}{3}(2\alpha_{\text{dt}}^m + \alpha_{\text{tr}}^m), \end{aligned} \quad (5.47)$$

then in combination with C , we may isolate single exponential decays originating from the trace and diagonal traceless sectors. Specifically the combinations

$$\begin{aligned} \Pi_Z - C &= \alpha_{\text{dt}}^m, \\ \Pi_Z + 2C &= \alpha_{\text{tr}}^m. \end{aligned} \quad (5.48)$$

5.5 Sampling strategy

This proposed RB method is simply scalable to multiple qubits by parallel application of the single-qubit Clifford gates. However, it is non-obvious if this method can achieve acceptable uncertainties of the estimated error per Clifford parameters, ε , in a reasonable number of experimental executions (samples) of the random circuits. In this section we verify the sample efficiency for a range of underlying ε affecting a pair of qubits. Furthermore, we determine a strategy for setting the sequence

lengths, m , that is sample efficient.

The Fisher information [133], $\mathcal{I}$, quantifies what we can learn from random observable, X , about a set of underlying unknown parameters, θ , that models the probability of observing X . It is particularly informative to consider the Cramér-Rao bound, which states that the inverse of $\mathcal{I}$ is a lower bound for the variance of any unbiased estimator of θ ,

$$\text{Cov}(\theta) \gtrsim \mathcal{I}(\theta)^{-1}. \quad (5.49)$$

The Fisher matrix is defined as,

$$\mathcal{I}_{ij}(\theta) = \mathbb{E}_X \left[\frac{\partial}{\partial \theta_i} \log f(X; \theta) \frac{\partial}{\partial \theta_j} \log f(X; \theta) \right], \quad (5.50)$$

where $f(X; \theta)$ is the likelihood function, and θ is a set of parameters. Here, we are interested in a multinomial probability function, with each shot resulting in a categorical outcome $X \in \{SS, SF, FS, FF\}$. The population functions given by equations 5.29, are then the probability of observing $X = x$ for a given set of parameters, $\theta = \{\varepsilon_L, \varepsilon_R, \varepsilon_{\text{corr}}\}$, yielding,

$$f(X = x; \theta) = P_x(\theta). \quad (5.51)$$

We now find the Fisher matrix to be,

$$\mathcal{I}_{ij}(\theta) = \mathbb{E}_X \left[\frac{1}{P_X^2} \frac{\partial P_X}{\partial \theta_i} \frac{\partial P_X}{\partial \theta_j} \right]. \quad (5.52)$$

For a discrete random variable, the expectation is,

$$\mathbb{E}[g(X)] = \sum_x P_x g(x), \quad (5.53)$$

giving our final expression for the single-shot multinomial Fisher matrix,

$$\mathcal{I}_{ij} = \sum_{x \in \{SS, SF, FS, FF\}} \frac{1}{P_x} \frac{\partial P_x}{\partial \theta_i} \frac{\partial P_x}{\partial \theta_j}. \quad (5.54)$$

The information gained in any one shot depends on the underlying error magnitudes, as well as the chosen sequence length, m . As the correlated profiles, P_x , appear as

multi-exponentials decays dependent on common parameters, θ , we must choose a range of probe sequence lengths that yield large information over all the characteristic decays. As the characteristic decays are unknown, a good heuristic strategy for fitting an unknown exponential is logarithmic sampling. However, the total extent of sequence lengths to test is non-obvious. It would be experimentally simple if the initial gradient of the correlated profiles successfully estimated the parameter θ , however, expanding the populations around $m = 0$, for small errors, we find that both purely uncorrelated and purely *common* noise decays are indistinguishable,

$$\frac{d\mathcal{P}_{SS,LR}}{dm}\big|_{m=0} \approx \frac{d\mathcal{P}_{SS,\text{corr}}}{dm}\big|_{m=0} \approx -2\varepsilon, \quad (5.55)$$

where $\mathcal{P}_{SS,LR}$ has errors $\varepsilon_L = \varepsilon_R = \varepsilon$ and $\varepsilon_{\text{corr}} = 0$, while $\mathcal{P}_{SS,\text{corr}}$ has errors $\varepsilon_L = \varepsilon_R = 0$ and $\varepsilon_{\text{corr}} = \varepsilon$. This suggests that we must collect data near the asymptotic structure of the decays to distinguish uncorrelated and *common* noise sources. To verify this, we set experimentally reasonable parameters for an underlying *common noise* dominated decay, and vary the maximum sequence length, m_{max} , while observing the expected uncertainty given by the Cramér-Rao bound of the Fisher information. We normalise by the magnitude of the parameter estimated to give the relative uncertainty, $\frac{\sigma_\varepsilon}{\varepsilon}$. We also find the minimum eigenvalue of the Fisher matrix, λ_{min} , which corresponds to how accurately the least constrained vector in parameter space can be measured. Figure 5.5 uses $N = 5000$ samples per sequence length (e.g. $s = 50$ shots, and $r = 100$ random sequences), and 10 total sequence lengths sampled logarithmically from $m_{\text{min}} > 2$ to $m_{\text{max}} = \gamma\tau$. The scalar γ is the relative probe depth and τ is the characteristic decay for the slow, α_{tr} mode,

$$\tau = \frac{-1}{\log(\alpha_{\text{tr}})} = \frac{-1}{\log(1 - 2(\varepsilon_L + \varepsilon_R))}. \quad (5.56)$$

We set error magnitudes of $\varepsilon_L = 1.0 \times 10^{-3}$, $\varepsilon_R = 1.2 \times 10^{-3}$, $\varepsilon_{\text{corr}} = 1.0 \times 10^{-2}$, which are similar magnitudes we observe experimentally (see § 5.6). It can be seen that the relative uncertainty for all θ are minimised when $\gamma \approx 1$. This length also corresponds to maximising λ_{min} — lowering the cost of estimating the least constrained vector of parameters, where the vector is written in the basis




Figure 5.6: (a) The Cramér-Rao derived relative uncertainty for a range of correlated and uncorrelated error ratios. The thin lines correspond to expected uncertainties using $\langle Z \rangle$ measurements from k -copy RB, while thick lines use both $\langle Z \rangle$ and parity measurements, as described in § 5.4. (b-c) The required number of shots, N , per sequence length, m , for unity fractional uncertainty, for a range of correlated and uncorrelated errors. In (b), only $\langle Z \rangle$ measurements are used, while in (c) both $\langle Z \rangle$ and parity measurements are used. A contour is shown for 5000 samples, which is typically used in our experiments.

$[\hat{\varepsilon}_L, \hat{\varepsilon}_R, \hat{\varepsilon}_{\text{corr}}]$. It is interesting to note that the least constrained parameter vector, when $\gamma \ll 1$, roughly corresponds to $\vec{\nu} \approx \frac{1}{\sqrt{3}}[1, 1, -1]$, confirming our suspicion that *common* and depolarising noise is indistinguishable at small m . In the case that $\gamma \gg 1$, we find $\vec{\nu} \approx \frac{1}{\sqrt{2}}[1, -1, 0]$, corresponding to the depolarisation rates on the left and right qubit are tending to be indistinguishable, however we note that even over a couple of orders of magnitude, the distinguishability of ε_L and ε_R does not dramatically drop. Here, we draw the conclusion that a good strategy for efficient information per shot, is to set m_{max} to the characteristic decay of the slowest mode.

We now explore the ability for k -copy RB to fit θ for a range of underlying error magnitudes. Figure 5.6(a) shows the relative uncertainty as a function of the error ratio, defined here as $\varepsilon_{\text{corr}}/(\varepsilon_L + \varepsilon_R)$. We fix $\varepsilon_L = \varepsilon_R = 1 \times 10^{-3}$, and again set $N = 5000$ samples per sequence length m , with $m_{\text{max}} = \tau$, and vary $\varepsilon_{\text{corr}}$ to vary the error ratio. The thin blue and red line show the expected relative uncertainty for ε_L and $\varepsilon_{\text{corr}}$ respectively. We can see that for the case that local noise dominates, $\varepsilon_L \gg \varepsilon_{\text{corr}}$, the sufficient information for accurately estimating the correlated noise is not available. This limitation was found to be experimentally important, and motivates the injection of fixed *common noise* offsets in § 5.6.4, and 5.6.5. Another

strategy to alleviate this problem is by including data from the coherence probes described in § 5.4. The thick blue and red lines correspond to measurements where $N = 2500$ samples are collected by the typical $\langle Z \rangle$ *k-copy* protocol, and $N = 2500$ samples are using the parity protocol, denoted as Π . We find that inclusion of the parity samples dramatically reduce the relative uncertainty of the fitted parameters when $\varepsilon_L \gg \varepsilon_{\text{corr}}$. Panels (b) and (c) show the required number of samples, N , for a relative uncertainty $\frac{\sigma_\varepsilon}{\varepsilon} = 1$, for a range of local and *common* noise magnitudes. We highlight the $N = 5000$ sample contour as this may be collected in reasonable time on our apparatus. Again, we see that including measurements using the parity protocol dramatically reduces the required number of samples in the regime of $\varepsilon_L \gg \varepsilon_{\text{corr}}$ (top left corner of plots).

Here we have shown that the Fisher information gathered in an experimentally realistic *k-copy* protocol is sufficient to accurately determine the underlying error per Cliffords, $\theta = \{\varepsilon_L, \varepsilon_R, \varepsilon_{\text{corr}}\}$, over many orders of magnitude of local and *common* errors. We shall show experimentally in § 5.6.4 that in the presence of experimental imperfections, and the fitting protocol (Levenberg-Marquadt), whose effects are not considered in this analysis, the proposed protocol still sufficiently recovers the magnitudes of the expected error per Cliffords with known noise injection.

5.6 Experimental Results

The results presented here constitute the first project-data output from the newly commissioned *FastGates* apparatus.

Experimentally demonstrating the *k-copy* RB protocol verifies the underlying assumptions on the error models we considered, and the efficacy of twirling on these channels. It provides evidence that this method is indeed simply executed on a physical platform and able to accurately estimate independent depolarisation and correlated errors. To test this accuracy we estimate ε when injecting errors of known magnitudes via random phase modulation, systematic detuning offsets, and intentional Stark shifts using the single-addressing system.

5.6.1 Physical implementation

We demonstrate *k*-copy RB using up to eight trapped $^{40}\text{Ca}^+$ ions separated by an average spacing of $\approx 5 \mu\text{m}$. The qubit is encoded in the optical transition $|4S_{1/2}, m_j = -\frac{1}{2}\rangle \leftrightarrow |3D_{5/2}, m_j = -\frac{5}{2}\rangle$ in a quantising field of $\approx 5 \text{ G}$. Single-qubit rotations are driven by 729-nm light, with details given in § 4.1.3. The two-qubit coherence and Stark shift measurements use the individual addressing system, detailed in § 3.6.5, while all other measurements use the global elliptical beam, detailed in § 3.6.6.

The single-qubit Clifford gate set is decomposed into native $\pi/2$ single-qubit rotations of the form $\hat{R}_x(\pi/2)$, $\hat{R}_y(\pi/2)$, and $\hat{R}_z(\pi/2)$ (defined in § 2.2.2). $\hat{R}_x$ and $\hat{R}_y$ rotations are performed natively on resonance with the qubit transition, with relative phase set by modulating the 729-nm beam via an acousto-optic modulator (AOM), while $\hat{R}_z$ rotations are performed virtually by updating the phase of subsequent $\hat{R}_\phi$ rotations. Two decomposition strategies are used in this work. Only the average numbers of physical $\hat{R}_x(\pi/2)$ and $\hat{R}_y(\pi/2)$ -gates that compose the Clifford are relevant here. The first uses an average of $\kappa = 2.17 \hat{R}_\phi(\pi/2)$ -gates [134]⁴, while the second uses $\kappa = 1.167$ gates⁵.

We begin by reproducing the multi-exponential decays seen in simulation, Figure 5.2, followed by parity measurements to observe the coherence of the intermediate symmetric state described in § 5.4. We then present data with known injected errors through phase and detuning offsets to show the ability for *k*-copy RB to discern correlated and uncorrelated errors for a range of magnitudes. We finish this section with evidence for scaling the protocol to more than two qubits. We show both data for three-qubit decay profiles, and pairwise analysis of an eight-ion chain with dominant correlated errors due to global drive field inhomogeneity.




Figure 5.7: Correlated two-qubit probabilities after *k*-copy RB. Panels (a,b) show the single-qubit success probabilities, which decay with a single exponential (solid line) toward $\mathcal{P}_S|_{m \rightarrow \infty} = 1/2$. Panel (c) shows the correlated two-qubit populations, which exhibit an intermediate $\mathcal{P}_{SS} = 1/3$ asymptote before relaxing toward $\mathcal{P}_{SS}|_{m \rightarrow \infty} = 1/4$. The extracted errors from the single- and two-qubit decays are self-consistent with $\varepsilon_c^{(i)} = \varepsilon_{\text{corr}} + \varepsilon_i$.

5.6.2 Z-basis measurements

The *k*-copy RB protocol, with circuit diagram shown in Figure 5.1, is executed on two qubits using a global field. To ensure a visually clear intermediate asymptote is seen, we set a fixed 10 kHz frequency detuning of the 729-nm laser compared to the bare qubit. As single qubits gates are calibrated for on-resonance pulses, this detuning will be a large *common noise* source. The Clifford sequence length is log-sampled up to $m = 3200$ total gates to capture both the intermediate asymptote due to this common unitary error, and the final decay to the fully depolarised state. The recovered multi-exponential decay from the correlated success and failure rates of the two qubits is shown in Figure 5.7(c). Each solid point is inferred through $r = 50$ randomisations and $s = 50$ shots per randomisation. For each random sequence j , the population $\mathcal{P}_{SS,j}$ is estimated from the s binary measurement outcomes, $\mathcal{P}_{SS,j} = \frac{1}{s} \sum_{i=1}^s X_{ij}$, where $X_{ij} = 1$ if the i th shot is measured in state SS , and $X_{ij} = 0$ otherwise. The recovered population is then the mean over the r random sequences, $\mathcal{P}_{SS} = \frac{1}{r} \sum_{j=1}^r \mathcal{P}_{SS,j}$. Error bars show the standard error of

⁴Decomposed into $\hat{R}_x(\pi/2), \hat{R}_y(\pi/2)$ gates.

⁵Decomposed into $\hat{R}_x(\pi/2), \hat{R}_y(\pi/2), \hat{R}_z(\pi/2)$ gates.

the mean across random sequences,

$$\text{SE}(\mathcal{P}_{SS}) = \sqrt{\frac{1}{r(r-1)} \sum_{j=1}^r (\mathcal{P}_{SS,j} - \mathcal{P}_{SS})^2}. \quad (5.57)$$

The correlated populations exhibit a clear intermediate asymptote at $m \approx 100$ Cliffords, corresponding to the initial state $|\downarrow\downarrow\rangle$ approaching $\hat{\rho}_{sym}$, equation ??, with $\mathcal{P}_{SS} \rightarrow \frac{1}{3}$ asymptote highlighted by the red dashed line. The dashed blue line shows the fully depolarised $\mathcal{P}_{SS} \rightarrow \frac{1}{4}$ asymptote, seen clearly by $m \approx 1000$ gates. The solid lines are found through modelling errors by a combination of local depolarising errors and *common noise*, given by equation 5.29 generalised to include independent state-preparation and measurement error (SPAM), ε_{sp} (see equation. 5.42 for example). The model is fit using the nonlinear Levenberg–Marquardt method, yielding extracted parameters $\varepsilon_{sp} = 0(5) \times 10^{-3}$, $\varepsilon_{\text{corr}} = 2.6(2) \times 10^{-2}$, $\varepsilon_L = 5(6) \times 10^{-4}$, and $\varepsilon_R = 0(6) \times 10^{-4}$, confirming that correlated errors dominate with this global field detuning.⁶

From the same shot data record we recover the single-qubit populations required for standard RB, as described in § 4.1.4. To match the naming of the relevant participating irreps, we label the two qubits L and R for left and right, respectively. The standard RB single exponential decay curves are shown in (a) and (b), with extracted single-qubit error per Cliffords of $\varepsilon_c^{(L)} = 2.6(2) \times 10^{-2}$ and $\varepsilon_c^{(R)} = 2.6(2) \times 10^{-2}$. As expected, *k*-copy RB provides evidence for this standard error per Clifford to be decomposed to local and correlated components, given by equation 5.35.

5.6.3 Coherence measurements

A partial-tomography parity measurement, given by the circuit seen in Figure 5.4(a), is used to independently confirm occupation of the trace sector, quantifying overlap of the final state with Bell state, $|\Psi^+\rangle$. As the final differential rotation pulse requires local addressing of the ions, the single addressing system is used. Each random

⁶The large relative uncertainty of both local depolarising rates is due to the large covariance between ε_L and ε_R . The fitting may be constrained to the sum of local errors, to find $\varepsilon_L + \varepsilon_R = 5(1) \times 10^{-4}$. However, if precision measurements of this local error is of interest, either r , and s could be increased, or the local error contributions could be isolated as discussed in § ??.




Figure 5.8: Parity contrast versus sequence length; (a) shows the standard Z measured correlated populations. (b) shows parity contrast — a signal proportional to occupation of the trace sector. Solid lines are multi-exponential decays given by equation 5.29 and equation 5.46 for (a) and (b) respectively. Dashed lines indicate the expected asymptotes. (c) and (d) show example parity fringes for sequences of $m = 25$ and $m = 1600$ Clifford gates respectively with solid lines given by equation 5.45.

sequence is repeated for varying final rotation angle, ϕ , to produce parity fringes, examples shown in (c) and (d) in Figure 5.8 for $m = 25$ and $m = 1600$ Clifford gates respectively. The contrast is extracted by equation 5.46, and a plot against sequence length is shown in (b). The contrast asymptote $C \rightarrow 1/3$ is observed for $m \in [10, 100]$, corresponding to the residual coherence of $\hat{\rho}_{\text{sym}}$, equation 5.37. The parity contrast data is fit (solid line) according to equation 5.46 to extract error per Cliffords of, $\varepsilon_{\text{corr}} = 2.1(2) \times 10^{-2}$ and $\varepsilon_L = \varepsilon_R = 3.9(3) \times 10^{-4}$. These extracted errors are of similar magnitude to those found by Z -measurements as demonstrated in the previous section. Importantly, this parity method cannot distinguish local errors on the left and right qubit, but rather reports the sum, $\varepsilon_L + \varepsilon_R$. We have explicitly assumed here that $\varepsilon_L = \varepsilon_R$. As discussed in § 5.5, it would be advantageous to combine both parity and Z -measurements to further improve the uncertainties in estimated parameters. Demonstrating this experimentally is left for future work.




Figure 5.9: Injecting correlated noise between two qubits via common drive phase noise. Points correspond to fitted error per Clifford from full *k*-copy RB experiments. The solid red line is a linear fit between measured and injected error.

5.6.4 Controlled noise injection

To verify the ability for *k*-copy RB to discern correlated and uncorrelated errors, we compare the magnitude of extracted errors to known injected errors. We limit the injected error method to those consisting of noise on common drive fields, and local depolarising type noise. Both of these noise models can be induced entirely through single qubit rotations, which are trivial to implement on our device.

Laser phase jitter

We first inject correlated errors via phase noise on a common drive field. This noise mechanism is prevalent in many platforms that use a common oscillator, such as laser, microwave, or radio-frequency (RF) fields to resonantly drive qubit transitions.

The injected error magnitude is inferred through considering the ideal and noisy unitary channel. A unitary qubit rotation is given by,

$$\hat{U} = e^{-i\frac{\theta}{2}\vec{n}\cdot\vec{\sigma}}, \quad (5.58)$$

where θ is the rotation angle, and $\vec{n} \cdot \vec{\sigma}$ is the rotation axis. We decompose the single-qubit Clifford gate set into a series of resonant $\theta = \pi/2$ qubit rotations around

axes $\hat{\sigma}_x$ and $\hat{\sigma}_z$. Practically, phase noise on either the qubit frequency or drive field will introduce some deviation of this rotation axis. We shall consider the effect of zero-mean phase noise on $\hat{\sigma}_x$ rotations, leading to an erroneous rotations axis,

$$\hat{\sigma}_{\delta\phi} = \cos(\delta\phi)\hat{\sigma}_x + \sin(\delta\phi)\hat{\sigma}_y, \quad (5.59)$$

where $\delta\phi$ is the standard deviation of the noise in radians. We may find the expected error of a single Clifford gate using,

$$\varepsilon_{\delta\phi} = \kappa (1 - \mathcal{F}(V, U)), \quad (5.60)$$

where κ is the average number of $\pi/2$ rotations in a Clifford gate, and $\mathcal{F}(V, U)$ is the average gate fidelity. The fidelity of two unitaries averaged over the Haar measure is given in Appendix equation A.30, reproduced here,

$$\mathcal{F}(V, U) = \frac{1}{d(d+1)} (|\text{Tr}(U^\dagger V)|^2 + d). \quad (A.30)$$

Considering $\theta = \pi/2$ rotations, we find,

$$\hat{U} = e^{-i\frac{\pi}{4}\hat{\sigma}_x} = \frac{1}{\sqrt{2}} (\hat{I} - i\hat{\sigma}_x), \quad (5.61)$$

$$\hat{V} = e^{-i\frac{\pi}{4}\hat{\sigma}_{\delta\phi}} = \frac{1}{\sqrt{2}} (\hat{I} - i\hat{\sigma}_{\delta\phi}). \quad (5.62)$$

The fidelity, given that dimension $d = 2$ for a single qubit, is found to be,

$$\mathcal{F}(V, U) = \frac{1}{6} ((1 + \cos(\delta\phi))^2 + 2). \quad (5.63)$$

For small phase noise standard deviations, $\delta\phi \ll 1$, we find an expected error per Clifford of,

$$\varepsilon_{\delta\phi} = \frac{\kappa}{3} (\delta\phi)^2. \quad (5.64)$$

A known phase error may be conveniently injected through phase offsets sampled from a normal distribution with zero mean and standard deviation $\delta\phi$ applied to the RF drive of a control AOM on the 729-nm global addressing path. The phase of the AOM RF is imprinted directly onto the 729-nm light leading to the expected drop in fidelity of the Clifford gates. In this experiment we use the Clifford decomposition

for which $\kappa = 2.17 \pi/2$ rotations. A full *k-copy* RB experiment is executed with logarithmic sampling up to $m = 512$ Clifford gates with $r = 50$ randomisations and $s = 50$ repeats. The resulting multi-exponential success and failure probability decay curves are fit to equation 5.29, again generalised to include state-preparation errors. The correlated and local error components are found for a range of injected phase noise standard deviations. The resulting measured errors are shown in Figure 5.9. The x -axis is the expected error injected $\varepsilon_{\delta\phi}$, while the y -axis is the extracted error magnitudes. The fit, shown by the solid red line, is a direct proportionality of injected error plus a baseline correlated error rate not due to phase injection,

$$\varepsilon_{\text{corr}} = \varepsilon_{\text{corr},0} + \varepsilon_{\delta\phi}. \quad (5.65)$$

The baseline correlated error is therefore the only parameter to fit. With $\varepsilon_{\text{corr},0} = 1.71(4) \times 10^{-3}$, we find good agreement between the model and data. The local depolarising error on both qubits remains relatively constant, with the left qubit having a larger local error compared to the right qubit. We believe this asymmetry in local error is due to the global 729-nm beam not being centrally aligned to both qubits. The uncertainty in extracted error is found through Levenberg-Marquadt nonlinear fit; it can be seen that the uncertainty in the depolarising errors increases substantially with the total correlated error. This may be explained by the available Fisher information gained by the chosen logarithmic sampling strategy, as explored in § 5.5. Figure 5.6 (a) supports that we expect an increase in depolarising parameter uncertainty for $\varepsilon_{\text{corr}} \gg (\varepsilon_L + \varepsilon_R)$.

Laser detuning scan

k-copy RB can isolate errors arising from shared control fields, making it useful both for device characterisation and for optimising global control parameters. A fixed detuning of the global drive, under the action of *k-copy* Clifford twirling, is well modelled by the common random unitary error channel described in § 5.3.3. We demonstrate this by scanning the 729-nm laser frequency used for single-qubit rotations away from the bare qubit resonance. As with the previous section, we

consider the fidelity of the error channel compared to the ideal $\pi/2$ $\hat{\sigma}_x$ single qubit rotation.

A coherent detuning error is written as,

$$V = \exp \left(-i \frac{\pi}{4} \left(\hat{\sigma}_x + \frac{\delta}{\Omega} \hat{\sigma}_z \right) \right), \quad (5.66)$$

where δ is the detuning from resonance, and Ω is the Rabi frequency. Assuming $\delta/\Omega \ll 1$, the fidelity is found to be,

$$\mathcal{F}(V, U) = 1 - \frac{1}{3} \left(\frac{\delta}{\Omega} \right)^2 + \mathcal{O} \left(\frac{\delta}{\Omega} \right)^4. \quad (5.67)$$

To convert this to an error contribution to a single qubit Clifford gate, we introduce the scalar κ_δ , to account for both the number of $\pi/2$ rotation in each Clifford gate, but also the duration and delays between pulses where phase tracking between qubit and laser also drift. The error contribution is then given by

$$\varepsilon_\delta = \kappa_\delta (1 - \mathcal{F}(V, U)) = \frac{\kappa_\delta}{3} \left(\frac{\delta}{\Omega} \right)^2 + \mathcal{O} \left(\frac{\delta}{\Omega} \right)^4. \quad (5.68)$$

Each data point is obtained from a complete *k*-copy RB experiment with 11 sequence lengths up to $m = 1024$, using $r = 50$ random sequences and $s = 50$ shots per sequence, with correlated populations fit to equation 5.29 to extract $\{\varepsilon_{\text{corr}}, \varepsilon_L, \varepsilon_R\}$ at each detuning. The upper panels of Figure 5.10 show example datasets for the *k*-copy RB sequences, while the lower panel shows the extracted error per Cliffords as a function of detuning. The correlated error is observed to increase quadratically with detuning, while the local depolarising errors remain approximately constant, consistent with detuning of a shared drive producing *common noise*. To fully parameterise the injected error, we include a background correlated error not originating from detuning, ε_0 , and a shift of the light-shifted qubit versus the bare qubit frequency, δ_0 . The solid red line is therefore fit to the model,

$$\varepsilon_{\text{corr}} = \varepsilon_0 + \frac{\kappa_\delta}{3} \left(\frac{\delta_L - \delta_0}{\Omega} \right)^2. \quad (5.69)$$

Here, $\Omega = 2\pi \times 100(1)$ kHz is the independently measured Rabi frequency. Fitting gives $\varepsilon_0 = 3.2(1) \times 10^{-3}$, $\kappa_\delta = 15.9(7)$, and $\delta_0 = 2\pi \times 0.80(5)$ kHz.




Figure 5.10: Correlated and single-qubit depolarising errors extracted as a function of shared-drive detuning. The correlated error (red points) rises quadratically with laser detuning, while the local depolarising errors (blue points) remain approximately constant. Error bars are the extracted standard deviations from a nonlinear Levenberg–Marquardt fit of the underlying correlated populations. The solid red line is fit according to equation 5.69. The red shaded region signifies the 68% confidence, generated through covariance-based parameter sampling.




Figure 5.11: Correlated and uncorrelated errors are injected via a single-addressed Stark-shift beam, and quantified through fitting the correlated success populations. The four addressing configurations are depicted visually above the plot. The wide beam signifies the on-resonance global drive which executes the parallel single-qubit rotations, while the narrow beams represent the off-resonant Stark shift beams addressing single-qubits (black circles). In (a) the errors are correlated between qubits, while panel (b) shows the extracted errors after independent errors are injected. Error bars are the extracted standard deviations from a nonlinear Levenberg–Marquardt fit of the underlying correlated populations. Panel (c) shows the calibration of Stark shift against total beam power used to infer the injected error.

Single addressed Stark shifts

Uncorrelated noise injection requires the introduction of noise with spatial inhomogeneity. This is implemented experimentally through an off-resonant single addressing beam inducing Stark-shifts between the on-resonant parallel single qubit rotations executing the *k*-copy twirling. The single addressing beam is detuned from the qubit resonance by -15 MHz, and the induced Stark shift with beam power, γ , is calibrated using an RPE Ramsey sequence (see § 4.1.3), with results shown in Figure 5.11 (c) for both the left and right qubit. From the frequency shift versus power, we extract $\gamma_L = 8.24(5) \times 10^7 \text{ Hz W}^{-1}$, $\gamma_R = 8.25(4) \times 10^7 \text{ Hz W}^{-1}$ for the left and right qubit respectively. The injected error, ε_{inj} , is then inferred by,

$$\varepsilon_{\text{inj}} = \frac{\kappa}{3} \phi^2, \quad (5.70)$$

where the phase error, given in radians is, $\phi = P\tau\gamma$, with P being the power in the single addressing beam, and τ the duration of the Stark-shift pulse. In these experiments, the duration of the Stark-shift pulse is fixed to $\tau = 20 \text{ }\mu\text{s}$, while the total power of the beam is varied to supply a known error.

Two experimental sequences were run, the first, shown in Figure 5.11(a), implements stochastic *common noise* through applying either two spots shifting both ions, or no Stark shift, with a 50 : 50 chance after each Clifford gate. The second experiment, shown in (b), applies local errors through a single spot either addressing the left or right qubit, again with 50 : 50 chance after each Clifford gate. These two arrangements are shown schematically in the upper panels. To ensure correlated errors dominate even with injected local errors, an additional constant correlated error offset of $\varepsilon_{\text{corr},0} \approx 2 \times 10^{-2}$ is applied through random phase fluctuations, as described in § 5.6.4. This reduces the number of samples required to accurately estimate the correlated and uncorrelated error contributions, which is of practical importance in collecting data in reasonable time (for reference, this data was collected over ~ 3 hours). It can be seen that $\varepsilon_{\text{corr}}$ increases only when injecting *common noise*, and remains constant when local errors are injected. Dashed lines are simple linear fits given by $\varepsilon_{\text{fit}} = \varepsilon_0 + \varepsilon_{\text{inj}}$, where ε_{fit} are the extracted errors, ε_0 is




Figure 5.12: Variation of $\hat{R}_x(\pi)$ -pulse gate duration along eight ion chain due to the Gaussian waist of the global addressing beam. Ion labels, i , are given on the upper axis, and equilibrium position given on the lower. The dotted line is fit using equation 5.71 to extract a waist radius of $27.7(9) \mu\text{m}$.




Figure 5.13: Correlated and local error per Cliffords measured for all qubit pairs after applying *k*-copy RB to an 8 ion chain. Noise is dominated by position dependent duration of $R_x(\pi/2)$ gates consistent with inhomogeneous illumination by the global drive. Panel (a) displays the extracted excess correlated errors, $\varepsilon_{\text{corr}}^{(i,j)}$ for each qubit pair, (i, j) . Panel (b) displays the sum of both independent depolarising errors, $(\varepsilon_L + \varepsilon_R)^{(i,j)}$. Panel (c) displays the ratio of excess correlated noise to the sum of local errors given in (b). This plot highlights the outer qubits are dominated by spatially correlated errors.

some residual error baseline. These straight line fits show good agreement between measured and injected error when accounting for some fit baseline error rates.

5.6.5 Scaling to a multi-ion register

Pairwise *k*-copy RB extends simply to > 2 qubits: correlated populations are fit to equation 5.29 independently for each qubit pair. For a register of k qubits, all $k(k-1)/2$ pairwise outcomes are extracted from the same shot record — no additional trials on the quantum register are required, although classical post-processing scales with the number of pair combinations. The protocol itself requires only synchronised application of identical single-qubit Clifford sequences and spatially resolved readout.

This scaling is demonstrated on an eight-ion chain, with parallel single-qubit rotations driven by an elliptical 729-nm beam. We expect to observe structured errors due to the Gaussian intensity profile of the beam. To verify the expected gate duration dependence due to the finite beam, we run parallel gate duration calibration experiments (see § 4.1.3 for details). The gate duration, τ , as a function of qubit position in the beam is expected to follow,

$$\tau(r) \propto \frac{1}{\sqrt{I(r)}}, \quad (5.71)$$

where $I(r)$ is a Gaussian distribution, $I = a \exp\left(\frac{-(r-r_0)^2}{2c^2}\right)$. Figure 5.12 shows the measured $\hat{R}_x(\pi)$ -pulse gate durations for each ion, with the fit given by the dashed line. From this, we find an expected major-axis waist radius of 27.7(9) μm aligned to the ion chain. The qubit spacing, on the lower axis of the calibration plot, is inferred through the axial confinement of 0.5 MHz using equilibrium positions of a linear chain given in [33]. At this confinement, the chain extent is 35 μm , giving an expected Rabi-frequency difference of 10.3(6)% between the outer and inner ions. We calibrate global single-qubit gates using the inner ion, $i = 4$, leading to an expected per Clifford error of (using fidelity definition from Appendix A.2),

$$\varepsilon_{c,max} = \frac{2}{3}\kappa \left(\alpha \frac{\pi}{4}\right)^2 = 5.08(2) \times 10^{-3}, \quad (5.72)$$

where, here the Clifford decomposition used gives $\kappa = 1.167$, and the fractional Rabi frequency difference is given by the above calibration curve, $\alpha = 0.103(6)$.

As with the single qubit Stark-shift dataset, we ensure that correlated errors dominate to improve the fitting sensitivity. This is done through a fixed phase noise error corresponding to $\varepsilon_{\text{corr},0} = 1.3 \times 10^{-2}$. In the following data, we subtract this baseline offset to highlight the excess correlated noise due to the beam structure, $\varepsilon_{\text{corr}}^> = \varepsilon_{\text{corr}} - \varepsilon_{\text{corr},0}$.

Figure 5.13 shows the extracted pairwise correlations. The excess correlated error is shown in (a), with the largest values seen for the outer ion pairs, agreeing with the expected maximum error of $\varepsilon_{\text{corr},\text{max}}^> = \varepsilon_{c,\text{max}}$. Panel (b) shows the sum of the local depolarisation error per Cliffords on both qubits, with the largest values now corresponding to pairs consisting of one central ion, and one outer ion. We plot the ratio of excess correlated error and depolarisation errors in panel (c) to further highlight the error distribution due to the beam shape.

5.7 Outlook

k-copy RB efficiently discerns correlated and uncorrelated error channels acting on a multi-qubit system. The symmetry enforced by twirling over the diagonal single-qubit Clifford representation $Cl_1^{\otimes k}$ (§ 5.3.2), causes structured errors to drive the decay of correlated populations into invariant subspaces. Observing transient populations in these invariant subspaces allows for error characterisation without full process reconstruction. Given a specified error model, here focussing on common unitary errors due to fluctuating global fields, these correlated decays can be used to directly quantify the magnitude of the individual error channels contributing to the total error per Clifford that is typically measured in standard RB.

In this chapter we proposed and experimentally demonstrated this method applied on up to an eight-qubit register through pairwise comparisons (§ 5.6.5). This approach permits highly efficient parallel data collection: the time required on the device is independent of qubit number, and the protocol requires exclusively local operations followed by classical post-processing of correlated survival rates.




Figure 5.14: Experimental data and analytic fit for 3-qubits subject to correlated random-unitary single-qubit noise injected via random phase jitter ($\sigma_\phi = 0.06$ turns). $\mathcal{P}_{[SSF]}$ and $\mathcal{P}_{[SFF]}$ represent the sum of the permutations of single- and two-qubit failures respectively. Explicitly, $\mathcal{P}_{[SSF]} = \mathcal{P}_{SSF} + \mathcal{P}_{SFS} + \mathcal{P}_{FSS}$. A clear intermediate asymptote of $\mathcal{P}_{SSS} = 1/4$ and final asymptote of $\mathcal{P}_{SSS} \rightarrow 1/8$ are observed, consistent with dominant *common noise*.

This work is particularly relevant to laser-free trapped-ion architectures [135, 136], where the use of global fields and gradients can give rise to correlated errors across multiple ions, providing a natural setting in which to apply this protocol.

5.7.1 Limitations

This chapter has analysed structured errors on two qubits, extending to more qubits only through pairwise comparisons. The current pairwise protocol described here does not determine correlated errors between >2 -qubits. These higher weight errors may be particularly detrimental to error correction circuits where the distance of the code, d , is less than the correlated error channel weight [105], therefore these errors would be advantageous to characterise. We expect the symmetry resolved dynamics revealed by pairwise k -copy RB to generalise to higher weight errors, however, without prior work on the sub-representation structure of the single-qubit Clifford group applied diagonally to >2 qubits, extending the analysis for k -copy RB is challenging. The number of relevant irreps does not increase, however, the multiplicity does increase dramatically. Some headway can be made if we consider the *common noise* error mode discussed throughout this chapter.




Figure 5.15: Circuit for probing time-dependent environmental errors by introducing a delay between random sequence execution on parallel copies in k -copy RB.

A simple relation can be found for the asymptotic states, namely, under *common noise* the state is driven to the maximally mixed state on the trace subspace, as this is the $SU(2)^{\otimes k}$ -invariant state on the symmetric irrep. This state is given by,

$$\hat{\rho}_{\text{sym}} = \frac{1}{k+1} \sum_{l=0}^k |D_l^k\rangle \langle D_l^k|, \quad (5.73)$$

where $|D_l^k\rangle$ is the Dicke state for l excitations (here successful recoveries of target state), out of k -qubits. From this expression we see that an asymptote of $\mathcal{P}_{SS\dots S}|_{m \rightarrow \infty} = 1/(k+1)$ is expected under *common noise*, while the maximally mixed state under local depolarisation leads to the familiar asymptote of $\mathcal{P}_{SS\dots S}|_{m \rightarrow \infty} = 1/(2^k)$. This result suggests that the presence of *common noise* affecting many qubits may be simply observed, however, without the explicit representation structure, it cannot be easily quantified.

We confirm this asymptote experimentally with data shown in Figure 5.14. For three-qubits, we leverage symbolic computation to diagonalise the resulting PTM representation of *common noise* and recover the analytic decay used in this fit. This direct diagonalisation approach is possible for this small system, and with the highly structured *common noise* error channel, however, for a full treatment of k -body error under $Cl_1^{\otimes k}$ twirling, we suspect the sub-representation structure will need to be found.

5.7.2 Future work

One avenue for future work is captured by the above limitations section: scaling the analysis to multi-body errors through studying the sub-representation structure of the single qubit Clifford group applied diagonally to k -qubits. We propose two

further dimensions in which to extend the *k-copy* analysis:

First, the subsystem on which the local twirling is applied could be extended to > 1 -qubit. As an example, the two-qubit Clifford group [137] or the SLeRB [124] protocol could be applied in parallel to k -copies of 2-qubit subsystems. This would reveal correlated errors and leakage that reduce the fidelity of two-qubit entangling gates.

Second, we may utilise this method to probe correlations in time. As the protocol uses exclusively single-qubit rotations, with no explicit coupling made between qubits by the ideal circuit, we could apply some time delay, Δt , between the parallel gate sequences. A proposed circuit is shown in Figure 5.15 for the two-copy case. In the case that Δt is much longer than the random sequence length, we are effectively probing the stability of the circuit over time (albeit on parallel subsystems) — specifically we are looking for deviations in circuit fidelity due to the external context in which the circuit is executed. This method is similar to the context dependent tests studied in reference [120]. However, as we execute parallel sequences on two copies, we can also set Δt to be significantly shorter than the random sequence length. In this regime, we could, for example, probe the frequency and coherence of noise on the common drive field used to execute the parallel gates. This would provide evidence and quantify violations of the Markovian QIP model, without full process tomography, as discussed in § 5.1.

6

Towards Hybrid System Control and Quantum Error Mitigation

Contents

5.1	Errors and Characterisation	98
5.2	Existing approaches to correlated-noise characterisation	101
5.3	The <i>k-copy</i> RB scheme	103
5.3.1	Protocol	104
5.3.2	Representations of $\mathcal{C}l_1^{\otimes 2}$	105
5.3.3	Common random unitary noise	106
5.3.4	Local depolarising noise	109
5.3.5	Combined model	110
5.3.6	State preparation errors	114
5.4	Parity contrast: an independent probe of coherence .	115
5.5	Sampling strategy	117
5.6	Experimental Results	121
5.6.1	Physical implementation	122
5.6.2	Z-basis measurements	123
5.6.3	Coherence measurements	124
5.6.4	Controlled noise injection	126
5.6.5	Scaling to a multi-ion register	133
5.7	Outlook	134
5.7.1	Limitations	135
5.7.2	Future work	136

Having established and characterised the experimental apparatus and demonstrated its application to *k-copy* randomised benchmarking, in this chapter we

consider future research goals utilising the multi-qubit arrays enabled through the *NPL* trap (§ 3.2), the individual-ion addressing system (§ 3.6.5), and the developed control over both the qubits, which encode discrete variables, and the motional modes, which encode continuous variables.

In the chapter we establish a set of unitary operations required for universal continuous-variable (CV) control of the ion motion and discuss the extensions to hybrid spin-oscillator systems. Following on from recent progress in spin-conditioned bosonic operations demonstrated within our group [42, 43, 45, 138–140], we are interested in developing the required addressing capabilities to scale to multiple qubits and motional modes. We outline the method we use to generate such spin-oscillator interactions experimentally, and demonstrate the generation of a single-mode squeezed state on the new apparatus. We then outline the next steps required for demonstrating the remaining single- and two-mode interactions, and future experiments this will enable.

In the second half of this chapter we consider an error mitigation technique applicable to both discrete and continuous systems. Developing this, and associated methods on our trapped ion platform will enable substantial reductions of certain errors, with significantly fewer resources required than full error correction. We give a short introduction to quantum error mitigation, followed by an outline of virtual distillation [91–93] — the specific method considered here. We provide a demonstration of this protocol applied to a discrete-variable only system. Finally, we explore the impact of error channels that induce multi-qubit correlations on the efficacy of virtual distillation, particularly focussing on correlated noise due to a common fluctuating control field as discussed in the Chapter 5 § 5.3.3. The *k-copy* protocol developed in that chapter is utilised to quantify this unfavourable error, and inform predictions on the performance of virtual distillation in this context. The section is concluded with a discussion on the application of virtual distillation to larger system sizes, and to hybrid systems, presenting a scalable method within reach of our apparatus.

6.1 Hybrid Spin-Oscillator Control in Trapped Ions

CV quantum computation [18] offers an alternative strategy for encoding quantum information, utilising quantum harmonic oscillators, termed *qumodes*, as the fundamental quantum object, rather than discrete variable (DV) systems, e.g. the two-level qubit. A single ideal qumode possesses an infinite dimensional Hilbert space, allowing for high (in this case arbitrarily high) density encoding of information in physical information carriers. This opens an alternative avenue for scaling QIPs, requiring dramatically fewer physical carriers compared to qubit-only systems. The inclusion of both CV and DV subsystems, termed the hybrid model [141], retains the convenient control and readout of discrete qubits while providing the larger state space and native interactions available to qumodes. The hybrid system is a promising avenue for quantum simulation, with the ability to encode spins with physical two-level systems, whilst encoding bosons efficiently using bosonic physical systems [45]. For computation the large state space of hybrid systems allows for highly resource efficient circuit decompositions. For example, reference [14] considers a factoring algorithm whose required physical resources are independent of the number to be factored. Remarkably, this scheme requires only a single qubit and three oscillators, however requires prohibitively large energy states of the oscillator. Practically, we will not be able to realise arbitrarily high energy oscillator states, but it is still an open question of what information density is possible in experimentally accessible quantum oscillators.

Hybrid spin-oscillator computation has been studied in the context of superconducting cavity-QED [142–146], neutral atom in optical tweezers [147–149], and in the trapped-ion platform [150–152]. We shall focus on the trapped-ion implementation, where the DV system is the collection of qubits embedded in the electron energy levels, and the qumodes are the shared secular motional modes of the ion crystal within the trapping pseudopotential.

Trapped ions benefit from multiple native methods to drive, and couple qubit and qumode: The qumodes may be directly driven through parametric trap potential

modulation [153–155], or through near mode-resonance RF fields (§ 2.3.3); the qubit may be driven either by optical fields (in the case of optical qubit encoding detailed in § 2.2.2), or through microwaves (in the case of ground state encoding) resonant with the energy-state transition; the qubit and qumode may be coupled through qubit-state dependent potential shifts [156, 157], or through optical fields and magnetic field gradients tuned near to motional-mode frequency sidebands of the qubit-carrier transition, for example through the Jaynes-Cummings interaction (§ 2.4.1), or the Mølmer-Sørensen (MS) spin-dependent force (SDF) [37] (§ 2.4.2). Taken together, these native control and coupling mechanisms, combined with the long coherence times of both the qubit and qumode, make the trapped-ion platform well-suited for exploring hybrid CV–DV quantum computation.

The apparatus developed in this thesis has demonstrated good qubit and radial motional coherence times, provides access to multiple qubits and qumodes, and includes a single-addressing system for selectively driving their coupling. These capabilities make it well-placed for studying methods for scaling hybrid-system control.

6.1.1 Theory

Gate set for a Hybrid System

For any computing paradigm, we must establish a set of interactions required for universal computing (i.e. Turing complete). In DV-only systems, quantum universality may be realised by a finite set of Clifford operations plus one non-Clifford gate [158, 159]. CV-only universality can analogously be realised by a finite set of Gaussian operations plus one non-Gaussian gate [17, 18]. The term Gaussian is understood in the position-momentum representation of the qumode, as discussed in § 2.3.2, where the bare oscillator Hamiltonian is given by,

$$\hat{H}_{\text{osc}} = \frac{1}{2m}\hat{p}^2 + \frac{1}{2}m\omega_x^2\hat{x}^2, \quad (6.1)$$

with $\hat{x}$ and $\hat{p}$ quadratures expressible in terms of creation and annihilation operators, as given in equation 2.25. Gaussian operators are then the set of interactions containing at most quadratic terms in $\hat{x}$ and $\hat{p}$, therefore mapping Gaussian input

Interaction name	Oscillator $\hat{H}_I^{(\text{osc})}$	Gauss.	N_{osc}	Two-SDF method		
				m	$\hat{H}_I^{(\text{s})}$	Ω_{eff}
Displacement	$\hat{a}e^{i\phi} + \hat{a}^\dagger e^{-i\phi}$	✓	1	-	$\hat{\sigma}_\alpha$	Ω_α
Parity	$(2\hat{a}^\dagger \hat{a} + 1) \sin(\phi)$	✓	1	+1	$\hat{\sigma}_\gamma$	Ω_2
Beam splitter	$\hat{a}\hat{b}^\dagger e^{i\phi} + \hat{a}^\dagger \hat{b}e^{-i\phi}$	✓	2	+1	$\hat{\sigma}_\gamma$	Ω_2
Squeezing	$\hat{a}^2 e^{i\phi} + \hat{a}^{\dagger 2} e^{-i\phi}$	✓	1	-1	$\hat{\sigma}_\gamma$	Ω_2
Trisqueezing	$\hat{a}^3 e^{i\phi} + \hat{a}^{\dagger 3} e^{-i\phi}$	✗	1	-2	$\hat{\sigma}_\beta$	Ω_3

Table 6.1: An example set of unitary interactions to compose a universal gate set for continuous variable computation. The name for the CV interaction, the bosonic operator component of the Hamiltonian, if it is Gaussian, and the number of oscillators it acts on are given. These CV interactions are accessible and may be extended to hybrid-interactions (CV-DV) by using the protocol presented in reference [150], here termed the “Two-SDF” method. The required relative detuning setting, m , the resulting spin conditioning, $\hat{\sigma}$, and the effective coupling strength Ω_{eff} are shown here. The first row is a single SDF with strength Ω_α . The other coupling strengths are found via the 2nd and 3rd terms of the Magnus expansion: $\Omega_2 = \Omega_\alpha \Omega_{\alpha'}/\Delta$, $\Omega_3 = \Omega_\alpha^2 \Omega_{\alpha'}/\Delta^2$.

states to Gaussian output states.

The condition for universality of a hybrid CV-DV is somewhat ambiguous as both the CV and DV subsystems are themselves sufficient for computational universality¹, however this presents an opportunity: the freedom in selecting the chosen gate-set for universality, combining both Clifford, non-Clifford, Gaussian, non-Gaussian, and spin-oscillator interactions together, allows for selecting efficient, native, and simply accessed interactions of the physical platform. We consider the over-complete instruction set presented in Table 6.1, consisting of spin-conditioned Gaussian and non-Gaussian interactions. The total Hamiltonian is given by,

$$\hat{H} = \hat{H}_s \otimes \hat{I}_{\text{osc}} + \hat{I}_s \otimes \hat{H}_{\text{osc}} + \hat{H}_I^{(\text{s})} \otimes \hat{H}_I^{(\text{osc})}, \quad (6.2)$$

where $\hat{H}_s$ and $\hat{H}_{\text{osc}}$ are the bare Hamiltonians describing the qubit and qumode, and the interaction Hamiltonian is split into spin-only, $\hat{H}_I^{(\text{s})}$, and motional-only, $\hat{H}_I^{(\text{osc})}$, acting components. These interaction components are provided in the table.

¹We note that “universality” is perhaps misleading terminology. We specifically mean here the ability to simulate any other physical system, i.e. Turing complete. However, we note that universal is also used in the context that the gate-set is sufficient to access any state in the corresponding Hilbert space to an arbitrary level of accuracy. These two criteria are not synonymous — the latter in fact encompasses Turing universality.

This set of interactions on the bosonic subsystems already enable computational universality, however the spin-conditioning provides a bridge between CV and DV subsystems, enabling readout and control to be implemented on auxiliary qubits. Specifically, we consider a cubic CV instruction set [18, 142, 160], taking the Gaussian operations and supplementing with a non-Gaussian “trisqueezing” gate, with Hamiltonian,

$$\hat{H}_T \propto \hat{a}^3 + \hat{a}^{\dagger 3}. \quad (6.3)$$

Spin-conditioned trisqueezing has been shown previously in our group [42], utilising the protocol given in reference [150], here termed the “two-SDF” method.

The remaining nonlinear Gaussian operations; squeezing, parity, and beam-splitter, may be generated through the same “two-SDF” mechanism, which we shall detail in the following section, and have also been previously demonstrated [42, 45, 138, 139].

The displacement operation can be realised either spin-independently by RF fields near resonance to the motional-mode frequencies (§ 2.3.1), or spin-conditioned using various SDF implementations, such as by the MS method using bichromatic fields (§ 2.4.2), or by static Stark-shift gradients due to an optical tweezer [157].

As highlighted in this section, the required unitaries for universal hybrid CV-DV control may all be accessed by a common protocol: The “two-SDF” method. This greatly simplifies experimental hardware and control complexity. We shall give an overview of this method in the following section, and provide a practical route and demonstration utilising four optical frequencies with fixed relative detunings and phases.

Gate Implementation

To implement the unitaries given in Table 6.1 with spin-dependent control, we use the “two-SDF” method proposed in reference [150]. The nonlinear bosonic

interactions are effectively driven through applying two simultaneous linear spin-dependent forces with non-commuting spin conditioning, on a single qubit. The two simultaneous SDFs, with Hamiltonians and parameters labelled by subscripts α , and β , are applied off-resonantly relative to motional modes with creation operators $\hat{a}_\alpha^\dagger$ and $\hat{a}_\beta^\dagger$, respectively. The resulting Hamiltonian is given generally by,

$$\hat{H}_{\text{2SDF}}(t) = \hat{H}_\alpha(t) + \hat{H}_\beta(t), \quad (6.4)$$

$$\begin{aligned} &= \frac{\hbar\Omega_\alpha}{2}\hat{\sigma}_\alpha \left(\hat{a}_\alpha e^{i(\Delta_\alpha t + \phi_\alpha)} + \hat{a}_\alpha^\dagger e^{-i(\Delta_\alpha t + \phi_\alpha)} \right) \\ &+ \frac{\hbar\Omega_\beta}{2}\hat{\sigma}_\beta \left(\hat{a}_\beta e^{i(\Delta_\beta t + \phi_\beta)} + \hat{a}_\beta^\dagger e^{-i(\Delta_\beta t + \phi_\beta)} \right), \end{aligned} \quad (6.5)$$

where the individual SDF strengths are Ω_i , the detuning from each mode is Δ_i , the motional phase is ϕ_i , and we ensure the spin-conditioning anticommute $\{\hat{\sigma}_\alpha, \hat{\sigma}_\beta\} = 0$, $[\hat{\sigma}_\alpha, \hat{\sigma}_\beta] = 2i\epsilon_{\alpha\beta\gamma}\hat{\sigma}_\gamma$, where $\epsilon_{\alpha\beta\gamma}$ is the Levi-Civita symbols. The unitary dynamics of $\hat{H}_{\text{2SDF}}$ may be evaluated by the Magnus expansion [161], with details given in references [138, 139, 150], and in Appendix A.1. To lend some intuition on the resulting physics, we highlight the main ideas here: The Magnus expansion expresses the generator of the unitary evolution as an infinite series of terms, with successive terms proportional to increasingly higher-order nested commutators of the Hamiltonian evaluated at different times. Due to the non-commuting spin conditioning, this series does not truncate, leading to arbitrarily high nonlinear bosonic interactions effectively driven via the application of only two linear SDFs. To utilise specific nonlinear terms within this infinite series, matching the desired interactions given in Table 6.1, we constrain the relative detunings of the two SDFs to $\Delta_\beta = m\Delta_\alpha$, where m is some integer given in the table. This immediately brings components of the Magnus terms matching the desired nonlinear $\hat{H}_{\text{osc}}$ into resonance. Any undesired higher-order terms that are also brought into resonance can be arbitrarily suppressed by increasing the detuning, as the effective coupling of the n th-order Magnus term scales as $\Omega_{\text{eff}} \propto \Omega_\alpha^n / \Delta^{n-1}$ (relabelling $\Delta_\alpha = \Delta$ and assuming $\Omega_\alpha = \Omega_\beta$). The remaining error terms are off-resonant, and may also be further suppressed by ensuring $\Delta \gg \Omega_\alpha$, and utilising appropriate smooth amplitude and phase envelopes for the pulses.




Figure 6.1: State tomography for a single mode through measurement of an ancilla qubit. (a) The Hadamard test method for measuring the expectation value of an arbitrary operator, $\hat{O}$, applied to the oscillator state, utilising an ancilliary qubit. (b) A Hadamard test circuit for measuring the characteristic function, $\chi(\beta)$, of the oscillator using native single qubit rotations and $\hat{\sigma}_x$ -conditioned spin-dependent forces (symmetric control represented by the diamond symbol) [162].

In this regime, the unitary dynamics of $\hat{H}_{\text{SDF}}$ are well approximated by the evolution of time-independent effective Hamiltonians. For example, the effective squeezing interaction is driven when $m = -1$, and both SDFs act on the same mode, $\hat{a}_\alpha^\dagger = \hat{a}_\beta^\dagger = \hat{a}^\dagger$ yielding,

$$\hat{H}_{\text{eff,sqz}} = \frac{\hbar\Omega_2}{2}\epsilon_{\alpha\beta\gamma}\hat{\sigma}_\gamma(\hat{a}^2e^{i\phi'_p} + \hat{a}^{\dagger 2}e^{-i\phi'_p}), \quad (6.6)$$

where $\Omega_2 = \frac{\Omega_\alpha\Omega_\beta}{\Delta}$ is an effective coupling strength, and phase $\phi'_p = \phi_\alpha + \phi_\beta + \frac{\pi}{2}$. The derivation and forms of the other spin-conditioned effective nonlinear bosonic interactions are given in Appendix A.1.

Frequency and spin-conditioning control of two SDFs therefore allow for the generation of the entire hybrid CV-DV gate-set proposed in Table 6.1. After a discussion on how we can characterise the oscillator state, we shall present the experimental setup for generating these interactions, and demonstrate the creation of a squeezed state by this method.

Readout

With the above methods we may generate non-classical bosonic states, now we provide a method for characterising them.

In contrast to a two-level system, the state of an oscillator may not be directly read out through state dependent fluorescence. This is due to the much larger state-space of the qumode, and the technical difficulty in direct motional readout on the ion-trap platform. To solve these issues, a modular readout scheme is used,

involving indirect readout of the oscillator state via a discrete collection of controlled unitaries subject to an ancillary qubit, and computational basis measurements of that ancilla.

The information is mapped to the ancilla by a Hadamard test [163], shown in Figure 6.1 (a). Here the expectation value of the operator $\hat{O}$ on the oscillator, $\langle \hat{O} \rangle = \text{Tr}(\hat{O}\hat{\rho}_{\text{osc}})$, is recovered through exclusive measurements of the ancilla qubit after qumode-dependent phase “kick-backs” are applied. The Hadamard test has the convenient relation between computational basis measurements of the ancilla, and the applied operator, $\langle \hat{\sigma}_z \rangle = \text{Re}[\text{Tr}(\hat{O}\hat{\rho}_{\text{osc}})]$.

Full state tomography may now be formulated by repeat measurements of the conditional displacement operation, $\hat{O} = e^{-i\theta}\hat{D}(\beta)$, to recover the characteristic function, defined as $\chi = \text{Tr}(\hat{D}(\beta)\hat{\rho}_{\text{osc}})$, with β being a complex parameter with displacement magnitude, $|\beta|$, and displacement axis, $\arg(\beta)$. Final computational state measurements of the Hadamard test circuit then yield,

$$\langle \hat{\sigma}_z \rangle = \text{Re}[e^{-i\theta}\text{Tr}(\hat{D}(\beta)\hat{\rho}_{\text{osc}})], \quad (6.7)$$

$$= \cos(\theta)\text{Re}[\chi] + \sin(\theta)\text{Im}[\chi], \quad (6.8)$$

where θ may be set to 0 or $\pi/2$ to recover the real or imaginary part of the characteristic function respectively. Each measurement provides only 1-bit of information on the (continuous) oscillator state, thus to approximate a useful reconstruction, many samples with a range of β values must be collected.

Once the characteristic function has been measure, it can be converted to other familiar representations of the oscillator, such as in the Fock-state basis, or Wigner representation [164–166]. In this work, we shall keep to the conveniently measured characteristic function.

To give two concrete examples, the vacuum state characteristic function is the circularly symmetric Gaussian state,

$$\chi_0 = \exp\left(\frac{-|\beta|^2}{2}\right), \quad (6.9)$$

while the squeezed state, generated by application of the Hamiltonian in equation 6.6 to the vacuum state, breaks this symmetry due to a contraction on one axis, and a corresponding expansion on the other [167],

$$\chi_s = \exp\left(\frac{-|\beta|^2}{2} \left(e^{2r} \cos^2(\phi) + e^{-2r} \sin^2(\phi)\right)\right), \quad (6.10)$$

where the complex squeezing parameter is $\xi = re^{i\phi_s} = \Omega_2 e^{i\phi'_p t}$, and ϕ is the angle between squeezing axis and probe axis, $\phi' = \arg(\beta) - \frac{\phi_s}{2}$.

Experimentally, $\hat{\sigma}_x$ conditioned SDFs are simply accessed by the bichromatic field method discussed in § 2.4.2. The Hadamard test circuit is then decomposed to native gates, shown in Figure 6.1 (b), first described in reference [162]. The proceeding I or $\hat{R}_y(\pi/2)$ set the quadrature to measure; $\text{Re}(\chi)$ or $\text{Im}(\chi)$ respectively. This circuit provides a practical method for characterising a single qumode, which we shall use in the following experimental demonstration. Similar to the constraints of state characterisation discussed in § 5.1, full tomographic reconstruction is intractable for multiple-qumodes due to the large state-space, even with reasonably truncated oscillator representations (e.g. imposing a Fock-state cut-off). In the case of characterising systems consisting of multiple qubits and qumodes, either reduced tomography [168], or simplified metrics of performance [169] must be used.

6.1.2 Experimental Demonstration

In our group, we have previously demonstrated higher order squeezing [42], superpositions of nonlinear bosonic states [43], beam-splitter [45], and two-mode squeezing interactions. However, due to the constraints of the previous experimental apparatus we were limited to global addressing fields, preventing the application of these hybrid interactions to more than two qubits. Here we demonstrate the creation and characterisation of a squeezed state on the newly constructed apparatus, described in Chapter 3, utilising the single-qubit addressing system. Single addressing both increases the power efficiency of the interactions, and allows the control to select which qubit is used to couple to the motion — this provides a scalable route to generating hybrid interactions on larger qubit, and qumode registers.




Figure 6.2: The quadchromatic field required for generating nonlinear motional interactions via the two-SDF method. (a) For single-mode interactions, both SDF fields are detuned around the sidebands of the same motional mode, ω_α . Integer m is varied to select the desired interaction. (b) For two-mode interactions, each SDF is detuned from a different motional mode.

Two Spin Dependent Forces

The individual SDFs are generated by the same bichromatic field method described in § 2.4.2 — i.e. simultaneous off-resonant blue- (BSB) and red-sidebands (RSB) of motional mode ω_α are applied. To generate two simultaneous SDFs, we then require four optical fields with, frequency, phase, and amplitude control. The four required frequencies of the “quadchromatic” field are shown schematically in Figure 6.2 with respect to the carrier frequency, ω_0 . Here, purple lines correspond to SDF_α , and orange lines to SDF_β . In this example we show the SDFs acting on a shared motional mode, but in general they may act on different modes. Similar to the generation of the bichromatic field, we may produce this quadchromatic field through the application of four RF tones onto a single-pass (SP) AOM. These tones are generated using an arbitrary waveform generator (AWG)². We cannot control the absolute optical phase at the ion, however we can set relative phases between the four fields by applying phase offsets between the RF tones with respect to some fixed time point. As discussed in § 2.4.2, the spin phase of the SDFs is set by the absolute phase of the RSB and BSB, $\hat{\sigma}_\phi = \cos(\phi_s)\hat{\sigma}_x + \sin(\phi_s)\hat{\sigma}_y$. To ensure the spin conditioning of the two SDFs fully anti-commute, we require $\pi/2$ phase difference between the two spin phases,

$$\phi_{\alpha,s} - \phi_{\beta,s} = \frac{\phi_{\alpha,\text{RSB}} + \phi_{\alpha,\text{BSB}}}{2} - \frac{\phi_{\beta,\text{RSB}} + \phi_{\beta,\text{BSB}}}{2} = \frac{\pi}{2}. \quad (6.11)$$

²*Sinara Phaser* board

In the following demonstration we set $\phi_{\alpha,\text{RSB}} = \phi_{\alpha,\text{BSB}} = \pi/2$, $\phi_{\beta,\text{RSB}} = \phi_{\beta,\text{BSB}} = 0$, however we note that for composing multiple hybrid operations into a circuit, the resulting motional phase of the SDFs can lead to further constraints on the chosen RSB and BSB phases.

The intensity of the quadchromatic field is ramped on and off with a rise envelope of $\sin^2(\pi t/(2t_{\text{ramp}}))$ applied via the AWG, where t_{ramp} is the ramp duration. This ramp suppresses excess off-resonant excitations of the A_1 phase loop term given in the previous section (further details in Appendix A.1), and so we must select ramp duration based on the characteristic scale, $t_{\text{ramp}} \gtrsim \frac{1}{\Delta}$.

We aim to balance the power in each of the 4 tones. This may be calibrated using an optical pick-off before addressing the ion, directed onto a fast photodiode, and measuring the beatnote contrast when applying any pair of tones. In the following squeezing demonstration, we balance the tones comprising the individual SDFs, but do not explicitly balance the power between SDFs, as the desired interaction scales as $\Omega_\alpha\Omega_\beta$, any imbalance leads to non-optimal interaction efficiency, but unlike imbalance of the RSB and BSB, does not increase coherent error terms. We observe that the duty cycle of the SP AOM can dramatically shift the power imbalance between all four tones, and so we take care ensuring that the total RMS power to drive one-, two-, and four-tones is the same as for the idle (no optical pulse) setting of the SP AOM, thus the only transient power drop is that of the ramp sequence.

Squeezing

We demonstrate the effective squeezing interaction on a single ion and radial motional mode. We are limited to using the radial modes due to the geometry of the single addressing system with respect to the trap. We use the upper radial mode with frequency $\omega_{\text{upper}} = 2\pi \times 2.87$ MHz, and motional coherence $\tau_{\text{upper}} = 18.0(8)$ ms given in § 4.2.4. The initial near-vacuum motional state is prepared by Doppler and sideband cooling, with a typical thermal fock state distribution with mean $\bar{n} = 0.03(1)$, as described in § 4.2.2. The spin state is initialised into the ground state

by optical pumping, as characterised in § 4.1.1. By setting the relative detuning between the SDFs to $m = -1$, and ensuring anti-commuting spin conditioning by setting the relative phases of the quadchromatic field according to equation 6.11, we realise a $\hat{\sigma}_z$ -conditioned squeezing interaction, given by equation 6.6. As the spin-state is prepared in an eigenstate of the spin-conditioned squeezing, we expect squeezing on the qumode, without any change to the qubit state. We use a total quadchromatic beam power of 53 μW , and set the base detuning, $\Delta = 2\pi \times 50$ kHz. We suppress excess excitation of the off-resonant error terms by using a ramp of $t_{\text{ramp}} = 40$ μs .

The state is characterised by the tomography method given in § 6.1.1, with a 35×35 sampling-grid of displacement values, β , and $N = 300$ repeat samples to reduce shot noise. The full dataset for one of the tomographic scans takes around 1 hour to collect, and so we interleave mode calibration experiments (§ ??) to run every 5 minutes.

The on-resonance splitting SDF used for tomography must have its strength, Ω_{SDF} , calibrated to infer the magnitude of β applied. We use a total bichromatic beam power of 40 μW , and through fitting the $\hat{\sigma}_z$ expectation for the time-dependent splitting of a vacuum state,

$$\langle \hat{\sigma}_z \rangle = e^{-2|\beta|^2}, \quad |\beta| = \frac{\Omega_{\text{split}} t}{2}, \quad (6.12)$$

we measure an SDF splitting strength of $\Omega_{\text{split}} = 2\pi \times 7.2(4)$ kHz.

With these settings, the characteristic function for a squeezed state is measured, and the resulting real part is shown in Figure 6.3 (a). We can see a contraction along one phase-space axis, and expansion along the orthogonal axis, together with a small phase-space tilt, likely due to imprecise phase tracking between the squeezing interaction and the displacement probe. Fitting the data to equation 6.10 we may extract both the squeezing axis, ϕ_s , and the strength, r , where r is given by,

$$r = |\Omega_{\text{eff,S}}| t = \frac{\Omega_\alpha \Omega_\beta}{\Delta} t. \quad (6.13)$$

An example fit is shown in panel (b), with extracted $r = 0.90(2)$, for a quadchromatic FWHM pulse duration of $t_{\text{sqz}} = 190$ μs . Panel (c) displays the squeezing strength as




Figure 6.3: Demonstration of the squeezing interaction through the two-SDF method. (a) The characteristic function, $\text{Re}(\chi(\beta))$, of the squeezed state measured using the state tomography method given by Figure 6.1. (b) The simulated state taking fit parameters inferred by equation 6.10. (c) The fitted squeezing parameter, r , as a function of two-SDF pulse duration.

a function of t_{sqz} , with fit values given by the points, and the black line is a linear fit. Assuming that the two SDF strengths are well balanced, $\Omega_\alpha = \Omega_\beta$, and taking the 53 μW used for the quadchromatic field compared to the SDF strength calibration at 40 μW , we find an expected squeezing strength of $\Omega_{\text{eff},S} = 2\pi \times 0.86(9)$ kHz. This is shown as the red dashed line and shaded region. The fitted strength of $\Omega_{\text{eff},S} = 2\pi \times 0.74(1)$ kHz is lower than the expected, suggesting that either we are not sufficiently balancing the strengths of the two simultaneous SDFs, the spin-conditioning is not fully anti-commuting, or that the power stabilisation routine for 4-tones is inaccurate. Further exploration is needed, including careful quantification of the individual SDF strengths constituting the quadchromatic field. However, this preliminary result confirms that the motional coherence, and stability of the single-addressing system is sufficient to create, and characterise non-classical motional states using the apparatus developed for this thesis.

6.1.3 Outlook

Hybrid systems offer access to large Hilbert spaces, even with modest ion numbers. In this section we outlined a convenient set of native interactions for universal hybrid qubit-qumode gates. We experimentally demonstrate the squeezing interaction

generated via the “two-SDF” method of reference [150]. This demonstration replicated previous work using the “two-SDF” method [42], however now shown on the newly built apparatus utilising a single-ion addressing system, and applied to the radial motional modes.

Enabled by the single-addressing, and improved motional coherence times of the new system, we are particularly interested in exploring short hybrid circuits composed of multiple interactions acting on an array of qubits and qumodes. We are interested in utilising the hybrid system with increased robustness against errors through the use of quantum error mitigation methods applicable to both spin and oscillator subsystems. We outline one such method in the next section, with a preliminary demonstration on the spin subsystem.

Limitations

Using a single multi-ion chain to access multiple qubits and qumodes for scaling hybrid QC on the trapped-ion platform will be ultimately limited by spectral crowding of the motional modes, i.e. the average frequency spacing between neighbouring modes decreases as the number of accessible modes increases[33, 170]. The “two-SDF” method requires frequency selectivity to drive bosonic interactions exclusively on the chosen mode. To selectively drive a single mode, the magnitude of the detuning Δ must follow $\Delta \ll \delta\omega$, where $\delta\omega$ is the frequency difference to the nearest neighbouring mode. As discussed in the “two-SDF” method section, the strength of each SDF must be less than the base detuning, $\Omega_\alpha \ll \Delta$, to well approximate the effective nonlinear unitary dynamics. As a result of these two requirements, the mode-frequency splitting directly limits the attainable strength of the nonlinear interaction in the adiabatic regime, resulting in increasing gate duration with increasing register size.

A possible mitigation would be to rather scale the hybrid system using a collection of fixed-length chains with adequate mode-frequency splitting, and to transport quantum information between these chains either through ion transport whilst

maintaining mode and spin coherence [171], generating remote entanglement through photonic interconnects [13], or through inducing collective motional interactions between separate ion-chains [172].

Another mitigation strategy could come from the pulse-shaping approach demonstrated for use in fast two-qubit entangling gates in the non-adiabatic regime [40], i.e. when $\Delta > \delta\omega$. Introducing extra degrees of freedom in the pulse amplitude envelope could allow for driving the desired interaction on the selected modes, while returning all spectator modes to their initial state by the end of the pulse. However, we note that the accumulation of non-global geometric phases, and possible non-commutativity of interactions on shared qubit-oscillator states could prohibit the use of such a pulse-shaping scheme. Ultimately, the desired interaction, mode frequency spectra, and target gate duration will determine the dominant coherent errors. We leave such an analysis for future work.

6.2 Virtual Distillation

With access to noisy, intermediate-scale quantum information processors, the natural question is whether useful computations may be executed, without requiring full fault-tolerant error correction codes. In this section, we explore a resource-efficient error mitigation technique, known as virtual distillation, that can dramatically suppress the effect of errors on the final measured outcome for expectation value problems.

This section starts with outlining the goals of quantum error mitigation (QEM), with the following section giving details on the virtual distillation method. We then present an experimental demonstration suppressing errors on two-copies of a single qubit system. We finish this section with a discussion on future plans, and the extension of virtual distillation to hybrid CV-DV systems.

6.2.1 Error Mitigation

Noise coupled to a quantum system results in unintended coherent, and incoherent evolution. Deviations of the realised state from the target state are known as errors, and, if left unchecked, lead to the failure of quantum algorithms. This failure can either be due to a reduction of signal to noise ratio, or a bias in the measured expectation value, resulting in misleading conclusions. Ultimately, the goal of quantum error correction (QEC) is to detect and correct these errors on the logical state level through embedding information in larger Hilbert spaces [173]. Full QEC will allow for the execution of arbitrarily long algorithms ensuring errors (of up to a certain distance) do not accumulate over time. However, current QEC methods require large overheads in both qubit number and circuit depth, and are therefore technically challenging on current quantum hardware.

Furthermore, QEC is typically discussed with respect to discrete “gate-based” quantum computers, which despite its promise to be universal, pay for that privilege in often inefficient decompositions of target unitaries into a discrete set of gates. This poses an opportunity for near term devices to show quantum advantage using either short depth, uncorrected circuits [174], or leveraging CV systems to access larger computational spaces [141]. However, the impact of errors on uncorrected computation must be carefully managed.

We consider algorithms that prepare some final state of interest and then take expectation-value measurements of some observable on that state. Uncertainty of the approximated expectation value due to averaging over a finite number of discrete samples, known as shot-noise, may be arbitrarily reduced by increasing the samples, N . However, due to imperfect execution of a given circuit, and coupling to the environment, there will be some systematic shift of the estimated expectation value compared to the ideal value of the noiseless state, $\hat{\rho}_{\text{id}}$. This shift, or bias, of the approximated observable may be reduced by the application of quantum error mitigation (QEM) techniques [28, 175].

QEM protocols are varied in nature, making it difficult to give an overarching




Figure 6.4: A schematic explanation of the goal and cost of QEM. Deviations of the expectation value of an operator $\hat{O}$ on a noisy state, $\hat{\rho}$, and on ideal state $\hat{\rho}_{\text{id}}$ are separated into shot-noise and bias components. Shot noise sets the variance of the distribution of estimated expectation values, while the bias shifts the center of the distribution (blue). Error mitigation reduces bias at the cost of increasing variance (green).

mechanism of operation. Reference [28] defines QEM as “... algorithmic schemes that reduce the noise-induced bias in the expectation value by postprocessing outputs from an ensemble of circuit runs, using circuits at the same noise level as the original unmitigated circuit or above.” This highlights the main distinction to QEC protocols: the physical noisy state is never corrected (often it is in fact degraded through the use of QEM) — rather the output from the noisy circuits is manipulated and postprocessed in such a way that we *effectively* recover an expectation value with suppressed deviation from the ideal value.

As a short example, consider a noisy state $\hat{\rho} = \lambda \hat{\rho}_{\text{id}} + (1 - \lambda) \hat{\rho}_{\text{err}}$, where we have some probability of preparing the ideal state, and some probability of preparing an error state, $\hat{\rho}_{\text{err}}$. We take N measurements of some observable on this mixed state, with the consecutive sample outcomes listed as $\vec{n} = (n_0, n_1, \dots, n_N)$. Now consider the (unphysical) perfect QEM, consisting of an oracle that “knows” if each sample outcome originated from a measurement of $\hat{\rho}_{\text{id}}$ or of $\hat{\rho}_{\text{err}}$, and flags them as such. In a postprocessing step, we may discard all samples associated with the error state, and retain the subset of samples of the ideal state, $N_{\text{id}} \subset N$. With this subset we may calculate the bias-free expectation value, however at the cost of increased shot noise, as $N_{\text{id}} \leq N$, and therefore increased variance of the estimated observable under many such experimental trials. This example, although fictitious, retains two important aspects of many QEM methods. First, we effectively “filter out” undesired information by some form of postprocessing, and second, this comes at a sample cost, where we effectively exchange a suppression of bias to an increase in

shot noise. This trade-off is shown schematically in Figure 6.4, for this oracle QEM.

Error mitigation protocols are diverse in their resource efficiency, assumptions on error-types, and scaling with system size. In this thesis, we consider the virtual distillation protocol [91–93] — a QEM method, applicable to both DV and CV systems, with few assumptions on the underlying noise model [104, 176]. It requires a fixed depth circuit appended to the state-preparation circuit under test, and a reasonable resource overhead of $k \times$ more physical qubits (or qumodes) for a k -exponential decrease in error bias.

Virtual distillation has been previously demonstrated experimentally on the superconducting platform to suppress errors in chemical simulations [177], and in neutral atoms to suppress residual thermal effects on observables of motional states [91]. The first example uses discrete-variable encoding, while the latter uses continuous-variable systems. Both examples used two-copies of the noisy state to approximate the purified state for measurement. The performance of the virtual distillation protocol is highly dependent on the presence of correlated errors between subsystems. Different experimental platforms have different dominant error channels. Therefore, it is important to verify the efficacy of virtual distillation on the ion-trap platform, especially in the presence of spatially correlated errors, which are prevalent when using global optical or microwave control fields.

With the apparatus constructed for this thesis, we may begin to explore error suppression using virtual distillation with multiple copies under realistic noise. The following section gives details on the mechanism of the protocol and practically implementable circuits.

6.2.2 Virtual Distillation Protocol

Virtual distillation is a purification-based QEM, utilising k -copies of a circuit to reduce the final bias of the measured expectation value of a noisy state, $\hat{\rho}$. The term *purification* is used to signify that the method effectively approximates the pure state closest to $\hat{\rho}$, with the underlying assumption that the dominant eigenvector

of $\hat{\rho}$ approximates the ideal pure state, $\hat{\rho}_{\text{id}} = |\psi_{\text{id}}\rangle\langle\psi_{\text{id}}|$. The validity of this approximation is explored in detail in reference [176], where they derive a bound for the deviation between the ideal state and dominant eigenvector, termed the “coherent mismatch”, c . This error provides the noise-floor for purification based techniques. Favourably, for many error models and circuit types, c is negligible to relevant applications, and in fact scales inversely proportionally to the number of qubits in a single copy, N , for fixed depth circuits.

The purified state may be accessed by,

$$\hat{\rho}_{\text{pur}}(k) = \frac{\hat{\rho}^k}{\text{Tr}[\hat{\rho}^k]}, \quad (6.14)$$

where the denominator acts as a normalisation factor, and for $k = 2$ is the usual definition for the purity of a mixed state. To elucidate the claim that this is the form of the purified state, we expand the spectrum of $\hat{\rho}$, assuming that the dominant eigenvector is the ideal state,

$$\hat{\rho} = \lambda_{\text{id}} |\psi_{\text{id}}\rangle\langle\psi_{\text{id}}| + \sum_{i=1}^{2^N-1} \lambda_i |\psi_i\rangle\langle\psi_i|, \quad (6.15)$$

with $\lambda_{\text{id}} > \lambda_i$ for all i , and from normalisation, $\lambda_{\text{id}} + \sum_{i=1}^{2^N-1} \lambda_i = 1$. Inserting this expansion into equation 6.14, we find,

$$\hat{\rho}_{\text{pur}}(k) = \frac{1}{\mathcal{N}} \left(\lambda_{\text{id}}^k |\psi_{\text{id}}\rangle\langle\psi_{\text{id}}| + \sum_{i=1}^{2^N-1} \lambda_i^k |\psi_i\rangle\langle\psi_i| \right), \quad (6.16)$$

with normalisation factor $\mathcal{N} = \lambda_{\text{id}}^k + \sum_{i=1}^{2^N-1} \lambda_i^k$. The exponentiation acts to relatively attenuate all error eigenvalues compared to λ_{id} , with the large k limit recovering the pure ideal state, $\lim_{k \rightarrow \infty} \hat{\rho}_{\text{pur}}(k) = |\psi_{\text{id}}\rangle\langle\psi_{\text{id}}|$.

The remaining question is: how can we experimentally access $\hat{\rho}_{\text{pur}}(k)$? In references [92, 93], it is shown that the expectation value of the purified state with respect to the applied operator $\hat{O}$ is,

$$\langle \hat{O} \rangle_{\text{pur}} = \text{Tr}[\hat{O} \hat{\rho}_{\text{pur}}(k)] = \frac{\text{Tr}[\hat{O} \hat{\rho}^k]}{\text{Tr}[\hat{\rho}^k]}, \quad (6.17)$$

and crucially may be estimated without ever physically preparing $\hat{\rho}_{\text{pur}}(k)$. Instead the numerator and denominator may be found through preparing k -copies of the

noisy state, applying a so-called “derangement” operation to entangle all copies together, and measuring the desired observable on one of the copies. This can be understood through the following identity³ providing the numerator,

$$\text{Tr}[\hat{O}\hat{\rho}^k] = \text{Tr}[\hat{O}^{(i)}\hat{S}^{(k)}\hat{\rho}^{\otimes k}], \quad (6.18)$$

where $\hat{S}^{(k)}$ is the derangement operator, which permutes the k copies of the state, and $\hat{O}^{(i)}$ is the operator acting on the i th copy. The denominator of the corrected expectation value is found by replacing the operator with the identity operator. There is some freedom in the form of the derangement, here we shall give it explicitly as the cyclic permutation of k objects,

$$\hat{S}^{(k)} |\psi_1\rangle \otimes |\psi_2\rangle \otimes \dots \otimes |\psi_k\rangle = |\psi_2\rangle \otimes |\psi_3\rangle \otimes \dots \otimes |\psi_k\rangle \otimes |\psi_1\rangle. \quad (6.19)$$

Practically, if we can perform this multi-copy derangement, $\hat{S}^{(k)}$, and single-copy operator $\hat{O}^{(i)}$, on k -copies of the noisy state, we can postprocess a collection of circuit repetitions to recover the k th purified expectation value $\langle \hat{O} \rangle_{\text{pur}}$.

In the following section we provide two possible approaches for implementing these unitaries and measurements, and give explicit circuit decompositions for both. In § 6.2.3 we explore the scheme experimentally, and particularly study the breakdown of error suppression due to correlated noise *between* copies. We give details of the impact on suppression when the underlying independence assumption of the errors on each noisy copy is violated.

Circuits

We discuss two methods of implementing virtual distillation. The first being a general approach able to act on arbitrary Hilbert spaces using Hadamard tests, and the second being more resource efficient, however not easily generalisable to arbitrary sizes of each copy, and to multi-qubit observables.

³Huggins [92] provides a proof of the identity, which we give an example for the case $k = 2$ in Appendix A.4.




Figure 6.5: Two circuits for implementing virtual distillation. (a) Hadamard test method using controlled derangement, $\hat{S}^{(k)}$, over k copies, and controlled operator, $\hat{O}$, applied on one copy. The ancilla qubit is measured to extract $\text{Tr}(\hat{O}\hat{\rho}^k)$. Repeating the circuit without the controlled operator measures normalisation factor, $\text{Tr}(\hat{\rho}^k)$, used to recover the virtually purified result, $\langle \hat{O} \rangle_{\text{pur}}$. (b) Two-copy implementation using a B -circuit, here decomposed into two entangling gates and single qubit phase gates. Measurement of both qubits in the computational basis, and classical postprocessing recovers $\text{Tr}(\hat{Z}\hat{\rho}^2)$ and purity, $\text{Tr}(\hat{\rho}^2)$, used to recover $\langle \hat{Z} \rangle_{\text{pur}}$.

We begin with the Hadamard test, to build further intuition on virtual distillation based on similar circuit topology discussed in § 6.1.1 for CV-tomography. The circuit for virtual distillation via an ancilla is shown in Figure 6.5 (a). The derangement and desired operator, $\hat{O}$, are applied conditioned on the ancilla qubit, with a final measurement only made on the ancilla. The circuit is repeated N_{shot} times and the outcomes are averaged to recover the estimated real part of the numerator in equation 6.17. The circuit is again repeated N_{shot} times with the controlled operator replaced with the identity operator. Subsequent averaging recovers the denominator, which may then be used to recover $\text{Re}[\langle \hat{O} \rangle_{\text{pur}}]$.

The controlled derangement operation may be decomposed into a series of controlled-swap gates [93], which for DV systems may be decomposed to native single and two-qubit gates accessible on the ion trap platform. For CV-systems, the derangement can be decomposed to a series of Fredkin gates [178, 179], which in turn may be accessible by either controlled-beam-splitters [151, 180] and/or controlled-parity operations [181]. These interactions are available in our system by the “two-SDF” method, as highlighted in Table 6.1.

The second method, known as the “ B -matrix” [92], is slightly more resource efficient, not requiring ancilla qubits, however does not simply scale for $k > 2$, and for non- $\hat{Z}$ measurements. The approach recognises that for $k = 2$ the symmetrised computational observable $\hat{Z}_2 = \frac{1}{2}(\hat{Z}\hat{I} + \hat{I}\hat{Z})$ commutes with the SWAP operator, $\hat{S}^{(2)}$, and consequently both observables required for virtual distillation commute,

$$[\hat{S}^{(2)}, \hat{Z}_2 \hat{S}^{(2)}] = 0. \quad (6.20)$$

Thus the two observables possess a common eigenbasis and may be simultaneously diagonalised by a single matrix⁴,

$$\hat{B} = \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} & 0 \\ 0 & -\frac{i}{\sqrt{2}} & \frac{i}{\sqrt{2}} & 0 \\ 0 & 0 & 0 & -i \end{pmatrix}. \quad (6.21)$$

In the computational eigenbasis, the observable is now given by,

$$\langle \hat{A} \rangle = \text{Tr}[\hat{B}^\dagger \hat{D} \hat{B} \hat{\rho}], \quad (6.22)$$

$$= \sum_{x \in \{0,1\}^2} d_x p(x), \quad (6.23)$$

where $\hat{A}$ is either $\hat{Z}_2$ or $\hat{S}^{(2)} \hat{Z}_2$, and the second line follows by the definition of the diagonalised operator $\hat{D} = \sum_x d_x |x\rangle \langle x|$, and $p(x) = \langle x | \hat{B} \hat{\rho} \hat{B}^\dagger | x \rangle$ is the probability of observing computational state $|x\rangle$ after the basis change given by $\hat{B}$.

As a result, both numerator and denominator from equation 6.17 may be simultaneously measured by computational basis measurements after applying a two-qubit basis rotation⁵. The values are related to computational expectation values by,

$$\text{Tr}[\hat{\rho}^2] = \frac{1}{2}(1 - \langle \hat{Z} \otimes \hat{I} \rangle + \langle \hat{I} \otimes \hat{Z} \rangle + \langle \hat{Z} \otimes \hat{Z} \rangle), \quad (6.24)$$

$$\text{Tr}[\hat{Z} \hat{\rho}^2] = \frac{1}{2}(\langle \hat{Z} \otimes \hat{I} \rangle + \langle \hat{I} \otimes \hat{Z} \rangle). \quad (6.25)$$

Figure 6.5 (b) shows the B -matrix decomposition into native gates to our platform. A demonstration of this simple two-copy example of virtual distillation is given in the following section.

⁴Note, we use a different convention than reference [92]. This is done for a simpler decomposition to native gates on our apparatus.

⁵Although not discussed here, this method scales simply to $\hat{Z}$ measurements on copies containing multiple qubits through parallel pairwise rotations [92].




Figure 6.6: Verification of the physically implemented B -circuit. (a-b) Preparing the initial state $|\downarrow\uparrow\rangle$, applying global $\hat{R}_x(\theta)$ rotations, and measuring the two-qubit joint probabilities provides a clear signal for the correct application of the B -matrix. (a) No B -circuit is applied. (b) B -circuit applied. (c) Globally applied Ramsey experiment followed by the B -circuit and subsequent postprocessing to recover the purity and numerator term given by equations 6.24.

6.2.3 Experimental Demonstration

This section begins with experimental details on implementing virtual distillation via the B -matrix decomposed to native two-qubit $\hat{R}_{xx}$ gates and single qubit phase shifts, $\hat{R}_z$. With this protocol we can then study the efficacy of virtual distillation in two different noise regimes — independent depolarisation, and correlated errors due to common random-unitaries (see § 5.3.3 for the noise model). This work provides evidence of the performance of virtual distillation under realistic structured errors originating from fluctuations of a common external field.

Virtual Distillation on Two Copies

We demonstrate virtual distillation via the B -circuit shown in Figure 6.5 (b) on two-copies of a single qubit. The qubit is encoded in the optical transition $|4S_{1/2}, m_j = -\frac{1}{2}\rangle \leftrightarrow |3D_{5/2}, m_j = -\frac{5}{2}\rangle$ of a $^{40}\text{Ca}^+$ ion in a quantising field of ~ 5 G. Gates are driven using 729-nm light focussed onto the ions by the individual addressing system detailed in § 3.6.5. Two-qubit entangling gates of the form $\hat{R}_{xx}$ are implemented using Walsh-1 MS gates on the in-phase upper radial mode, as characterised in § 4.3.2. For this decomposition, we specifically require $\hat{R}_{xx}(\frac{\pi}{4})$ -gates,

as opposed to the $\hat{R}_{xx}(\pi/2)$ maximally-entangling two-qubit gates detailed in § 4.3.2. To calibrate these gates, two consecutive $\hat{R}_{xx}(\frac{\pi}{4})$ pulses are applied, and a partial state-tomography protocol [97] is utilised (as summarised by equation 4.14).

After calibration, the Walsh-1 $\hat{R}_{xx}(\frac{\pi}{4})$ gate settings are the following: 21 μW total 729-nm beam power, -15 kHz detuning from in-phase upper radial motional mode, 10 μs ramp duration (ramped on and off for both arms of the Walsh), and a 67.5 μs FWHM pulse duration for each arm. The total pulse duration including ramp durations is then 155 μs . Stark-shifts must be carefully calibrated to ensure consecutive MS gates have spin-phases set relative to the qubit. The Stark-shift for this power is calibrated using an SDF on a single ion (described in § 2.4.2), finding a shift of 2.49 kHz. This phase is accounted for by updating the RF-tones on the SP AOM for subsequent pulses. With these settings we measure a fidelity for two-consecutive $\hat{R}_{xx}(\frac{\pi}{4})$ gates of $\mathcal{F}_2 = 96.0(8)\%$, with population component of $P_{\downarrow\uparrow} + P_{\uparrow\downarrow} = 0.010(3)$, and parity contrast of $C = 0.93(2)$. The inferred single $\hat{R}_{xx}(\frac{\pi}{4})$ gate fidelity, assuming depolarising errors, is then $\mathcal{F} = 98.0(4)\%$.

Single qubit phase shifts, $\hat{R}_z$, are executed virtually. As we require individual $\hat{R}_z$ phase shifts on the two qubits, we update the relative phases of the two pairs of RF tones applied to the acousto-optic deflectors, used for generating the addressed spots. As these single-qubit gates are virtual, we expect negligible error contributions.

The final computational-basis measurements are executed by state-dependent fluorescence, described in § 4.1.1. From the single-shot outcomes, the joint probabilities P_{ij} , with $i, j \in \{\downarrow, \uparrow\}$ denoting the bright and dark states of ions 1 and 2 are recovered, respectively. The required expectation values given in equations 6.24 are then,

$$\langle \hat{Z} \otimes \hat{I} \rangle = P_{\downarrow\downarrow} + P_{\downarrow\uparrow} - P_{\uparrow\downarrow} - P_{\uparrow\uparrow}, \quad (6.26)$$

$$\langle \hat{I} \otimes \hat{Z} \rangle = P_{\downarrow\downarrow} - P_{\downarrow\uparrow} + P_{\uparrow\downarrow} - P_{\uparrow\uparrow}, \quad (6.27)$$

$$\langle \hat{Z} \otimes \hat{Z} \rangle = P_{\downarrow\downarrow} - P_{\downarrow\uparrow} - P_{\uparrow\downarrow} + P_{\uparrow\uparrow}. \quad (6.28)$$

To initially verify the performance of the physical implementation of the B -circuit, two heuristic tests are used⁶. First, we want a clear indication that the transformation due to the applied B -circuit matches the convention given by the B -matrix in equation 6.21. Inspecting the B -matrix, we see the central block acts distinctly on the $\{|\downarrow\uparrow\rangle, |\uparrow\downarrow\rangle\}$ subspace. As such, a method that preferentially probes this subspace would give a clear signal for conformity between the applied B -circuit and the desired B -matrix. To access this subspace, a singly addressed $\hat{R}_x(\pi)$ -pulse is applied to prepare the initial state $|\downarrow\uparrow\rangle$. A global $\hat{R}_x(\theta)$ is then applied with variable rotation angle, θ . Figure 6.6 (a) shows the measured two-qubit populations as a function of θ without the B -circuit applied. Panel (b) shows the populations with the global rotation and the B -circuit applied. The solid lines correspond to the expected populations inferred from equation 6.21, while points are experimental data with $N = 150$ samples per probed rotation angle. We can see good qualitative agreement between the observed data, and the desired behaviour suggesting the phase relationships enforced by the B -circuit indeed conform with the ideal B -matrix.

The second heuristic check verifies the conventions for recovering the numerator and denominator in equations 6.24, from the qubits measured joint probabilities, given in equations 6.26. To test this, a simple zero-delay Ramsey experiment is performed on both qubits simultaneously. The qubits are initially prepared in $|\downarrow\downarrow\rangle$. A global pulse $\hat{R}_x(\frac{\pi}{2})$ is then applied, followed by a final analysis pulse $\hat{R}_\phi(\frac{\pi}{2})$, applied with variable spin-phase, ϕ . After these pulses, the B -circuit is applied, and the two-qubit populations are measured. As the two qubits remain in identical pure-states to one another (up to errors in the parallel single-qubit gates), we expect the inferred purity to remain unity, $\text{Tr}[\hat{\rho}^2] = 1$. Given this, we expect the recovered numerator term to be equivalent to typical Ramsey oscillations on a single qubit,

$$\text{Tr}[\hat{Z}\hat{\rho}^2] = \langle \hat{Z} \rangle_\phi = -\cos(\phi). \quad (6.29)$$

⁶Process-tomography would be ideal for this two-qubit implementation, however we have not yet implemented this on our apparatus.




Figure 6.7: Verifying virtual distillation under injected local depolarising and two-qubit correlated errors. Random Clifford sequences of length m are followed by the B -circuit. Local errors are generated using independent Cliffords on each copy, while correlated errors are generated by setting $C_{i,L} = C_{i,R}$.

Figure 6.6 (c) displays the post-processed purity and numerator terms extracted as a function of analysis pulse phase, ϕ , with solid lines showing the expected behaviour of the ideal B -circuit. Again we see excellent agreement between experimental and desired behaviour, giving evidence that we are successfully applying virtual-distillation.

In the next section we use virtual distillation for its intended purpose — to suppress the impact of errors. We will discuss the efficacy of virtual distillation under two different noise models, providing experimental evidence for both.

Virtual Distillation under Correlated Errors

To quantify the achieved suppression of errors through the use of virtual distillation, we must characterise and inject noise with a known structure. Applying randomised Clifford sequences, the same as those used in standard randomised benchmarking (RB) (details in § 4.1.4), is a convenient method for enforcing a simple structure on injected errors. We frame the task of quantifying the observed error suppression in the context of single-qubit RB by comparing the exponential decay due to some error-per-Clifford, ε_{raw} , to the effective decay profile of the error-mitigated observable, $\langle \hat{Z} \rangle_{\text{pur}}(m)$, where m is the Clifford sequence length.

Figure 6.7 shows the circuit implemented for controlled error injection. Random Clifford sequences with a final inverting gate are applied in parallel to both qubit copies (now labelled “left”, L , and “right”, R), effectively driving depolarising channels. The total circuit error rate is set by an error-per-Clifford, ε , and the sequence length, m . The B -circuit is then applied followed by $\hat{Z}$ measurements. Postprocessing of the two-qubit joint probabilities recover the error-suppressed expectation value, $\langle \hat{Z} \rangle_{\text{pur}}$.

The RB protocol is repeated without the applied B -circuit to recover the standard single-qubit exponential decays, $\langle \hat{Z} \rangle_{\text{raw}}$, and the corresponding error-per-Clifford, ε_{raw} .

Crucially, with this circuit we may explore two different error regimes: first, we may apply independent single-qubit Clifford sequences to each copy, which locally twirls the error channel acting on each qubit. This decorrelates errors between the two copies [79, 182] leading to favourable performance of virtual distillation.

In contrast, with the second regime we purposefully apply the same Clifford sequence to each qubit, i.e. by enforcing $C_{i,L} = C_{i,R}$. In this regime, any common-mode errors, for example due to noisy global control fields, or systematic coherent control offsets, will not be decorrelated by the parallel twirling (see *k-copy* twirling in Chapter 5). These correlated errors can significantly degrade the performance of virtual distillation. We shall spend some time now to make this connection concrete. Consider the k -copy mixed state, $\hat{\rho}^{(k)}$, for a common-mode error channel,

$$\hat{\rho}^{(k)} = \sum_i \kappa_i \hat{\rho}_i^{\otimes k}. \quad (6.30)$$

Applying virtual distillation, using equation 6.17, and the identity equation 6.18, the effective measured observable is,

$$\langle \hat{O} \rangle_{\text{pur,c}} = \frac{\sum_i \kappa_i \text{Tr}[\hat{O} \hat{\rho}_i^k]}{\sum_i \kappa_i \text{Tr}[\hat{\rho}_i^k]}, \quad (6.31)$$

where crucially the trace operator is now within the sum. Clearly, the components κ are not exponentiated in the same manner as the λ coefficients of the previously purified density matrix, $\hat{\rho}_{\text{pur}}$, in equation 6.16 — as such the dominant state is no longer effectively amplified through derangement. In the specific case of common random unitary noise, discussed in § 5.3.3, $\hat{\rho}_i$ remains a pure state, $\hat{\rho}_i = U_i |\psi_{\text{id}}\rangle \langle \psi_{\text{id}}| U_i^\dagger$, as such the effective observable reduces to $\langle \hat{O} \rangle_{\text{pur,c}} = \text{Tr}[\hat{O} \hat{\rho}] = \langle \hat{O} \rangle_{\text{raw}}$, where $\langle \hat{O} \rangle_{\text{raw}}$ is the expectation value of a single copy, $\hat{\rho} = \sum_i \kappa_i \hat{\rho}_i$. For this particular correlated error channel, virtual distillation is entirely ineffective.

Quantifying the proportion of correlated and independent error components for a pair of qubits is exactly the task *k-copy* randomised-benchmarking was designed for,

Twirling type	$\varepsilon_L/10^{-3}$	$\varepsilon_R/10^{-3}$	$\varepsilon_{\text{corr}}/10^{-3}$	$\varepsilon_{\text{raw}}/10^{-3}$
Independent	7(1)	7(1)	4(1)	11(2)
Identical	2.0(4)	1.4(4)	9.3(5)	11.2(6)

Table 6.2: Local depolarising ($\varepsilon_L, \varepsilon_R$) and correlated error ($\varepsilon_{\text{corr}}$) contributions of both the independently and identically twirled channels on two copies of a single-qubit. These error-per-Cliffords are characterised by the *k-copy* RB method described in § 5.3. The raw error (ε_{raw}) is the error-per-Clifford determined through standard single-qubit RB, described in § 4.1.4.

details in Chapter 5. The independent depolarising error components, labelled as $\varepsilon_L, \varepsilon_R$, and the two-qubit correlated component, $\varepsilon_{\text{corr}}$, may be estimated for both the “independent” and “identical” twirled channels. These error-per-Cliffords are then used to predict the performance of virtual distillation. Table 6.2 summarises the measured error-per-Clifford gates for both twirling schemes. The interleaved gate sequences are implemented by alternating the position of the single-ion addressing beam after each Clifford, requiring a delay of ≈ 10 μs between gates. To ensure a baseline correlated unitary error is present, the single-qubit gates comprising the Clifford sequence are driven off-resonantly by $\delta = 2\pi \times 1$ kHz. The different twirling methods change the distribution of errors between correlating and depolarising, while retaining a consistent total single-qubit error rate. The residual estimated correlated error of $\varepsilon_{\text{corr}} = 4(1) \times 10^{-3}$ in the case of independent twirling suggests that the twirling circuit depth was not sufficient to entirely decorrelate the errors between the two qubits.

The analytic decay of the mitigated observable $\langle \hat{Z} \rangle_{\text{pur}}$ when the two-copies are subject to a shared error channel, Λ_T , consisting of *k-copy* twirled common random unitaries (parameterised by $\varepsilon_{\text{corr}}$) and local depolarising errors on the two copies (parameterised by ε_L and ε_R) is given by,

$$\langle \hat{Z} \rangle_{\text{pur}} = (1 - 2\varepsilon_{\text{corr}})^m \frac{\cosh(m(\varepsilon_L - \varepsilon_R))}{\cosh(m(\varepsilon_L + \varepsilon_R))}, \quad (6.32)$$

where we have assumed small errors. This is simply calculated by extending equation 5.28 to include the action of the B -matrix, here using the superoperator form,

$$P_{LR} = \langle \langle L, R | B \Lambda_T^m | \downarrow\downarrow \rangle \rangle, \quad (6.33)$$




Figure 6.8: Verifying the performance of virtual distillation under either dominant local depolarising or correlated errors. The effective suppression of errors in the single-qubit $\hat{Z}$ observable using virtual distillation on two-copies. The green points correspond to the raw single-qubit decay due to the physical errors, with the solid green line being an exponential fit. The dotted lines correspond to the expected error mitigated single qubit decays inferred by the model given in equation 6.32, and the error-per-Cliffords measured by *k-copy* RB. The red and blue points correspond to the error mitigated $\hat{Z}$ observables for in the regime of dominant correlated errors and local depolarising errors respectively.

where P_{LR} is the joint probability for the two qubits, with $L, R \in \{\downarrow, \uparrow\}$. The effective observable $\langle \hat{Z} \rangle_{\text{pur}}$ is then recovered by equations 6.26.

For small m , and assuming small errors, this may be expanded to second order as,

$$\langle \hat{Z} \rangle_{\text{pur}} \approx 1 - 2m\varepsilon_{\text{corr}} - 2m^2\varepsilon_L\varepsilon_R, \quad (6.34)$$

where we can see that correlated errors maintain a linear decay, while depolarising errors are suppressed to a quadratic decay. We can also see that in the ideal case of exclusively depolarising type errors of equal magnitude, the effective observable decays as $\langle \hat{Z} \rangle_{\text{pur}} = \text{sech}(2m\varepsilon_L)$.

With this model, and the quantified error components, we plot the expected effective error-mitigated decay profiles in Figure 6.8. The green solid line is the raw decay of a single qubit as found by standard RB, with an error-per-Clifford given by ε_{raw} . The blue dotted line corresponds to the model in equation 6.32 for the “independent” twirled error components given in Table 6.2. The red dotted line corresponds to the same model using the “identical” twirled error components.

The blue and red points correspond to the postprocessed $\langle \hat{Z} \rangle_{\text{pur}}$ recovered through measurements after the applied B -circuit. We can see that there is good agreement between the expected mitigated decay profiles and the observed data. The final decay to $\langle Z \rangle < 0$ suggests we either have a systematic readout error, or the B -circuit contributes a systematic coherent error.

An effective error-per-Clifford for small total circuit errors may be calculated by the survival decay gradient at the limit $m = 0$. The inset plot in Figure 6.8 shows this initial difference in gradient of the data and fitted models. Analytically this is given by the first order decay of equation 6.34, which is dependent exclusively on the correlated error component. For $\langle Z \rangle_{\text{pur}}$ with independent Clifford twirling sequences, an effective error suppression of $g = 2.8(7)$ is found, where g is the ratio of gradients,

$$g = \left(\frac{d\langle Z \rangle_{\text{raw}}}{dm} \Big|_{m=0} \right) / \left(\frac{d\langle Z \rangle_{\text{pur}}}{dm} \Big|_{m=0} \right) = \frac{\varepsilon_{\text{raw}}}{\varepsilon_{\text{corr}}}. \quad (6.35)$$

For identical twirling sequences a small suppression of $g = 1.20(9)$ is estimated, highlighting the negative impact of correlated errors on the performance of virtual distillation, as well as the importance of sufficiently decorrelating errors, which we do here by local Clifford twirling.

6.2.4 Outlook

Quantum error mitigation is accessible to near-term quantum computers, under some circumstances allowing for resource-efficient, and substantial reductions in measurement bias. In this section we have provided some background on quantum error mitigation via virtual distillation, highlighting two possible implementations: a Hadamard test, and the B -circuit. We demonstrate the B -circuit on the trapped-ion platform, and use the protocol to explore the ability to suppress errors under unfavourable correlated error channels. These demonstrations provide the groundwork for two future research directions — exploring the practical achievable advantage when extending to more than two copies; and extending to error-mitigation acting on a hybrid system.

Copy Num, k	Total $\hat{R}_{xx}(\frac{\pi}{4})$	Total $\hat{R}_x(\frac{\pi}{2})$	Expected Faults
$2(B)$	2	0	0.02
2	22	7	0.117
3	33	14	0.234
4	44	21	0.351
5	55	28	0.468

Table 6.3: Required number of $\hat{R}_{xx}(\frac{\pi}{4})$ and $\hat{R}_x(\frac{\pi}{2})$ gates to implement the derangement circuit for virtual distillation on multiple copies. Apart from $n = 2(B)$, which uses the B -circuit, all other rows are using the ancilla-based Hadamard test circuit, implemented by controlled-SWAP gates. We give the expected total faults assuming fidelities of 99% for two-qubit operations, and 99.9% for single qubit operations.

Extending to > 2 copies

From equation 6.16, we can see that increasing the copy number should exponentially approximate the dominant pure state. To our knowledge virtual distillation with > 2 -copies has not been experimentally demonstrated. We want to scale past two copies, and either verify this exponential suppression, or study why practically this suppression breaks down. We predict that twirling the error channels incident on each copy will be crucial to mitigate multi-qubit correlations induced by structured error channels.

For more than two-copies we will use the ancilla-based Hadamard circuit, as shown in Figure 6.5 (a). This method simplifies the readout procedure and decomposes the derangement straightforwardly into controlled-SWAP gates.

For n copies of noisy state, $\hat{\rho}$, the controlled derangement can be decomposed to a minimum of $n - 1$ controlled-SWAP gates [183]. Using accessible native gates, these controlled-SWAP are decomposed to $n_{2Q} = 11$ $\hat{R}_{xx}(\frac{\pi}{4})$ gates and at most $n_{1Q} = 7$ single qubit $\hat{R}_x(\frac{\pi}{2})$ rotations. Assuming that we can implement $\hat{R}_{xx}(\frac{\pi}{4})$ with fidelity, $\mathcal{F} = 99\%^7$, and single qubit gates with fidelity, $\mathcal{F} = 99.9\%$, we may calculate the expected total faults per circuit for a given copy size. Here, we assume the error on each gate is independent, and define expected fault number as $N_f = n_i(1 - \mathcal{F}_i)$. Table 6.3 summarises the gate scaling and expected fault number

⁷We have currently demonstrated $\mathcal{F} \approx 98\%$ gates, however we believe this can be further improved through systematic calibration, and an improved PDH lock of the 729-nm laser.




Figure 6.9: Proposed circuit for applying virtual distillation to a hybrid system. Using the ancilla method for measuring $\text{Tr}(\hat{O}\hat{\rho}^k)$, we may decompose the controlled derangement to the transversal application of Fredkin gates acting on either qubit or qumode subsystems.

for the derangement operation. The first row gives the two-copy method via the B -circuit for reference.

To effectively suppress errors, we (heuristically) desire that the derangement operation contributes fewer faults than the state-preparation of the noisy copy. Further, we desire that the total circuit has at most $N_f \approx 1$. From the table, we see that although it might be possible to demonstrate up to $k = 5$ with our expected gate errors, this might be currently out of reach due to the large error contribution from the derangement. The difficulty in predicting the efficacy of the noisy derangement is that the impact of the errors on the performance of virtual distillation highly depends on their structure. Practically observing the suppression of circuit errors with copy number k , will thus provide evidence of the impact of individual gate errors on the derangement. It may also be useful to further characterise the individual gate errors using tomographic methods to predict the performance of the noisy-derangement.

Virtual Distillation on Hybrid Systems

Hybrid systems, as discussed in § 6.1, offer advantages in computation, due to the large accessible state space, with low numbers of physical information

carriers. Quantum harmonic oscillators (qumodes) also have applications in quantum simulation [184] and sensing [185, 186]. However, due to environmental and control noise, the performance and reliability of these systems has strong limitations. Virtual distillation offers a resource efficient method for suppressing errors in CV systems if a pure state is targeted. The proposal for virtual distillation originated from measuring an observable on a thermal qumode state effectively prepared at a lower temperature than the constituent noisy copies [91].

We want to extend this protocol to a register of qubits and qumodes, applying transversal derangement, as seen in Figure 6.9, to suppress errors in hybrid systems. Particularly, the mitigation of dephasing and Gaussian errors, such as coherent displacements, on the qumode is an under-explored topic, and poses an opportunity for greatly reducing the bias in quantum simulation tasks [187].

To access this, we require a Fredkin (controlled-SWAP) operation acting between an ancilla qubit and two qumodes. For certain Fock-state distributions, this may be realised by a single controlled-beamsplitter interaction [151], which we have simple access to via the “two-SDF” method in § 6.1.1. However, for the general case, we require both parity and beamsplitter operations, with control on either gate [188]. Demonstrating the Fredkin on a hybrid system is also of particular interest due to its reversible nature, and the practical advantage that it conserves the total excitation number of the system [178].

Limitations

Error mitigation is resource efficient, but only allows access to measurements of an *effectively* error suppressed state, $\hat{\rho}_{\text{pur}}$ — entropy is never physically removed from the system, as is the case for QEC. The required samples for a fixed variance of the measured observable scales exponentially with the expected total number of faults in the circuit, N_f [28]. As such, QEM is only practically advantageous up to $N_f \approx 1$ (although this is somewhat a heuristic bound, and depends on many implementation details). QEM is thus suited for near-term applications on non-fault

tolerant, short-depth circuits.

As shown in § 6.2.3, correlated and structured errors between copies can greatly reduce the efficacy of virtual distillation. This may limit the practicality of virtual distillation on some experimental platforms relying on global addressing fields. This may be mitigated either through appropriate independent twirling of the copies or through random circuit compositions [182]. Furthermore, physically separating the copies using networked modules can ensure that errors originating from external environmental factors are no longer common-mode between copies [183].

7

Conclusion

7.1 Summary

In this thesis we constructed an experimental apparatus to enable access to a multi-qubit and qumode register. We characterised this apparatus, and have fully demonstrated the required elementary control for the qubit subsystem. We developed a novel and scalable method to characterise noise that induces correlated error on multi-qubit registers, and explored this experimentally with the apparatus. We provided an outline for fully utilising the qumodes in a hybrid qubit-qumode system using spin-conditioned bosonic interactions and experimentally verified the underlying method to drive these interactions in a scalable way using single-ion-addressing optical beams. Finally, we identify virtual-distillation as an error mitigation technique amenable to hybrid qubit-qumode registers, demonstrate its efficacy on the qubit subsystem under both depolarising and correlated errors, and give a future perspective on its implementation on the full hybrid system.

The Experimental Platform

The constructed apparatus was initially designed to enable the study of fast two-qubit entangling gates on multi-ion chains, which on a previous apparatus was

limited by the attainable optical field intensity, and the inability to selectively drive only two qubits in a multi-qubit register. As such the new apparatus features: a modern segmented Paul RF trap with large-solid-angle optical access, dual high-NA objective lenses for providing near-diffraction-limited high-intensity optical fields at the ion plane, and a crossed-AOD single-ion addressing system for spatial selectivity. The construction of the apparatus began in earnest at the start of 2023, with first ions trapped in December that year. Commissioning the apparatus lasted until mid-2025, when the single-addressing system was installed. This marked a transition from build-up and characterisation, to investigating novel methods and demonstrations on the apparatus. At the time of writing, there still remains one considerable apparatus upgrade: the optical standing waves.

The work in this thesis has utilised travelling-wave single-addressing beams. The next major hardware extension, planned for the end of 2026, is installation of the counter-propagating addressing system. With active phase stabilisation, using a similar scheme to the previous apparatus [138], this will enable an array of single-ion addressed optical standing waves. These standing waves allow for two major capabilities: first, we will have access to sharp electric field and intensity spatial gradients. These sharp features can strongly drive spin-dependent forces, enabling fast spin-boson-gates and boson-mediated two-qubit entangling-gates. Second, the ion may be placed in the structured optical field to suppress unwanted interactions. For example, sitting on the maximum curvature of the optical field suppressed the undesired carrier-interaction that was limiting gate-speed in two-qubit MS gates [41]. Further, the spatial Stark-shift profiles due to the structured optical field intensity can be used to directly access spin-mediated bosonic interactions, which have been previously proposed and shown with the radial gradients of tightly focussed beams [157], and by the axial structure of phase-stable standing waves [156, 189]. Completion of this hardware upgrade will place our apparatus as a versatile, and modern platform for studying both fast interactions on multi-qubit chains, and the control of hybrid systems with multiple qubits and qumodes.

Scaling Characterisation

With access and control of ≈ 10 qubits, full tomography is already unfeasible on our system. We considered the problem of characterising structured errors on a multi-qubit register, and developed *k-copy RB* to provide accessible metrics of commonplace correlated noise sources. This method has practical utility in benchmarking scaled systems, and verifying the efficacy of noise tailoring techniques, as we found when studying virtual-distillation. We have demonstrated this method on the trapped-ion platform under common-mode noise due to fluctuating globally acting fields, however this method may be applied to all quantum computing platforms with access to qubit-resolved readout. The method is also applicable to other correlating noise-sources, such as always on ZZ -interactions relevant to superconducting qubits, and to certain non-Markovian error channels leading to correlations over time. This work contributes a useful and scalable tool for benchmarking, and categorising the errors impacting experimental quantum computers.

Scaling Resources

Trapped ions provide two possible quantum information carriers, the discrete electronic energy levels, and the continuous collective motional modes. We considered a strategy to fully utilise both these resources, and access a considerable sized, and controllable Hilbert space. The coherence of both subsystems was characterised on the constructed apparatus, a practical gate-set for universal hybrid control was discussed, and the experimental method for driving the spin-conditioned bosonic gates was demonstrated. This work was a continuation of previous demonstrations in our group [42, 43, 45], now utilising the improved motional coherence times, and the single-ion addressing system of the new apparatus. By establishing this control on the new apparatus, we enable further research into hybrid systems on the trapped-ion platform. Recent work by our group has identified multiple protocols and tools utilising spin-conditioned nonlinear bosonic interactions applied to multiple qubits [140]. These include squeezing-mediated geometric phase gates, bosonic state preparation, and effective N -body spin interactions. With this new

platform, we have the ability to explore these proposals, and develop the further experimental control for their implementation.

Scaling Imperfect Systems

To suppress the impact of errors, we experimentally investigated an implementation of virtual distillation. We demonstrated that under common-mode noise, which is typical when using qubits that are spatially near to one another or when using global control fields, virtual distillation is rendered ineffective. This evidence was further supported through *k-copy* RB characterisation and through structured error injection. This constitutes a demonstration for the need of noise-tailoring techniques [182], to enforce the underlying assumption of independent, or low-weight errors required in many error mitigation and correction strategies. Furthermore, we discussed the application of virtual distillation to more than two copies, and to hybrid CV-DV systems. Error mitigation on the large accessible Hilbert space of multi-qubit and qumode present on our apparatus will enable short hybrid computations, and simulations. These extended resources can allow for further development of experimentally simulated lattice gauge theories [44, 45], computing Feynman diagrams through phase-space measurements [187], and resource efficient simulation of chemical systems [184, 190, 191].



This thesis followed the build-up, characterisation, and first experimental demonstrations of a new ion-trap apparatus. The resulting platform now supports several distinct research programmes, and is intended to enable experiments extending well beyond the work presented here.

Appendices

A

Appendix

A.1 Magnus expansion

The unitary evolution due to a time dependent Hamiltonian may be written in the ordered, short-time limit as,

$$U(t) = \lim_{N \rightarrow \infty} e^{\frac{-i}{\hbar} \hat{H}(t_N) \Delta t} \dots e^{\frac{-i}{\hbar} \hat{H}(t_2) \Delta t} e^{\frac{-i}{\hbar} \hat{H}(t_1) \Delta t} = e^{A(t)}, \quad (\text{A.1})$$

where $\Delta t = t/N$, and $t_n = n\Delta t$, and $A(t)$ is a time dependent exponent, which is expanded in the Magnus construction to [161],

$$A(t) = \sum_{k=1}^{\infty} A_k(t). \quad (\text{A.2})$$

The first three $A_k(t)$ terms are given by,

$$A_1(t) = \frac{-i}{\hbar} \int_0^t dt_1 \hat{H}(t_1) \quad (\text{A.3})$$

$$A_2(t) = -\frac{1}{2\hbar^2} \int_0^t \int_0^{t_1} dt_1 dt_2 [\hat{H}(t_1), \hat{H}(t_2)] \quad (\text{A.4})$$

$$A_3(t) = \frac{i}{6\hbar^3} \int_0^t \int_0^{t_1} \int_0^{t_2} dt_1 dt_2 dt_3 \left([\hat{H}(t_1), [\hat{H}(t_2), \hat{H}(t_3)]] + [\hat{H}(t_3), [\hat{H}(t_2), \hat{H}(t_1)]] \right), \quad (\text{A.5})$$

where we can see that if the Hamiltonian fails to commute with itself at different times, the higher-order Magnus terms become relevant. This may be understood as the continuous analogy of the Baker-Campbell-Hausdorff formula. It is particularly

useful in this thesis to understand the origin of the geometric phase in the MS gate scheme, and the nonlinear bosonic interactions generated in the “two-SDF” protocol given in § 6.1.1.

A.1.1 AC Stark Shift

Here we analyse AC Stark shifts due to far off-resonant optical interactions. The interaction between a spin and an oscillating electric field is given by equation 2.8 (in the interaction picture of the spin, and neglecting fast oscillating terms), reproduced here as,

$$\tilde{H}_I \approx \frac{\hbar\Omega}{2} \hat{\sigma}_+ e^{i(\vec{k} \cdot \vec{R} - \delta t + \phi)} + \text{h.c.} \quad (2.8)$$

Neglecting the motional term, taking $\phi = 0$, and setting δ to some large detuning $|\Delta| \gg \Omega$, the Hamiltonian is,

$$\tilde{H}_I \approx \frac{\hbar\Omega}{2} (\cos(\Delta t) \hat{\sigma}_x - \sin(\Delta t) \hat{\sigma}_y). \quad (A.6)$$

The first two Magnus terms are,

$$A_1(t) = -\frac{i\Omega}{2\Delta} (\sin(\Delta t) \hat{\sigma}_x - (1 - \cos(\Delta t) \hat{\sigma}_y)), \quad (A.7)$$

$$A_2(t) = -\frac{i\Omega^2}{4\Delta} \hat{\sigma}_z \left(t - \frac{\sin(\Delta t)}{\Delta} \right). \quad (A.8)$$

The first term is bounded in time contributing to negligible population transfer of order $\left(\frac{\Omega}{\Delta}\right)^2$. The second term contains both a bounded oscillating and unbounded secular term in time. Neglecting the oscillating $\sin(\Delta t)$ term, and taking only the linearly increasing phase term, we find the effective Hamiltonian corresponding to rotations around $\hat{\sigma}_z$,

$$\hat{H}_{\text{eff}} = \frac{\hbar\delta_{\text{AC}}}{2} \hat{\sigma}_z, \quad \delta_{\text{AC}} = \frac{\Omega^2}{2\Delta}, \quad (A.9)$$

with rotation angle $\theta = \delta_{\text{AC}} t$.

A.1.2 Geometric Phase

Here we analyse the off-resonant motional displacement, described by the Hamiltonian given in equation 2.32, where we are in the interaction picture with respect to $\hat{H}_{\text{osc}}$, and have assumed the fast oscillating terms are negligible, reproduced here as,

$$\tilde{H}_I \approx -\frac{\hbar\Omega}{2}(\hat{a}^\dagger e^{-i(\delta t+\phi)} + \hat{a}e^{i(\delta t+\phi)}), \quad (2.32)$$

The first two Magnus terms are,

$$A_1(t) = \frac{i\Omega}{\delta} \sin\left(\frac{\delta t}{2}\right) (\hat{a}^\dagger e^{-i(\delta t/2+\phi)} + \hat{a}e^{i(\delta t/2+\phi)}), \quad (A.10)$$

$$A_2(t) = -\frac{i\Omega^2}{4\delta} \left(t - \frac{\sin(\delta t)}{\delta}\right) \hat{I}. \quad (A.11)$$

The dynamics of A_1 describe the familiar displacement operator, $\hat{D}(\beta)$ with $\beta = \frac{\Omega}{\delta} \sin\left(\frac{\delta t}{2}\right) ie^{-i(\delta t/2+\phi)}$. The maximum displacement from the origin is $\beta_{\text{max}} = \frac{\Omega}{|\delta|}$. The trajectory traces a circle of radius $\frac{\Omega}{2|\delta|}$ which returns to the initial motional state with periodicity $T_n = \frac{2\pi n}{|\delta|}$, for integer n . The second term, A_2 , contains both a bounded oscillating and unbounded secular term in time. The term acts entirely as the accumulation of a phase over time. Considering the application of the off-resonant displacement for $t = T_n$, then the dynamics reduces entirely into a single accumulating phase of,

$$\Phi(T_n) = \frac{\Omega^2}{4\delta} T_n. \quad (A.12)$$

For a single loop the phase accumulated is thus,

$$\Phi(T_1) = \frac{\pi\Omega^2}{2\delta^2}. \quad (A.13)$$

This geometric phase is utilised in the Mølmer-Sørensen scheme to perform two-qubit entangling gates.

A.1.3 “Two-SDF” Dynamics

Here we analyse the “two-SDF” Hamiltonian given in equation 6.4 using the Magnus expansion. We recover effective Hamiltonians for both Gaussian interactions, as

well as for tri-squeezing.

The “two-SDF” Hamiltonian is reproduced for convenience here,

$$\hat{H}_{2\text{SDF}}(t) = \hat{H}_\alpha(t) + \hat{H}_\beta(t) \quad (\text{A.14})$$

$$= \frac{\hbar\Omega_\alpha}{2} \hat{\sigma}_\alpha \left(\hat{a}_\alpha e^{i(\Delta_\alpha t + \phi_\alpha)} + \hat{a}_\alpha^\dagger e^{-i(\Delta_\alpha t + \phi_\alpha)} \right) \quad (\text{A.15})$$

$$+ \frac{\hbar\Omega_\beta}{2} \hat{\sigma}_\beta \left(\hat{a}_\beta e^{i(\Delta_\beta t + \phi_\beta)} + \hat{a}_\beta^\dagger e^{-i(\Delta_\beta t + \phi_\beta)} \right). \quad (\text{A.16})$$

The first three terms of the corresponding Magnus expansion are found by inserting $\hat{H}_{2\text{SDF}}$ into equations A.3. We shall later rationalise that with certain values of SDF magnitudes, detunings, and phases, we may arbitrarily suppress the higher order terms.

The A_1 term appears as separable fast oscillations in phase space due to each SDF,

$$A_1(t) = -i \sum_{i \in \{\alpha, \beta\}} \left(\frac{\Omega_i \hat{\sigma}_i}{\Delta_i} \sin \left(\frac{\Delta_i t}{2} \right) \left(\hat{a}_i e^{i \left(\frac{\Delta_i t}{2} + \phi_i \right)} + \hat{a}_i^\dagger e^{-i \left(\frac{\Delta_i t}{2} + \phi_i \right)} \right) \right). \quad (\text{A.17})$$

This is effectively the first-order only dynamics of the displacement operator leading to excursions in phase-space. The effects of this term may be entirely suppressed at certain gate durations, $A_1(T) = 0$, where T is an integer multiple of both $\tau_\alpha = 2\pi/\Delta_\alpha$ and $\tau_\beta = 2\pi/\Delta_\beta$. Excursions may also be made arbitrarily small by ensuring $\Omega_i \ll \Delta_i$.

The commutator in the A_2 term may be expanded to reveal four relevant components,

$$A_2(t) = S_\alpha(t) + S_\beta(t) + X_{\text{pair}}(t) + X_{\text{exch}}(t), \quad (\text{A.18})$$

where S denotes “self-terms”, and X “cross-terms” of the two simultaneously applied SDFs. We label the first cross-term “pair” due to the pair creation and annihilation of excitations, and the second group, “exch” is labelled due to driving the exchange of excitation.

The self-terms, with the form,

$$S_i(t) = \frac{i\Omega_i^2}{4} \left(\frac{\sin(\Delta_i t)}{\Delta_i^2} - \frac{t}{\Delta_i} \right) \quad (\text{A.19})$$

are recognisable from § A.1.2 as geometric phase factors. Taking the assumption that $t \gg \frac{1}{\Delta_i}$, we get a linearly growing geometric phase term, $S_i(t) \approx \frac{-i\Omega_i^2 t}{4\Delta_i}$. In the

case that this geometric phase acts globally on the system, they may be ignored. If the accumulated phase is spin, or oscillator state dependent, then interesting dynamics may be observed, however these effects are out of the scope of this thesis. The remaining cross-terms contain the 2nd order nonlinear bosonic interactions we are interested in emulating. We set the two detunings relative to one another by, $\Delta_\beta = m\Delta_\alpha$. Expanding the commutator of the cross terms we find for general m the following interactions, grouped by their effect on the oscillator state,

$$X_{\text{pair}} = \frac{i\epsilon_{\alpha\beta\gamma}\Omega_\alpha\Omega_\beta}{4\Delta}\hat{\sigma}_\gamma\left(\hat{a}_\alpha\hat{a}_\beta e^{i\phi_p}J(m,t) + \hat{a}_\alpha^\dagger\hat{a}_\beta^\dagger e^{-i\phi_p}J(m,t)^\dagger\right), \quad (\text{A.20})$$

$$X_{\text{exch}} = \frac{i\epsilon_{\alpha\beta\gamma}\Omega_\alpha\Omega_\beta}{8\Delta}\left(\{\hat{a}_\alpha, \hat{a}_\beta^\dagger\}e^{i\phi_m}F(m,t) + \{\hat{a}_\alpha^\dagger, \hat{a}_\beta\}e^{-i\phi_m}F(m,t)^\dagger\right), \quad (\text{A.21})$$

with $\phi_p = \phi_\alpha + \phi_\beta$, and $\phi_m = \phi_\alpha - \phi_\beta$, and J and F containing resonance like terms,

$$J(m,t) = \frac{1-m}{m}\left(\frac{e^{i(m+1)\Delta t} - 1}{(m+1)\Delta t}\right)t, \quad (\text{A.22})$$

$$F(m,t) = \left(\frac{m+1}{m}\frac{e^{-i(m-1)\Delta t} - 1}{(m-1)\Delta t} + \frac{e^{i\Delta t} - e^{-im\Delta t}}{m\Delta t}\right)t. \quad (\text{A.23})$$

We now may find the effective Hamiltonians governing the dynamic evolution for certain resonance conditions. We can convert from the generator form, $A(T)$, to familiar effective time-independent Hamiltonian, $\hat{H}_{\text{eff}}$ for a given time, T , by $\hat{H}_{\text{eff}} = \frac{i\hbar}{T}A(T)$.

Taking the limit, $m = -1$, we find $F(-1,t) = 0$ and $J(-1,t) = -2it$, leading to a time-linear term X_{pair} , with effective Hamiltonian,

$$\hat{H}_{\text{eff,pair}} = \frac{\epsilon_{\alpha\beta\gamma}\hbar\Omega_2}{2}\hat{\sigma}_\gamma\left(\hat{a}_\alpha\hat{a}_\beta e^{i\phi'_p} + \hat{a}_\alpha^\dagger\hat{a}_\beta^\dagger e^{-i\phi'_p}\right), \quad (\text{A.24})$$

where $\Omega_2 = \frac{\Omega_\alpha\Omega_\beta}{\Delta}$, and $\phi'_p = \phi_\alpha + \phi_\beta + \frac{\pi}{2}$. For $a_\alpha^\dagger = a_\beta^\dagger$ this is the squeezing interaction, whereas for $a_\alpha^\dagger \neq a_\beta^\dagger$ this is two-mode squeezing.

Taking the limit, $m = +1$, we find $J(1,t) = 0$ and $F(1,t) = -2it(1 - \text{sinc}(\Delta t))$. If we consider the secular term only, we find the time-linear term X_{exch} . Now considering the anticommutation relationships between ladder operators we find for $a_\alpha^\dagger = a_\beta^\dagger$, an effective Hamiltonian,

$$\hat{H}_{\text{eff,parity}} = \frac{\epsilon_{\alpha\beta\gamma}\hbar\Omega_2}{2}\hat{\sigma}_\gamma\sin(\phi_m)(2\hat{a}_\alpha^\dagger\hat{a}_\alpha + 1), \quad (\text{A.25})$$

which is the parity, or in this case “dispersive” Hamiltonian due to the spin-conditioning.

Finally, considering $a_\alpha^\dagger \neq a_\beta^\dagger$, we recover the beamsplitter interaction, with effective Hamiltonian,

$$\hat{H}_{\text{eff,bs}} = \frac{\varepsilon_{\alpha\beta\gamma}\hbar\Omega_2}{2}\hat{\sigma}_\gamma \left(\hat{a}_\alpha\hat{a}_\beta^\dagger e^{i\phi'_m} + \hat{a}_\alpha^\dagger\hat{a}_\beta e^{-i\phi'_m} \right), \quad (\text{A.26})$$

with $\phi'_m = \phi_\alpha - \phi_\beta + \frac{\pi}{2}$.

Similar arguments hold for the higher order terms, where, for example, we recover the tri-squeezing interaction for $m = -2$, with effective Hamiltonian of the form,

$$\hat{H}_{\text{eff,tri}} = \frac{\hbar\Omega_3}{2}\hat{\sigma}_\beta \left(\hat{a}_\alpha^3 e^{i\phi} + \hat{a}_\alpha^{\dagger 3} e^{-i\phi} \right), \quad (\text{A.27})$$

with $\Omega_3 = \frac{\Omega_\alpha^2\Omega_\beta}{2\Delta^2}$.

For further details on the use of these higher order effective interactions refer to the following two theses:[138, 139].

A.2 Gate Fidelity

We derive the fidelity of a noisy gate implementation with respect to the ideal unitary.

A quantum channel,

$$\mathcal{E} : \hat{\rho} \mapsto \hat{\rho}, \quad (\text{A.28})$$

which executes ideal unitary, U , with some error, has an average operation fidelity given by,

$$F(\mathcal{E}, U) = \int d\psi \langle \psi | U^\dagger \mathcal{E}(|\psi\rangle \langle \psi|) U | \psi \rangle, \quad (\text{A.29})$$

where the integral is taken over the uniform Haar measure on the space of pure states [192, 193]. We shall see that the language used for this analysis of single qubit gate fidelities closely resembles the group theory approach describing randomised benchmarking. The Haar measure is relevant here to consider the fidelity of the operation independent of the initial state (assuming it is pure), which is exactly

the enforced criteria through Clifford twirling during RB. For a single qubit, this can be expressed as an integral over the Bloch sphere. If we further assume that the error channel is unitary, as is the case of a detuned channel considered here, we shall find that the fidelity follows

$$F(V, U) = \frac{1}{d(d+1)} \left(|\text{Tr}(U^\dagger V)|^2 + d \right), \quad (\text{A.30})$$

where V is the error unitary, and d is the dimension of the Hilbert space.

Here we shall give a derivation of this result, by considering the symmetry imposed by the Haar measure. The standard fidelity of two unitaries following from equation A.29 is given by

$$\begin{aligned} F(V, U) &= \int d\psi \langle \psi | U^\dagger V | \psi \rangle \langle \psi | V^\dagger U | \psi \rangle, \\ &= \int d\psi |\langle \psi | W | \psi \rangle|^2, \end{aligned} \quad (\text{A.31})$$

where $W = U^\dagger V$. Using the identity $\text{Tr}(A)\text{Tr}(B) = \text{Tr}(A \otimes B)$, we can rewrite the fidelity as a two-copy expectation value,

$$F(V, U) = \text{Tr} \left((W \otimes W^\dagger) \int d\psi (|\psi\rangle \langle \psi|)^{\otimes 2} \right), \quad (\text{A.32})$$

where we have now separated the fidelity into a part dependent on the error channel, and a part equivalent to the second moment of the Haar measure, explicitly rewritten as

$$\mathcal{M}_2 = \int d\psi (|\psi\rangle \langle \psi|)^{\otimes 2}. \quad (\text{A.33})$$

We now consider two apparent symmetries of $\mathcal{M}_2$: first, the Haar measure is invariant under unitary transformations, and so our expression must commute with $U \otimes U$ for any unitary U ; second, the two states in the tensor product are identical, and so $\mathcal{M}_2$ must be invariant under the swap operator S which exchanges the two copies. These two symmetries are sufficient to show that $\mathcal{M}_2$ is proportional to the projector onto the symmetric subspace of the two-copy Hilbert space, which is given by

$$\Pi_{\text{sym}} = \frac{1}{2}(\mathbb{I} + S). \quad (\text{A.34})$$

To normalise the moment (as the Haar measure is a probability distribution), we enforce $\text{Tr}(\mathcal{M}_2) = 1$. In this two-copy expression, it is apparant that $\text{Tr}(I) = d^2$ and $\text{Tr}(S) = d$, resulting in,

$$\mathcal{M}_2 = \frac{1}{d(d+1)} (\mathbb{I} + S). \quad (\text{A.35})$$

Incorporating this result into the fidelity expression, we find

$$F(V, U) = \frac{1}{d(d+1)} \text{Tr} \left[(W \otimes W^\dagger)(\mathbb{I} + S) \right], \quad (\text{A.36})$$

where we can evaluate the two terms using the identities $\text{Tr}(A \otimes B) = \text{Tr}(A)\text{Tr}(B)$ and $\text{Tr}((A \otimes B)S) = \text{Tr}(AB)$, the latter of which is highly relevant to the mechanism of virtual distillation, described in chapter ???. This results in the final expression for the fidelity of a unitary error channel with respect to a target unitary, given above by equation A.30.

A.3 Irreducible Representations of Cl_1

In this section we provide some background group theory useful in describing how *k-copy* RB simplifies general CPTP error channels. We start with an overview of twirling, with background applicable to all RB methods. We then provide further details on the particular symmetries enforced by twirling over $Cl_1^{\otimes 2}$.

A.3.1 Twirling

Twirling a channel over a group enforces the group symmetry onto that channel — a useful method in simplifying error channels for characterisation, or to lessen the impact of structured errors on correction and mitigation protocols.

Consider a group $\mathcal{G}$ and a representation $r : \mathcal{G} \rightarrow GL(\mathcal{L})$, where GL is the general linear group (consisting of the set of invertible matrices and the matrix multiplication operation), and $\mathcal{L}$ denotes the Liouville (vectorised operator) space

Class	Representation	Member
Identity	e	I
Swap	(12)	H
Two-swap	(12)(34)	X
Three-cycle	(123)	HS
Four-cycle	(1234)	S

Table A.1: Conjugacy classes for the single qubit Clifford group.

of two qubits. Twirling over the elements g in $\mathcal{G}$ enforces symmetry on a general error channel, Λ , given in Liouville form,

$$\tilde{\Lambda} = \frac{1}{|\mathcal{G}|} \sum_{g \in \mathcal{G}} r(g) \Lambda r(g)^\dagger, \quad (\text{A.37})$$

where $|\mathcal{G}|$ is the order of the finite group $\mathcal{G}$. The twirling enforces commutation with the group action,

$$r(g) \tilde{\Lambda} r(g)^\dagger = \tilde{\Lambda} \quad \forall g \in \mathcal{G}. \quad (\text{A.38})$$

The representation of the space, $\mathcal{L}$ decomposes to irreps λ as,

$$\mathcal{L} = \bigoplus_{\lambda} (V_{\lambda} \otimes \mathbb{C}^{m_{\lambda}}), \quad (\text{A.39})$$

where V_{λ} is the irrep space, and $\mathbb{C}^{m_{\lambda}}$ is a complex valued multiplicity space with dimension of the multiplicity m_{λ} . From Schur's lemma, equation A.38 implies the twirled error channel has the form,

$$\tilde{\Lambda} = \bigoplus_{\lambda} (I_{\lambda} \otimes \mathcal{M}_{\lambda}), \quad (\text{A.40})$$

where $\mathcal{M}_{\lambda}$ is the multiplicity matrix of dimensions $m_{\lambda} \times m_{\lambda}$.

A.3.2 Character Table of Cl_1

Irreps are uniquely determined by their character on the conjugacy classes of the Clifford group. Two group elements $g, h \in G$ are said to be in the same conjugacy class if they are related by conjugation,

$$g = khk^{-1}, \quad (\text{A.41})$$

irrep	e	(12)	(12)(34)	(123)	(1234)
Identity/Trace	1	1	1	1	1
Left/Right/Antisymmetric	3	-1	-1	0	1
Diagonal traceless	2	0	2	-1	0
Symmetric off-diagonal	3	1	-1	0	-1

Table A.2: Character table for the Clifford group given the irreps in Table 5.1, and with conjugacy classes: identity, transposition, double transposition, 3-cycle, and 4-cycle. The “sign” irrep of the single-qubit Clifford group has zero multiplicity in the two-copy case, and so is excluded here.

for all $k \in G$.

With this definition, the 24 single qubit Clifford elements are divided into five permutation classes, shown in Table A.1, with a representative name and example member following each class. These classes make obvious the known isomorphism between Cl_1 and the symmetric group S_4 (modulo phase). Character tables [194], are useful tools in representation theory, to uniquely describe irreps by the traces of the representative matrices for group elements of each conjugacy class. The table is presented as a 2D matrix with rows corresponding to group representations and columns to group conjugacy classes. The character, ξ , is defined by,

$$\xi_\lambda(g) = \text{Tr}[r_\lambda(g)], \quad (\text{A.42})$$

where r_λ is the irreducible component of the representation of G acting on some vector space,

$$r_\lambda(g)_{i,j} = \text{Tr}[B_i^\dagger g B_j g^\dagger], \quad (\text{A.43})$$

where B_i are the orthonormal basis vectors for irrep λ . Taking the two-qubit sub-representations reproduced in Table 5.1, the character table calculated is shown in Table A.2. We can see that the left, right, and antisymmetric sub-representations form a multiplicity of the 3D-irrep labelled 3, while the trace sub-representation is a multiplicity of the 1D trivial irrep.

A.3.3 Projectors of $Cl_1^{\otimes 2}$

We outline the method of recovering the α parameters described in § 5.3.3 using projectors onto the irrep spaces. A special thanks to Iver Romvig Øvergaard who performed this analysis applied to the *common noise* channel and gave permission for reproducing the results here.

The decay due to a single rotation acting identically on both qubits which, in the Pauli-transfer matrix (PTM) representation, can be written as $\Lambda^{\mathcal{P}} = \text{diag}(1, R, R, R \otimes R)$, where superscript $\mathcal{P}$ represents the block diagonal matrix is ordered by the two-qubit Pauli basis $\mathcal{P} = \{I \otimes I, \vec{\sigma} \otimes I, I \otimes \vec{\sigma}, \vec{\sigma} \otimes \vec{\sigma}\}$, with $\vec{\sigma} = \{\hat{\sigma}_x, \hat{\sigma}_y, \hat{\sigma}_z\}$, and $R = R(\vec{n}, \theta) \in SO(3)$ describes a rotation by θ around $\vec{n}$. Finding the vector $\vec{\alpha}$ can now be done using Schur's lemma for the case of no mixing between irrep spaces:

$$\tilde{\Lambda}^{\mathcal{P}} = \sum_{\lambda} \alpha_{\lambda} \Pi_{\lambda}, \quad \alpha_{\lambda} = \frac{1}{d_{\lambda}} \text{Tr}(\Pi_{\lambda} \Lambda^{\mathcal{P}}). \quad (\text{A.44})$$

Here, Π_{λ} denotes the PTM projector onto the irrep space labelled λ and with dimension d_{λ} . It follows immediately that the identity sector is fixed, while the left- and right-qubit sectors produce $\alpha_L = \alpha_R = \text{Tr}(R)/3$. In principle, α_{λ} for the two-qubit sector can be evaluated by explicitly constructing the projectors from the basis operators in Table 5.1. However, in this case the symmetries of these irrep spaces allow the projectors to be expressed more simply in terms of operators whose traces are easily evaluated:

$$\Pi_{\text{tr}} = T, \quad (\text{A.45})$$

$$\Pi_{\text{dt}} = D - T, \quad (\text{A.46})$$

$$\Pi_{\text{so}} = \frac{1}{2}(I_9 + S) - D, \quad (\text{A.47})$$

$$\Pi_{\text{as}} = \frac{1}{2}(I_9 - S), \quad (\text{A.48})$$

where $T = 1/3 \sum_{ij} |ii\rangle \langle jj|$, $D = \sum_i |ii\rangle \langle ii|$, and $S = \sum_{ij} |ji\rangle \langle ij|$ are the trace projector, diagonal projector, and swap operator, respectively. The required traces are

$$\text{Tr}(I_3 R) = \text{Tr}(R) = (1 + 2c), \quad (\text{A.49})$$

$$\text{Tr}(I_9(R \otimes R)) = (\text{Tr}(R))^2 = (1 + 2c)^2, \quad (\text{A.50})$$

$$\text{Tr}(T(R \otimes R)) = \frac{1}{3} \text{Tr}(R^T R) = 1, \quad (\text{A.51})$$

$$\text{Tr}(D(R \otimes R)) = \sum_i R_{ii}^2 = c^2 + 2c + (1 - c)^2 \kappa, \quad (\text{A.52})$$

$$\text{Tr}(S(R \otimes R)) = \text{Tr}(R^2) = 4c^2 - 1, \quad (\text{A.53})$$

where $c = \cos(\theta)$, and $\kappa = n_x^4 + n_y^4 + n_z^4$ is an anisotropy parameter reflecting the fact that Clifford twirling enforces symmetry under permutations of the Pauli axes, rather than under arbitrary rotations. Substituting these traces into equation A.44 and using equation A.45 yields

$$\alpha_L = \alpha_R = \alpha_{\text{as}} = \frac{1}{3}(1 + 2c), \quad (\text{A.54})$$

$$\alpha_{\text{tr}} = 1, \quad (\text{A.55})$$

$$\alpha_{\text{dt}} = \frac{1}{2}(c^2 + 2c - 1 + (1 - c)^2 \kappa), \quad (\text{A.56})$$

$$\alpha_{\text{so}} = \frac{1}{3}(3c^2 - (1 - c)^2 \kappa). \quad (\text{A.57})$$

A.3.4 ZZ-type noise

In the main text we explore a separable error channel that induces correlations between qubits in a register. Here we shall explore entangling type noise, which although not prevalent on the ion-trap qubit platform, is a dominant error source in superconducting qubits.

The always on ZZ-type error common to superconducting qubits is described by the Hamiltonian,

$$H_{ZZ} = \frac{J_{ZZ}}{2} Z \otimes Z, \quad (\text{A.58})$$

where J_{ZZ} is the coupling strength. Generally, H_{ZZ} will couple the subrepresentations of the 3-irrep, given in Table 5.1. We can show the qualitative difference




Figure A.1: Liouville operator representation of error channels after twirling over $Cl_1^{\otimes 2}$, ordered by the Schur basis. (a) *Common noise*, (b) *Continuous ZZ*, (c) *Stochastic ZZ*. Orange blocks represent an entry of 1, green represent 0, and red represent a decay with dependence on m .

in the subspace dynamics by plotting the Liouville representation of the twirled *common noise*, and the *ZZ*-type noise. Specifically we plot $\tilde{\Lambda}^m$, where m is the sequence length. Figure 5.5 plots the maps in the Schur basis ordered so isotypic representations are adjacent to one another. Orange blocks have entries of 1, light green blocks are 0, and red blocks represent some decay parameters dependent on m . Panel (a) shows the fully diagonal *common noise* transfer matrix, where we can see the invariant trace subspace that was noted in the main text. Panel (b) shows the general case for twirled *ZZ*-type noise, given by equation A.58. The off-diagonal entries exactly couple the three subspaces of the 3-irrep. We note that the resulting block diagonal form is inconvenient to express as a general multi-exponential decay profile as given in the main text for *common* and local depolarising type noise. However, the resulting block diagonal Liouville form may be used directly with experimental data to fit these parameters of interest. Further, we may still explore the expected asymptotic decays due to this model, and find that under dominant *ZZ*-type noise, we expect an asymptote of $\mathcal{P}_{SS}|_{m \rightarrow \infty} = 1/2$.

We may recover a simple diagonal Liouville form in the case of stochastic

ZZ noise, given by,

$$\mathcal{E}_{ZZ}(\hat{\rho}) = (1 - p_{ZZ})\hat{\rho} + p_{ZZ}(Z \otimes Z)\hat{\rho}(Z \otimes Z), \quad (\text{A.59})$$

where p_{ZZ} is the probability of a correlated phase flip error on both qubits. In the Shur basis the following sector decays are found for the stochastic ZZ error,

$$\begin{aligned} \alpha_L = \alpha_R = \alpha_{\text{so}} = \alpha_{\text{as}} &= 1 - 4p_{ZZ}/3, \\ \alpha_{\text{tr}} = \alpha_{\text{dt}} &= 1. \end{aligned} \quad (\text{A.60})$$

We show this dependence in Figure A.1 (c). The same intermediate asymptote of $\mathcal{P}_{SS}|_{m \rightarrow \infty} = 0.5$ is expected.

The difference in correlated decay structure due to *common*, depolarising, and ZZ -type errors suggests that twirling over $Cl_1^{\otimes 2}$ both simplifies a generic error map to be measurable via RB, but also, retains some noise-dependent structure of the correlated decay profiles. Designing optimal initial and final states, as well as optimal sampling strategies for extracting a larger set of error parameters covering multiple physical error mechanisms is left for future work. A possible strategy to avoid introducing arbitrary error channels, is to be agnostic of the physical error mechanisms, and rather directly extract the $\vec{\alpha}$ and $\vec{m}$ parameters described in § A.3. Analogous protocols to this is described by character RB [125], and tomography via RB [195].

A.4 Derangement for $k = 2$

Huggins [92] provides a proof of the identity given in equation 6.19, which I reproduce here for the simple example of $k = 2$. The derangement operator for two copies is simply the swap operator, $S^{(2)} = \text{SWAP}$, and we will assume the observable acts

only on the first copy, giving $\hat{O}^1 = \hat{O} \otimes \mathbb{I}$. The identity is then,

$$\text{Tr}[\hat{O}^1 S^{(2)} \hat{\rho}^{\otimes 2}] \quad (\text{A.61})$$

$$= \sum_{\substack{i_1, i_2 \\ j_1, j_2 \\ k_1, k_2}} \langle i_1 | \hat{O} | k_1 \rangle \delta_{i_2, k_2} \langle k_1, k_2 | S^{(2)} | j_1, j_2 \rangle \langle j_1 | \hat{\rho} | i_1 \rangle \langle j_2 | \hat{\rho} | i_2 \rangle \quad (\text{A.62})$$

$$= \sum_{\substack{i_1, i_2 \\ j_1, j_2 \\ k_1, k_2}} \hat{O}_{i_1, k_1} \delta_{i_2, k_2} \delta_{k_1, j_2} \delta_{k_2, j_1} \rho_{j_1, i_1} \rho_{j_2, i_2} \quad (\text{A.63})$$

and contracting the Kronecker deltas, we find,

$$\text{Tr}[\hat{O}^1 S^{(2)} \hat{\rho}^{\otimes 2}] \quad (\text{A.64})$$

$$= \sum_{i_1, j_1, j_2} \hat{O}_{i_1, j_2} \rho_{j_1, i_1} \rho_{j_2, j_1} \quad (\text{A.65})$$

$$= \text{Tr}[\hat{O} \hat{\rho}^2], \quad (\text{A.66})$$

with the last equality following from the definition of matrix multiplication.

References

- [1] Paul Benioff. “The computer as a physical system: A microscopic quantum mechanical Hamiltonian model of computers as represented by Turing machines”. en. In: *J Stat Phys* 22.5 (1980), pp. 563–591.
- [2] Richard P. Feynman. “Simulating physics with computers”. en. In: *Int J Theor Phys* 21.6 (1982), pp. 467–488.
- [3] A. M. Turing. “On Computable Numbers, with an Application to the Entscheidungsproblem”. en. In: *Proceedings of the London Mathematical Society* s2-42.1 (1937), pp. 230–265.
- [4] Alonzo Church. “An Unsolvable Problem of Elementary Number Theory”. In: *American Journal of Mathematics* 58.2 (1936), pp. 345–363.
- [5] David Deutsch. “Quantum theory, the Church–Turing principle and the universal quantum computer”. In: *Proc. R. Soc. A* 400.1818 (1985), pp. 97–117.
- [6] Charles H. Bennett and Gilles Brassard. “Quantum cryptography: Public key distribution and coin tossing”. In: *Theoretical Computer Science. Theoretical Aspects of Quantum Cryptography – celebrating 30 years of BB84* 560 (2014), pp. 7–11.
- [7] P. Drmota et al. “Experimental Quantum Advantage in the Odd-Cycle Game”. In: *Phys. Rev. Lett.* 134.7 (2025), p. 070201.
- [8] David P. DiVincenzo. “The Physical Implementation of Quantum Computation”. en. In: *Fortschritte der Physik* 48.9-11 (2000), pp. 771–783.
- [9] Peter W. Shor. “Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer”. In: *SIAM J. Comput.* 26.5 (1997), pp. 1484–1509.
- [10] Madelyn Cain et al. Shor’s algorithm is possible with as few as 10,000 reconfigurable atomic qubits. en. arXiv:2603.28627 [quant-ph]. 2026.
- [11] Paul Webster et al. The Pinnacle Architecture: Reducing the cost of breaking RSA-2048 to 100 000 physical qubits using quantum LDPC codes. arXiv:2602.11457 [quant-ph]. 2026.
- [12] D. Kielpinski, C. Monroe, and D. J. Wineland. “Architecture for a large-scale ion-trap quantum computer”. en. In: *Nature* 417.6890 (2002). Number: 6890, pp. 709–711.
- [13] L. Luo et al. “Protocols and techniques for a scalable atom–photon quantum network”. en. In: *Fortschritte der Physik* 57.11-12 (2009), pp. 1133–1152.

- [14] Lukas Brenner et al. Factoring an integer with three oscillators and a qubit. en. arXiv:2412.13164 [quant-ph]. 2024.
- [15] David P. DiVincenzo. “Two-bit gates are universal for quantum computation”. In: *Phys. Rev. A* 51.2 (1995), pp. 1015–1022.
- [16] A. Yu Kitaev. “Quantum computations: algorithms and error correction”. en. In: *Russ. Math. Surv.* 52.6 (1997), p. 1191.
- [17] Samuel L. Braunstein and Peter van Loock. “Quantum information with continuous variables”. In: *Rev. Mod. Phys.* 77.2 (2005), pp. 513–577.
- [18] Seth Lloyd and Samuel L. Braunstein. “Quantum Computation over Continuous Variables”. In: *Phys. Rev. Lett.* 82.8 (1999), pp. 1784–1787.
- [19] M. Malinowski, D.T.C. Allcock, and C.J. Ballance. “How to Wire a \$1000\$-Qubit Trapped-Ion Quantum Computer”. In: *PRX Quantum* 4.4 (2023), p. 040313.
- [20] Karan K. Mehta et al. “Integrated optical multi-ion quantum logic”. en. In: *Nature* 586.7830 (2020), pp. 533–537.
- [21] Jens Eisert et al. “Quantum certification and benchmarking”. en. In: *Nat Rev Phys* 2.7 (2020), pp. 382–390.
- [22] Akel Hashim et al. “Benchmarking quantum logic operations relative to thresholds for fault tolerance”. en. In: *npj Quantum Inf* 9.1 (2023), p. 109.
- [23] Isaac L. Chuang and M. A. Nielsen. “Prescription for experimental determination of the dynamics of a quantum black box”. In: *Journal of Modern Optics* 44.11-12 (1997), pp. 2455–2467.
- [24] Alexander Erhard et al. “Characterizing large-scale quantum computers via cycle benchmarking”. en. In: *Nat Commun* 10.1 (2019), p. 5347.
- [25] David C. McKay et al. “Benchmarking Quantum Processor Performance at Scale”. In: (2023). _eprint: 2311.05933.
- [26] Dorit Aharonov, Alexei Kitaev, and John Preskill. “Fault-Tolerant Quantum Computation with Long-Range Correlated Noise”. In: *Phys. Rev. Lett.* 96.5 (2006), p. 050504.
- [27] A. M. Steane. “Overhead and noise threshold of fault-tolerant quantum error correction”. English. In: *Phys. Rev. A* 68.4 (2003). Place: College Pk WOS:000187004700041, p. 042322.
- [28] Zhenyu Cai et al. “Quantum error mitigation”. In: *Rev. Mod. Phys.* 95.4 (2023), p. 045005.
- [29] Colin D. Bruzewicz et al. “Trapped-Ion Quantum Computing: Progress and Challenges”. en. In: *Applied Physics Reviews* 6.2 (2019). arXiv:1904.04178 [quant-ph], p. 021314.
- [30] C. Langer et al. “Long-Lived Qubit Memory Using Atomic Ions”. In: *Phys. Rev. Lett.* 95.6 (2005), p. 060502.

- [31] T. P. Harty et al. “High-Fidelity Preparation, Gates, Memory, and Readout of a Trapped-Ion Quantum Bit”. In: *Phys. Rev. Lett.* 113.22 (2014), p. 220501.
- [32] P A Barton et al. “Measurement of the lifetime of the 3d 2D5 $\tilde{O}$ 2 state in 40Ca $\dot{i}$ ”. en. In: *PHYSICAL REVIEW A* ().
- [33] D.F.V. James. “Quantum dynamics of cold trapped ions with application to quantum computation”. en. In: *Applied Physics B: Lasers and Optics* 66.2 (1998), pp. 181–190.
- [34] Wolfgang Paul. “Electromagnetic traps for charged and neutral particles”. en. In: *Rev. Mod. Phys.* 62.3 (1990), pp. 531–540.
- [35] Hans Dehmelt. “Experiments with an Isolated Subatomic Particle at Rest (Nobel Lecture)”. In: *Angewandte Chemie International Edition in English* 29.7 (1990), pp. 734–738.
- [36] D. Leibfried et al. “Quantum dynamics of single trapped ions”. In: *Rev. Mod. Phys.* 75.1 (2003), pp. 281–324.
- [37] Anders Sorensen and Klaus Molmer. “Entanglement and quantum computation with ions in thermal motion”. en. In: *Phys. Rev. A* 62.2 (2000). arXiv:quant-ph/0002024, p. 022311.
- [38] Klaus Mølmer and Anders Sørensen. “Multiparticle Entanglement of Hot Trapped Ions”. In: *Phys. Rev. Lett.* 82.9 (1999), pp. 1835–1838.
- [39] Roei Ozeri. “Tutorial: The trapped-ion qubit toolbox”. en. In: *Contemporary Physics* 52.6 (2011). arXiv:1106.1190 [quant-ph], pp. 531–550.
- [40] V. M. Schäfer et al. “Fast quantum logic gates with trapped-ion qubits”. en. In: *Nature* 555.7694 (2018). Number: 7694, pp. 75–78.
- [41] S. Saner et al. “Breaking the Entangling Gate Speed Limit for Trapped-Ion Qubits Using a Phase-Stable Standing Wave”. In: *Phys. Rev. Lett.* 131.22 (2023), p. 220601.
- [42] O. Băzăvan et al. Squeezing, trisqueezing, and quadsqueezing in a spin-oscillator system. en. arXiv:2403.05471 [physics, physics:quant-ph]. 2024.
- [43] S. Saner et al. Generating arbitrary superpositions of nonclassical quantum harmonic oscillator states. arXiv:2409.03482 [quant-ph]. 2024.
- [44] O. Băzăvan et al. Synthetic $\mathbb{Z}_2$ gauge theories based on parametric excitations of trapped ions. en. arXiv:2305.08700 [cond-mat, physics:hep-lat, physics:quant-ph]. 2023.
- [45] S. Saner et al. Real-Time Observation of Aharonov-Bohm Interference in a $\mathbb{Z}_2$ Lattice Gauge Theory on a Hybrid Qubit-Oscillator Quantum Computer. en. arXiv:2507.19588 [quant-ph]. 2025.
- [46] Kaushal Choonee, Guido Wilpers, and Alastair G. Sinclair. “Silicon microfabricated linear segmented ion traps for quantum technologies”. In: *2017 19th International Conference on Solid-State Sensors, Actuators and Microsystems (TRANSDUCERS)*. 2017, pp. 615–618.

- [47] D. Allcock. “Surface-electrode ion traps for scalable quantum computing”. English. <http://purl.org/dc/dcmitype/Text>. Oxford University, UK, 2011.
- [48] V. Kaushal et al. “Shuttling-based trapped-ion quantum information processing”. In: *AVS Quantum Sci.* 2.1 (2020), p. 014101.
- [49] Jonathan D. Sterk et al. “Closed-loop optimization of fast trapped-ion shuttling with sub-quanta excitation”. en. In: *npj Quantum Inf* 8.1 (2022), p. 68.
- [50] J P Home et al. “Normal modes of trapped ions in the presence of anharmonic trap potentials”. en. In: *New J. Phys.* 13.7 (2011), p. 073026.
- [51] D. J. Wineland et al. “Experimental Issues in Coherent Quantum-State Manipulation of Trapped Atomic Ions”. In: *J Res Natl Inst Stand Technol* 103.3 (1998), pp. 259–328.
- [52] Q. A. Turchette et al. “Heating of trapped ions from the quantum ground state”. In: *Phys. Rev. A* 61.6 (2000), p. 063418.
- [53] M. Brownnutt et al. “Ion-trap measurements of electric-field noise near surfaces”. en. In: *Rev. Mod. Phys.* 87.4 (2015), pp. 1419–1482.
- [54] Ivan A. Boldin, Alexander Kraft, and Christof Wunderlich. “Measuring Anomalous Heating in a Planar Ion Trap with Variable Ion-Surface Separation”. In: *Phys. Rev. Lett.* 120.2 (2018), p. 023201.
- [55] Patrick See et al. “Fabrication of a Monolithic Array of Three Dimensional Si-based Ion Traps”. In: *Journal of Microelectromechanical Systems* 22.5 (2013). Conference Name: Journal of Microelectromechanical Systems, pp. 1180–1189.
- [56] Guido Wilpers et al. “A monolithic array of three-dimensional ion traps fabricated with conventional semiconductor technology”. en. In: *Nature Nanotech* 7.9 (2012). Number: 9, pp. 572–576.
- [57] Guido Wilpers et al. “A compact UHV package for microfabricated ion-trap arrays with direct electronic air-side access”. en. In: *Appl. Phys. B* 111.1 (2013), pp. 21–28.
- [58] Grzegorz Kasprawicz et al. “Urukul – open-source frequency synthesizer module for quantum physics”. en. In: *International Journal of Electronics and Telecommunication* 68.1 (2022), pp. 123–128.
- [59] Grzegorz Kasprawicz et al. “ARTIQ and Sinara: Open Software and Hardware Stacks for Quantum Physics”. EN. In: *OSA Quantum 2.0 Conference (2020), paper QTu8B.14*. Optica Publishing Group, 2020, QTu8B.14.
- [60] T. Harty and Thomas Peter Harty. “High-fidelity microwave-driven quantum logic in intermediate-field $^{43}\text{Ca}^+$ ”. English. <http://purl.org/dc/dcmitype/Text>. Oxford University, UK, 2013.
- [61] Magnus Danielson. AM to PM conversion of linear filters. arXiv:1805.07224 [eess.SP]. 2018.
- [62] K M Birnbaum. “Ultra-High Vacuum Chambers”. en. In: (2005).

- [63] J. Wolf. “Cryogenic, near-field quantum logic chips with passive field nulling on $^{43}\text{Ca}^+$ ”. English. <http://purl.org/dc/dcmitype/Text>. University of Oxford, 2019.
- [64] D. M. Lucas et al. “Isotope-selective photoionization for calcium ion trapping”. In: *Phys. Rev. A* 69.1 (2004), p. 012711.
- [65] R. W. P. Drever et al. “Laser phase and frequency stabilization using an optical resonator”. en. In: *Appl. Phys. B* 31.2 (1983), pp. 97–105.
- [66] Robin Oswald. “Characterization and control of a cryogenic ion trap apparatus and laser systems for quantum computing”. en. Accepted: 2023-03-17T14:31:09Z. Doctoral Thesis. ETH Zurich, 2022.
- [67] Eric D. Black. “An introduction to Pound–Drever–Hall laser frequency stabilization”. In: *Am. J. Phys.* 69.1 (2001), pp. 79–87.
- [68] Long-Sheng Ma et al. “Delivering the same optical frequency at two places: accurate cancellation of phase noise introduced by an optical fiber or other time-varying path”. en. In: *Opt. Lett.* 19.21 (1994), p. 1777.
- [69] Iver Oevergaard. “Individual ion addressing in a trapped-ion quantum processor”. en. Master Thesis. University of Oxford, 2025.
- [70] A. S. Sotirova et al. High-fidelity heralded quantum state preparation and measurement. arXiv:2409.05805 [quant-ph]. 2024.
- [71] S. Crain et al. “Individual addressing of trapped $^{171}\text{Yb}^+$ ion qubits using a microelectromechanical systems-based beam steering system”. en. In: *Appl. Phys. Lett.* 105.18 (2014), p. 181115.
- [72] Ye Wang et al. “High-Fidelity Two-Qubit Gates Using a Microelectromechanical-System-Based Beam Steering System for Individual Qubit Addressing”. In: *Phys. Rev. Lett.* 125.15 (2020), p. 150505.
- [73] M. Mazzanti et al. Alignment and Optimisation of Optical Tweezers on Trapped Ions. en. arXiv:2406.06721 [quant-ph]. 2024.
- [74] I. Pogorelov et al. “Compact Ion-Trap Quantum Computing Demonstrator”. en. In: *PRX Quantum* 2.2 (2021), p. 020343.
- [75] Lukas Perntaler. “Single Ion Addressing”. en. PhD thesis. University of Innsbruck, 2019.
- [76] Shelby Kimmel, Guang Hao Low, and Theodore J. Yoder. “Robust calibration of a universal single-qubit gate set via robust phase estimation”. In: *Phys. Rev. A* 92.6 (2015), p. 062315.
- [77] Akel Hashim et al. “A Practical Introduction to Benchmarking and Characterization of Quantum Computers”. en. In: *PRX Quantum* 6.3 (2025). arXiv:2408.12064 [quant-ph], p. 030202.
- [78] E. Knill et al. “Randomized benchmarking of quantum gates”. In: *Phys. Rev. A* 77.1 (2008), p. 012307.

- [79] Christoph Dankert et al. “Exact and approximate unitary 2-designs and their application to fidelity estimation”. en. In: *Phys. Rev. A* 80.1 (2009), p. 012304.
- [80] Chao Fang et al. “Crosstalk Suppression in Individually Addressed Two-Qubit Gates in a Trapped-Ion Quantum Computer”. In: *Phys. Rev. Lett.* 129.24 (2022), p. 240504.
- [81] D. J. Wineland and Wayne M. Itano. “Laser cooling of atoms”. en. In: *Phys. Rev. A* 20.4 (1979), pp. 1521–1540.
- [82] Michael Ramm et al. “Precision Measurement Method for Branching Fractions of Excited $\{P\}_{1/2}$ States Applied to $^{40}\mathrm{Ca}^+\{\mathbf{+}\}$ ”. In: *Phys. Rev. Lett.* 111.2 (2013), p. 023004.
- [83] Donald C. Morton. “Atomic Data for Resonance Absorption Lines. III. Wavelengths Longward of the Lyman Limit for the Elements Hydrogen to Gallium”. en. In: *ApJS* 149.1 (2003), p. 205.
- [84] F. Diedrich et al. “Laser Cooling to the Zero-Point Energy of Motion”. In: *Phys. Rev. Lett.* 62.4 (1989), pp. 403–406.
- [85] D. M. Meekhof et al. “Generation of Nonclassical Motional States of a Trapped Atom”. en. In: *Phys. Rev. Lett.* 76.11 (1996), pp. 1796–1799.
- [86] Alistair Robertson Milne. “Construction of a linear ion trap and engineering controlled spin-motional interactions”. en. In: ().
- [87] Q. A. Turchette et al. “Decoherence and decay of motional quantum states of a trapped atom coupled to engineered reservoirs”. en. In: *Phys. Rev. A* 62.5 (2000), p. 053807.
- [88] C. Monroe et al. “A Schrödinger Cat” Superposition State of an Atom”. en. In: *Science* 272 (1996). Company: C Monroe, D M. Meekhof, B E. King, David J. Wineland Distributor: C Monroe, D M. Meekhof, B E. King, David J. Wineland Institution: C Monroe, D M. Meekhof, B E. King, David J. Wineland Label: C Monroe, D M. Meekhof, B E. King, David J. Wineland Last Modified: 2021-10-12T11:10-04:00, pp. 1131–1136.
- [89] P. J. J. O’Malley et al. “Qubit Metrology of Ultralow Phase Noise Using Randomized Benchmarking”. In: *Phys. Rev. Appl.* 3.4 (2015), p. 044009.
- [90] Daniel Kienzler. “Quantum Harmonic Oscillator State Synthesis by Reservoir Engineering”. PhD thesis. ETH Zurich, 2015.
- [91] Jordan Cotler et al. “Quantum Virtual Cooling”. In: *Phys. Rev. X* 9.3 (2019). arXiv:1812.02175 [quant-ph], p. 031013.
- [92] William J. Huggins et al. “Virtual Distillation for Quantum Error Mitigation”. In: *Phys. Rev. X* 11.4 (2021), p. 041036.
- [93] Bálint Koczor. “Exponential Error Suppression for Near-Term Quantum Devices”. In: *Phys. Rev. X* 11.3 (2021), p. 031057.
- [94] Shaun C Burd. “Squeezing and amplification of trapped-ion motion”. en. In: ().

- [95] D. Hayes et al. “Coherent Error Suppression in Multiqubit Entangling Gates”. en. In: *Phys. Rev. Lett.* 109.2 (2012), p. 020503.
- [96] J. P. Gaebler et al. “Randomized Benchmarking of Multiqubit Gates”. In: *Phys. Rev. Lett.* 108.26 (2012), p. 260503.
- [97] C. A. Sackett et al. “Experimental entanglement of four particles”. en. In: *Nature* 404.6775 (2000), pp. 256–259.
- [98] P.W. Shor. “Fault-tolerant quantum computation”. In: *Proceedings of 37th Conference on Foundations of Computer Science*. ISSN: 0272-5428. 1996, pp. 56–65.
- [99] Dorit Aharonov and Michael Ben-Or. “Fault-Tolerant Quantum Computation with Constant Error Rate”. In: *SIAM J. Comput.* 38.4 (2008), pp. 1207–1282.
- [100] Ji Zou, Stefano Bosco, and Daniel Loss. “Spatially correlated classical and quantum noise in driven qubits”. en. In: *npj Quantum Inf* 10.1 (2024), p. 46.
- [101] Mohan Sarovar et al. “Detecting crosstalk errors in quantum information processors”. en. In: *Quantum* 4 (2020). arXiv:1908.09855 [quant-ph], p. 321.
- [102] David Layden, Mo Chen, and Paola Cappellaro. “Efficient Quantum Error Correction of Dephasing Induced by a Common Fluctuator”. en. In: *Phys. Rev. Lett.* 124.2 (2020), p. 020504.
- [103] John Preskill. “Sufficient condition on noise correlations for scalable quantum computing”. en. In: *QIC* 13.3&4 (2013), pp. 181–194.
- [104] Jonathan Foldager and Bálint Koczor. “Can shallow quantum circuits scramble local noise into global white noise?” en. In: *J. Phys. A: Math. Theor.* 57.1 (2024), p. 015306.
- [105] James P. Clemens, Shabnam Siddiqui, and Julio Gea-Banacloche. “Quantum error correction against correlated noise”. In: *Phys. Rev. A* 69.6 (2004), p. 062313.
- [106] Easwar Magesan, J. M. Gambetta, and Joseph Emerson. “Scalable and Robust Randomized Benchmarking of Quantum Processes”. In: *Phys. Rev. Lett.* 106.18 (2011), p. 180504.
- [107] Erik Nielsen et al. “Gate Set Tomography”. en. In: *Quantum* 5 (2021). arXiv:2009.07301 [quant-ph], p. 557.
- [108] J. F. Poyatos, J. I. Cirac, and P. Zoller. “Complete Characterization of a Quantum Process: The Two-Bit Quantum Gate”. In: *Phys. Rev. Lett.* 78.2 (1997), pp. 390–393.
- [109] Andrew M. Childs, Isaac L. Chuang, and Debbie W. Leung. “Realization of quantum process tomography in NMR”. en. In: *Phys. Rev. A* 64.1 (2001), p. 012314.
- [110] Steven T. Flammia and Joel J. Wallman. “Efficient estimation of Pauli channels”. en. In: *ACM Transactions on Quantum Computing* 1.1 (2020). arXiv:1907.12976 [quant-ph], pp. 1–32.

- [111] Robin Harper, Steven T. Flammia, and Joel J. Wallman. “Efficient learning of quantum noise”. In: *Nat. Phys.* 16.12 (2020). arXiv:1907.13022 [quant-ph], pp. 1184–1188.
- [112] Easwar Magesan, Jay M. Gambetta, and Joseph Emerson. “Characterizing quantum gates via randomized benchmarking”. In: *Phys. Rev. A* 85.4 (2012), p. 042311.
- [113] Kenneth Rudinger et al. “Experimental Characterization of Crosstalk Errors with Simultaneous Gate Set Tomography”. en. In: *PRX Quantum* 2.4 (2021), p. 040338.
- [114] Robin Harper and Steven T. Flammia. “Learning Correlated Noise in a 39-Qubit Quantum Processor”. en. In: *PRX Quantum* 4.4 (2023), p. 040311.
- [115] Erik Nielsen et al. “Efficient flexible characterization of quantum processors with nested error models”. en. In: *New J. Phys.* 23.9 (2021), p. 093020.
- [116] Jay M. Gambetta et al. “Characterization of Addressability by Simultaneous Randomized Benchmarking”. In: *Phys. Rev. Lett.* 109.24 (2012), p. 240504.
- [117] David C. McKay et al. Correlated Randomized Benchmarking. arXiv:2003.02354 [quant-ph]. 2020.
- [118] Joel J. Wallman et al. “Estimating the Coherence of Noise”. en. In: *New J. Phys.* 17.11 (2015). arXiv:1503.07865 [quant-ph], p. 113020.
- [119] Matthew Girling, Cristina Cîrstoiu, and David Jennings. “Estimation of correlations and nonseparability in quantum channels via unitarity benchmarking”. In: *Phys. Rev. Res.* 4.2 (2022), p. 023041.
- [120] Kenneth Rudinger et al. “Probing Context-Dependent Errors in Quantum Processors”. en. In: *Phys. Rev. X* 9.2 (2019), p. 021045.
- [121] Balázs Gulácsi, Joris Kattemölle, and Guido Burkard. “Characterizing correlated noise with single-qubit operations”. In: *Phys. Rev. Res.* 7.4 (2025), p. L042067.
- [122] Ethan R. Clements et al. “Lifetime-Limited Interrogation of Two Independent $\{27\}\{\mathrm{Al}\}^{\{+\}}$ Clocks Using Correlation Spectroscopy”. In: *Phys. Rev. Lett.* 125.24 (2020), p. 243602.
- [123] H. Hainzer et al. “Correlation Spectroscopy with Multiqubit-Enhanced Phase Estimation”. In: *Phys. Rev. X* 14.1 (2024), p. 011033.
- [124] R. T. Sutherland et al. Subspace Leakage Error Randomized Benchmarking of Mølmer-Sørensen Gates. en. arXiv:2510.09508 [quant-ph]. 2025.
- [125] Jonas Helsen et al. “A new class of efficient randomized benchmarking protocols”. en. In: *npj Quantum Inf* 5.1 (2019), p. 71.
- [126] Joel J Wallman and Steven T Flammia. “Randomized benchmarking with confidence”. en. In: *New J. Phys.* 16.10 (2014), p. 103032.
- [127] Jonas Helsen, Joel J. Wallman, and Stephanie Wehner. “Representations of the multi-qubit Clifford group”. en. In: *Journal of Mathematical Physics* 59.7 (2018). arXiv:1609.08188 [quant-ph], p. 072201.

- [128] Peng Zhao et al. “Quantum Crosstalk Analysis for Simultaneous Gate Operations on Superconducting Qubits”. en. In: *PRX Quantum* 3.2 (2022), p. 020301.
- [129] Simon Pettersson Fors, Jorge Fernández-Pendás, and Anton Frisk Kockum. Comprehensive explanation of ZZ coupling in superconducting qubits. en. arXiv:2408.15402 [quant-ph]. 2024.
- [130] Mohammed Alghadeer et al. Crosstalk Dispersion and Spatial Scaling in Superconducting Qubit Arrays. en. arXiv:2512.18148 [quant-ph]. 2025.
- [131] Ángel Rivas and Markus Müller. “Quantifying spatial correlations of general quantum dynamics”. en. In: *New J. Phys.* 17.6 (2015), p. 062001.
- [132] Alexei Gilchrist, Daniel R. Terno, and Christopher J. Wood. Vectorization of quantum operations and its use. en. 2009.
- [133] R. A. Fisher. “XXI.—On the Dominance Ratio”. en. In: *Proceedings of the Royal Society of Edinburgh* 42 (1923), pp. 321–341.
- [134] Jeffrey M. Epstein et al. “Investigating the limits of randomized benchmarking protocols”. In: *Phys. Rev. A* 89.6 (2014), p. 062321.
- [135] C.M. Löschnauer et al. “Scalable, High-Fidelity All-Electronic Control of Trapped-Ion Qubits”. In: *PRX Quantum* 6.4 (2025), p. 040313.
- [136] A. C. Hughes et al. Trapped-ion two-qubit gates with >99.99% fidelity without ground-state cooling. arXiv:2510.17286 [quant-ph]. 2025.
- [137] A. D. Córcoles et al. “Process verification of two-qubit quantum gates by randomized benchmarking”. In: *Phys. Rev. A* 87.3 (2013), p. 030301.
- [138] Oana Bazavan. “Easy as XYZ: Exploring non-commuting coherent interactions in trapped ions”. en. In: ().
- [139] Sebastian Saner. Quantum Control of Hybrid Spin-Oscillator Systems. en. Springer Theses. Cham: Springer Nature Switzerland, 2026.
- [140] Kai Shinbrough et al. Applications of Spin-Dependent Generalized Squeezing in Hybrid Spin-Oscillator Quantum Processors. arXiv:2608.04197 [quant-ph]. 2026.
- [141] Yuan Liu et al. “Hybrid Oscillator-Qubit Quantum Processors: Instruction Set Architectures, Abstract Machine Models, and Applications”. In: *PRX Quantum* 7.1 (2026), p. 010201.
- [142] Timo Hillmann et al. “Universal Gate Set for Continuous-Variable Quantum Computation with Microwave Circuits”. In: *Phys. Rev. Lett.* 125.16 (2020), p. 160501.
- [143] Alexandre Blais et al. “Circuit quantum electrodynamics”. In: *Rev. Mod. Phys.* 93.2 (2021), p. 025005.
- [144] J. Majer et al. “Coupling superconducting qubits via a cavity bus”. en. In: *Nature* 449.7161 (2007), pp. 443–447.

- [145] Xiaozhou Pan, Pengtao Song, and Yvonne Y. Gao. “Continuous-Variable Quantum Computation in Circuit QED”. en. In: *Chinese Phys. Lett.* 40.11 (2023), p. 110303.
- [146] Alec Eickbusch et al. “Fast universal control of an oscillator with weak dispersive coupling to a qubit”. en. In: *Nat. Phys.* 18.12 (2022), pp. 1464–1469.
- [147] A. M. Kaufman, B. J. Lester, and C. A. Regal. “Cooling a Single Atom in an Optical Tweezer to Its Quantum Ground State”. In: *Phys. Rev. X* 2.4 (2012), p. 041014.
- [148] Vincent Lienhard et al. “Generation of Motional Squeezed States for Neutral Atoms in Optical Tweezers”. In: *Phys. Rev. Lett.* 135.25 (2025), p. 253404.
- [149] Jaeyong Hwang et al. “Hybrid qubit-oscillator module from motional states of two interacting atoms”. In: *Phys. Rev. Lett.* 137.11 (2026). arXiv:2512.06429 [quant-ph], p. 113403.
- [150] R. T. Sutherland and R. Srinivas. “Universal hybrid quantum computing in trapped ions”. en. In: *Phys. Rev. A* 104.3 (2021). arXiv:2105.05768 [physics, physics:quant-ph], p. 032609.
- [151] H. C. J. Gan et al. “Hybrid Quantum Computing with Conditional Beam Splitter Gate in Trapped Ion System”. en. In: *Phys. Rev. Lett.* 124.17 (2020), p. 170502.
- [152] C. Flühmann et al. “Encoding a qubit in a trapped-ion mechanical oscillator”. en. In: *Nature* 566.7745 (2019), pp. 513–517.
- [153] D. J. Heinzen and D. J. Wineland. “Quantum-limited cooling and detection of radio-frequency oscillations by laser-cooled ions”. In: *Phys. Rev. A* 42.5 (1990), pp. 2977–2994.
- [154] S. C. Burd et al. “Quantum amplification of mechanical oscillator motion”. en. In: *Science* 364.6446 (2019), pp. 1163–1165.
- [155] S. C. Burd et al. “Quantum amplification of boson-mediated interactions”. en. In: *Nat. Phys.* 17.8 (2021), pp. 898–902.
- [156] Martín Drechsler et al. “State-dependent motional squeezing of a trapped ion: Proposed method and applications”. In: *Phys. Rev. A* 101.5 (2020), p. 052331.
- [157] Philip Leindecker et al. State-dependent Gaussian gate set using an optical tweezer for trapped ions. arXiv:2606.31864 [quant-ph]. 2026.
- [158] Michael A. Nielsen and Isaac L. Chuang. Quantum Computation and Quantum Information: 10th Anniversary Edition. en. ISBN: 9780511976667. 2010.
- [159] Scott Aaronson and Daniel Gottesman. “Improved simulation of stabilizer circuits”. In: *Phys. Rev. A* 70.5 (2004), p. 052328.
- [160] Axel M. Eriksson et al. “Universal control of a bosonic mode via drive-activated native cubic interactions”. en. In: *Nat Commun* 15.1 (2024), p. 2512.

- [161] Wilhelm Magnus. “On the exponential solution of differential equations for a linear operator”. en. In: *Communications on Pure and Applied Mathematics* 7.4 (1954), pp. 649–673.
- [162] Christa Flühmann and Jonathan P. Home. “Direct characteristic-function tomography of quantum states of the trapped-ion motional oscillator”. en. In: *Phys. Rev. Lett.* 125.4 (2020). arXiv:1907.06478 [physics, physics:quant-ph], p. 043602.
- [163] R. Cleve et al. “Quantum algorithms revisited”. In: *Proc. R. Soc. A* 454.1969 (1998), pp. 339–354.
- [164] E. Wigner. “On the Quantum Correction For Thermodynamic Equilibrium”. In: *Phys. Rev.* 40.5 (1932), pp. 749–759.
- [165] K. E. Cahill and R. J. Glauber. “Density Operators and Quasiprobability Distributions”. In: *Phys. Rev.* 177.5 (1969), pp. 1882–1902.
- [166] K. Vogel and H. Risken. “Determination of quasiprobability distributions in terms of probability distributions for the rotated quadrature phase”. In: *Phys. Rev. A* 40.5 (1989), pp. 2847–2849.
- [167] Hsiang-Yu Lo et al. “Spin–motion entanglement and state diagnosis with squeezed oscillator wavepackets”. en. In: *Nature* 521.7552 (2015). Number: 7552, pp. 336–339.
- [168] Kevin He et al. “Efficient multimode Wigner tomography”. en. In: *Nat Commun* 15.1 (2024), p. 4138.
- [169] Christophe H. Valahu et al. “Benchmarking Bosonic Modes for Quantum Information with Randomized Displacements”. In: *PRX Quantum* 5.4 (2024), p. 040337.
- [170] David Schwerdt et al. “Scalable Architecture for Trapped-Ion Quantum Computing Using rf Traps and Dynamic Optical Potentials”. In: *Phys. Rev. X* 14.4 (2024), p. 041017.
- [171] A. Walther et al. “Controlling Fast Transport of Cold Trapped Ions”. In: *Phys. Rev. Lett.* 109.8 (2012), p. 080501.
- [172] M. Valentini et al. “Demonstration of Two-Dimensional Connectivity for a Scalable Error-Corrected Ion-Trap Quantum Processor Architecture”. In: *Phys. Rev. X* 15.4 (2025), p. 041023.
- [173] Barbara M. Terhal. “Quantum error correction for quantum memories”. In: *Rev. Mod. Phys.* 87.2 (2015), pp. 307–346.
- [174] John Preskill. “Quantum Computing in the NISQ era and beyond”. en. In: *Quantum* 2 (2018). arXiv:1801.00862 [cond-mat, physics:quant-ph], p. 79.
- [175] Kristan Temme, Sergey Bravyi, and Jay M. Gambetta. “Error Mitigation for Short-Depth Quantum Circuits”. In: *Phys. Rev. Lett.* 119.18 (2017), p. 180509.
- [176] Bálint Koczor. “The Dominant Eigenvector of a Noisy Quantum State”. In: *New J. Phys.* 23.12 (2021). arXiv:2104.00608 [quant-ph], p. 123047.

- [177] T. E. O’Brien et al. Purification-based quantum error mitigation of pair-correlated electron simulations. en. 2022.
- [178] Edward Fredkin and Tommaso Toffoli. “Conservative logic”. en. In: *Int J Theor Phys* 21.3 (1982), pp. 219–253.
- [179] Yvonne Y. Gao et al. “Entanglement of bosonic modes through an engineered exchange interaction”. en. In: *Nature* 566.7745 (2019), pp. 509–512.
- [180] Honggi Jeon et al. “Two-Mode Bosonic State Tomography with Single-Shot Joint-Parity Measurement of a Trapped Ion”. en. In: *PRX Quantum* 6.4 (2025), p. 040352.
- [181] Wonhyeong Choi et al. Multimode Fock-State Measurements using Dispersive Shifts in a Trapped Ion. en. arXiv:2601.04591 [quant-ph]. 2026.
- [182] Joel J. Wallman and Joseph Emerson. “Noise tailoring for scalable quantum computation via randomized compiling”. In: *Phys. Rev. A* 94.5 (2016), p. 052325.
- [183] Tenzan Araki, Joseph F. Goodwin, and Bálint Koczor. Space-time tradeoff in networked virtual distillation. arXiv:2503.19245 [quant-ph]. 2025.
- [184] Mingyu Kang et al. “Seeking a quantum advantage with trapped-ion quantum simulations of condensed-phase chemical dynamics”. en. In: *Nat Rev Chem* 8.5 (2024), pp. 340–358.
- [185] C. L. Degen, F. Reinhard, and P. Cappellaro. “Quantum sensing”. In: *Rev. Mod. Phys.* 89.3 (2017), p. 035002.
- [186] Christophe H. Valahu et al. “Quantum-enhanced multiparameter sensing in a single mode”. In: 11.39 (2025).
- [187] S. Varona et al. Towards quantum computing Feynman diagrams in hybrid qubit-oscillator devices. arXiv:2411.05092 [quant-ph]. 2024.
- [188] Radim Filip. “Overlap and entanglement-witness measurements”. In: *Phys. Rev. A* 65.6 (2002), p. 062320.
- [189] Alfredo Ricci Vasquez et al. “Control of an Atomic Quadrupole Transition in a Phase-Stable Standing Wave”. In: *Phys. Rev. Lett.* 130.13 (2023), p. 133201.
- [190] Tomas Navickas et al. “Experimental Quantum Simulation of Chemical Dynamics”. In: *J. Am. Chem. Soc.* 147.27 (2025), pp. 23566–23573.
- [191] Vanessa C. Olaya-Agudelo et al. “Simulating open-system molecular dynamics on analog quantum computers”. en. In: *Phys. Rev. Research* 7.2 (2025), p. 023215.
- [192] M. D. Bowdrey et al. “Fidelity of Single Qubit Maps”. en. In: *Physics Letters A* 294.5-6 (2002). arXiv:quant-ph/0201106, pp. 258–260.
- [193] Michael A Nielsen. “A simple formula for the average gate fidelity of a quantum dynamical operation”. In: *Physics Letters A* 303.4 (2002), pp. 249–252.

- [194] Chin-Yen Lee et al. “On character table of Clifford groups”. en. In: *Journal of Mathematical Physics* 66.2 (2025). arXiv:2309.14850 [math.RT], p. 022201.
- [195] Shelby Kimmel et al. “Robust Extraction of Tomographic Information via Randomized Benchmarking”. en. In: *Phys. Rev. X* 4.1 (2014), p. 011050.